

Structured Writing for Content Quality

Structured Writing for Content Quality:

Table of Contents

1. Foreword	1
2. Why structure writing?	4
3. How ideas become content	7
From ideas to dots	7
The three domains	11
I. The Structured Writing Domains	14
4. Writing in the Media Domain	16
Factoring out font information	19
Using constraints to improve efficiency	20
Enabling the media domain	21
5. Writing in the Document Domain	24
Extensibility	26
Context and Structure	27
Constraining document structure	29
Backsliding into the media domain	30
Reversing the backsliding	31
Extending the document domain	34
6. Writing in the Subject Domain	37
A key difference between subject and document domains	42
Using subjects to establish context	43
7. The Management Domain: an Intrusion	45
Example: Including boilerplate content	45
An alternative approach in the subject domain	46
Hybrid approaches	48
Ad hoc management decisions	49
The management domain and content management	50
II. Algorithms	52
8. Quality in Structured Writing	57
Robots that read	57
Dumbing it down for the robots	57
Making humans better writers	58
We write better for robots than we do for humans	58
Structure, art, and science	60
Contra-structured content	61
Until the robots take over	61
9. Authoring	63
Functional lucidity	65
Granularity	66
Simplicity and Clarity	66
10. Separating Content from Formatting	69
Separate out style instructions	69
Separate out formatting characters	70

Name your abstractions correctly	71
Make sure you have the right set of abstractions	71
Create containers to provide context	73
Move the abstractions to the containers	74
Separate out abstract formatting	75
11. Processing Structured Text	78
Two into one: reversing the factoring out of invariants	78
Adding back style information	78
Rules based on structures	80
The order of the rules does not matter	80
Applying rules in the document domain	81
Processing based on context	82
Processing container structures	84
Restoring factored-out text	84
Processing in multiple steps	86
Query-based processing	88
12. Single Sourcing	89
Basic single sourcing	89
Differential single sourcing	92
Differential organization and presentation	95
Conditional differential design	97
Primary and secondary media	98
Responsive Design	99
13. Reuse	100
Fitting pieces of content together	100
Common into variable	101
Variable into common	102
Variable into variable	103
Using IDs	104
Using keys	105
Common with conditions	108
Factor out the common	110
Factor out the variable	111
Assemble from pieces	112
Combining multiple techniques	116
Content reuse is not a panacea	116
Quality traps	116
Cost traps	117
14. Composition	119
Fundamental composability	119
Structural composability	119
Stylistic composability	120
Narrative composability	121
15. Single source of truth	123
Single vs Findable	125

16. Publishing	128
The Rendering Algorithm	128
The Encoding Algorithm	129
The Formatting Algorithm	131
The Presentation Algorithm	132
Differential presentation algorithms	133
Presentation sub-algorithms	134
The linking algorithm	134
The navigation algorithm	134
The public metadata algorithm	135
The document structure normalization algorithm	135
The Synthesis Algorithm	136
Differential synthesis	136
Synthesis sub-algorithms	136
The inclusion algorithm	137
The filtering algorithm	137
Coordinating inclusion and filtering	137
The extraction algorithm	138
The merge algorithm	138
The resolve algorithm	138
Deferred synthesis	139
Combining algorithms	139
The consistency challenge	140
Functional lucidity and a layered architecture	141
17. Linking	143
Deflection in the media domain	144
Deflection in the document domain	145
Deflection in the management domain	146
The problem with IDs and Keys	147
Relationship tables	147
The problem with link tables	148
Conditional linking	149
Deflection in the subject domain	150
Finding resources to link to	153
Deferred Deflection	155
Different domain, different algorithm	157
18. Conformance	158
Completeness	160
Consistency	160
Accuracy	161
Semantic constraints	161
Entry validation constraints	164
Referential integrity constraints	165
Conformance to external sources	165
Conformance and change	167

Design for conformance	167
19. Auditing	170
Correctness of the definition of the content set	171
Ensuring the content set remains uncontaminated	172
Ensuring that the content set is well integrated	173
Making content auditable	174
Performance auditing	175
20. Relevance	176
Being relevant	176
Showing relevance	177
Showing relevance to algorithms	178
21. Active content	180
Media Domain	182
Document Domain	182
Subject Domain	183
22. Extract	185
Tapping external sources of content	185
Information created for other purposes	186
The Extraction Algorithm	188
The diversity of sources	189
23. Merge	190
24. Information Architecture	192
Top-down vs. bottom-up information architecture	193
Categorization	193
Linking	194
Tables of Contents	195
Lists	196
Dynamic content	198
25. Change management	199
26. Content Management	204
Metadata is the foundation of management	204
The location of metadata	205
Managing the process	207
Conflicting constraints	209
Content management and reuse	211
Creating manageable content	212
27. Collaboration	213
Bridging silos	215
Expanding your pool of authors	216
28. Timeliness	217
29. Repeatability	223
30. Conversion	225
31. Exchange	228
32. Translation	231
XLIFF	231

III. Structures	232
33. Rhetorical Structure	234
Irreducible rhetorical models	236
Rhetorical metamodels	238
Meta-models vs generic models	240
Making the rhetorical model explicit	242
Structure and randomness	242
34. Mechanical Structure	244
Flat vs. nested structures	245
The role of syntax	249
Agreeing on boxes, names, and labels	250
Constraint and annotation	252
Secondary structures of interpretation	254
Elements vs. attributes	256
35. Metadata	261
The recursive nature of metadata	262
Where should metadata live?	263
Taxonomy and controlled vocabulary	266
Top-down vs. bottom-up terminology control	267
Local and global names	271
Ontology	272
36. Blocks, fragments, paragraphs, and phrases	274
Semantic blocks	274
Functional blocks	277
Narrative blocks	278
Fragments	279
Paragraphs and phrases	280
37. Wide Structures	283
Tables	283
Alternatives to tables	287
Alternate presentation	288
Semantic structure	288
Record data as data	289
Don't use tables for layout	289
Code	290
Pictures and graphics	291
Inline graphics	297
38. Subject structures	299
39. Patterns	305
IV. Languages	310
40. Markup	312
Why markup?	312
Markup vs. regular text	313
Markup languages	315
Concrete markup languages	315

Abstract markup languages	316
Instances of abstract markup languages	317
Concrete languages in abstract clothing	318
The ability to extend	319
The ability to constrain	319
Showing and hiding structure	320
Hybrid languages	321
Instances of hybrid markup languages	323
41. Lightweight Languages	325
Markdown	325
Wiki markup	326
reStructuredText	327
ASCIIDoc	328
LaTex	329
Subject Domain Languages	332
SAM	333
42. Heavyweight markup languages	334
DITA	335
DocBook	339
S1000D	340
HTML	340
Subject domain languages	341
43. Extensible and Constrainable Languages	345
XML	345
DITA	346
Specializing between domains	350
DocBook	351
RestructuredText	352
TeX	353
Sam	353
SPFE	355
44. Constraint Languages	357
V. Choices	360
45. Choosing	362
Complexity	362
Scale	364
Tool driven vs. structure driven decision making	365
Index	369
A. Copyright and Legal Notices	371

List of Figures

4.1. Objects vs. dots	18
4.2. A vector graphics text object	19
4.3. Merging text and font information	20
4.4. Printer's markup	22
5.1. One document domain object, two media domain styles	28
5.2. Markdown Editor	34
37.1. Broken table formatting	284

Chapter 1. Foreword

This is a book about structured writing. It proposes a theory of structured writing, by which I mean a breakdown of the ways in which structured writing is useful for getting quality writing done efficiently.

I feel the need right off the bat to make a distinction between a theory of structured writing and a theory of markup design. Most forms of structured writing involve the use of markup languages (though markup languages are not the only way to express the structure of content). Most of the study of markup languages has not focused on their role in facilitating the writing process. Markup design theory largely concerns itself with the representation of the structure and/or meaning of texts. Its intended use is to make the texts available for further use or study. From a theoretical standpoint, it is concerned with whether and how markup truly reflects important aspects of the content of the text. In other words, markup is seen principally as a means to annotate the text, not to shape it.

The annotation of existing texts is a matter of significant concern in the humanities, where the Text Encoding Initiative¹ (TEI) plays an important role in determining standards for representing texts for study. But TEI is not about the writing process. It is not about improving the quality of text. In fact, it tries not to influence or change the text, since the text is the object being studied.

Structured writing, on the other hand, is an attempt to improve what is written. As such, its first concern is to constrain the text, to control what can and cannot be written and how it can be expressed. This sometimes means factoring out part of the text itself. The published document that results from the structured writing process may look and read differently from the structured source from which it was produced. It is a means to shape text, and only incidentally to annotate it.

The biggest consequence of this is that ease of authoring becomes a key design goal. Structured writing aims to make content better. Markup that is hard to use or to understand does not make content better because it distracts the author from the writing task. It may even disqualify the most appropriate author from writing the content at all, if they are not willing or able to learn the system.

At the same time, structured writing does impose constraints on the author. Markup that annotates without constraining does nothing to improve content quality because it provides no guidance and does nothing to detect mistakes or to make content easier to validate and audit. The right constraints help the writer create better content, but they do impose on the author's attention as they are writing. It is important to make sure that the constraints are congenial if they are not to do more harm than good.

Many existing markup languages show some balance of constraint and annotation in their design. This book aims to provide a systematic way of thinking about structured

¹<http://www.tei-c.org>

writing as a system of constraints on the writing process. It will treat markup design as a discipline supporting the expression of such constraints. Markup, for this purpose, is not an end, but a means to an end, which is the creation and delivery of high quality content.

This does not mean that it is unconcerned with markup as annotation, since content delivered as annotated markup is an increasingly important part of publishing on the Web. Highly annotated content may therefore be an important output of a modern structured writing process. This book, however, is not concerned with the theory behind what annotation is desirable in published content, but with how we can constrain the writing process to produce that annotation with the greatest consistency and quality.

I am using the term “structured writing”, even though the term “structure content” is more commonly used today. This is deliberate. The focus of this book is on writing, on the point at which, and the process by which, ideas get written down.

Many advocates of structured content talk about content primarily as a stored resource, a corporate content asset. They talk about how structure unlocks the value of that asset by allowing it to be used in multiple ways to generate more revenue. (The term “intelligent content” is increasingly used to describe this approach.) Unlocking the value of content as a corporate asset a good thing, but if we only think of structure this way we can end up designing structures for the sake of repositories and processing algorithms without thinking enough about how the author is going to actually write content in those formats.

Not thinking about the writing process can result in the definition of structures that are difficult for authors to create. This divides the author’s attention, which reduces the quality of the writing they do. It also fails to take advantage of the capacity of structured writing techniques to actively improve the quality of content. A system that focuses only on preserving and unlocking the value of content can often get in the way of creating valuable content in the first place.

As this book will show, creating high quality content and making its value more available and exploitable are not incompatible aims. By focusing on structured writing as a tool for content creation, and for content quality enhancement, we can supply content that meets all of the downstream value extraction needs as well. However, the focus in this book is on the writing process itself, and how we can use structured writing techniques to make authoring easier while improving content quality.

Some readers may detect something of a split personality in this book. On the one hand, I have very deliberately tried to cover the field of structured writing in an even handed way, and to show both the benefits and costs of the various structured writing techniques. At the same time, I have definite opinions on many topics in structured writing and I cannot pretend that these never influence the things I choose to talk about or the way I characterize them. There is definitely an element of advocacy in this work. I will, from time to time, declare my prejudices as a warning to the reader that the opinions of others on a subject may diverge from my own. I may not do this as often as I should.

Others have their prejudices as well, of course. Much of what is written about structured writing and structured writing systems comes from people who have a vested interest in one particular system or another, either because they are associated with a vendor or consultant whose business is built on a particular model, or because they have training or experience in a particular model and want those skills to continue to be relevant. That does not invalidate what they say, but the reader should be aware there may be a bigger picture.

I hope that this book will serve to provide an independent language and set of ideas in which the various approaches to structured writing can be discussed in neutral terms, including discussion of, and disagreement with, my own prejudices in the field.

Chapter 2. Why structure writing?

Creating content is a complex business. You have to research your subject matter and your audience. You have to figure out what needs to be said and how to say it to achieve your desired effect. You then have to record the content, manage it, and publish it, potentially to many different audiences and many different media. If you are part of an organization that creates a lot of content you have to coordinate what you are doing with what everyone else is doing and make sure it all serves the organization's goals. Each of these things are complicated tasks in themselves and they become all the more complex when you have to coordinate them all. Not everyone has the skills or the knowledge to do all these things well, and if too much of the complexity falls on one person's shoulders, it will limit their ability to perform well or deliver on time. In particular, writing is an activity that requires the whole of the writer's attention. Any part of the complexity of the content creation process that impinges on the writer's attention divides their attention and makes it more difficult for them to write quality content.

You can't reduce the fundamental complexity of content creation. All that stuff has to get done, and it is all difficult in one way or another. All you can do is to make sure that the inherent complexity is handled by the people or processes that are best equipped to deal with it. Managing a content creation process is all about distributing the complexity of content creation in the right way.

What happens if the inherent complexity of content creation is not managed properly? The complexity cannot be destroyed, so it has to go somewhere. If no one in the publishing organization handles it, it gets dumped on the reader. If a piece of content is missing information the reader needs, then the burden of finding the missing information falls on the reader. If the reader can't understand the term used, or read the font chosen, or find content to explain an unfamiliar concept, or get the information they need when they need it, or tell current information from outdated, that is complexity distributed to the reader. Any part of the complexity of content creation that is not handled in the organization falls down to the reader, who must then handle it themselves or abandon the content.

Another name for this is poor quality content. Poor quality content is the direct result of the complexity of content creation being distributed to the reader instead of being handled in the organization. The way for an organization to consistently create quality content, therefore, is to make sure that the inherent complexity of content creation is directed to a competent and adequately resourced person or process in the organization itself.

Structured writing is one of the most powerful tools for redistributing the complexity of content creation to the right people or processes. A lot of this has to do with taking tasks that are being performed by people and having them performed by algorithms instead.

Most complex tasks have repeatable elements in them. A design question that has been settled once can be used over and over again without having to redo the design work.

If one pieces of content is formatted one way, chances are many other pieces can be formatted the same way. Algorithms are great at doing the same task the same way over and over. Human are lousy at this.

Writing an algorithm to do a repetitive task redistributes complexity from the person who used to do the task the algorithm, but also to the person who writes and maintains the algorithm. For example, using a CSS stylesheet to format lists redistributes the task of designing list formatting from the writers, whose job is to know what they are talking about and how to say it, to a publications designer, whose job it is to know how to attractively format lists and how to code good CSS. Now the writer does not have to know about formatting or design, and the designer does not have to know about writing or the subject matter. By redistributing the complexity of formatting away from the writer, we simplify the writer's life, allowing them to focus on their part of the complexity of content delivery, composition, and we distribute it towards someone better qualified to deal with that complexity, the designer.

The problem is, content, in its natural state, is not in a format that algorithms can work with. If we are going to redistribute some aspects of complexity from writers to algorithms, we have to get the content into a form that algorithms can read. This, of course, introduces new complexity into the writing process.

All tools introduce new complexity into the processes they serve. People have to design the tools and build the tool and learn to use the tool, and maintain the tool. All this is new complexity that did not exist before. We accept this new complexity into our processes because a good tool allows us to distribute our existing complexity more efficiently, so that less of the complexity falls on the people and more on the tools, and so that each person in the system only has to deal with the complexity that they are able to cope with.

Thus introducing CSS into your web publishing process increased the total complexity in the system. There is now a new language that someone has to learn, and possibly tools that go with that language. But introducing CSS allows you to better manage how complexity is handled in your organization by directing non-writing-related complexity away from writers and by directing design complexity to designers, who are better qualified to handle it.

Adding complexity to the system in order to distribute complexity better is pervasive today. The complexity of the technology stack that lets you send an email or a text message is staggering, but its directs the complexity of communicating over distance away from the users towards systems and the people who maintain them for a huge net gain in productivity. Structured writing will undoubtedly add complexity to your content creation stack, but if you use it correctly is will help you better manage the complexity of content creation, improving the efficiency of your process and the quality of your output.

This book is about how to use structured writing technique to distribute the complexity of content creation in the way that best suits your organization. As we will see, this is not a one-size-fits-all proposition. Different organizations have different kinds of complexity

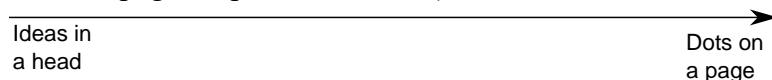
Why structure writing?

to deal with an need to distribute it in different ways to maximize the quality of the content they produce. This books will examine the principles of structured writing and show how specific structured writing techniques work to distribute specific kinds of complexity within the organization. Finding the right combination of techniques for your organization is up to you. Hopefully, though, this book will give your the information you need to figure out what will work best for you.

Chapter 3. How ideas become content

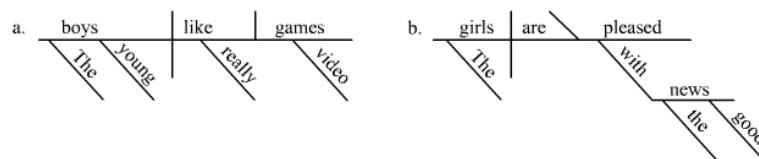
From ideas to dots

The process of creating and delivering content consists of translating ideas (stuff someone thinks or knows) into concrete physical form that can be read (dots or lines of ink on a page, or pixels on screen).



The writing and publishing process is all about how we get from ideas in a head to dots on a page. Structured writing is about applying a structured methodology to that process. It is a long road from ideas to dot, and structured writing techniques can be applied at many points along that road. Almost all writing done today uses structured writing techniques to one extent of another. As this book will show, the basic principles of structured writing apply across the spectrum, from the tools and techniques used in most offices today, to most sophisticated structured writing systems. It is not a matter of unstructured vs structured approaches, but the degree of structure used.

All writing has structure in the literary sense of the word. Every comprehensible sentence has a grammatical structure. You may even have learned to diagram that structure in school.¹



Many texts also have deliberate rhetorical structure: a particular way of assembling words and other media to effectively inform or persuade an audience. Sometimes that rhetorical structure is ad hoc, invented by the writer for a unique individual piece. But often it is repeated for many similar pieces of content. Thus newspapers stories have a repeated rhetorical structure known as the pyramid model, in which the most important facts are promoted to the opening sentences of the pieces. Recipes have a very familiar structure, with an introduction, a list of ingredients, and a list of preparation steps.

Hard Boiled Eggs

A hard boiled egg is simple and nutritious. Prep time, 15 minutes. Serves 6.

¹By Tjo3ya - Own work, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=18312612>

Ingredients

- 12xeggs
- 2qt water

Preparation

1. Place eggs in pan and cover with water.
2. Bring water to a boil.
3. Remove from heat and cover for 12 minutes.
4. Place eggs in cold water to stop cooking.
5. Peel and serve.

As a writer, you don't have to stop to think about the rhetorical structure of a recipe: you follow the established pattern. Many other subjects have similarly well known rhetorical patterns used to describe them. These are all examples of structured writing in the rhetorical sense, and it is well known that adherence to sound rhetorical models greatly improves the quality of writing. They do this by constraining writers to follow the established rhetorical pattern.

This idea of constraints is fundamental to structured writing. Whenever you specify a structure to be followed, you are imposing a constraint. Constraints define what is and is not required or acceptable in a piece of writing. If you have ever been asked to write following a template or a style guide, you have been asked to follow a set of constraints.

However, the term structured writing today generally means more than following a style guide or template. It usually means the use of computer data formats (most commonly XML) to capture written content so that it can be processed by algorithms. These data structures are also constraints. They require the writer to write content to fit into a set of specific data structures. These mechanical constraints may or may not correspond to the rhetorical constraints of a particular piece of writing. In some cases, the mechanical constraints are there to serve a publishing or content management purpose and may have nothing to do with the rhetorical structure of the writing. In some cases they may actually interfere with it.

However, it is also possible to use mechanical structures like this to help improve the rhetorical structure of content. This is the kind of structured writing that this book is about. From this point on, therefore, I am going to use structured writing to mean writing in data structures. As we go on, we will show how those data structures can give greater and greater support to the literary and rhetorical structure of content, and that this

approach can be be hugely beneficial in managing how complexity is distributed in your content creation process.

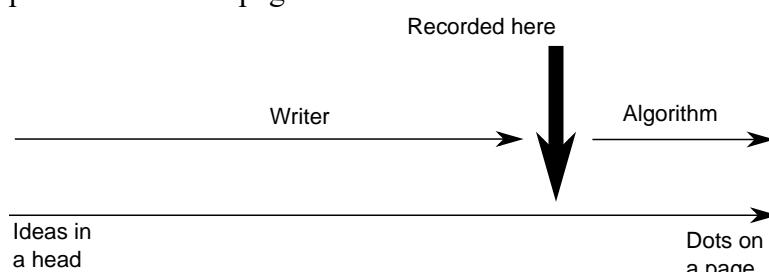
Just as all writing has at least a basic grammatical structure, all writing on a computer involves creating basic data structures. Thus the only case in which no structured writing techniques are involved in the writing process is when an author writes down their ideas with pen and paper, and gives that paper directly to the reader. In this case, the entire writing process, from an idea in the writer's head to words on paper takes place in the writer's head and any constraints that are imposed on the content are captured and imposed by the writer themselves as they inscribe their words with the pen.



It is rare for writers to record their ideas directly in the final physical form these day. For instance, the writer may write in a word processor, edit the text on screen, and press Print to sent the content to a printer to create the final form that the reader receives. Or they may press Send and have the final form rendered on someone else's monitor. Either way, recording the content happens before the final form is created, and the way the content looks when you write may be different from how it looks when the reader sees it in final form. The point at which the content is recorded, in the journey from ideas to dots, has been moved back just a little bit.



Word processing, desktop publishing software, and various approaches to structured writing all establish a point in this process from ideas to dots where the content will be recorded by the author, and then provide algorithms to complete the movement from that point to dots on a page.



I'm going to use the word algorithm a lot in this book. An algorithm is a formalized and consistent way to do something. Basically, if you give an algorithm the same inputs, it should always produce the same outputs every time it is run. Computer software is an encoding of an algorithm that a computer can execute. But an algorithms is not the same thing as a program. An algorithm is a method. A program describes the algorithm to a computer. You do not need to be a programmer to create an algorithm. Nor do you need a computer to execute an algorithm. Human being can execute algorithms as well.

The reasons I am going to talk about algorithms and not software is that algorithms are fundamental to structured writing. Algorithms and structures work together and you can't design one without the other. The reason you add structure is to enable algorithms. The heart of this book is a description of the principal structured writing algorithms and the structures that support them. The goal of structured writing is to make content better, and it is the algorithms, not the structures that do that. The structures exist to support the algorithms.

Constraints play an important role for algorithms. First, they help assure the quality of the data that the algorithms consume by controlling what authors create. Secondly, they constrain the interpretation of the data. In other words, the structures that authors create in the content tell the algorithms what each piece of content means. The more precisely the structures constrain the interpretation of the data, the more the algorithms can do with that data.

We use structured writing algorithms to help us with the writing process. We use them either to take over or to support parts of the writing task. This usually involves moving the point at which content is recorded earlier in the process from idea to dots, allowing the algorithms more scope to assist in the process.

This leads us to a working definition of structured writing for purposes of this book:

Structured writing is writing that follows a set of constraints and records that it has done so.

For instance, lets say that we have a constraint that every section must start with a title. We can follow this constraint in just about any tool we choose. We could even follow in when writing by hand on paper. But if we choose a structured writing format, such as DocBook, we can not only follow this constraint but record that we have followed it:

```
<section>
  <title>The quick brown fox</title>
  <para>The quick brown fox jumps over the lazy dog.</para>
</section>
```

This markup records that we have followed the constraint by identifying the title as the first thing in the section. This constrains the interpretation of the string "The quick brown fox" by algorithms. Algorithms can now do things such as:

- Check that all sections start with titles (that is, check that the constraint has been followed)
- Format all section titles the same way
- Create a list of all the section titles in the document

All of these redirect tasks away from the author, redistributing those bits of complexity to the algorithms and the people who write them.

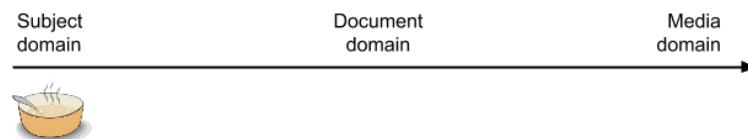
Following constraints can make content better, but it is in recording the constraints that we have followed that we constrain the interpretation of the text and therefore make it available to algorithms. And it is in making it available to algorithms that we gain the ability to redistribute the complexity of the content creation process.

Generally speaking, the more we constrain the creation and interpretation of content with structure, the more accessible it becomes to algorithms, and the more opportunities we gain to distribute complexity effectively. This means recording the content at a point nearer to ideas in the head and further from dots on a page. The nearer we are to ideas in a head, the more we know about what the content says as opposed to what a page looks like, and the more effectively we can constrain the creation and interpretation of the text.

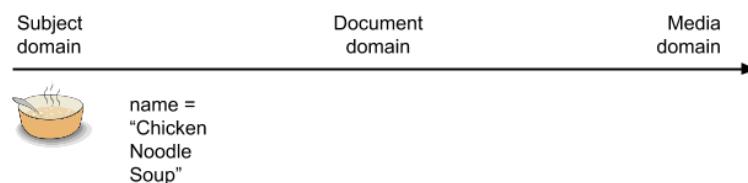
We can describe this process of earlier recording in terms of three domains, each domain reflecting a stage in the progress from ideas to dots. The domains are the media domain (which is concerned with lines and dots on paper or screen), the document domain (which is concerned with the expression and organization of documents), and the subject domain (which is concerned with the ideas that we write about).

The three domains

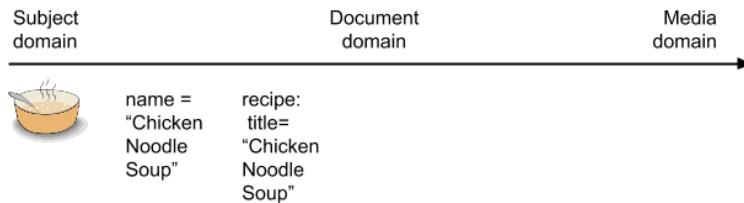
Let's suppose that an author is planning to write a recipe for chicken noodle soup. They start out with the idea of a soup made with chicken and noodles. This is an idea about the subject matter and not yet any form of content.



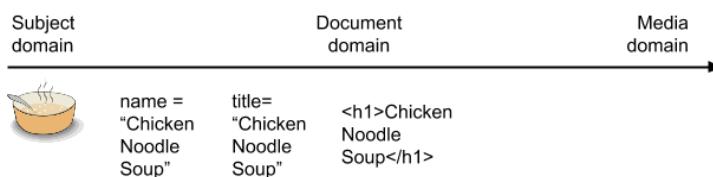
They then decide to give the dish the name “Chicken Noodle Soup.” That name is content. However, it is not yet part of a document. It is a piece of data in the subject domain.



Then the author composes a recipe for Chicken Noodle Soup. They are now composing a document that has structures like a title, a list of ingredients, and a preparation procedure. They use the name of the dish as the title of the recipe. That title is a piece of data in the document domain, but because it is the title of a recipe, it still has a strong association with the subject domain (a recipe is a document with a very specific subject).

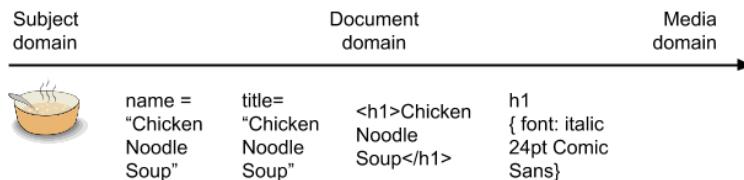


As the process continues, some format for publishing the recipe is chosen. This might be HTML. HTML is a markup language in the document domain. That is, it contains markup for typical document structures such as titles, paragraphs, and lists. To express the title of the recipe in HTML, we turn the declaration of the title of the recipe into a declaration about an HTML heading level:



This breaks the association of the title with the subject domain. At this stage we are recording that “Chicken Noodle Soup” is the title of a web page, not that it is the title of a recipe, and not that it is the name of a dish. However, we have now specified what type of document this is: a web page.

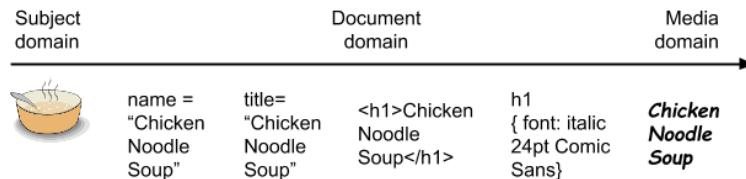
If you follow modern practice, your HTML should not include any specific information about how the document will be rendered -- what fonts will be used, how big the margins will be, etc. These things should be specified by a CSS stylesheet.



Adding the CSS moves the entire package further into the media domain. Notice, however, that in this case, all of the document domain information in the HTML format has been retained. The movement from the subject domain through the document domain into the media domain does not have to mean that information from an earlier domain has to be thrown away as you add information from the later domains. For instance, HTML5 microformats allow you to retain subject domain information all the way into the user’s browser. Traditional publishing techniques tended to throw away subject domain

information as the process advanced. Today, that information is more likely to be retained as long as possible. (One of the implications of the term “intelligent content”.)

When the resulting page is loaded into a browser, dots are painted on the screen in the appropriate shape.



This process involves translating the document and media domain information in the HTML and CSS into the graphics primitives of the platform on which the content is displayed -- basically lighting up dots on a screen or printing dots on paper. This final step will destroy all document and subject domain information, but since the browser retains the HTML source, any information that has reached that point is available to code running in the browser and to other services running on the Web.

All content passes through the three domains. Content always begins with the author thinking about subjects in the real world. They then decide to express ideas about those subjects in words. They collect their ideas together and determine an order and structure to express them. Finally, they decide how they will be formatted in a particular media. The question is, where in this process that all content goes through does the author start recording the content.

Where did the author in this scenario actually record their content? They could have recorded it in the subject domain and let an algorithm take it from there. Since the chosen output format is HTML, which is a document domain language, they could not have delayed recording any longer than that. Actually, that's not entirely true. If they used a graphical HTML editor that allowed them to add styles by hand, they could actually have done their work in the media domain, even if a document domain tagging language was being created under the surface. Of course, if they had been intended to publish the recipe on paper, perhaps just making one copy to give to a friend, they could have delayed recording it until the very end of the media domain. Wherever they recorded it, however, their thought process in creating the finished recipe would have had to go through all three domains.

An content must pass through all three domains to get from ideas in a head to dots on a page. But the author can record that content in the media domain, the document domain, or the subject domain. In the next three chapters we will look at what it is like to write in each domain.

Part I. The Structured Writing Domains

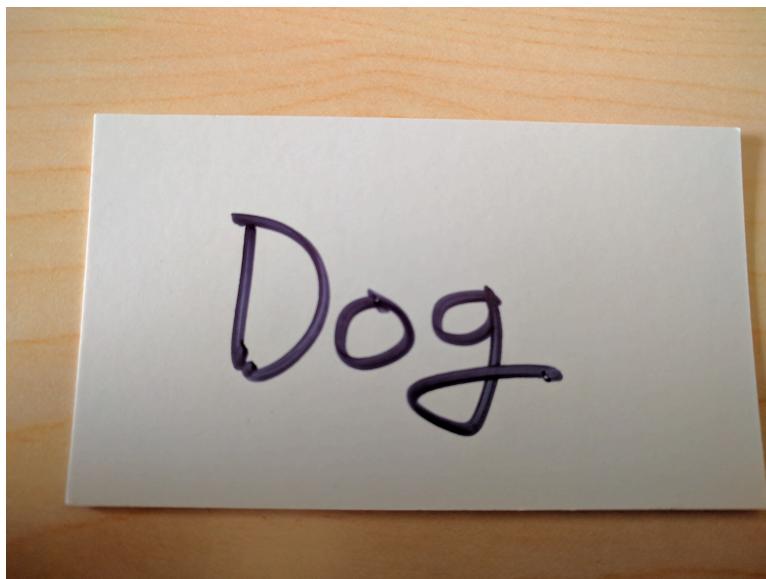
Table of Contents

4. Writing in the Media Domain	16
Factoring out font information	19
Using constraints to improve efficiency	20
Enabling the media domain	21
5. Writing in the Document Domain	24
Extensibility	26
Context and Structure	27
Constraining document structure	29
Backsliding into the media domain	30
Reversing the backsliding	31
Extending the document domain	34
6. Writing in the Subject Domain	37
A key difference between subject and document domains	42
Using subjects to establish context	43
7. The Management Domain: an Intrusion	45
Example: Including boilerplate content	45
An alternative approach in the subject domain	46
Hybrid approaches	48
Ad hoc management decisions	49
The management domain and content management	50

Chapter 4. Writing in the Media Domain

The media domain is the structured writing domain in which the structures relate to the media in which the content is displayed. Such content is often considered “unstructured”, but all content has structure, and we can actually find all the patterns and techniques of all forms of structured writing in the media domain. This makes it a good place to study the fundamentals of structured writing.

At its most basic, a hand guiding the pen over paper or chisel into stone is working in the media domain through direct physical interaction with the media.



The closest you can get to pen and paper in the computer world is to use a paint program to directly place dots on the screen. You can select the pen tool and use your mouse or a stylus to write your text. This will record the text as a matrix of dots.



There is very little structure here. We are recording a pattern of dots. Those patterns of dots are text characters only in the sense that the patterns are recognizable as characters to the human eye. The computer has no idea they are characters.

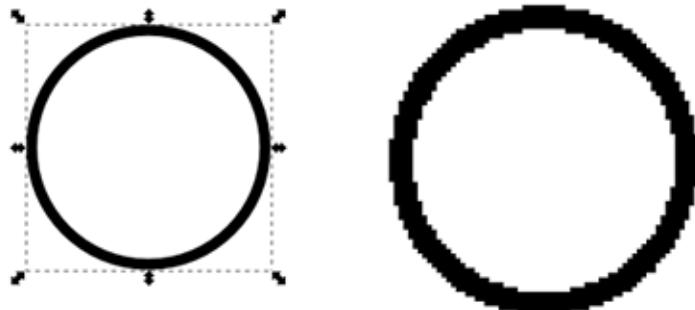
This is a pretty inefficient way to write. You can work faster if you use the paint program's text tool.



This allows you to type letters on the keyboard. However, those letters are still recorded as a set of dots, not as characters, so you can't go back and edit your text as text, only as dots. The computer is providing a slightly more structured way to enter the text, but is not maintaining any more structure in recording the text. (This is an important lesson: Some tools record less structure than you may think. For instance, some template-based authoring solutions don't record which template was used, so the template does not guide subsequent editing or use of the content. Structure that is used but not recorded has no downstream value.)

To edit more efficiently, you need to move away from dots and start working in a program that records characters as characters. You could go to a text editor, but a text editor does not keep any formatting information (unless you create markup -- but that would be getting ahead of ourselves). For most publishing purposes, plain text is inadequate. We need to maintain the ability to format the document.

One type of program that lets us record text as text and also lets us attach formatting to the text is a vector graphic program. A vector graphics (or “draw”) program creates graphics as a collection of objects (“vector” meaning the mathematical representation of a shape or line). For example, it allows you to create a circle as a shape, described mathematically in computer memory, rather than as a set of dots. Rather than recording an actual circle, the program records an abstraction of a circle: the essential properties needed to reproduce an actual circle on a media, such as its center, diameter, and line weight. The computer then lets you manipulate that abstraction as an object, only rendering it as actual dots when the graphic is displayed on screen or paper.

Figure 4.1. Objects vs. dots

A circle as an object, displayed in a vector graphics program (Inkscape), left, vs. a circle as a set of dots in a raster graphics program (Microsoft Paint), right.

In a typical vector graphics program, a shape is rendered into dots on screen instantly as you draw or edit the shape. Nonetheless, the computer is storing data describing the shape, not a circular pattern of dots, as it would in a paint or photo editing program. This is an instance of what in structured writing is called “separating content from formatting”. The mathematical abstraction of a circle is the content; the dots that represent it on screen are the formatting, or rather, the result of applying formatting to the object.

To give you some idea of how a circle might be represented as an object, here is a sample piece of Java code for creating a circle:

```
Circle circle = new Circle();
circle.setCenterX(100.0f);
circle.setCenterY(100.0f);
circle.setRadius(50.0f);
```

This code creates a new circle object and specifies three values that you would need to draw a basic circle, an X and Y coordinate for the center of the circle, and the radius. When a vector graphics file is displayed or printed, a rendering algorithm turns those objects into a matrix of dots on the current media, using more dots for higher resolution media and fewer dots for low resolution media.

If representing the circle as an object creates an abstraction of “circle” in the media domain, this bit of PostScript code takes that abstraction (a circle as a set of coordinates) and (in concert with the underlying graphics driver) makes it concrete again (dots on a media).

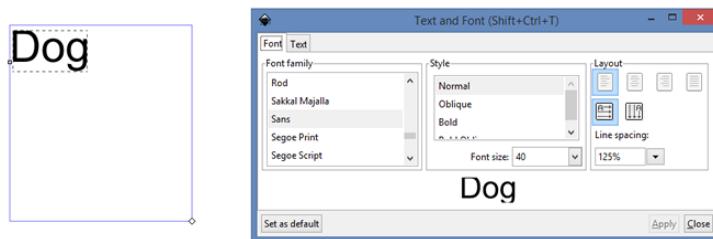
The point here is not to understand the code, but to see that by creating a more abstract form of the content we can make it easier to work with, and that we can use algorithms to render it back into concrete form. All of structured writing comes down to finding better

ways to represent content for a particular purpose (more efficient, more versatile, more verifiable, more precisely constrained) and ways to transform those representations, into dots on a screen or paper so that the content can be read.

This is a pattern we will see as we look at structured writing across the three domains. Rather than storing the image of the document to be printed, the computer stores an abstraction of the document which consists of raw text combined with (or, in some cases, replaced by) additional pieces of information which are commonly called metadata. This metadata can then be used to drive algorithms to convert it back into a concrete media domain representation that can be displayed to readers (and also to do a number of other useful things that we will look at later).

Just as a vector graphics program represents a circle as a circle object, it represents text as a text object. A text object is a rectangular area that contains characters. It has numerous media domain properties, such as margins, background and foreground colors, the text string, and the font face, size, and weight used to display that text, as in this example from Inkscape:

Figure 4.2. A vector graphics text object



A text object in a vector graphics program, with object properties shown.

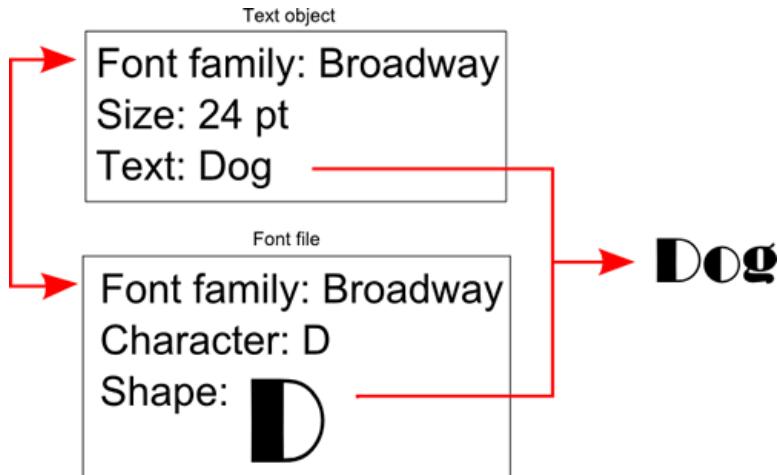
Factoring out font information

A vector graphics program displays text in a chosen font. If you change the value of the text object's font attribute, it will immediately redraw the text in the new font. The shapes of the individual letters in the font are required information for rendering the text object in the media domain. However, they are not stored as part of the text object. The representation of the text in the paint program includes the shape of the letters. In the vector graphics program it does not. That information has been factored out.

The shape of the characters (technically, “glyphs”) that make up the font are stored separately in font files. Font files consist of a set of shape objects that describe each glyph, together with metadata such as the name of the font and the name of each glyph. To actually display the text block on screen, the graphics program (or rather the graphics system API to which it delegates this task) combines information from the font file with

information from the text object by matching the metadata to find the right font and character, and then drawing the appropriate glyph on the current media.

Figure 4.3. Merging text and font information



The vector graphic text object factors out letter shapes to a separate font file.

This is a pattern we will see over and over again in structured writing. In order to simplify the objects that we create to store our content, we take part of the information needed to do the final rendering, and move it to a separate file. By factoring out information that is constant for a given application (the shape of a capital D is the same for all capital D's for text in a given font), we simplify the format of the information we are preparing and keep the downstream presentation more consistent.

Adding more structure to content means adding more metadata to it. But if we just kept adding metadata, it would very quickly overwhelm the content. So whenever we can we carve some metadata off into separate files and create rules for pulling it back in when it is time to publish the content.

Whether it is, “the capital D will always be this shape” or “the list of ingredients will always have the ingredient name aligned left and the quantity aligned right”, filtering out these invariant rules into a separate file is a key part of structuring content.

This means that designing a content structure, regardless of the domain you choose to work in, essentially consists of identifying the places in the content where we can separate out these invariant properties into separate structures: separating variants from invariants using metadata.

Using constraints to improve efficiency

Writing a document in a vector graphics program is certainly better than in a paint program, but it is hardly ideal for writing long documents. This is why most of us use tools that are designed specifically for writing documents.

A vector graphics program works purely in the media domain, and pretty much lets you put shapes and text boxes anywhere you like. This gives you enormous freedom, but it also makes you do a lot of extra work if what you want to create is a typical document that is basically one long text flow with some headings and graphics thrown in.

Word processors and desktop publishing programs make it easier to create documents by introducing some document domain constraints, as well as some higher-level tools for managing large text flows. A document is made up of a series of pages that have margins and contain text flows. Text flows are made up of blocks (paragraphs, headings) inside of which text can flow, even from one page to the next. Common features like tables are supported as objects than can exist in text flows. New pages are created automatically as text expands.

Pages, paragraphs, headings, and tables, are all document domain objects. Rather than working on a blank slate, as you do in a graphics program, you are now working within the constraints of these document domain objects. These constraints remove or constrain decisions about positioning of elements, which makes creating documents faster and more consistent. Structured writing is about making content that obeys constraints, and these basic document domain constraints are the next step in that journey.

These constraints are not without their negative consequences. You always give up something when you impose a constraint. There are certain page layout effects that are difficult or impossible to achieve in Word or FrameMaker because you have given up some of the liberty of a vector graphics program. (You also gave up some liberty in moving from raster to vector graphics, which is why photo editing, which requires adjusting individual pixels, is done in raster rather than vector format.)

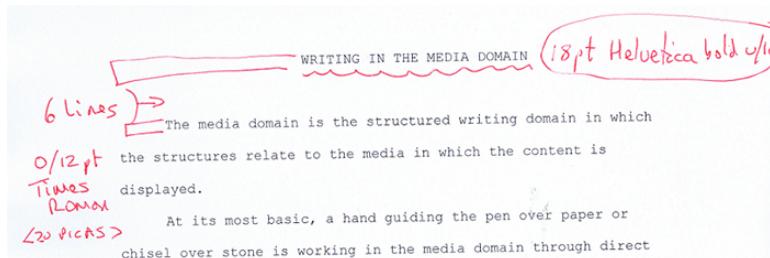
But while they make authoring easier by introducing document domain constraints into the program, both word processors and desktop publishing programs still have one foot firmly planted in the media domain. While there are some basic document domain objects being created under the surface, the author sees and interacts with the media domain representation of those objects on the screen. And the way that these programs allow writers to distinguishing one block of text from another is almost always be applying formatting styles. The document domain objects they provide are just enough to give the author something to hang media domain styles on.

Enabling the media domain

Providing the ability for the author to work in the media domain was at the heart of the desktop publishing revolution. For centuries, scribes worked directly in the media domain, using pen and ink to inscribe words and pictures on paper or velum. With the printing press, however, authors no longer worked directly in the media domain. While authors were still directly placing ink on paper, at first by pen and then by typewriter, they were no longer preparing the final visual form of the content. That would be created later by the typesetter.

To tell the typesetter how to create the final visual form, document designers had to add additional instructions (metadata) to the author's manuscript. The designers did this using typesetter's marks, and the process was called "marking up" the document. We still use "marking up" to describe how structured writing is done today.

Figure 4.4. Printer's markup



The writer preparing a manuscript for typesetting was working in the document domain, indicating basic document structures like paragraphs, lists, and titles, without any indication of how they should look in print. The designer then wrote a set of instructions for applying formatting to those structures -- a formatting algorithm. Then the typesetter executed that algorithm by setting the type which the printer then used to print final output.

This is pretty much exactly what we do today when we create an HTML page and specify a CSS stylesheet to supply the formatting instructions. Those instructions are then executed by the browser to render the content on screen or paper.

Actually, we are getting ahead of ourselves here. A better analogy to old style typesetter's marks is an HTML page with the styles specified inline.

```
<p style="font-family: serif; font-weight: bold; font-size: 12pt">
```

You can see that this markup is very very similar to the old typesetters marks in Figure 4.4, "Printer's markup".

The intent of the desktop publishing movement was to combine the job of designer and writer into one, effectively moving the author from working in the document domain to working in the media domain, formatting content as it was written.

The inefficiency of applying formatting directly to every part of the document soon became apparent, however. Stylesheet mechanisms were devised to separate the media domain elements of a document from the document domain elements.

Using styles rather than direct formatting does not mean that you have moved from working in the media domain to working in the document domain. It just means that you have moved from working in an unstructured way in the media domain to working in a structured way in the media domain.

The set of document domain objects in a typical word processor is so small that it does not even distinguish between a heading and a list item. Both are just paragraphs with different styles applied. What distinguishes them is not what type of structure they are, but what styles are applied to them.

Relying on just a few document domain structures and distinguishing them with styles makes word processing simple, but it can lead to problems. Microsoft Word's difficulties with maintaining list numbering are well documented, for instance. Because Word does not formally record lists as distinct structures, it can easily lose track of where they begin and end.

Word processing and desktop publishing tend to roughly straddle the boundary between the document and media domains, with authors thinking partially in document domain terms and partially in media domain terms, and the applications creating and storing structures from both domains.

While we can clearly divide the structured writing spectrum into the three domains, many actual writing tools and markup languages include structures from more than one domain. We will see this over and over again in this book. But next we will look at what it is like to work more squarely in the document domain.

Chapter 5. Writing in the Document Domain

Word processors and desktop publishing applications tend to straddle the divide between the media domain and the document domain. While they are built on a basic set of document domain objects -- pages, paragraphs, tables, etc. -- they use a WYSIWYG display to keep the author working and thinking mostly in terms of styles and formatting -- the concerns of the media domain. This makes it difficult to apply meaningful document domain constraints to the author's work, or to record which constraints the author has followed. This makes it impossible to redistribute formatting complexity away from the author. For that we need to move to the document domain.

The simplest reason for moving to the document domain is actually to enforce media domain constraints that are hard to enforce in the media domain itself. (In fact, the reason for moving to the next domain is often to enforce -- or factor out -- constraints in the previous domain. This is one of the consistent patterns we find in structured writing.)

Consider a list. You may want to impose a constraint that the spacing above the first item of a list must be different from the spacing between other items of the list. This is a media domain constraint -- it is about formatting, not the structure of the document -- but it is hard to enforce in the media domain. Most media domain applications create lists by applying styles to ordinary paragraphs. The usual way to apply the extra space above the first item is to create two different styles, which we can call first-item-of-list style, and following-item-of-list style. The first-item-of-list style is defined with more spacing above.

This creates a constraint that the author is required to follow. They are required to apply first-item-of-list style to the first item of a list and following-item-of-list style to the rest. Nothing in the word processor enforces this constraint. The author simply has to remember it. This creates two problems:

1. It makes authoring a little bit harder (and all the little bits add up).
2. The author can get it wrong without anyone noticing. They may forget to apply first-item-of-list style or they may add a new first item to the list and not realize that the second item in the list now has first-item-of-list style.

As we noted before, structured writing works by factoring out invariants. Most constraints are invariants -- that is, they are rules that apply to all instances of the same content structure (such as all lists must have extra space before the first item). The easiest way to enforce a constraint is not to enforce it on the author, but to factor it out altogether.

To do this, we create a list structure -- not a styled paragraph, but a structure that is specifically a list. The word processor version of a list is structured like this, as a flat list of styled paragraphs:

```
p{first-item-of-list}: Carrots
p{following-item-of-list}: Celery
p{following-item-of-list}: Onions
```

An explicit list structure looks like this:

```
list:
  list-item: Carrots
  list-item: Celery
  list-item: Onions
```

List structures like this belong to the document domain because they are a common rhetorical tool that is not specific to any one subject area and can be formatted in a wide variety of ways.

Once we have a list object, we can create rules -- in a separate file -- about how lists are formatted. We are familiar with this from HTML. Here's that same structure in HTML (actually, this is a slightly more specific structure, but we'll get to that):

```
<ol>
  <li>Carrots</li>
  <li>Celery</li>
  <li>Onions</li>
</ol>
```

Now that we have a distinct list object, we can factor out our invariant list formatting rule into a separate file. For HTML, this is usually done with a CSS stylesheet:

```
li:first-child
{
  padding-top: 5pt;
}
```

Now the correct spacing above lists is not something the author has to think about. In fact, it is not something they can manipulate even if they want to. Authors just create list structures. Media domain list formatting rules have been factored out of the author's world. The media domain constraint about spacing above a list will now be followed automatically and reliably by algorithms. Authoring gets a little bit simpler again (the author is already thinking "list" as they write) and the possibility of error is removed.

But wait! That's fine if all lists are formatted exactly the same way, but we know that is not true. At very least, some lists are bulleted and some are numbered. And then there

are nested lists, which are formatted differently from their parents, and specialized lists, like lists of ingredients, definitions, or function parameters. If we are going to create list structures in the document domain rather than applying list styles in the media domain, how do we make sure each of these types of lists gets formatted appropriately?

Extensibility

At this point it is worth looking at a very important feature of all structured writing systems: extensibility. In media-domain word processing and desktop publishing programs, authors may need many different styles to format their documents, and these applications do not attempt to anticipate or provide all the styles every author might need. Some, like Word, come with a basic set of styles that may meet some basic needs, but all these programs let authors define new styles as well. The set of document domain objects in these programs is small and fixed, but the set of media domain styles is extensible -- you can create as many as you need.

A word processor or desktop publishing application that supports the definition of styles is essentially creating an extensible media domain environment. Styles are media domain structures that factor out a set of style metadata that you can attach to a block of text to specify how it will be displayed. Every time you create a new style you are extending your set of media domain structures.

This need for extensibility is another common pattern in structured writing. If you are working in the media domain, you may need to extend your set of styles. If you are working in the document domain, you may need to extend your set of document domain structures.

If we need more than one type of list structure in the document domain, we need to extend our document domain language with new lists types. But how many types do we need?

One obvious formatting difference between lists is that some are numbered and some are bulleted. How does a formatting algorithm tell whether to use bullets or numbers to format a given list? One way would be to add a style attribute to specify bullets or numbers, but then the author would be working in the media domain again. To keep the author in the document domain, we need to create document domain structures that contain the metadata we need to make those decisions at the formatting stage.

The common way to handle bullets vs numbers is to create two different list structures, the ordered list and the unordered list. Different markup languages call them by different names -- `ol` and `ul` in HTML, `orderedlist` and `itemized-list` in DocBook, for example -- but they are conceptually the same thing. Thus the HTML example above is a little more specific than just being a list structure. It is an ordered list structure.

The choice of the terms “unordered” and “ordered” is important, because it focuses on the document-domain properties of a list -- whether its order matters -- rather than on

its media domain properties -- bullets or numbers. Whether an ordered list should be formatted with numbers or letters or Roman numerals, belongs entirely to the media domain. It has been factored out of the document domain structures.

Context and Structure

Does the need for separate ordered and unordered list objects imply that we will need a separate document domain list structure for every possible way a list could be formatted in the media domain? No. In fact, that would really just be working in the media domain by proxy. When we work in the document domain we are specifically thinking in terms of document structures, not formatting, and so each document domain object we create needs to make sense in document domain terms, not media domain terms.

For example, consider nested lists. Items in a nested list are formatted differently from the list that contains them. At minimum they are indented more and usually they have different number and bullet styles as well. In the media domain, we would need a different style for each level. However, we don't need a separate nested list structure in the document domain. Instead, we express the fact that a list is nested by actually nesting it inside its parent list. For instance, we can nest one ordered list inside another ordered list:

```
<ol>
  <li>
    <p>Dogs</p>
    <ol>
      <li>Spot</li>
      <li>Rover</li>
      <li>Fang</li>
      <li>Fluffy</li>
    </ol>
  </li>
  <li>
    <p>Cats</p>
    <ol>
      <li>Mittens</li>
      <li>Tobermory</li>
    </ol>
  </li>
</ol>
```

In the document domain the inner and outer list are both identical `ol/li` structures. In the media domain, one will be formatted with Arabic numerals and the other with letters.

Figure 5.1. One document domain object, two media domain styles

1. Dogs

- a. Spot
- b. Rover
- c. Fang
- d. Fluffy

2. Cats

- a. Mittens
- b. Tobermory

Both the inner and outer lists are ordered list items in the document domain, but in the media domain they are formatted differently based on context.

In this case, the algorithm that formats the page distinguishes the inner and outer lists by looking at their parentage. For instance in CSS:

```
ol>li>ol>li
{
    list-style-type: lower-alpha;
}
```

This ability to distinguish structures by context is vital to structured writing. It enables us to reduce the number of structures we need to fully describe our content, particularly in the document and subject domains. It also allows us to name structures more logically and intuitively, since we can name them for what they are, not how they are to be formatted or for where they reside in the hierarchy of the document as a whole.

It also points out another important difference between the way media domain and document domain writing is usually implemented. The media domain almost always uses a flat structure with paragraphs, tables, etc. following one after the other in sequence. For instance, a nested list in Word is constructed as a flat sequence of paragraphs with

different styles. Inner and outer lists are expressed purely by the indent applied to the paragraphs. (Word tries to maintain auto-numbering across such listed nest structured, but does not always get it right.)

In the document domain, document structures are usually implemented hierarchically. List items are *inside* lists. Nested lists are *inside* list items. Sections are *inside* chapters. Subsections are *inside* sections. Where the media domain typically only has before and after relationships (except in tables), the document domain typically adds inside/outside relationships to the mix. This use of nested, rather than flat, structures helps to create context, which helps to reduce the number of different structures you need.

Some document domain languages are more hierarchical than others. HTML is relatively flat in structure. For example, it has six different heading styles H1 through H6. Docbook, a widely used document domain structured writing language, is much more hierarchical in structure and has only one element for the same purpose: `title`. But DocBook's `title` element can occur inside 84 different elements, and therefore can potentially be formatted in 84 different ways based on context. In fact, it can potentially be formatted in more ways than that, since some of the elements that contain it can also be nested in other elements, creating even more contexts.

There is a balance to be struck here, however. Nested structures are harder to create and can be harder to understand. Often they require the writer to find just the right place in a hierarchical structure to insert a new piece of content, which is more difficult than simply starting a new paragraph in a word processor.

Constraining document structure

There are other reasons for working in the document domain beside abstracting out formatting rules. One of the main ones is to constrain how documents are structured. For example, let's say that you want to make sure that all graphics inserted into your documents have a figure number, a title, and a caption. This is a document domain constraint rather than a media domain constraint. The requirement for a graphic to have a figure number, title, and caption is one of document structure and organization and does not say anything about how the title or caption should be formatted, (which is a media domain concern).

In the media domain, you can make styles available for figure-numbers, titles, and captions, but you can't enforce a rule that says all graphics must have these elements (which is, by its very nature, a document domain rule). In the document domain, you can express these constraints. You can literally make sure that the only way to include a graphic is to make it a figure and give it a title and a caption by making it illegal to place an `image` element anywhere else in the document structure. A structure that implemented this constraint might look like this:

`figure:`

```
title: Cute kitty
caption: This is a cute kitten.
image: images/cute.jpg
```

If the only way to include an image is to use the image element, and the only place where the image element is allowed is inside the figure element, and if the title and the caption elements are required and must have content, then there is no way for an author to add a graphic without a figure, title, and caption. A document that lacked these elements would be rejected by the conformance algorithm (see Chapter 18, *Conformance*).

This is an example of a constraint being enforced rather than factored out. With the list formatting we could factor out the formatting constraint completely by separating the formatting information from the list structure. Here we enforce the image constraint by providing a specific structure that must be followed to insert images.

(The figure number would be generated automatically, of course, just like the numbers in an ordered list. That constraint has been factored out rather than enforced.)

There are many ways in which you might want to constrain the structure of a document. In a typical media domain application, there is no restrictions on the order in which paragraph styles can be applied. If you want to put a level two heading between two steps in a procedure, nothing other than common sense will stand in the way of your doing so. In a document domain language, however, that kind of thing will usually not be allowed.

Instead, the document domain language will have a set of constraints on how procedures are structured. Procedure structures will have step structures nested within them. Step structures will only be allowed to appear inside procedure structures. Only certain text elements -- such as paragraphs, lists, or code blocks -- will be allowed to occur inside a step. There will be no way to place a second level heading inside a procedure.

Constraints like these are important to document domain languages. If you want to control how procedure are written, or how graphics are labeled, you need to create specific document domain structures for these things, and constrain them to avoid them being misused. Without such constraints, it is easy for a language to slip back into the media domain, something that has happened to HTML.

Backsliding into the media domain

HTML was originally designed for sharing scientific papers. It was not designed to strictly control the organization and presentation of scientific papers -- it was designed more to accommodate requirements than to constrain them -- but it does have features that betray its origins. One example is definition lists. Definition lists are a common feature of scientific papers, which need to precisely define how they will use key terms. But as HTML has come to be adopted as the generic language of web pages, the definition list (`dl`) structure has come to be treated as a generic labeled list structure, used for all kinds of things other than definition lists.

This highlights one of the challenges of structured writing, which is to make sure that structures do not get used for purposes other than what they were intended for. This often happens when people look for an easy way to create an effect in the media domain. If the writer wants a piece of text formatted in a particular way, but the only document domain structure that formats that way is intended for something different, it is easy for them to use the structure incorrectly to get the formatting effect they are after.

But this means that the document domain structure that is being misused no longer expresses the constraints it was designed to express. This means you lose functionality. For instance, you can't find all the definitions in a set of HTML documents by looking for all the `d1` elements. You will get all kinds of things that are not definitions. You might also miss a lot of definitions that were not created using the `d1` structure.

This highlights something that is clearly true about HTML, and that can potentially affect any document domain language. It can slip back into the media domain if people start using it for the media domain effects produced the default transformation algorithms rather than adhering to its document domain structures. Today, structured writing advocates often dismiss HTML as an unstructured language. They will point to other languages, such as DocBook or DITA, as being structured by contrast, despite the fact that all three languages are document domain languages at heart with many similar structures between them.

And while it was not always so, HTML has largely become a set of basic document domain structures that authors can hang styles on (using CSS). In other words, it has come to be used much like traditional media-domain word processing and desktop publishing applications. When people write in HTML, they largely do so in WYSIWYG environments, using style-oriented tools that mimic traditional word processing very closely. The result is often an HTML document that formats more or less correctly, but that is coded very inconsistently from the point of view of the document domain, and which is thus very hard to work with -- either to edit by hand or manipulate with algorithms.

Reversing the backslicing

One way to get back to writing in the document domain has emerged recently in the form of a new syntax called MarkDown. The idea behind MarkDown is to represent the major document domain structures of HTML using the kind of formatting people use in text-only email messages. This approach removes the difficulties associated with typing raw HTML. Here's a sample of Markdown (courtesy of Wikipedia):

Heading

=====

Sub-heading

Another deeper heading

Paragraphs are separated
by a blank line.

Leave 2 spaces at the end of a line to do a
line break

Text attributes **italic**, ****bold****,
``monospace``, ~~~~strikethrough~~~~ .

A [link](<http://example.com>).
[28]

Shopping list:

- * apples
- * oranges
- * pears

Numbered list:

1. apples
2. oranges
3. pears

The rain---not the reign---in
Spain.

It translates into the following HTML (again, courtesy of Wikipedia):

```
<h1>Heading</h1>

<h2>Sub-heading</h2>

<h3>Another deeper heading</h3>

<p>Paragraphs are separated
by a blank line.</p>

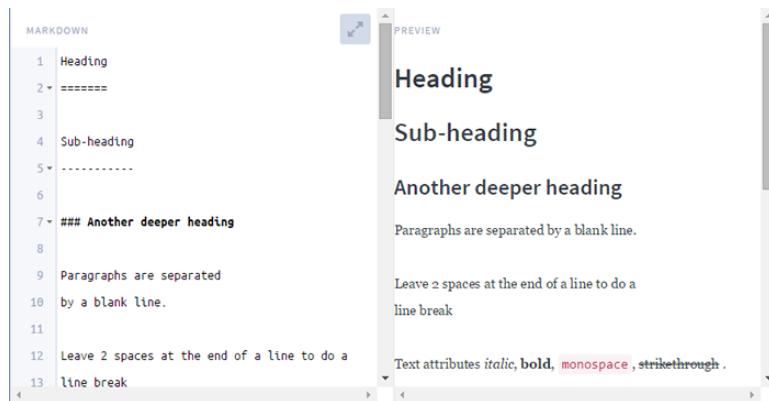
<p>Leave 2 spaces at the end of a line to do a<br />
line break</p>
```

```
<p>Text attributes <em>italic</em>, <strong>bold</strong>,  
<code>monospace</code>, <s>strikethrough</s>. </p>  
  
<p>A <a href="http://example.com">link</a>. </p>  
  
<p>Shopping list:</p>  
  
<ul>  
<li>apples</li>  
<li>oranges</li>  
<li>pears</li>  
</ul>  
  
<p>Numbered list:</p>  
  
<ol>  
<li>apples</li>  
<li>oranges</li>  
<li>pears</li>  
</ol>  
  
<p>The rain&mdash;not the  
reign&mdash;in Spain.</p>
```

Markdown was not designed to be a pure document domain language. It was simply designed to let you write HTML quickly in a text editor. But the effect of using Markdown is that it gets you out of working in a WYSIWYG view and lets you see the structure of the document you are creating.

Many markdown editors use a split screen view that shows the formatted version in one pane as the writer writes markdown syntax in the other. But even here, the writer is still working in the document domain because they still see the structure in the view they are working on. A Markdown editor is never going to produce the kind of messy HTML that a pure WYSIWYG HTML editor can produce.

Figure 5.2. Markdown Editor



A detail of the Dillinger Markdown editor showing the Markdown and browser views side by side.

Something else interesting is at work here. A list in markdown is just a sequence of paragraphs that start with asterisk characters. On the face of it, this is just like a word processor creating lists by styling paragraphs. But if you look at the resulting HTML, you will see that it creates a proper list wrapper element around the list. The markdown interpreter is inferring the hierarchical structure of the document domain from the essentially flat Markdown syntax.

The author is working in something that looks and feels quite like the media domain, thought they have no actual styles and cannot change the formatting at all. But they use abstract formatting notation (underlines for headings, asterisks for unordered lists) to create document domain objects. The beauty of this is that the document domain constraints are preserved, while the author can work in a simple format that is easy to type, and reasonably easy to read. This is an example of functional lucidity, which we will look at in depth in Chapter 9, *Authoring*.

This is an important reminder that XML and languages based on it are not the only route to structured writing. In fact, there are many other ways to create structured texts that obey the appropriate constraints for a particular use case. We will look at some of them in later chapters.

Extending the document domain

Another factor in HTML's slide into the media domain is that it provides only a fairly basic set of document domain structures. As we have seen, enforcing or factoring out media domain constraints requires specific document domain structures. But the possible list of such structures is quite large. There are a few basic features that are common to all documents, such as paragraph, lists, and titles. But these structures alone are not enough to hang meaningful and useful document domain constraints on, which is why, as we noted above, extensibility is an important part of all structured writing domains.

For example, think about a bibliography. A bibliography is a document structure for listing works cited in or recommended by a document. It generally consists of a heading “Bibliography” followed by a set of paragraphs listing the cited works. In the media domain, it is not a particularly complicated structure to create. It is just a sequence of paragraphs with some bold and italic formatting for author names, book titles, etc.

In your media domain stylesheet, you may have some character styles defined that arguably belong to the document or subject domains, such as author-name or book-title. You may even have a specific paragraph style for bibliography entries, but it is unlikely to be more complicated than that.

But these few media domain styles don’t really cover all the rules for creating bibliographies that your institution or publisher is likely to insist on. There are rules for the presentation of a bibliography entry which go into detail about how each work and its authors are listed and how the listings are presented. These are constraints on the writing of the bibliography that the author has to follow, but which are not modeled by the media domain styles they are working with. The authors have to learn and follow these constraints for themselves, and when they have finished writing, these constraints will not be recorded in the content in a machine-readable way. If the constraints are not recorded, you won’t be able to write an algorithm to pull information out of bibliography entries because all the algorithm can see is a bunch of paragraphs with some bold and italic formatting.

If you want to control how bibliographic information is presented, and enable algorithms to find and extract data from bibliographic entries, you need a document domain structure for a bibliography. DocBook has one:

```
<biblioentry id="bib.xsltrec">
  <abbrev id="bib.xsltrec.abbrev">REC-XSLT</abbrev>
  <editor>
    <firstname>James</firstname>
    <surname>Clark</surname>
  </editor>
  <title>
    <ulink url="http://www.w3.org/TR/xslt">XSL Transformations
      (XSLT) Version 1.0</ulink>
  </title>
  <publishername>W3C Recommendation</publishername>
  <pubdate>16 November 1999</pubdate>
</biblioentry>
```

The example is in XML, which can be hard to read, so here is the same structure the SAM markup notation that I have used in earlier examples:

```
biblioentry:(#bib.xslltrec)
    abbrev:(#bib.xslltrec.abbrev) REC-XSLT
    editor:
        firstname: James
        surname: Clark
    title: XSL Transformations (XSLT) Version 1.0
    publishername: W3C Recommendation
    pubdate: 16 November 1999
```

This structure does not just constrain how bibliography entries are presented and formatted, it actually factors out many of those constraints by breaking down the components of a bibliography entry into separate labeled fields. Given a `biblioentry` structure like this, you could create an algorithm to present and format a bibliography entry almost any way you wanted to. Not only that, you could write an algorithm to extract bibliography information from a document by looking for `biblioentry` structures and selecting the desired information from them. For instance, if you want to build a list of authors cited in the document, you could do so by searching the `biblioentry` records and extracting the name of the authors from the structures that record them in the bibliography structure.

This is another important way in which we can cut down the number of document domain structures we need. If we capture the individual pieces of information that make up a bibliography entry, we only need one bibliography entry structure even if we want to present bibliography entries differently in different publications.

Chapter 6. Writing in the Subject Domain

Structured writing in the subject domain means creating structures that are specific to the subject matter being discussed. Sometimes the use of the subject domain is a minor part of a primarily document domain language. For example, DocBook includes subject domain tags like `GUILabel` for references to part of a computer display. Sometimes it takes the form of a document type that is specific to one subject such as a recipe, a used car review, a tide table, or a travel itinerary.

When a document type is specific to one subject, it may be either rhetorically focused or data-focused. If the document type is designed to constrain the way the information on a particular subject is presented, it is rhetorically focused. A rhetorically-focused subject domain document type is still mainly concerned with document-domain structures, but it specifies exactly which document domain structures are supposed to be used in which order in order to present a specific subject.

Alternatively, a subject-domain document type may choose to factor out the presentation details and just capture the pieces of information required to cover a subject. This makes it possible to apply different rhetorical strategies when presenting the document to the reader.

A recipe is a useful example for examining these differences. Notice that the word recipe itself has two meanings. It can mean the formula for a dish. It can also mean the document that records that dish. Recording the formula for a dish is a data oriented approach. Recording the desired presentation of a recipe is the rhetorical approach. We will examine this distinction in this chapter. To begin with, though, we will look at what a recipe looks like in a generic document domain language (one this is not specific to the subject of food preparation).

Here is a recipe written in reStructuredText, a lightweight general purpose markup language:

```
Hard Boiled Eggs
=====
A hard boiled egg is simple and nutritious.
Prep time, 15 minutes. Serves 6.
```

```
Ingredients
-----
===== =====
Item    Quantity
```

```
===== =====
eggs    12
water   2qt
===== =====

Preparation
-----
1. Place eggs in pan and cover with water.
2. Bring water to a boil.
3. Remove from heat and cover for 12 minutes.
4. Place eggs in cold water to stop cooking.
5. Peel and serve.
```

In reStructuredText, a line underlined with equals signs is a major heading and one underlined with dashes is a minor heading. A table is created by using equals signs to mark the beginning and end of the table and the boundary between the table head and the table body. Rows are separated by new lines and columns by spaces in the rows of equals signs. Numbered lists are created by putting numbers in front of lines of text. The equivalent HTML document would look like this:

```
<html>
<h1>Hard Boiled Eggs</h1>

<p>A hard boiled egg is simple and nutritious.  
Prep time, 15 minutes. Serves 6.</p>

<h2>Ingredients</h2>
<table>
  <thead>
    <tr>
      <th>Item</th>
      <th>Quantity</th>
    </tr>
  </thead>
  <tbody>
    <tr>
      <td>eggs</td>
      <td>12</td>
    </tr>
    <tr>
      <td>water</td>
      <td>2qt</td>
    </tr>
  </tbody>
</table>
```

```

</table>
<h2>Preparation</h2>
<ol>
    <li>Place eggs in pan and cover with water.</li>
    <li>Bring water to a boil.</li>
    <li>Remove from heat and cover for 12 minutes.</li>
    <li>Place eggs in cold water to stop cooking.</li>
    <li>Peel and serve.</li>
</ol>
</html>

```

This document follows the normal rhetorical pattern of a recipe. That is, it has the things a recipe normally has, in the order they normally occur in a recipe: introduction, list of ingredients, preparation steps. However, it does not record the fact that it follows this rhetorical pattern. There is nothing in the markup to say that this is not a novel, a car manual, or knitting pattern. Nor would the markup constrain an author to follow the normal rhetorical pattern of a recipe while writing.

Moving this document to the subject domain would allow us to impose these rhetorical constraints, and to record that we have done so. Neither reStructuredText nor HTML give us a way to do that, so we will need a different markup language. Here's what that might look like:

```

recipe: Hard Boiled Egg
introduction:
    A hard boiled egg is simple and nutritious.
    Prep time, 15 minutes. Serves 6.
ingredients:
    * 12 eggs
    * 2qt water
preparation:
    1. Place eggs in pan and cover with water.
    2. Bring water to a boil.
    3. Remove from heat and cover for 12 minutes.
    4. Place eggs in cold water to stop cooking.
    5. Peel and serve.

```

This structure breaks the document up into a collection of named structures. Those structures are “introduction”, “ingredients”, and “preparation” and they are contained in an overall structure called “recipe”. This is the basic rhetorical structure of a recipe. This markup make it explicit that this is a recipe (not a novel, a car manual, or a knitting pattern). The author is explicitly guided to follow this pattern. They are also constrained to present the ingredients as an unordered list and the preparation as a numbered list. (Chapter 18, *Conformance* will look at how such constraints are expressed and enforced.)

One of the common patterns of structured writing is the factoring out of invariants. One of the invariants of the recipe pattern is that it has sections titled “Ingredients” and “Preparation” (or words to that effect). Notice that these titles been factored out here. The titles are part of the rhetorical structure of a recipe, and since the markup now models the rhetorical structures, we can factor out the titles themselves by creating structures called `ingredients` and `preparation`. Since we need those sections to record the rhetorical pattern we have followed, the titles are redundant in our source and can be factored out and added back into the content by an algorithm at publishing time.

By factoring out titles, we factor out the constraint on what those titles must be. An author no longer has to remember the standard titles to use, and a whimsical author can no longer decide to title these sections “Stuff you need” and “Stuff your do”, or any other variant of the standard that the organization has chosen.

If your organization publishes a lot of recipes, you probably have a lot more constraints on the rhetorical structure of your recipes. For instance, you might have a constraint that every recipe must state its preparation time and the number of people it serves. In our subject domain markup, we can enforce and record that constraint by moving the information from the introduction section to separate fields:

```
recipe: Hard Boiled Egg
introduction:
    A hard boiled egg is simple and nutritious.
ingredients:
    * 12 eggs
    * 2qt water
preparation:
    1. Place eggs in pan and cover with water.
    2. Bring water to a boil.
    3. Remove from heat and cover for 12 minutes.
    4. Place eggs in cold water to stop cooking.
    5. Peel and serve.
prep-time: 15 minutes
serves: 1
```

Does this mean that the preparation time will now be displayed as separate fields in the output, rather than in-line? Not necessarily. It might be a good idea to call it out in separate fields so that readers can find the information more easily, but if you really wanted that information at the end of the introduction in every recipe, it would be a simple matter for the presentation algorithm (see Chapter 16, *Publishing*) to construct the sentences “Prep time, 15 minutes. Serves 6.” from the `prep-time` and `serves` field values.

You will always have a rhetorical constraint about how this information is to be presented, but the markup does not enforce the how or the where. It merely enforces that

the information must be captured, leaving it to the presentation algorithm to implement the current governing constraint on how it is to be presented.

So, something interesting has happened here. In order to enforce a rhetorical constraint -- that prep time and number of servings must be specified -- we have actually moved away from markup that specifies presentation to markup that merely records data. In other words, `prep-time` and `serves` are data-oriented subject domain structures that do not specify presentation at all.

This is a recurring pattern in structured writing, and one of the most important things to understand about how structured writing works. It is almost always better to factor out a constraint than to enforce it. This can be a difficult idea to adjust to. If we have a particular form of presentation we want to achieve, our first instinct naturally is to specify it in detail. But this is not always easy to do, especially if you want to specify that a paragraph should always contain certain pieces of information. Nor is it flexible if you want to vary the presentation for any reason (we will look at some reasons why you might in Chapter 12, *Single Sourcing*). It is important to condition yourself so that when you look at these kinds of problems the first question you ask yourself is, is there a way I can factor out this constraint. Only resort to trying to impose a constraint if it is not possible to factor it out.

By the way, using data-oriented subject-domain markup like this also offers some interesting publishing possibilities. For instance, with this markup in place, you could easily query your set of recipes to create a cookbook of recipes you can make in 30 minutes or less because you can query your set of recipes and pull out just those in which the `prep-time` field has a value of 30 minutes or less.

Are there other elements of presentation that we can factor out of the recipe structure? As we noted, the reStructuredText version above specifies a table presentation for the ingredients. Our recipe structure currently specifies a simple list. The block that contains it is labeled “ingredients”, but the contents of the block is just an ordinary unordered list.

There is a relevant constraint here, one which authors need to follow, but one that our markup does not yet impose or record: how to express ingredients. The specification of ingredients in a recipe generally requires three pieces of information, the name of the ingredient, the quantity, and the unit of measure used to express this quantity. These can be presented as a list or a table. To factor out the presentation choice, we can create an ingredient structure that calls out each piece of information separately:

```
ingredients:  
    ingredient:  
        name: eggs  
        quantity: 12  
        unit: each  
    ingredient:
```

```
name: water
quantity: 2
unit: qt
```

There are some shortcuts we can take to make this markup less verbose. (This is a markup syntax named SAM that I will talk about later):

```
ingredients:: ingredient, quantity, unit
eggs, 12, each
water, 2, qt
```

By adding and recording these constraints, we get similar benefits as before. We can better enforce any constraints we have about how ingredient lists are structured and formatted, and we gain access to the specific data involved, meaning, for example, that we could write an algorithm to convert our units from imperial to metric for publication in markets where metric units are preferred.

A key difference between subject and document domains

At this point we can see that with subject domain markup, we have a lot of choices about how documents are constructed. This highlights an important difference between document domain and subject domain markup. A document domain markup language specifies the content and order of a document. We expect that the document domain markup will specify exactly what content is to appear on the page and in what order. This is necessary in the document domain because the document domain does not record any information about the specific subject matter of individual pieces of text. We can't write an algorithm to publish certain pieces of a document domain file, because the markup does not record which are which. (More on this in the next chapter.)

In pure document domain structured writing, the structures created when writing express the author's intent about the rhetorical presentation of the content. But if the author was complying with any specific rhetorical constraints, they did so without any guidance or constraint from the document domain language. The language may have imposed general rhetorical constraints, like how to structure a procedure, but noting specific to the subject matter.

When we move to rhetorically-focused subject domain structure, the a specific rhetorical approach to the subject matter is enforced on the author and expressed by the structures they create. With these subject-specific rhetorical structure in place, you could begin to select different blocks of content to be included or excluded. As we factor out more of the presentation, we can do more and more to manipulate the presentation algorithmically. The recorded content no longer specifies exactly what content is to appear on the eventual

rendered page and in what order. Rather it is a collection of identifiable pieces of content that you can select from or reorder for publication.

Let's suppose we run a publishing company that publishes a number of magazines. We want to create a common store of recipes for use in all the magazines. But different magazines have different requirements. *Wine Weenie* magazine needs to have a wine match with every recipe. *The Teetotaler's Trumpet*, naturally, wants a non-alcoholic suggestion.

Here is how that might be handled in the document domain:

```
<section publication="Wine Weenie">
    <title>Wine match</title>
    <p>Pinot Noir</p>
</section>
<section publication="The Teetotaler's Trumpet">
    <title>Suggested beverage</title>
    <p>Lemonade</p>
</section>
```

This is an example of what we call conditional text. The publication attribute on the section element says, display this text only in this publication. (This makes it management metadata, which we will talk about in Chapter 7, *The Management Domain: an Intrusion*.)

By contrast, this is how this might be handled in the data-oriented subject domain structures:

```
<wine-match>Pinot Noir</wine-match>
<beverage-match>Lemonade</beverage-match>
```

This markup says nothing about which documents should contain either of these pieces of information. Nor does it contain the subheadings what would introduce either of them in the appropriate publication. All these decisions are now left to algorithms. This allows us to do far more things with this content without having to rewrite the source files in any way.

Using subjects to establish context

In Chapter 5, *Writing in the Document Domain*, we noted that we can use context to identify the role that certain structures play in a document, which allows us to get away with fewer structures. For instance, we can use a single title tag for all titles because we can tell what kind of title each one is from the context in which it occurs. The same

is true with subject domain structures. They can provide context that allows us to treat basic text structures differently.

Consider our markup language for recipes:

```
recipe: Hard Boiled Egg
introduction:
    A hard boiled egg is simple and nutritious.
ingredients:: ingredient, quantity, unit
    eggs, 12, each
    water, 2, qt
preparation:
    1. Place eggs in pan and cover with water.
    2. Bring water to a boil.
    3. Remove from heat and cover for 12 minutes.
    4. Place eggs in cold water to stop cooking.
    5. Peel and serve.
prep-time: 15 minutes
serves: 6
```

With the ingredients, we saw that we needed to add additional structure to factor out whether the ingredients would be presented as a list or a table. For the preparation, the steps are currently marked up as a numbered list. Suppose we want to present the steps as steps, rather than just as a generic numbered list (for instance, by labeling them as **Step 1.**, etc, rather than just **1.**). Do we need to create an additional step structure to do this? Not necessarily. In this case we can tell the difference between an ordinary ordered list and a set of preparation steps based on context. We can write a rule in the presentation algorithm that creates special formatting just for ordered lists that are the children of preparation elements that are children of recipe elements (just as we could create rules to format a nested list differently from its parent list based on its context in Chapter 5, *Writing in the Document Domain*).

Chapter 7. The Management Domain: an Intrusion

So far I have talked about three domains that content passes through and in which it can be recorded: the media domain, the document domain, and the subject domain. But there is a forth domain that intrudes into this picture: the management domain.

Why do I call the management domain an intrusion? Because while the subject, document, and media domains are all about recording the content itself, the management domain is not about the content, but about the process of managing it. Here is the example we looked at in Chapter 6, *Writing in the Subject Domain*:

```
<section publication="Wine Weenie">
    <title>Wine match</title>
    <p>Pinot Noir</p>
</section>
<section publication="The Teetotaler's Trumpet">
    <title>Suggested beverage</title>
    <p>Lemonade</p>
</section>
```

Here the publication attribute on the section element is management domain metadata. It does not specify the formatting, organization, or subject matter of the document. It specifies which publication the content should appear in, which is a management decision.

This idea of placing management metadata in the content itself is something that is unique to structured writing. When you manage unstructured content, the content file contains just the content.¹ Any management metadata related to the content is contained in a separate file or some kind of content management system. But this approach limits how fine-grained your management of the content can be. You can only manage content at the file level. Because structured writing makes it possible to address, and therefore manage, individual content structures within a file, you can extend management structures and management metadata into the content itself.

Example: Including boilerplate content

For example, let's say you have a standard warning statement that you are required to include in a document wherever you have a dangerous procedure. Structured writing is

¹As we noted in Chapter 3, *How ideas become content*, all content is structured. What matters is, what purpose is it structured for. So “unstructured” here means unstructured relative to your requirements. If you don’t create or choose a format that supports the management metadata you need at the content level, you can only manage that content at the file level.

all about factoring out invariants, and the invariant here is that this warning statement must appear whenever you describe a dangerous procedure.

Just as we extracted formatting information into a separate file when we moved content from the media domain to the document domain, we now extract the invariant warning from the document and place it in a separate file. Any place we want this warning to occur, we insert an instruction to include the contents of the file at that location.

```
procedure: Blow stuff up
    >>>(files/shared/admonitions/danger)
    step: Plant dynamite.
    step: Insert detonator.
    step: Run away.
    step: Press the big red button.
```

In the SAM markup above the `>>>` is an insert command. It inserts the content of the file located at `files/shared/admonitions/danger`.

Why is this operation part of the management domain, rather than the document domain? Because it deals with a system operation: locating a file in the system and loading its contents. If we were purely in the document domain, the author would be the one performing this operation: finding the file with the warning in it, opening it, and copying the contents into the document. The insertion instruction is just that: an instruction. It is not a declaration about the subject matter or structure of a document, such as we find in subject domain or document domain markup. It is an instruction to a machine to perform an operation. The management domain consists of instructions or the declaration of data required to perform management functions.

Different structured writing systems have different instruction sets for handling the situation described above. In DITA, for instance, this use case is handled using something called a `conref` or a `conkeyref`. In Docbook it can be handled using a generic XML facility called `XInclude`. We will examine this case in greater detail in Chapter 13, *Reuse*.

An alternative approach in the subject domain

There is another way to handle this situation, this time using the subject domain. As we saw in Chapter 6, *Writing in the Subject Domain*, factoring out invariant text is a feature of the subject domain. To understand the subject domain approach to this problem, remember what the invariant rule is here: A dangerous procedure must have a standard warning.

The management domain approach to this is to allow authors to insert the standard warning so that it is only stored once instead of being repeated in every procedure

(something that is often called content reuse) an algorithms we will look at in Chapter 13, *Reuse*. Notice that the management domain markup does not encapsulate our invariant rule that dangerous procedures must have a standard warning. It just provides a generic mechanism for inserting content as a reference to a file. It leaves it entirely up to the author to remember and enforce the rule about dangerous procedures.

The subject domain approach, on the other hand, is all about the invariant rule itself. Specifically, it expresses the aspect of the subject domain that triggers the rule: whether a procedure is dangerous or not:

```
procedure: Blow stuff up
    is-it-dangerous: yes
    step: Plant dynamite.
    step: Insert detonator.
    step: Run away.
    step: Press the big red button.
```

This markup simply records that this procedure is dangerous, a fact about the subject matter. This retains the information on which our invariant rule is based, but factors out the action to be taken. Rather than asking authors to remember to include the file (and how to included it, and how to find it) we delegate that responsibility to the presentation algorithm. It is now the algorithm, not the writer, that needs to remember to include the material in `files/shared/admonitions/danger` whenever the `is-it-dangerous` field of a procedure structure is set to “yes”. This is the sort of task that algorithms are much better at than humans.

Of course, the human writer does still have a job to do here. They have to remember set `is-it-dangerous` to “yes”. But we can make remembering to do this much easier if we make `is-it-dangerous` a required field in the procedure structure. In other words, we set up our structured writing language in such a way that an error will occur if `is-it-dangerous` is not specified. Now the writer is forced to answer the question “yes” or “no” for every procedure they write.

This approach makes the writer’s job much easier because they not only get reminded of the need to address the question of danger with every procedure, they are also asked it in a way that does not require them to know anything about how the content management system works, what warning text is required, or where it is located. They are recording a fact about the subject, not giving an instruction or interpreting a style guide.

One the other hand, this approach only factors out the reuse of one particular piece of content -- the warning for dangerous procedures. If you had multiple such invariant rules about different kinds of subject matter you would need separate subject domain structures for each of them, whereas a single management domain include instructions would let authors handle them all.

On the other other hand, if you have many such invariant rules, and you expect all of them to be enforced by authors from memory, your are going to limit your pool of authors to a few highly trained individuals, and even then they are still likely to miss some instances. The cost of ensuring full compliance with all these rules without subject domain markup to enforce the constraints could be quite high.

Hybrid approaches

It is not always an either/or decision to use pure management domain or pure subject domain approaches. Management domain structures tend to be used in generic document domain languages, since such languages are not designed to be specific to any particular subject matter. Nonetheless, such languages often have roots in particular fields and sometimes include subject-domain structures from those fields. Both DocBook and DITA, for instance, originated in the field of software documentation and both include structures related to the subject of software, such as code blocks and structures for describing user interface elements.

In some cases, such languages can mix subject domain structures into their management structures. One example is the `product` attribute, which is part of DITA's conditional text processing system.

In DITA, you can add the `product` attribute to a wide variety of elements. You can then set a value for products in the build systems and any element with the `product` attribute will only be included in the final output if it matches one of the product values specified in the build.

```
<p>The car seats
<ph product="CX-5">5</ph><ph product="CX-9">7</ph>
</p>
```

DITA can afford to use this bit of subject domain markup for products because product variations are an extremely common reason for using conditional text processing in technical communication, the area for which DITA was created. (Through a process called “specialization”, DITA can add other subject domain attributes for conditional processing in other subject areas.)

The reason I call this a hybrid approach is that the DITA `product` attribute does not exist merely to declare that a piece of text applies to a particular product. It is specifically a conditional processing attribute. That is, it is an instruction, even though it is phrased as a subject domain declaration.

To appreciate the difference, consider that there is another approach to documenting multiple versions of a product. Rather than generating a separate document for each product variant, you could create a single document that covered all product variants

and highlighted the differences between them. A pure subject domain approach would support either approach by simply recording the data for each variant:

```
seats:  
CX-5: 5  
CX-9: 7
```

This information could be presented as data similar to its source format or it could be used to algorithmically construct a sentence like this:

The CX-5 seats 5 and the CX-7 seats 7.

That is not something that the product attribute supports:

```
<p>The car seats  
<ph product="CX-5">5</ph><ph product="CX-9">7</ph>  
</p>
```

This markup is only designed to produce a CX-5 or CX-9 specific document. It is not designed to support the production of a document that covers both cars at once because it does not specify that the values 5 and 7 are numbers of seats. That information is in the text, but not in a form that a publishing algorithm could reliably locate and act on.

Also, creating a single document covering both cars is not the expectation that goes with creating the markup. It is not what the author is told the markup means. The markup is not a simple declaration of facts about each car. It is conditional text markup, and therefore an instruction.

Really, is it a contraction of the more explicitly imperative form (not actually used in DITA):

```
<ph condition="product=CX-5">5</ph>
```

While the introduction of subject domain names into management domain structures is an appropriate bit of semantic sugar for authors, this hybrid approach really remains firmly in the management domain.

Ad hoc management decisions

So far we have contrasted management domain and subject domain approaches to handling invariant rules. Sometimes the management decisions are being made ad hoc by writers as they write, not based on any invariant properties of the subject matter or document structure, and if the decisions affect only part of the content in a file, then the

only way to record those decisions so that the publishing algorithms can act on them is to include management domain metadata in the content.

An example of this is content reuse. The safety warning example was a case of an invariant rule for including a fixed piece of content, which was, therefore, being reused. But there are other situations in which the same text, or substantially the same text, may appear in different places for ad hoc reasons, or for reasons where any rule would apply to so few cases that it would not be worth defining and modeling.

The obvious and traditional way of handling such cases is either to write the text again (if you are not aware that other instances exist, or where they are) or to cut and paste the text. The downside of this is that the text now exists in multiple places, which creates management headaches if you ever need to change it. It also costs more to research and write the same content over and over again.

You can minimize these problems by storing the content once and inserting it wherever it is needed. This approach is not without its own downsides, however, which we will look at in Chapter 13, *Reuse*.

The management domain and content management

So far I have talked about the management domain as an intrusion into structured writing. But it is worth looking at where that intrusion comes from. Management domain structures in structured writing are an intrusion from the content management process.

So far in this book I have talked about structured writing as a way to make content better. In other words, I have talked about it having a fundamentally rhetorical purpose: it is there to improve the text. And since we now live in a world of dynamic content delivery options, it is there to help us build better rhetorical structures that might be difficult to build and manage by hand.

The world of content management, on the other hand, exists mainly to manage the vast collection of content that many organizations own today, and the process by which it is created. Content management is a significant business problem quite apart from the rhetorical properties of the content. Many organizations adopt structured writing not for any rhetorical purpose, but as an enabler of content management processes.

Many content management processes involve finding, identifying, combining and publishing content from diverse sources. These processes often require that content meets certain constraints, which means that structured writing -- the application of constraints to writing -- is a useful content management tool.

Tools and processes that are designed primarily to tackle the content management problem, rather than the rhetorical problem, almost invariably use document domain

The Management Domain: an Intrusion

structures with a significant injection of the management domain, though in these cases it might be more accurate to say that the document domain structures are being used as an extension of the management domain.

The focus of this book is the use of structured writing as a tool for rhetoric -- for making content better. As such, it will talk a lot more about the subject domain than would a text focused on content management. Nonetheless, as shown briefly here, the subject domain can have powerful content management features as well. I will look more at this in later chapters.

Part II. Algorithms

Table of Contents

8.	Quality in Structured Writing	57
	Robots that read	57
	Dumbing it down for the robots	57
	Making humans better writers	58
	We write better for robots than we do for humans	58
	Structure, art, and science	60
	Contra-structured content	61
	Until the robots take over	61
9.	Authoring	63
	Functional lucidity	65
	Granularity	66
	Simplicity and Clarity	66
10.	Separating Content from Formatting	69
	Separate out style instructions	69
	Separate out formatting characters	70
	Name your abstractions correctly	71
	Make sure you have the right set of abstractions	71
	Create containers to provide context	73
	Move the abstractions to the containers	74
	Separate out abstract formatting	75
11.	Processing Structured Text	78
	Two into one: reversing the factoring out of invariants	78
	Adding back style information	78
	Rules based on structures	80
	The order of the rules does not matter	80
	Applying rules in the document domain	81
	Processing based on context	82
	Processing container structures	84
	Restoring factored-out text	84
	Processing in multiple steps	86
	Query-based processing	88
12.	Single Sourcing	89
	Basic single sourcing	89
	Differential single sourcing	92
	Differential organization and presentation	95
	Conditional differential design	97
	Primary and secondary media	98
	Responsive Design	99
13.	Reuse	100
	Fitting pieces of content together	100
	Common into variable	101
	Variable into common	102

Variable into variable	103
Using IDs	104
Using keys	105
Common with conditions	108
Factor out the common	110
Factor out the variable	111
Assemble from pieces	112
Combining multiple techniques	116
Content reuse is not a panacea	116
Quality traps	116
Cost traps	117
14. Composition	119
Fundamental composability	119
Structural composability	119
Stylistic composability	120
Narrative composability	121
15. Single source of truth	123
Single vs Findable	125
16. Publishing	128
The Rendering Algorithm	128
The Encoding Algorithm	129
The Formatting Algorithm	131
The Presentation Algorithm	132
Differential presentation algorithms	133
Presentation sub-algorithms	134
The linking algorithm	134
The navigation algorithm	134
The public metadata algorithm	135
The document structure normalization algorithm	135
The Synthesis Algorithm	136
Differential synthesis	136
Synthesis sub-algorithms	136
The inclusion algorithm	137
The filtering algorithm	137
Coordinating inclusion and filtering	137
The extraction algorithm	138
The merge algorithm	138
The resolve algorithm	138
Deferred synthesis	139
Combining algorithms	139
The consistency challenge	140
Functional lucidity and a layered architecture	141
17. Linking	143
Deflection in the media domain	144
Deflection in the document domain	145

Deflection in the management domain	146
The problem with IDs and Keys	147
Relationship tables	147
The problem with link tables	148
Conditional linking	149
Deflection in the subject domain	150
Finding resources to link to	153
Deferred Deflection	155
Different domain, different algorithm	157
18. Conformance	158
Completeness	160
Consistency	160
Accuracy	161
Semantic constraints	161
Entry validation constraints	164
Referential integrity constraints	165
Conformance to external sources	165
Conformance and change	167
Design for conformance	167
19. Auditing	170
Correctness of the definition of the content set	171
Ensuring the content set remains uncontaminated	172
Ensuring that the content set is well integrated	173
Making content auditable	174
Performance auditing	175
20. Relevance	176
Being relevant	176
Showing relevance	177
Showing relevance to algorithms	178
21. Active content	180
Media Domain	182
Document Domain	182
Subject Domain	183
22. Extract	185
Tapping external sources of content	185
Information created for other purposes	186
The Extraction Algorithm	188
The diversity of sources	189
23. Merge	190
24. Information Architecture	192
Top-down vs. bottom-up information architecture	193
Categorization	193
Linking	194
Tables of Contents	195
Lists	196

Dynamic content	198
25. Change management	199
26. Content Management	204
Metadata is the foundation of management	204
The location of metadata	205
Managing the process	207
Conflicting constraints	209
Content management and reuse	211
Creating manageable content	212
27. Collaboration	213
Bridging silos	215
Expanding your pool of authors	216
28. Timeliness	217
29. Repeatability	223
30. Conversion	225
31. Exchange	228
32. Translation	231
XLIFF	231

Chapter 8. Quality in Structured Writing

When I talk to programmers about what I do, they often ask my why structured content is important any more. Machines are getting so good at reading human language, they argue, that semantic markup to assist the machine is increasingly becoming pointless.

Robots that read

Indeed, machines are getting better and better at understanding human language. An approach called Deep Learning is increasingly becoming a key technology for companies like Facebook, Google, and Baidu for both language comprehension¹ and speech recognition²).

The semantic web initiative has long sought to create a Web that is not just people talking to people but also machines talking to machines. This has traditionally involved an essentially separate communication channel -- semantic markup embedded in texts but not presented to the human reader. It has also involved the creation of specialized semantic data stores with query language to match, to teach computers to understand relationships that humans would express in ordinary language. Content management systems have implemented metadata schemes, often involving elaborate taxonomies, in an attempt to make content more findable when regular text-based searches don't work as well as they would like.

But this two-channel approach -- one text for the human, another for the machine -- only makes sense if we assume that the machine cannot read human language. If machine and human can both read the same text and understand it with the same level of sophistication (or if the machine's understanding is actually more sophisticated than the human's) then we shouldn't need two channels. The human Web becomes the semantic Web.

After all, the human text always was semantic. Semantics is simply the study of meaning. All meaningful texts have semantics. It is just that it has been difficult to build algorithms that could read and understand like humans do. Semantic technologies are about dumbing it down for the machine because the machine is not bright enough to read the regular semantics.

Dumbing it down for the robots

This dumbing down necessarily involves omitting a great deal of the semantics of the text. Fully expressing all the meaning and implications of even the simplest text in RDF triples would be daunting, for instance. This has always created a problem for semantic technologies: which semantics do you select to dumb down to the machine's level, and for

¹<http://www.technologyreview.com/featuredstory/540001/teaching-machines-to-understand-us/>

²<http://www.technologyreview.com/news/544651/baidus-deep-learning-system-rivals-people-at-speech-recognition/>

what purpose? This is why there is no universal approach to structured writing that works for all purposes and all subject matter. You can only represent a fraction of the human semantics to the machine, and which you choose depends on what specific functions you want the machine to perform.

But if the machine could read the text as well as you can, then these limitation vanish. Deep learning is moving us in that direction.

Why then should we bother with structured writing? Quite simply because while machines are rapidly learning to read human text , that text is still written by humans, and most humans are not good writers.

Making humans better writers

By that I don't just mean that they use poor grammar or spelling or that they create run-on sentences or use the passive voice too much, though all those things may be true, and annoying. I mean something more fundamental than that: they don't say the right things in the right way for the right audience. They leave out stuff that needs to be said, or they say it in a way that is hard to understand.

We all suffer from a malady called The Curse of Knowledge which makes it difficult for us to understand what it is like not to understand something we know. We take shortcuts, we make assumptions, we say things in obscure ways, and we just plain leave stuff out.

This is not a result of mere carelessness. The efficiency of human communication rests on our ability to assume that the person we are communicating with shares a huge collection of experiences, ideas, and vocabulary in common with us.³ Laboriously stating the obvious is as much a writing fault as omitting the necessary. Yet what is obvious to one reader may be obscure to another. The curse of knowledge is that as soon as something becomes obvious to us, we can no longer imagine it being obscure to someone else.

Thus much of human to human communication fails. The recipient of the communication simply does not understand it or does not receive the information they need because the writer left it out. Machines may learn to be better readers than we are, but even machines are not going to learn to read information that just isn't there.

We write better for robots than we do for humans

Actually, one of the advantages of the relative stupidity of robots is that it forces us to be very careful in how we create and structure data for machines to act on. The computer

³<http://everypageispageone.com/2015/08/04/the-economy-of-language-or-why-we-argue-about-words/>

science community quickly hit on the phrase “garbage in, garbage out”, because the machines were, and to a large extent still are, too stupid to know when the information they were taking in was garbage, and did not have the capacity, like human beings, to seek clarification or consult other sources. They just spit out garbage.

This meant that we had to put a huge emphasis on improving the quality and precision of the data going in. We diligently worked out its structures and put elaborate audit mechanisms in place to make sure that it was complete and correct before we fed it to the machine.

We have never been as diligent at improving the quality of the content that we have fed to human beings. Faced with poor content, human beings do not halt and catch fire; they either lose interest or do more research. Given our adaptability as researchers and our tenacity in pursuing things that really matter to us, we often manage to muddle through bad content, though at considerable economic cost. And the distance that often separates writers from readers means that the writers often have no notion of what the poor reader is going through. If readers did halt and catch fire, we might put more effort and attention into content quality.

Even today, when a huge emphasis is being placed on enterprise content management and the ability to make the store of corporate knowledge available to all employees, most of the emphasis is on making content easier to find, not on making it more worth finding. (This despite the fact that the best thing you can do to make content easy to find is to make it more worth finding.) People trying to build the semantic web spend a lot of time trying to make the data they prepare for machines correct, precise, and complete. We don't do nearly as much for humans. Until we do, deep learning alone may not be enough to make the human web the semantic web.

Part of the problem has always been that improving content quality runs up against the curse of knowledge. Both the authors who create the content and most of the subject matter experts who review it suffer from the curse, meaning that there are few effective ways to audit written content. Style guides and templates can help remind authors of what is needed, but they are difficult to remember and to audit, meaning there is little feedback for an author who strays.

Structured writing provides a way to guide and audit content for quality. My reply to the people who ask me whether structured writing is relevant, therefore, is “garbage in, garbage out”. Structured writing enables us to write in the subject domain (wholly or partially) and this allows us to guide and audit in ways not previously possible. It also allows us to factor out many constraints, simplifying the author's task, and therefore allowing them to give more of their attention to the writing task. These things make content better, whether that content is going to be read by people or robots.

That the structure we add to writing also makes it available to other algorithms for downstream use is a substantial and welcome bonus which, among other things, help pay for the tools and the structures that we use to make content better.

Structure, art, and science

To many writers, the idea that imposing constraints can improve quality is controversial. Many see quality writing as a uniquely human and individual act, an art, not a science, something immune to the encroachment of algorithms and robots. But I would suggest that the use of structures and algorithms as tools does not diminish the human and artistic aspects of writing. Rather, it supplements and enhances them.

And I would suggest that this is a pattern we see in all the arts. Music has always depended on the making and the perfecting of instruments as tools of the musician. Similarly the mathematics of musical theory gave us well tempered tuning, on which modern Western music is based.

Computer programming is widely regarded as an art⁴ among its practitioners, but the use of sound structures is recognized as an inseparable part of that art. Art lies not in the rejection of structure but in its mature and creative use. As noted computer scientist Donald Knuth observes in his essay, *Computer Programming as an Art*, most fields are not either an art or a science, but a mixture of both.

Apparently most authors who examine such a question come to this same conclusion, that their subject is both a science and an art, whatever their subject is. I found a book about elementary photography, written in 1893, which stated that “the development of the photographic image is both an art and a science”. In fact, when I first picked up a dictionary in order to study the words “art” and “science,” I happened to glance at the editor’s preface, which began by saying, “The making of a dictionary is both a science and an art.”

—<http://dl.acm.org/citation.cfm?id=361612>

As writers we can use structures, patterns, and algorithms as aids to art, just like every other profession.

Of course, few writers would claim that there is no structure involved in writing. We have long recognized the importance of grammatical structure and literary structure in enhancing communication. The question is, can the type of structures the structured writing proposes improve our writing, and if so, in what areas? Traditional poetry is highly structured, but it is doubtful that using an XML schema would help you write a better sonnet. On the other hand, it is clear that following the accepted pattern of a recipe would help you write a better cookbook, and using structured writing to create your recipes can help you both improve the consistency of your recipes and to produce them more efficiently.

The question then becomes, how much of our work is like recipes and would benefit from structured writing, and how much is like sonnets and would not. The answer, I believe,

⁴<http://ruben.verborgh.org/blog/2013/02/21/programming-is-an-art/>

is that a great deal of business and technical communication, at least, can benefit greatly. If you look at much of that communication and see no obvious structure, I would suggest that this is not evidence that structure is inappropriate, but that appropriate structure has not been developed and applied to the content.

Contra-structured content

We must also acknowledge that many writers have had a bad experience with structured writing. In many of these cases, the structured writing system was not chosen or designed by the writers to enhance their art; it was imposed externally for some other purpose, such as to facilitate the operation of a content management system or make it easier to reuse content. In some cases, these systems actively interfere with the author's art and directly hinder the production of quality content.

Writers who have had bad experiences with structured content have usually been faced with structures that were not designed for the writer's purpose. But such content is not merely unstructured for these author's purposes, it is actually contra-structured. It has an enforced structure that actively gets in the way of the author doing their best work.

I talk to authors all the time who show me page designs and layouts that make no sense, lamenting that the system does not give them any other choices. Content structure is not generic, and you cannot expect to simply install the flavor of the month CMS or structured writing system and get a good outcome.

Properly applied, however, as a means to guide and enhance the work of authors, structured writing can substantially improve content quality. In upcoming chapters, we will look at the algorithms of structured content, many of which relate directly to the enhancement of content quality.

Until the robots take over

Of course, this all supposes that the machines are not becoming better writers than us as well. Companies like Narrative Science are working on that, but I don't think they are nearly as far along that path as the deep learning folks are in teaching computers to read.

Do robots suffer from the curse of knowledge? Maybe not. But current writing robots certainly work with highly structured data, so structured writing is still key to quality content even when the robots do come for our keyboards.

Actually, according to James Bessen's recent article in *The Atlantic*, The Automation Paradox⁵, automation does not decimate white collar jobs the way we have been told to fear. By reducing costs, it increases demand, resulting in net growth of jobs, at least for people who learn to use the new technology effectively.

⁵<http://www.theatlantic.com/business/archive/2016/01/automation-paradox/424437/>

That said, all the semantic technology and content management in the world is not going to make the difference it should until we improve the quality of content on a consistent basis. Structured writing, particularly structured writing in the subject domain, is one of our best tools for doing that.

Chapter 9. Authoring

It may not be immediately obvious that authoring is an algorithm. It is not something done to content, rather it is the process that creates content. The creation of content is obviously a process and activity, but it is (for the most part) performed by people rather than machines, and it is not obvious that it is done using a set of fixed processing rules, which would define it as an algorithm.

Creating structured content is an activity over and above the pure act of writing. The only way we get content in the right structures is if authors create those structures as they write. Our structured writing system can only be as good as the structured content our writers create. Getting the best possible structure from our writers, therefore is key to all of the algorithms and all of the benefits of structured writing. This is not something we can afford to leave to chance. We need to be systematic about it. We need an authoring algorithm.

All forms of structured writing, even in the media domain require authors to do something other than simply write. Since writing is an intellectually challenging activity that requires full attention, adding structured writing requirements into the mix necessarily takes away some of the attention that should otherwise be given to content, which obviously has the potential to reduce the quality of the content. Clearly, therefore, we need to make sure that the intrusion of structured writing requirements onto the writing process is as minimal as possible.

But the equation is not quite so simple as this. Structured writing may be an additional requirement, but quality writing does not result from an author simply spilling words onto a page or screen in stream of consciousness fashion. Writing is a design activity. It creates a structure of words that conveys complex ideas and information about the real world. It very much matters that the author says the right things using the right words in the right order. If structured writing techniques can help with this literary design work, they can lessen the intellectual burden on the author, and thus potentially improve the quality of the content.

Of course, structured writing can improve content quality in other ways. Most of the algorithms we will look at pertain to content quality in one way or another, such as improved linking or better management of terminology. Still, these algorithms work on the structures that the authors create. If those structures are weak, there is a limit to what downstream algorithms can do to improve quality. It really all begins with getting authors to reliably create the required structures as they write.

In media-domain systems like word processor and desktop publishing systems, the writer is asked to think about formatting structures while writing. One of the traditional arguments for structured writing is to relieve the writer of the burden of thinking about (and manipulating) formatting so they can focus on writing.

This means moving to the document domain. But in the document domain, the writer has a new set of structures to think about: document domain structures. Is it easier on the writer to think about and manipulate document domain structures rather than media domain structures? In some cases, yes. For instance, writing a blog post or a web page in MarkDown may be less cumbersome for some writers than using a WYSIWYG HTML editor.

However, Markdown does not contain enough structure, or enough constraints on its structure, to enable many of the structured writing algorithms we will be looking at. If we want to support these algorithms, we will need something more structured, and this can easily mean something that requires more of the author's attention. If we are proposing to implement management-intensive algorithms, such as reuse, it can mean that authors need to learn and manipulate an entire management system and the management policies that the organization puts in place around it. Depending on how complex these policies are and how foreign they are to the author's experience, this can create a burden far greater than that of creating and manipulating formatting according to a style guide.

We could look at this and say, okay, yes, authoring is now more difficult and more complex than it was before because of all this additional structure and all that authors need to learn to apply that structure, but we are getting additional advantages as well, so it is worth it overall. The problem is, attention is a finite resource. Writing is a task that takes full attention, so any requirement added to the task takes attention away from writing. As writings get more difficult, writers do all of the component tasks less well. The more of the writer's attention is required on structure, the less is available for writing, and the quality of the writing suffers. And as the quality of the writing suffers, the writer becomes frustrated with the system, and becomes more interested in getting their ideas down than in obeying the onerous structural rules that are getting in their way. When that happens, the quality of the structure suffers as well. And if both the quality of the writing and the quality of the structures decline, all of your algorithms become less reliable, compromising all of benefits you hoped to obtain.

To look at this another way, the more complex your system becomes, and the more algorithms you are attempting to support, the more important conformance to its constraints becomes. But conformance is fundamentally a human activity. It is that author who must conform. Good conformance results from creating structures that are easy to conform to. It all begins with authoring. Unfortunately, authoring is often the last thing people think about in designing content management and structured writing systems

One of the most familiar tropes of the content management industry is that problems with content management systems are not technology problems, they are human problems. The solution, this trope suggests, lies in better change management and more training. The presumption here is that the tools work fine if only you give them correct input. If the input is not correct, that is the fault of the humans who created the input. But this is an argument we would not accept for any other kind of system. For any other kind of system we would say, "this system is too hard to use", not "the problem is everybody needs to be better trained and more accepting of change". But the real fault here is poor

system design. If humans cannot conform to the structures that the system requires, the fault is in the system design. The structures should be redesigned to be easier to conform to. (We will look at this in more detail in Chapter 18, *Conformance*.

How is designing structures for ease of authoring consistent with designing them to match the specific constraints that we want to impose for the sake of quality and efficient processing? As we have seen, moving from the media domain to the document domain allows us to factor out or impose certain structural constraints, but often require the introduction of the management domain to impose content management constraints. These complexities detract from ease of authoring. But we have also seen that by moving to the subject domain we can factor out many of the document domain and management domain structures that drive the publishing process and reintroduce them with algorithms when it is time to publish. Designing structures for authoring, therefore, often consists of factoring out complex publishing and content management structures using subject-domain structures that are lucid for authors.

One of the most important consequences of this, both for ease of authoring and reliability of data, is that in the subject domain, you are not asking the author to think, and to structure content, in terms of algorithms. In this sense, the move to the subject domain not only factors out specific constraints from the author, it factors out the need to think in algorithms at all, leaving the author free to think in terms of subject matter. This freedom to focus on content is a property I call functional lucidity.

Functional lucidity

Functional lucidity means the way that you actually use language when you are writing, which is to say the way that you use language when you are in the throes of figuring out what you want to say and how you want to say it. If you are asked to add markup to your content as you write, if you are asked to shape your content according to the constraints that a structured writing language dictates, then the lucidity of that markup and the structures it defines are vital to your success. The names of the structures, the order in which they occur should spring into your mind as readily (if not more so) than the words and phrases and ideas you are trying to record.

As anyone who has struggled to write even a paragraph in a language they are only beginning to learn can attest, writing in a language in which you are not fluent is painful. The effort of finding words and correct grammatical structures takes all of the attention that should be reserved for what you are trying to say. Writing in a tagging language where the tags don't make intuitive sense, where the structures don't seem to fit the thoughts you are trying to express, is very much like this. Lucidity is essential to avoid having the markup absorb all of the attention that should be focused on the content.

Functional lucidity is not an absolute property, of course. What is lucid for one writer may be opaque to another. In particular, professional technical writers who have been used to

writing in structured document domain templates in applications like FrameMaker may find a markup language like DocBook functionally lucid, whereas someone not used to thinking in these terms would find it difficult and distracting. On the other hand, those writers used to FrameMaker often find DITA's structure difficult to get used to because they do not find its approach to topics lucid. To still others it seems very natural.

But while different writers may have different degrees of experience and familiarity with abstract document structures, all writers should have familiarity with the subject matter they are writing about. Thus a well-designed subject domain language tends to naturally have functional lucidity for everyone who is likely to use it. (Though writers can sometimes disagree about what needs to be said on a subject and how it should be said.)

Granularity

There can be a conflict between ease of authoring and ease of content management. Content management may want to manage content down to a fine level of granularity, especially for purposes of content reuse. This content management algorithm may be best served by managing fairly small chunks of content -- semantic units rather than narrative units. But for the writer, something less than a narrative unit can be difficult. It can be difficult for the author to get a sense of how the semantic block they are writing will meet the reader's needs when they don't see the narrative block it will fit into. It is hard to create parts rather than wholes unless the parts are really well defined. A writer might carry the whole of an essay in their head, for instance, and be able to structure it well on that basis. But if they are making only parts and cannot see the wholes that will be created, it is hard to correctly structure a part without very clear and explicit guidance.

Simplicity and Clarity

One of the biggest benefits of subject domain markup for authors is a much higher degree of functional lucidity compared with a typical document domain language.

While a general document domain language like DocBook needs to have structures for a wide range of document structures, a recipe markup language such as we have developed in this chapter, has only a few simple elements. Better still, there are very few permutations of those elements. Because subject domain languages do not specify document order, we don't need to allow for many possible document orderings in the language, thus reducing the permutations we have to allow for and deal with. The synthesis algorithm can take the named structures of the subject domain markup and order them in any way you like.

Because subject domain structure describe the subject matter they contain, they are also much clearer to authors, who may not understand complex document structures (or, more often, the subtle distinctions between several similar document structures), but who do (we hope) understand their subject matter.

The combination of simplicity and clarity mean that in many cases you can get authors to create subject-domain structured content with little or no training. For instance, even if we add some additional fields to our recipe markup, you could still hand a sample like the one below to an author and ask them to follow it as a template, without giving them any training or any special tools.

```
recipe: Hard Boiled Egg
introduction:
    A hard boiled egg is simple and nutritious.
ingredients:: ingredient, quantity
    eggs, 12
    water, 2qt
preparation:
    1. Place eggs in pan and cover with water.
    2. Bring water to a boil.
    3. Remove from heat and cover for 12 minutes.
    4. Place eggs in cold water to stop cooking.
    5. Peel and serve.
prep-time: 15 minutes
serves: 6
wine-match: champagne and orange juice
beverage-match: orange juice
nutrition:
    serving: 1 large (50 g)
    calories: 78
    total-fat: 5 g
    saturated-fat: 0.7 g
    polyunsaturated-fat: 0.7 g
    monounsaturated-fat: 2 g
    cholesterol: 186.5 mg
    sodium: 62 mg
    potassium: 63 mg
    total-carbohydrate: 0.6 g
    dietary-fiber: 0 g
    sugar: 0.6 g
    protein: 6 g
```

Of course, the downside is that recipe markup is only good for one thing: recipes. A general document domain language can be used to write all kinds of documents. It will not enforce or record nearly as many constraints, or enable nearly as many options for validation or publishing, and it won't be nearly as clear and simple for authors to use. But neither will it require you to create subject domain languages for each of the subjects you write about. At first glance, that may seem like a slam dunk case for sticking with the document domain, as the idea of inventing subject domain languages and the synthesis

and presentation algorithms to go with them may seem daunting. But as we will see, the decision is not so clear cut, as sticking with the document domain comes with a lot of complexity, and sometimes custom development, that may not be apparent at first. We will look at this trade off in more depth as we look at the other algorithms.

Chapter 10. Separating Content from Formatting

If there is one phrase that most people associate with structured writing, it is “separating content from formatting”. We structure content so that we can process it with algorithms. An algorithm is a repeatable method for doing something. If we want our processing algorithms to be repeatable, we need the structure of the content to be consistent. That means that the separation of content from formatting has to be done consistently and repeatably. That makes it the first algorithm of structured writing.

Separate out style instructions

Let’s start with a simple bit of content that includes a description of its format. I’m going to use CSS syntax to describe the format, but that is not significant. This is just to show what we are doing when we do the separation. I’m also going to represent certain characters by their names in square brackets, just so we can see exactly where everything is going:

```
{font: 10pt "Open Sans"}The box contains:  
{font: 10pt "Open Sans"}[bullet][tab]Sand  
{font: 10pt "Open Sans"}[bullet][tab]Eggs  
{font: 10pt "Open Sans"}[bullet][tab]Gold
```

This file contains content and formatting, so let’s separate the two. Of course, when we remove the formatting from the content we are going to have to add something in its place so we can add the formatting back later. The algorithm is not really “separate content from formatting” but “separate formatting from content and replace it with something else”.

The simplest thing to replace it with is a style:

```
{style: paragraph}The box contains:  
{style: paragraph}[bullet][tab]Sand  
{style: paragraph}[bullet][tab]Eggs  
{style: paragraph}[bullet][tab]Gold
```

Then, of course, we need to record the style (we are separating it, not eliminating it altogether):

```
paragraph = {font: 10pt "Open Sans"}
```

Now that they are separated, we have the choice of substituting different formatting by changing the definition of the style, while leaving the content alone:

```
paragraph = {font: 12pt "Century Schoolbook"}
```

Separate out formatting characters

Cool, but suppose we would like to change the style of the bullet we use for lists. The style of bullet used is certainly part of what we would consider “formatting”, but bullets are text characters. To change them you don’t just have to change the font applied to the characters, you have to change the characters themselves.

Sometimes the typed characters in your text are part of the content, and sometimes they are part of the formatting. So now we need to extend our idea of a style to include content.

```
paragraph = {font: 12pt "Century Schoolbook"}  
bullet-paragraph = {font: 12pt "Century Schoolbook"}[bullet]
```

Now our content looks like this:

```
{style: paragraph}The box contains:  
{style: bullet-paragraph}[tab]Sand  
{style: bullet-paragraph}[tab]Eggs  
{style: bullet-paragraph}[tab]Gold
```

Except that now the writer will be starting the bulleted lines with a tab, which is awkward and probably error prone, so we move that character to the style as well.

```
paragraph = {font: 12pt "Century Schoolbook"}  
bullet-paragraph = {font: 12pt "Century Schoolbook"}[bullet][tab]
```

Now our content looks like this:

```
{style: paragraph}The box contains:  
{style: bullet-paragraph}Sand  
{style: bullet-paragraph}Eggs  
{style: bullet-paragraph}Gold
```

And now you can change the bullet style:

```
bullet-paragraph = {font: 12pt "Century Schoolbook"}[em dash][tab]
```

And then we maybe realize that “bullet-paragraph” is not the best name any more, because the style is now a dash, not a bullet. In other words we discover that we had not done as good a job as we thought of separating content from formatting, because the content actually still contains formatting information in the form of a style that is named for a particular piece of formatting.

Name your abstractions correctly

When we separate formatting from content, we have to insert something in its place, and it matters what that something is and what it is called. If we call it the wrong thing we set up a false expectation, and that will lead to authors using it incorrectly, which will mean we can't format it reliably.

So the first lesson about the algorithm of separating content from formatting is that it matters what you call things. When you do this, you are creating an abstraction, and you need to figure out what that abstraction is and name it appropriately.

So what is the abstraction here? It is a list, of course. So maybe we do this:

```
{style: paragraph}The box contains:  
{style: list-item}Sand  
{style: list-item}Eggs  
{style: list-item}Gold
```

and

```
list-item = {font: 12pt "Century Schoolbook"} [em dash][tab]
```

Make sure you have the right set of abstractions

But then, of course, we run into this problem:

```
{style: paragraph}To wash hair:  
{style: list-item}Lather  
{style: list-item}Rinse  
{style: list-item}Repeat
```

Here our list should have numbers, not dashes or bullets. So we realize that the abstraction we are after is not so broad as all list items. We look at the differences between the

Separating Content from Formatting

different kinds of list items we use and try to group them into abstract types and come up with names for those types. Maybe we come up with “ordered-list-item” and “unordered-list-item”. Then we have:

```
{style: paragraph}The box contains:  
{style: unordered-list-item}Sand  
{style: unordered-list-item}Eggs  
{style: unordered-list-item}Gold
```

and

```
{style: paragraph}To wash hair:  
{style: ordered-list-item}Lather  
{style: ordered-list-item}Rinse  
{style: ordered-list-item}Repeat
```

And the style for ordered-list-items now looks something like this:

```
ordered-list-item = {font: 12pt "Century Schoolbook" }<count>.[tab]
```

And then we realize that we need a way to increment the count and to reset it to 1 for a new list. So we have:

```
{style: paragraph}To wash hair:  
{style: first-ordered-list-item}Lather  
{style: ordered-list-item}Rinse  
{style: ordered-list-item}Repeat
```

and

```
first-ordered-list-item = {font: 12pt "Century Schoolbook" }<count=1>  
ordered-list-item = {font: 12pt "Century Schoolbook" }<++count>.[tab]
```

(`++count` here means add one to `count` and then display it.)

And this is pretty much how you do lists in FrameMaker or Microsoft Word today, as well as a number of other tools. But the reason for going through it in such detail is to point out what is involved in even this simple bit of separation. We began by simply removing formatting commands, but then started to remove characters as well, which forced us to include characters in our style definitions, and then to be able to actually calculate characters in our style definitions. And we saw that in performing these separations, we

were creating abstractions, and that it was important to consider all the cases we might run into and create the appropriate abstractions to handle them.

Create containers to provide context

As we noted in Chapter 5, *Writing in the Document Domain*, one problem with this approach is that the writer has to remember to apply a different style to the first item of a list. It would be better if they could use the same style for each list item and have the numbering just work. But this is hard to do because there is nothing in the content to say where one numbered list ends and the next begins. For this we need a new abstraction. So far we have abstractions for two kinds of list items: ordered and unordered list items. But we don't have an abstraction for lists themselves.

So far, we have been separating content from formatting purely in the media domain. We have replaced direct formatting definitions with indirect definitions through styles. The only thing that abstracts any of this beyond the media domain is the names that we have given to the styles that we have created. But now we start to venture into the document domain, creating the abstract idea of a list and inserting that abstract idea into our content.

```
paragraph: To wash hair:  
list:  
    ordered-list-item:Lather  
    ordered-list-item:Rinse  
    ordered-list-item:Repeat
```

There are a number of significant changes here. First, our structure is no longer flat. We have introduced the idea of a container. A list is a container for list items. In creating this container we have added something to the content that was not there before. Previously it was a series of paragraphs with different styles attached. Now we have a container, which, as far as the formatting is concerned, simply never existed in the original. The writer and reader knew that the sequence of bulleted paragraphs formed a list, but that was an interpretation of the formatting. Now we have taken that interpretation and instantiated it in the content itself. Those lines starting with bullets constitute a list, and now we have made it explicit.

By creating the idea of a list, we are able to further separate list formatting from the content of the list -- because now an algorithm, one I will call the formatting algorithm, can recognize it as a list and can make formatting decisions based on that knowledge.

The second important thing that has happened here is that the content no longer contains references to style names. Instead we have structures. `list` is a structure and so are `paragraph` and `numbered-list-items`.

We have replaced styles with structures because the same structure may get a different style depending on where it is in the document. The formatting algorithm is responsible

for determining if an ordered-list-item is the first one inside a list and formatting it accordingly. (Which is just how list formatting works in CSS¹.)

Now authors no longer apply styles to content, even ones with abstract names. Rather they place content in structures and allow the formatting algorithm to apply styles appropriately. The content is separated even more from the formatting.

Move the abstractions to the containers

But there is an obvious problem here. What if an author inadvertently does this:

```
paragraph: To wash hair:  
list:  
    ordered-list-item:Lather  
    unordered-list-item:Rinse  
    ordered-list-item:Repeat
```

To avoid this, we move the abstraction outward. Instead of ordered and unordered list items, we create ordered and unordered lists:

```
paragraph: To wash hair:  
ordered-list:  
    list-item:Lather  
    list-item:Rinse  
    list-item:Repeat
```

and

```
paragraph: The box contains:  
unordered-list:  
    list-item:Sand  
    list-item:Eggs  
    list-item:Gold
```

And, of course, the `list-item` structure can be used in either an unordered list or an ordered list, because it is a list item in either case, and the formatting algorithm can tell the difference based on which type of list it belongs to. The name `list-item` describes its role in the document in a way that is entirely separated from how it will be styled.

¹<https://css-tricks.com/numbering-in-style/>

Moving the abstraction out to the container is an important part of the algorithm of separating content from formatting. It helps keep things consistent and reduces the number of things authors have to remember.

Creating containers and abstracting out the differences between their contents is an important piece of separating content for formatting. For example, HTML and Markdown both provide six different levels of headings. But content under an H2 or an H5 heading is not in any container. The content simply comes after the heading. This means that it is perfectly possible and legal in these languages to place different heading elements in any order you want. Writers have to pay attention to which heading level they are creating and how it fits in the structure of the document they are creating.

By contrast, in DocBook, we have a `section` structure. Like a list, a section is part of the writer's interpretation of what they are creating in the document, but it is only implied, not instantiated by the formatting. The `section` structure, instantiates the concept of a section. And once we have the instantiation of a section, we don't need six levels of heading. We can have one structure called `title`. Sections can be nested inside other sections, and the formatting algorithm can apply the correct style to the title based on context:

```
section:  
  title:  
  paragraph:  
  section:  
    title:
```

This will eliminate incorrect heading element choices, ensuring that the headings in the output consistently reflect the section and subsection structure of the document.

(Now it must be said that not everyone holds to the view that headings in a text do or should reflect a hierarchy of sections. Sometimes they may be simply signposts along the way, and like any signpost, the size of the sign reflects the size of the town, not a strict hierarchy of sign sizes. So if that is how you look at document structures, you should choose a different way to separate content from formatting in your content.)

Separate out abstract formatting

We noted that in the case of ordered and unordered lists, separating content from formatting actually involved separating out some of the content as well. Or rather, that it involved separating out some of the characters. In other words, the distinction between what is represented in a document using character codes and what is represented in other data structures is not necessarily the same as the logical distinction between content and formatting from a structured writing point of view.

Consider a structure that we might call a labeled list:

Separating Content from Formatting

Street	123 Elm Street
Town	Smallville
Country	USA
Code	12345

The generic structure of a labeled list might look like this:

```
labeled-list:  
    list-item:  
        label: Street  
        content: 123 Elm Street  
    list-item:  
        label: Town  
        contents: Smallville  
    list-item:  
        label: Country  
        contents: 123 USA  
    list-item:  
        label: Code  
        contents: 12345
```

But what if you have hundreds of addresses, all with the same labels. In this case, are the labels really content, or are they formatting? Since they don't change from one list to another, we could look at them as being part of how the content is presented, rather than being part of the content itself. So we look for a way to separate them from the content.

As always, when we separate something from our content, we have to replace it with something else, and that something is generally an appropriately named structure. So that gets us a structure like this:

```
address:  
    street: 123 Elm Street  
    town: Smallville  
    country: USA  
    code: 12345
```

Now, of course, we have moved our content into the subject domain. At this point we should make a distinction between formatting and presentation. Though these two words are often used interchangeably, we need to make a distinction between them for structured writing purposes. For purposes of this discussion, the formatting of text is the precise details of its appearance: the font chosen, the width of the text column, the size

Separating Content from Formatting

of the characters, the spacing between line, the size and shape of the bullet characters. Presentation, on the other hand, is about how we organize the content. Deciding to use a list to present a certain piece of information is a presentation decision, and it is independent of the formatting details we apply to lists. When we move content from the media domain to the document domain, we separate the formatting of the content from the presentation of the content. The decision to present the information as a list remains; the decision about what that list will look like has been separated out.

When we move the content from the document domain to the subject domain, however, we separate the information from the presentation. The subject domain structure above is not a labeled list. It is a data record that could be turned into many different forms of presentation.

The job of turning such structures into a specific form of presentation is the job of the presentation algorithm, which we will look at in more detail in Chapter 16, *Publishing*. The presentation algorithm could turn it into a labeled list, a table, a paragraph (with the fields separated by commas), or the address label for an envelope.

In the subject domain, with the content separated from both formatting and presentation, we also gain the ability to query and reorganize the content in various interesting and useful ways (which we will explore in further chapters).

This is as far as we can go in separating content from formatting, and we can't separate all content from formatting to quite this extent. It should be clear at this point that separating content from format is not a binary thing. There are various stages of separation that we can achieve for various reasons. It is important to understand exactly which degree of separation will best serve your needs.

Chapter 11. Processing Structured Text

As I have said from the beginning, structured writing involves the use of algorithms to assist in writing better content. Now it is time to take a look at what basic structured writing algorithms look like and how they are expressed. There are two parts to every structured writing algorithm: the analysis and imposition of structure, and the processing of that structure. Since the most basic of structured writing algorithms is separating content from formatting, the exploration of the algorithms is not complete until we look at how you reunite the separated pieces to produce output. Understanding the basics of these algorithms is important to understanding structured writing even if you don't intend to code the algorithms yourself. So now let's look at the algorithm for putting formatting and content back together.

Two into one: reversing the factoring out of invariants

Moving content from the media domain to the document domain or the subject domain involves progressively factoring out invariants in the content. Each step in this process creates two artifacts, the structured content and the invariant piece that was factored out.

Thus processing structured text is about putting the pieces back together: combining the structured content with the invariants that were factored out. If factoring out the invariants moves content toward the document or subject domains, recombining the content with the invariants moves it in the opposite direction, toward the media domain. This could mean moving the content from the subject domain to the document domain or from the document domain to the media domain, or simply from a more abstract form in the media domain to a more concrete form, which will be our first example.

Adding back style information

The first example of separating content from formatting that we looked at involved factoring out the style information from this structure:

```
{font: 10pt "Open Sans"}The box contains:  
{font: 10pt "Open Sans"}[bullet][tab]Sand  
{font: 10pt "Open Sans"}[bullet][tab]Eggs  
{font: 10pt "Open Sans"}[bullet][tab]Gold
```

We replaced the style information with style names:

```
{style: paragraph}The box contains:
```

```
{style: bullet-paragraph}Sand
{style: bullet-paragraph}Eggs
{style: bullet-paragraph}Gold
```

And then we defined the styles:

```
paragraph = {font: 10pt "Open Sans"}
bullet-paragraph = {font: 10pt "Open Sans"}[bullet][tab]
```

To unite the styles with the appropriate paragraphs, we can write as set of simple search and replace rules:

```
find {style: paragraph}
      replace {font: 10pt "Open Sans"}

find {style: bullet-paragraph}
      replace {font: 10pt "Open Sans"}[bullet][tab]
```

I said at the beginning that the basic processing algorithm was to combine two source of information to create a new one. Where are these two sources? The first source is the structured text. The second source is the style definitions, and they are embedded in the rules themselves. This is how it is usually done. In some cases, though, the rules may pull content from a separate file. We will see examples of this later.

The result of applying these rules is that we get back the original content:

```
{font: 10pt "Open Sans"}The box contains:
{font: 10pt "Open Sans"}[bullet][tab]Sand
{font: 10pt "Open Sans"}[bullet][tab]Eggs
{font: 10pt "Open Sans"}[bullet][tab]Gold
```

If we want to change the styles, we can apply a different set of rules:

```
find {style: paragraph}
      replace {font: 12pt "Century Schoolbook"}

find {style: bullet-paragraph}
      replace {font: 12pt "Century Schoolbook"} [em dash][tab]
```

Applying these rules will result in the a change in the formatting of the original content:

```
{font: 12pt "Century Schoolbook"}The box contains:
{font: 12pt "Century Schoolbook"} [em dash][tab]Sand
```

```
{font: 12pt "Century Schoolbook"}[em dash][tab]Eggs  
{font: 12pt "Century Schoolbook"}[em dash][tab]Gold
```

Rules based on structures

The tools that do this sort of processing do not literally use search and replace like this. Rather, they parse the source document to pull out the structures and allow you to specify your processing rules by referring to the structures.

We are not concerned here with the actual mechanism by which a processing tool recognizes structures. We are concerned with what to do when a structure is found. So let's rewrite our rules to match structures rather than find literal strings in the text:

```
match paragraph  
    apply style {font: 12pt "Century Schoolbook"}  
  
match bullet-paragraph  
    apply style {font: 12pt "Century Schoolbook"}  
    output "[em dash][tab]"
```

The way I have written these rules is an example of what is called pseudocode. It is a way for human being to sketch out an algorithm to make sure that they understand what they are trying to do before they write actual code. There is no formal grammar or syntax to pseudocode. It is intended for humans, not computers, and you can use whatever approach you like as long as it is clear to your intended audience. But pseudocode should clearly lay out a set of logical steps for accomplishing something. It should make clear exactly how the pieces go together.

Writing algorithms in pseudocode is a great way to make sure that we understand the algorithms we are creating without worrying about the details of code -- or even learning how to code. They are also a great way to communicate to actual coders about what we need a program to do.

The result of applying these rules is just the same as before:

```
{font: 12pt "Century Schoolbook"}The box contains:  
{font: 12pt "Century Schoolbook"}[em dash][tab]Sand  
{font: 12pt "Century Schoolbook"}[em dash][tab]Eggs  
{font: 12pt "Century Schoolbook"}[em dash][tab]Gold
```

The order of the rules does not matter

You may have noticed that what these rules are doing is pretty much exactly what style sheets do in an application like Word or FrameMaker. In fact, it is exactly what a style

sheet does. If you understand style sheets, you understand a good deal of how structured writing algorithms work.

One important thing to notice is that when you create a style sheet in Word or FrameMaker, you don't specify the order in which styles will be applied to the document. The same is true when you create a CSS stylesheet for the Web. The style sheet is just a flat list of rules. The order in which the rules are applied to the document depends entirely on the order in which the various structures occur in the document.

This may seem very obvious, but it is key to understanding how structured text is usually processed. It is a subject that is sometimes quite confusing to people who have been trained to write procedural computer programs, which is why I am making a point of calling it out.

Things get a tiny bit more complex when we move into processing the nested structures of the document domain and subject domain, but the basic pattern of a set of unordered rules to describe a transformation algorithm still applies.

Applying rules in the document domain

Suppose we have a piece of document domain structured text that contains this `title` structure:

```
title: Moby Dick
```

We want to transform this document into HTML (which kind of sits on the boundary between the media domain and the document domain). When our rule matches a structure in the source document, it outputs the equivalent HTML structure. Here is the pseudocode for this rule (it is in a slightly different format from the pseudocode above):

```
match title
    create h1
        continue
```

This says, when you see a `title` structure in the source, create an `h1` structure in the output, and then continue applying rules to the content of the `title` structure.

The `continue` instruction is indented under the `create h1` instruction to indicate that the results of continuing will appear inside the `h1` structure.

In our pseudocode, we are assuming that the text content of each structure will be output automatically (as is the case in many tools), so the output of this rule (expressed in HTML) is:

```
<h1>Moby Dick</h1>
```

But suppose that there is another structure inside the title in our source. In this case it is an annotation of part of the title text:

```
title: Review of {Rio Bravo} (movie)
```

Here the annotated text is set off with curly braces and the annotation itself in parentheses immediately after it. (This is a feature of the SAM markup syntax that I am using for most of the examples in this book.) So the annotation says that the words “Rio Bravo” refer to a movie. The annotation is a content structure just like the title structure, and is nested inside the text of the title.

So what do we do with our rule for processing titles to make it deal with `movie` annotations embedded in the title text? Absolutely nothing. Instead, we write a separate rule for handling `movie` annotations no matter where they occur:

```
match movie
    create i
        continue
```

When the processor hits `continue` in the `title` rule, it processes the content of the title structure. In doing so, it encounters the `movie` structure and executes the `movie` rule. The result is output that looks like this:

```
<h1>Review of <i>Rio Bravo</i></h1>
```

The `continue` instruction is really all we need to add to our rules to allow them to deal with nested structures. They remain an unordered collection of rules, just like a stylesheet. (In fact, XSLT, a language that implements this model, calls a set of processing rules a “stylesheet”.)

Processing based on context

When we move to the document domain, we use context to reduce the number of structures that we need. For example, where HTML has six different heading structures, H1 through H6), DocBook has only one: `title`, which can occur in many different contexts. So how do we apply the right formatting to a title based on its context? We create different rules for the `title` structure in each of its contexts. We express the context by listing the parent structure names separated by slashes:

```
match book/title
```

```

create h1
    continue

match chapter/title
    create h2
    continue

match section/title
    create h3
    continue

match figure/title
    create h4
    continue

```

Now here is the clever bit. You don't have to change the `movie` rule to work with any of these versions of the `title` rule. Suppose our title is the title of a section, like this:

```

section:
    title: Review of {Rio Bravo}(movie)</title>

```

When we process this with our rules, the `section/title` rule will be executed to deal with the title structure, and the `movie` rule will be executed when the `movie` structure occurs in the course of processing the content of the `title` structure, with the following result:

```
<h3>Review of <i>Rio Bravo</i></h3>
```

This is the basic pattern for most structured writing algorithms. An algorithm consists of a set of rules.

- For each structure, you create a rule that says how to transform that structure into the structure you want.
- Each rule specifies the new structures to create and where to place the content and any nested structures.
- In each rule, you specify where the processor should process any nested structures and apply any rules that apply to them.
- If you want a different rule for a structure occurring in different contexts, write a separate rule for each context.

Why is it important for you to understand this? Because when you are going through the process of abstracting out invariants to move content to the document domain or subject

domain, it is really useful to understand how those invariants will be factored back in. In fact, understanding how this works can help you recognize invariants in your source and give you the confidence to factor them out. Writing down the pseudocode for processing the structure you are creating can help you validate that you have factored things out correctly and that the structures you are creating will be easy to process and that the processing rules will be clear, consistent, and reliable.

Obviously there is more involved in a complete processing system, and we are not going to get into the gritty details here, but let's look at few a cases that come up frequently.

Processing container structures

When we move content to the document domain or the subject domain, we often create container structures to provide context. These container structures don't have any analog in the media domain, so what do we do with them when it is time to publish? We obviously use them to provide context for the rest of our processing rules, but what to we do with the containers themselves?

In the previous example the content was contained in a `section` structure. So how does the `section` structure get processed?

```
match section
    continue
```

Yes, it's that simple. Just don't output any new structure in its place. The `section` container has done its work at this point so we simple discard it. We still want the stuff inside it though, so we use the `continue` instruction to make sure the contents get processed. In short, the container is a box. We unpack the contents and discard the box.

Restoring factored-out text

Sometimes when we factor out the invariants in content, we are not only factoring out styles, we are also factoring out text. To process the content we need to put the text back (obviously we can put back different text depending on our needs -- which was why we factored it out in the first place).

As we saw, a simple example of factoring out text is numbered and bulleted lists, where we factor out the text of the numbers and bullets. Let's look at how we create rules to put them back.

Suppose we have a document that contain these two different kinds of lists:

```

paragraph: To wash hair:
ordered-list:
    list-item:Lather
    list-item:Rinse
    list-item:Repeat

paragraph: The box contains:
unordered-list:
    list-item:Sand
    list-item:Eggs
    list-item:Gold

```

Let's write a set of rules to deal with this document. Converting this to HTML lists won't tell us much, since HTML handles list numbering and bullets itself, so we'll create instructions for printing on paper. We won't use real printing instructions (they get tediously detailed). Instead we will use the same style specification shorthand we used above. The paragraph rule is simple enough:

```

match paragraph
    apply style {font: 10pt "Century Schoolbook" }
    continue

```

Now let's deal with the `ordered-list`. The ordered list structure is just a container, so we don't need to create an output structure for it. But because this is an ordered list, we need to start a count to number the items in the list. That means we need a variable to store the current count. We will use a \$ prefix to indicate that we are creating a variable:

```

match ordered-list
    $count=1
    continue

```

Then the rule for each ordered list item will output the value of the variable and increment it by one:

```

match ordered-list/list-item
    apply style {font: 12pt "Century Schoolbook" }
    output $count
    output "[tab]"
    $count=$count+1
    continue

```

Every time the `ordered-list/list-item` rule is fired, the count will increase by one, resulting in the list items being numbered sequentially.

If a new numbered list is encountered, the `ordered-list` rule will be fired, resetting the count to 1.

This rule will not match `list-item` elements that are children of an `unordered-list` element, so we need a separate set of rules of unordered lists. Because `unordered-list` is just a container and does not produce any formatted output, its rule has nothing to do:

```
match unordered-list
    continue

match ordered-list/list-item
    apply style {font: 12pt "Century Schoolbook"}
    output "[em dash][tab]"
    continue
```

Applying these rules will produce output like this:

```
{font: 10pt "Century Schoolbook"}To wash hair:
{font: 10pt "Century Schoolbook"}1.[tab]Lather
{font: 10pt "Century Schoolbook"}2.[tab]Rinse
{font: 10pt "Century Schoolbook"}3.[tab]Repeat
{font: 10pt "Century Schoolbook"}The box contains:
{font: 10pt "Century Schoolbook"}[em dash][tab]Sand
{font: 10pt "Century Schoolbook"}[em dash][tab]Eggs
{font: 10pt "Century Schoolbook"}[em dash][tab]Gold
```

Note how the structure has been flattened and all of the abstractions of document structure have been removed. We are back in the media domain, with a flat structure that specifies formatting and text.

Processing in multiple steps

We do not always want to apply final formatting to our content in a single step. When we separated content from formatting, we did the separation in several stages. It is often desirable to put them back together in several stages. Not only are the algorithms involved easier to write and maintain if they only do one step of the process, we can often reuse some of the downstream steps (nearer the media domain) for many different types of document domain and subject domain content.

So far we have looked at examples from the media domain and the document domain. Let's look at one from the subject domain. We used an example of completing the separation of content from formatting by moving a labeled list from the document domain to the subject domain.

```
address:  
  street: 123 Elm Street  
  town: Smallville  
  country: USA  
  code: 12345
```

Now let's look at the algorithm (the set of rules) for getting it back to the document domain, where it should look like this:

```
labeled-list:  
  list-item:  
    label: Street  
    contents: 123 Elm Street  
  list-item:  
    label: Town  
    contents: Smallville  
  list-item:  
    label: Country  
    contents: 123 USA  
  list-item:  
    label: Code  
    contents: 12345
```

Here is the set of rules to accomplish this transformation:

```
match address  
  create labeled-list  
  continue  
  
match street  
  create list-item  
  create label  
    output "Street"  
  create contents  
  continue  
  
match town  
  create list-item  
  create label  
    output "Town"  
  create contents  
  continue
```

```
match country
  create list-item
    create label
      output "Country"
  create contents
  continue

match code
  create list-item
    create label
      output "Code"
  create contents
  continue
```

Notice that the text of the labels, which we factored out when we moved to the subject domain, are being factored back in here, and are specified in the processing rules. As we moved the content from the media domain to the document domain to the subject domain, we first factored out invariant formatting and then invariant text. In the algorithms, we put back the text and the formatting, each at a different processing stage.

Processing content in multiple steps can save us a lot of time. The subject domain address structure is specific to a single subject and we might have many similar structures in our subject domain markup. But it is presented as a `labeled-list` structure. A labeled list is a document domain structure that can be used to present all kinds of information, and that can be formatted for many different media. By transforming the address structure into a `labeled-list` structure, we avoid having to write any code to format the address structure directly. We can format the address correctly for multiple media using the existing `labeled-list` formatting rules.

Query-based processing

The rule-based approach shown here is not the only way to process structured writing. There is another approach which we could call the query-based approach.¹ In this approach, you write a query expression that reaches into the structure of a document and pulls out a structure or a set of structures from the middle of the document.

This is a useful technique if you want to radically rearrange the content of a document or if you want to pull content out of one document to use in another. We will look at algorithms that use the query-based approach in later chapters.

¹The rule-based and query-based approaches are often called “push” and “pull” methods respectively, but I sometimes find it hard to remember which is which. I find rule-based and query-based more descriptive.

Chapter 12. Single Sourcing

Single sourcing was one of the earliest motivations for structured writing. However, the term “single sourcing” gets used to mean different things, all of which involve a single source in one way or another, but which use different approaches and achieve different ends. To make life easier, I will distinguish three main meanings of “single sourcing” as follows:

Single sourcing	Producing the same document in different media.
Content reuse	Using the same content to create different documents.
Single source of truth	Ensuring that each piece of information is recorded only once.

In this chapter we will look at single sourcing as defined above. In Chapter 13, *Reuse* we will look at content reuse and in Chapter 15, *Single source of truth* we will look at single source of truth.

Basic single sourcing

The basic single sourcing algorithm is straightforward and we have covered most of it already in Chapter 11, *Processing Structured Text*.

Basic single sourcing involves taking a piece of content in the document domain and processing its document domain structures into different media domain structures for each of the target media.

Suppose we have a recipe recorded in the document domain:

```
page:  
    title: Hard Boiled Eggs  
  
    A hard boiled egg is simple and nutritious.  
    Prep time, 15 minutes. Serves 6.  
  
    section: Ingredients  
        ul:  
            li: 12 eggs  
            li: 2qt water  
  
    section: Preparation  
        ol:  
            li: Place eggs in pan and cover with water.
```

```
    li: Bring water to a boil.  
    li: Remove from heat and cover for 12 minutes.  
    li: Place eggs in cold water to stop cooking.  
    li: Peel and serve.
```

We can output this recipe to two different media by applying two different formatting algorithms. First we output to the Web by creating HTML.

```
match page  
  create html  
    stylesheet www.example.com/style.css  
    continue  
  
match title  
  create h1  
  continue  
  
match p  
  copy  
  continue  
  
match section  
  continue  
  
match section/title  
  create h2  
  continue  
  
match ul  
  copy  
  continue  
  
match ol  
  copy  
  continue  
  
match li  
  copy  
  continue
```

In the code above, paragraph and list structures have the same names in the source format as they do in the output format (HTML) so we just copy the structures rather than recreating them. This is a common pattern in structured writing algorithms.

The above algorithm transforms our source into HTML that looks like the following:

```
<html>
  <head>
    <link rel="stylesheet" type="text/css" href="//www.apache.org/
  </head>

  <h1>Hard Boiled Eggs</h1>

  <p>A hard boiled egg is simple and nutritious. Prep time, 15 minu
  <h2>Ingredients</h2>

  <ul>
    <li>12 eggs</li>
    <li>2qt water</li>
  </ul>

  <h2>Preparation</h2>

  <ol>
    <li>Place eggs in pan and cover with water.</li>
    <li>Bring water to a boil.</li>
    <li>Remove from heat and cover for 12 minutes.</li>
    <li>Place eggs in cold water to stop cooking.</li>
    <li>Peel and serve.</li>
  </ol>
</html>
```

Outputting to paper (or to PDF, which is a kind of virtual paper) is more complex. On the Web, you output to a screen which is of flexible width and infinite length. The browser generally takes care of wrapping lines of text to the screen size (unless formatting commands tell it to do otherwise) and there is no issue with breaking text from one page to another. For paper, though, you have to format for a fixed size page. This means that formatting for paper involves fitting the content into a set of fixed size pages.

This leads to a number of formatting problem, such as where to break each line of text, how to avoid a heading appearing at the bottom of a page or the last line of a paragraph appearing as the first line of a page. It also creates issues with references. For instance, a reference to content on a particular page cannot be known until the pages are paginated by the algorithm.

Because of issues like this, you don't write a formatting algorithm for paper directly, the way you would write an algorithm to output HTML. Rather, you use an intermediate typesetting system which already knows how to handle things like inserting page number references and handling line and page breaks. Rather than handling these things yourself,

you tell the typesetting system how you would like it to handle them and then let it do its job.

One such typesetting system is XSL-FO (Extensible Stylesheet Language - Formatting Objects). XSL-FO is a typesetting language written in XML. To format your content using XSL-FO, you transform your source content into XSL-FO markup, just the way you transform it into HTML for the Web. But then you run the XSL-FO markup through an XSL-FO processor to produce your final output, such as PDF. (I call this the encoding algorithm.)

Here is a small example of XSL-FO markup:

```
<fo:block space-after="4pt">
    <fo:wrapper font-size="14pt" font-weight="bold">
        Hard Boiled Eggs
    </fo:wrapper>
</fo:block>
```

As you can see, the XSL-FO code contains a lot of specific media domain instructions for spacing and font choices. The division between HTML for basic structures and CSS for specific formatting does not exist here. Also note that as a pure media-domain language, XSL-FO does not have document domain structures like paragraphs and titles. From its point of view, a document consists simply of a set of blocks with specific formatting properties attached to them.

Because of all this detail, I am going to show the literal XSL-FO markup in the pseudocode of the algorithm, and I am not going to show the algorithm for the entire recipe. (The point is not for you to learn XSL-FO here, but to understand how the single-sourcing algorithm works.)

```
match title
    output '<fo:block space-after="4pt">'
        output '<fo:wrapper font-size="14pt" font-weight="bold">'
            continue
        output '</fo:wrapper>'
    output '</fo:block>'
```

Other typesetting systems you can use for print output include TeX and later versions of CSS.

Differential single sourcing

Basic single sourcing outputs the same document presentation to different media. But each media is different, and what works well in one media does not always work as well

in another. For example, online media generally support hypertext links, while paper does not. Let's suppose that we have a piece of content that includes a link.

```
In Rio Bravo, {the Duke}(link "http://JohnWayne.com")  
plays an ex-Union colonel out for revenge.
```

In SAM syntax, the piece of markup represented by "http://JohnWayne.com" specifies the address to link to. In the algorithm examples below, this markup is referred to as the "specifically" attribute using the notation @specifically.

In HTML we want this output as a link using the HTML a element, so we write the algorithm like this:

```
match p  
    copy  
        continue  
  
match link  
    create a  
        attribute href = @specifically  
        continue
```

The result of this algorithm is:

```
<p>In Rio Bravo, <a href="http://JohnWayne.com">The Duke</a>  
plays and ex-Union colonel out for revenge.  
</p>
```

But suppose we want to output this same content to paper. If we output it to PDF, we could still create a link just like we do in HTML, but if that PDF is printed, all that will be left of the link will be a slight color change in the text and maybe an underline. It will not be possible for the reader to follow the link or see where it leads.

Paper can't have active links but it can print the value of a URL so that reader can type it into a browser if they want to. An algorithm could do this by printing the link inline or as a footnote. Here is the algorithm for doing it inline. (We'll dispense with the complexity of XSL-FO syntax this time.)

```
match p  
    create fo:block  
        continue  
  
match link
```

```
continue
output " (see: "
output @specifically
output ") "
```

This will produce:

```
<fo:block>
In Rio Bravo, the Duke (see: http://JohnWayne.com)
plays an ex-Union colonel out for revenge.
</fo:block>
```

This works, but we should note that the effect is not exactly the same in each media. Online, the link to JohnWayne.com serves to disambiguate the phrase “The Duke” for those readers who do not recognize it. A simple click on the link will explain who “the Duke” is. But in the paper example, such disambiguation exists only incidentally, because the words “JohnWayne” happen to appear in the URL. This is not how we would disambiguate “The Duke” if we were writing for paper. We would be more likely to do something like this:

The Duke (John Wayne) plays an ex-Union colonel.

This provides the reader with less information, in the sense that it does not give them access to all the information on JohnWayne.com, but it does the disambiguation better and in a more paper-like way. The loss of the reference to JohnWayne.com is probably not an issue here. Following that link by typing it into a browser is a lot more work than simply clicking on it on a Web page. If someone reading on paper wants more information on John Wayne they are far more likely to type “John Wayne” into Google than type “JohnWayne.com” into the address bar of their browser.

With the source content written as it is, though, there is no easy way to produce this preferred form for paper. The content is in the document domain and the document domain specifies the presentation of the content. It is therefore not a neutral format between different media that require different forms of presentation. The choice to specify a link gives it a strong bias towards the Web and online media rather than paper. A document domain approach that favored paper would similarly lead to a poorer online presentation that omitted the link.

What we need to address the problem is a differential approach to single sourcing, one that allows us to differ not only the formatting but the presentation of the content for different media.

One way to accomplish this differential single sourcing is to record the content in the subject domain, thus removing the prejudice of the document domain presentation for one class of media or another. Here is how this might look:

{The Duke} (actor "John Wayne") plays an ex-Union colonel.

In this example, the phrase “The Duke” is annotated with a subject domain annotation that clarifies exactly what the text refers to. That annotation says that “the Duke” is the name of an actor, specifically “John Wayne”.

Our document domain examples attempted to clarify “the Duke” for readers, but did so in media-dependent ways. This subject domain example clarifies the meaning of “The Duke” in a presentation-neutral way by explicitly marking up what the phrase refers to in a formal and unambiguous way. This disambiguation is not written for readers, however. It is written for algorithms. This allows an algorithm to generate a media-appropriate clarification for the reader in the target media.

For paper:

```
match actor
    continue
    output " ( "
    output @specifically
    output " ) "
```

For the Web:

```
match actor
    create link
        $href = get link for actor named @specifically
        attribute href = $href
    continue
```

This supposes the existence of a system that can respond to the `get link` instruction and look up pages to link to based on the type and a name of a subject. We will look at how a system like that works in Chapter 17, *Linking*.

Differential organization and presentation

Differences in presentation between media can be broader than this. Paper documents sometimes use complex tables and elaborate page layouts that often don’t translate well to online media. Effective table layout depends on knowing the width of the page you have available, and online you don’t know that. A table that looks great on paper may be unreadable on a mobile device, for instance.

And this is more than a layout issue. Sometimes the things that paper does in a static way should be done in a dynamic way in online media. For example, airline or train schedules have traditionally been printed as timetables on paper, but you will virtually never see them presented that way online. Rather, there will be an interactive travel planner that lets you choose your starting point, destination, and desired travel times and then presents you with the best schedule, including when and where to make connections.

Single sourcing your timetable to print and PDF will not produce the kind of online presentation of your schedule that people expect, and that can have a direct impact on your business.

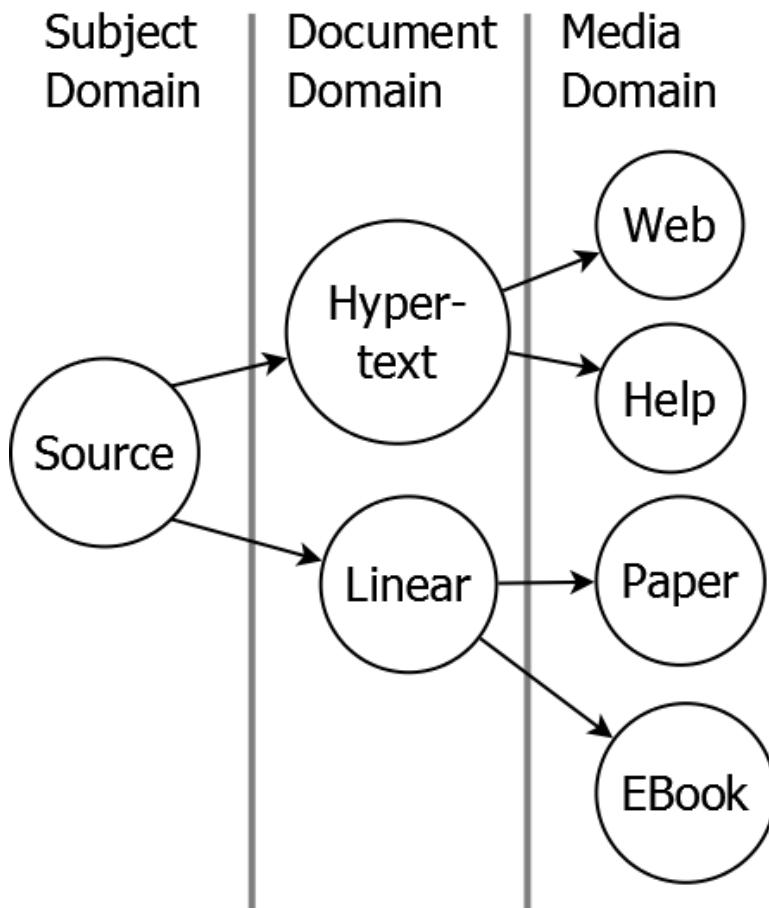
To single source schedule information to paper and online, you can't maintain that content in a document domain table structure. You need to maintain it in a timetable database structure (which is subject domain, but really looks like a database, not a document at all).

An algorithm (which I will call the synthesis algorithm) can then read the database to generate a document domain table for print publication. For the Web, however, you will create a web application that queries the database dynamically to calculate routes and schedules for individual travelers.

Differences in linking between media can also go much deeper than how the links are presented. Links are not simply a piece of formatting like bold or italics. Links connect pieces of content together. On paper, documents are designed linearly, with one section or chapter after another. But online you can organize information into a hypertext¹ with links that allow the reader to navigate and read in many different sequences.

The difference between linear information design and hypertext information design is not a media domain distinction but a document domain distinction. But if you are thinking about single sourcing your content it is one that you have to take into consideration. In other words, single sourcing is not just about one document domain source with many media domain outputs. It can also be about a single subject domain source with multiple document domain outputs expressing different information designs, and outputting them to different media.

¹<http://everypageispageone.com/2013/03/19/web-organization-is-not-like-book-organization/>



More radical forms of differential single sourcing start to look a lot like reusing the same content to build quite different documents (albeit on the same subject) and therefore start to use the techniques of content reuse, which we will deal with in Chapter 13, *Reuse*.

Conditional differential design

You can also do differential single sourcing by using conditional (management domain) structures in the document domain.

For instance, if you are writing a manual that you intend to single source to a help system, you might want to add context setting information to the start of a section when it appears in the help system. The manual may be designed to be read sequentially, meaning that the context of individual sections is established by what came before. But help systems are always accessed randomly, meaning that the context of a particular help topic may not be clear if it was single sourced from a manual. To accommodate this, you could include a context setting paragraph that is conditionalized to appear only in help output:

```
section: Wrangling left-handed widgets
```

```
~~~(?help-only)
```

Left-handed widgets are used when wrangling counter-clockwise

To wrangle a left handed widget:

1. Loosen the doohickey using a medium thingamabob.
2. Spin three times around under a full moon.
3. Touch the sky.

In the SAM markup above, the `~~~` creates a “fragment” structure to which conditional tokens can be applied. Content indented under the fragment marker is part of the fragment.

To output a manual, we suppress the help-only content:

```
match fragment where conditions = help-only
    ignore
```

To output help, we include it:

```
match fragment where conditions = help-only
    continue
```

Primary and secondary media

While there is a lot you can do in the way of differential single sourcing to successfully output documents that work well in multiple media, there are limits to how far this approach can take you.

In the end, linear and hypertext approaches a fundamentally different ways of writing which invite fundamentally different ways of navigating and using information. Even moving content to the subject domain as much as possible will not entirely factor out these fundamental differences of approach.

When single sourcing content to both linear paper-like media and hypertext web-like media, you will generally have to choose a primary media to write for. Single sourcing that content to the other media will be on a best-effort basis. It may be good enough for a particular purpose, but it will never be quite as good as it could have been had you designed for that media.

Many of the tools used for single sourcing have an built in bias towards one media or another. Desktop-publishing tools like FrameMaker, for instance, were designed for linear media. Online collaborative tools like wikis were designed for hypertext media.

It is usually a good idea to pick a tool that was designed for the media you choose as your primary.

In many cases, the choice of primary media is made implicitly based on the tools a group has traditionally been using. This usually means that the primary media is paper, and it often continues to be so even after the organization had stopped producing paper and their readers are primarily using online formats.

Some organizations seem to feel that they should only switch to tools that are designed primarily for online content when they have entirely abandoned the production of paper and paper-like formats such as PDF. This certainly does not need to be the case. It is perfectly possible to switch to an online-primary tools and still produce linear media as a secondary output format.

Changing your information design practices from linear paper based designs to hypertext Every Page is Page One (???) designs is non-trivial, but such designs better suit the way many people use and access content today. Don't expect single sourcing to successfully turn document-oriented design into effective hypertext by themselves. To best serve modern readers it will usually be much more effective to adopt an Every Page is Page One approach to information design and use structured writing techniques to do a best-effort single sourcing to linear media for those of your readers who still need paper or paper-like formats -- or for those documents where linear organization still makes sense.

Responsive Design

Responsive design, an approach to web design in which a page automatically adjust to the viewport it is displayed on, is a form of differential single sourcing. Many attempts at responsive design are not as responsive as we would like them to be because they content is not stored in a source format that allows a sufficient degree of differential presentation. If you want to achieve effective and consistent responsive design in your output, you will find that the use of subject domain structures in your content will be enormously helpful.

It is very difficult to expect writers to encode the presentation structures of a responsive design into their files as they write. Not only that, the very idea of responsive design should take into account the possibility of having to respond to new media in the future. Every sound responsive design decision is based on the appropriate organization of subject matter in a particular view port. The rules of effective responsive design are going to be express in content terms. The best way to create content that will presented using responsive design, therefore, is in the subject domain where the subject matter divisions on which responsive design rules are expressed are made fully explicit in the content.

Chapter 13. Reuse

The reuse of content in different contexts has become one of the main drivers of structured writing, particularly in the form of widespread adoption of DITA. Content reuse is not, in itself, a technique that promotes content quality (except perhaps if it succeeds in freeing up time to focus on creating new and better content with the time saved by reusing existing content). However, implementing content reuse without thinking about content quality can do serious damage to content quality. Therefore any consideration of structured writing for content quality has to look seriously at the content reuse algorithm and how it can be implemented without hurting content quality.

Content reuse is not one technique, but a collection of many techniques, each of which requires specific content structures in the subject, document, and management domains. The simplest content reuse technique is cutting and pasting content from one source to another. This approach is quick and easy to implement, but it creates a host of management problems. When people talk about content reuse they generally mean any and every means of reusing content other than cutting and pasting.

This means that reusing content really means storing a piece of content in one place and inserting it into more than one publication by reference. “Reusing” can suggest that this activity is somewhat akin to rummaging through that jar of old nuts and bolts you have in the garage looking for one that is the right size to fix your lawnmower. While you can do it that way, that approach is neither efficient nor reliable. The efficient and reliable approach involves deliberately creating content for use in multiple locations. This means that you need to place constraints on the content to be reused and the content that reuses it, and that means you are in the realm of structured writing.

Fitting pieces of content together

If you are going to create one piece of content that can be used in many outputs, you have to make sure it fits in each of those outputs.

If you cut and paste, this is not a concern. You can cut any text you like, paste it in anywhere, and edit it to fit if you need to. But if the content you want to use is used in other places, you can’t edit it to fit because that might cause it to no longer fit in the other places. For reuse to work, the content must be written to fit in multiple places. In other words, it has to meet a set of constraints that will allow it to fit in multiple places. We will look at this in more detail in Chapter 14, *Composition*. In this chapter we will focus on the algorithms for fitting the pieces together.

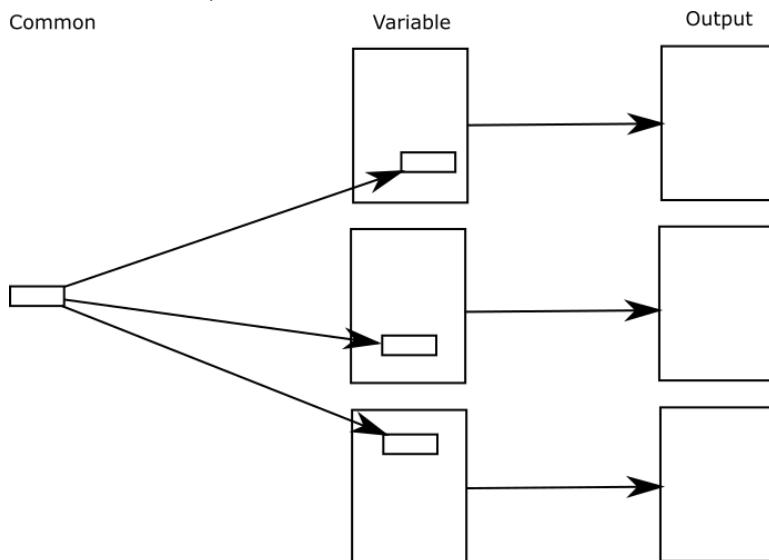
There are seven basic models for fitting pieces of content together:

- Common into variable

- Variable into common
- Variable into variable
- Common with conditions
- Factor out the common
- Factor out the variable
- Assemble from pieces

Common into variable

In the common into variable case, you have a common piece of content that occurs in many places. This could mean it occurs in many documents or in many places in the same document, or both.



For instance, if you are writing procedures and there is a common safety warning that must appear on all dangerous procedures, each individual procedure is the variable part and the standard warning is the common part.

We looked at an example of this in Chapter 7, *The Management Domain: an Intrusion*.

```
procedure: Blow stuff up
    >>>(files/shared/admonitions/danger)
    step: Plant dynamite.
    step: Insert detonator.
    step: Run away.
```

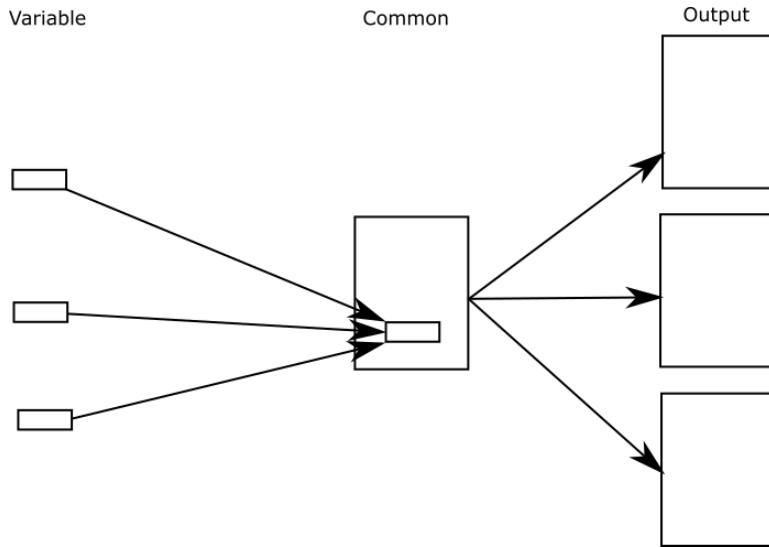
step: Press the big red button.

To ensure that the included content will always fit, you need to make sure that there is a clear division of responsibilities between the common content and each of the documents it will be inserted into. The inserted content should give the safety warning, the whole safety warning, and nothing but the safety warning. Each document that includes it should include it in the required place in the procedure structure.

Of course, you can also use the subject-domain approach to common into variable that we looked at in Chapter 7, *The Management Domain: an Intrusion*.

Variable into common

In the variable into common case, you have a single document that will be output in many different ways by inserting variable content at certain locations.



For instance, if you are writing a manual to cover a number of car models you can factor out the number of seats each model has.

The vehicle seats >(\$seats) people.

This is the fixed content that will occur in all manuals, with the number of seats pulled in from an external source. Let's say we have a collection of vehicle data that is stored in a structure like this:

```
vehicles:  
  vehicle: compact  
  seats: four
```

```
colors: red, green, blue, white, black
transmissions: manual, CVT
doors: four
horsepower: 120
torque: 110 @ 3500 RPM
vehicle: midsize
seats: five
colors: red, green, blue, white, black
transmissions: CVT
doors: four
horsepower: 180
torque: 160 @ 3500 RPM
```

Then we write the algorithm to process the insert so that it queries this structure.

```
match insert where variable = $seats
$number_of_seats = vehicles/vehicle[$model]/seats
output $number_of_seats
```

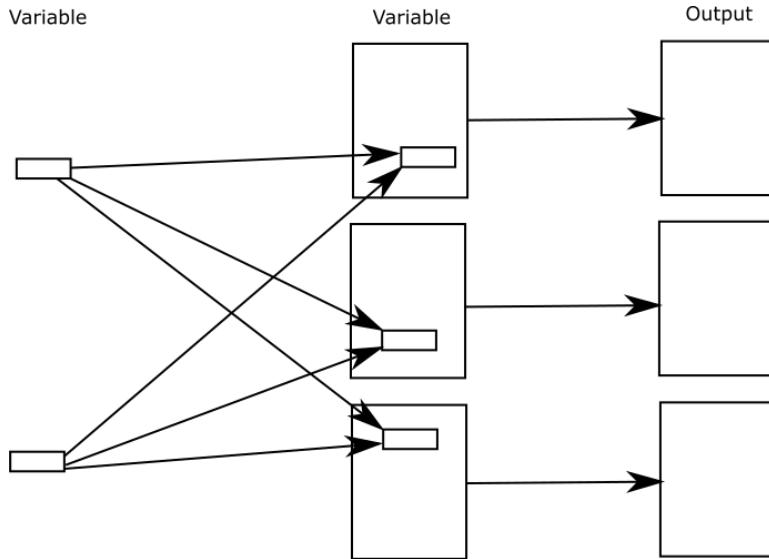
All these insert and query mechanisms are pseudocode, of course. Exactly how things work and exactly how you delineate, identify, and insert content vary from system to system.

With the variable into common technique, you are creating a common source by factoring out all the parts of the different outputs that are not common. This is, in some ways, the inverse of the usual pattern of factoring out invariants: we are actually factoring out the variants. But really, it amounts to the same thing. We are factoring variants from invariants. The only real difference between this and the common into variable is whether the common parts are embedded in the variable parts or vice versa. Either way, we still end up with two artifacts: the variable piece or pieces and the common piece or pieces.

Variable into variable

Variable into variable is a variation on common into variable in which you can make a wholesale change of the common elements that you are pulling into a set of variable documents.

For example, suppose you decide to market your product line to a new market. The new market has different safety regulation which means you need to insert a different standard warning into all your manuals. In this case, you want to swap out the common elements used in your home market and substitute the common elements for the foreign market.



Here we need to talk about how we identify the content to be inserted. In the common into variable example, we inserted the content of a file that contains a standard warning. But this approach is fragile. You can't reorganize your files without breaking reuse, and you almost always need to reorganize your files eventually. Plus, it forces you to have every piece of reusable content in a separate file, which may not be efficient.

But for variable into variable this approach simply does not work. Variable into variable requires loading a different file, which is difficult when the content specifies a particular file name to import.

As always in structured writing, we look for a way to factor out the problematic content. So here we look for a way to factor out the file name and replace it with something else.

Using IDs

The most basic way to factor out the file name is to give the content of the file an ID and use the ID to identify the content in a location independent way. Here is the warning file with the ID `#warn_danger` added:

```
warning: (#warn_danger)
title: Danger
```

Be very very careful. This could kill you.

We can then insert the warning into our procedure by referring to that ID.

```
procedure: Blow stuff up
>>>(#warn_danger)
```

```
step: Plant dynamite.  
step: Insert detonator.  
step: Run away.  
step: Press the big red button.
```

The responsibility for locating the warning has now been shifted from the content to the algorithm.

```
match insert with ID  
  $insert_content = find ID in $content_set  
  output $insert_content
```

This is a constant pattern in structured writing. When it comes to locating resources, you want to move that responsibility from the content to the algorithm. This makes it easier to update the locations, but it also gives you far more options for storing and managing your content, since algorithms can interact with a variety of systems in sophisticated ways, rather than just storing a static address. It also means you can make wholesale changes in how your content is stored without having to edit the content itself.

This means that the synthesis algorithm needs some way to resolve the ID and find the content to include. In many cases, a content management system is used to resolve the ID. In other cases it is as simple as the algorithm searching through a set of files to find the ID or building a catalog that points to the files that contain IDs.

To do variable into variable reuse in a system that uses IDs, you simply point the synthesis algorithm at a different set of files that contain the same IDs, but attached to different content. So if your foreign market requires a different warning, you can create a file like this:

```
warning: (#warn_danger)  
        title: Look out!
```

Pay close attention. You could really hurt yourself.

By telling the build to search this file for IDs rather than the file with the domestic market warning, you automatically get the the foreign warning rather than the domestic one.

Using keys

Another way to do this is with keys. Instead of assigning an ID directly to the content, the keys approach use an intermediate lookup table to resolve keys to particular resources.

So in this case we have the warning in a file called files/shared/admonitions/domestic/danger with the following content (no ID):

```
warning:  
    title: Danger
```

Be very very careful. This could kill you.

And we have the procedure which includes the warning via a key:

```
procedure: Blow stuff up  
    >>>(%warn_danger)  
    step: Plant dynamite.  
    step: Insert detonator.  
    step: Run away.  
    step: Press the big red button.
```

(I am using # to denote IDs and % to denote keys. This is purely arbitrary and has nothing to do with how they work. Different systems will denote IDs and keys in different ways.)

To connect the key to the warning file, we then create a key lookup table:

```
keys:  
    key:  
        name: warn_danger  
        resource: files/shared/admonitions/domestic/danger
```

When the synthesis algorithm processes the procedure, it sees the key reference %warn_danger and looks it up in the key lookup table. The key lookup table tells the algorithm that the key resolves to the resource files/shared/admonitions/domestic/danger. The algorithm then loads that file and inserts the contents into the output.

```
match insert with key  
    $resource = find key in lookup-table  
    output $resource
```

To output your content for the foreign market, you simply prepare a new key lookup table:

```
keys:  
    key:  
        name: warn_danger
```

```
resource: files/shared/admonitions/foreign/danger
```

You then tell the synthesis algorithm to use this lookup table instead.

Using keys is not necessarily better than using IDs. What it comes down to is that you need some kind of bridge between the citation of an identifier in the source file and the location of a resource with that identifier in the content store. This bridge can be created by a key lookup table, by remapping file URLs, or by modifying a query to a content repository.

One feature of the key approach is that, because it does not attach the key directly to the content, it can be used to identify resources that do not have IDs, which may include resources that you do not control.

One downside of a strictly external approach to keys is that they can only point to a whole resource. This could force you to keep all your reusable units in separate files. To avoid this, you can combine keys with IDs. The following example combines the foreign and domestic danger warnings into one file and gives each an ID:

```
warnings:
```

```
  warning: (#warn_danger_domestic)
    title: Danger
```

Be very very careful. This could kill you.

```
  warning: (#warn_danger_foreign)
    title: Look out!
```

Pay close attention. You could really hurt yourself.

Now we can rewrite our key lookup tables to use the IDs to pull the right warning out of this common file. For the domestic build we would use a key lookup table like this:

```
keys:
```

```
  key:
```

```
    name: warn_danger
    resource: files/shared/warnings#warn_danger_domestic
```

And for the foreign build, one that looks like this:

```
keys:
```

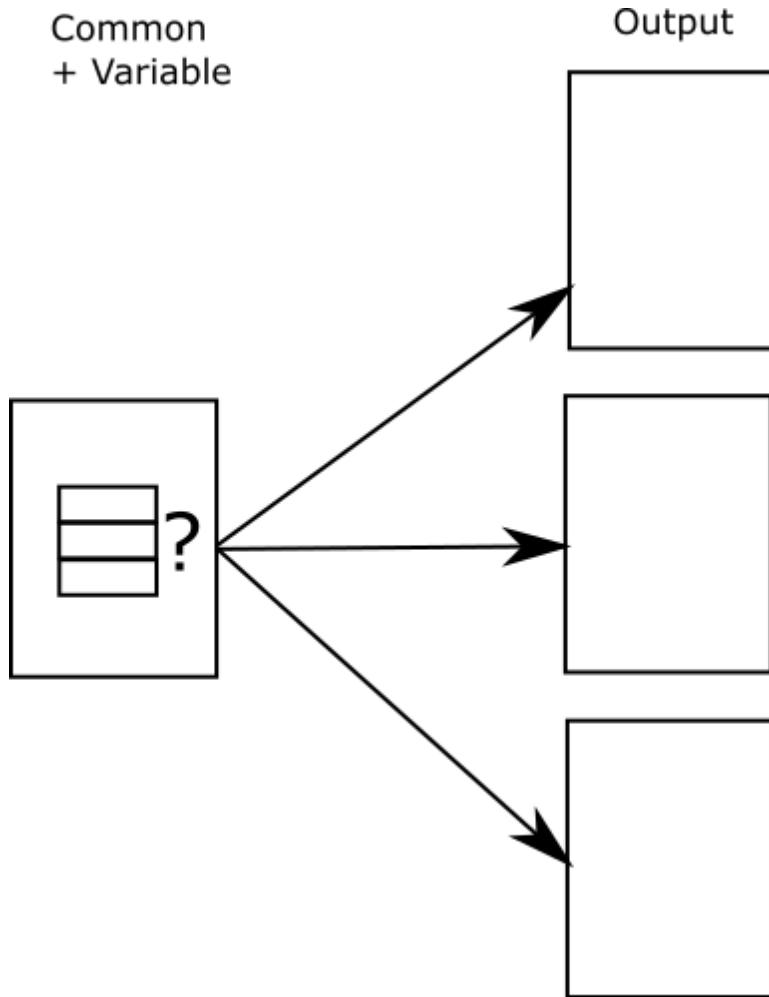
```
  key:
```

```
    name: warn_danger
```

```
resource: files/shared/warnings#warn_danger_foreign
```

Common with conditions

In some cases of variable into common, the variant pieces is not actually factored out into a separate file. Rather, each of the possible alternatives is included in the file conditionally.



For instance in content for a car manual you might have conditional text for the number of people the car seats.

```
The vehicle seats {four} (?compact) {five} (?midsize) {seven} (?van).
```

Here the main text is the fixed piece and the variable pieces are the words “four”, “five”, and “seven”. Which of these will be included in the output depends on which condition is applied during the build. If the condition `midsize` is applied, then the output text will be “five” and the other alternatives will be suppressed.

```
match phrase with condition
    if condition in $build_conditions
        continue
    else
        ignore
```

The upside of the conditional approach is that it keeps all the variants in one file, so your algorithm does not have to know where to go to find the external content.

But there are a number of downsides to this approach:

- It gets very cumbersome to read the source if there are many different conditions applied.
- When the subject matter changes, you have to find all the places the conditions occur and update them.
- If the same data point (the number of seats) is mentioned in many different documents, that information is still being duplicated all over the content, which makes it hard to maintain and verify, and hard to change if the compact seats five in the next model year.

Common with conditions is not limited to cases where there are alternate values, however. In some cases, content may simply be inserted or omitted for certain outputs.

The main features of the car are:

```
ol:
    li: Wheels
    li: Steering wheel
    li:(?deluxe) Leather seats
    li: Mud flaps
```

In this case, the list item “Leather seats” would only be published if the condition `deluxe` was specified in the build. It would be omitted for all other builds. These are the kinds of cases where it is harder to get away from the use of conditionals as a reuse mechanism.

This approach to reuse is often called filtering or profiling. Some systems have elaborate ways of specifying filtering or profiling of the content. The net effect is the same as the simple condition tokens shown here, but they may allow for more sophisticated or elaborate conditions than shown here.

Because common with conditions is essentially a form of variable into common where the variable content is contained inside the common source, it can technically be replaced

by a variable into common approach in all cases. In practice, the use of conditions tends to occur when:

- The number of variations is small and thought to be fixed or to change infrequently.
- The variable pieces are eccentric or contextually dependent.
- The writer or organization wishes to avoid managing multiple files.
- The current tools don't support variable into common.

How successful a common with conditions approach will be also depends on what you choose for your conditional expressions. Generally, subject domain conditions will be much more stable and manageable than document domain conditions. For instance, conditions that relate to different vehicles (subject domain) are based in the real world and are therefore objectively true as long as the subject matter remains the same. Conditions that relate to different publications or different media, on the other hand, are not objectively true and can't be verified independently. They only way to verify them is to build the different documents or media and see if you got the content you expected. This makes maintaining such conditions cumbersome and error prone.

Factor out the common

In Chapter 7, *The Management Domain: an Intrusion*, we noted that the subject domain alternative to using an insertion instruction for the warning text was to specify which procedures were dangerous, thus factoring out the constraint that the warning must appear. In effect, this factors out the common content as well.

```
procedure: Blow stuff up
    is-it-dangerous: yes
    step: Plant dynamite.
    step: Insert detonator.
    step: Run away.
    step: Press the big red button.
```

In this case, the author does not have to identify the material to be included, either directly by file name or indirectly through an ID or a key. Instead, it is up to the algorithm to include it:

```
match procedure/is-it-dangerous
    if is-it-dangerous = 'yes'
        output files/shared/warnings#warn_danger_domestic
```

To produce the foreign market version of the documentation, you simply edit the rule:

```
match procedure/is-it-dangerous
  if is-it-dangerous = 'yes'
    output files/shared/warnings#warn_danger_foreign
```

The beauty of this approach is that the content is entirely neutral as to what kind of reuse may be going on or how dangerous procedures may be treated. Because the content itself contains only objective information about the procedure itself, you can implement any algorithm you like to publish or reuse the content in any way you like at any time based on this information. By making the content not specific to any form of reuse or any reuse mechanism, we effectively make it much more reusable.

We are also making the content much easier to write, since this approach does not require the writer to know how the reuse mechanism works, how to identify reusable content, or even that reuse is occurring at all. All they have to do is answer a simple question about the content -- one to which they should already know the answer.

Structured writing is about factoring invariants and complexities out of content and this approach enables the widest range of reuse possibilities while factoring all the complexities of reuse out of the content. The downside of this is that it is not as general. The `is-it-dangerous` metadata applies only to dangerous procedures. It does not address the inclusion of reusable content in other places. You would need to factor out other interesting reuse cases in a similar way to create a complete subject-domain solution.

Factor out the variable

You can also factor out the variable content. For example, in the case of the different models of a car, rather than conditionalizing the list of features in the document, like this:

The main features of the car are:

```
ol:
  li: Wheels
  li: Steering wheel
  li:(?deluxe) Leather seats
  li: Mud flaps
```

You can factor out the list entirely:

The main features of the car are:

```
>>>(%main_features)
```

You can then maintain the features list in a database. The organization probably already has a database of features for each vehicle, so we don't need to create anything new. We simply query the existing database. (After all, this is about reusing what already exists rather than recreating it!)

So now our algorithm looks something like this:

```
match insert with key
    $resource = lookup key in lookup-table
    output $resource
```

We then have a key lookup table where the resource is identified by a query on the database

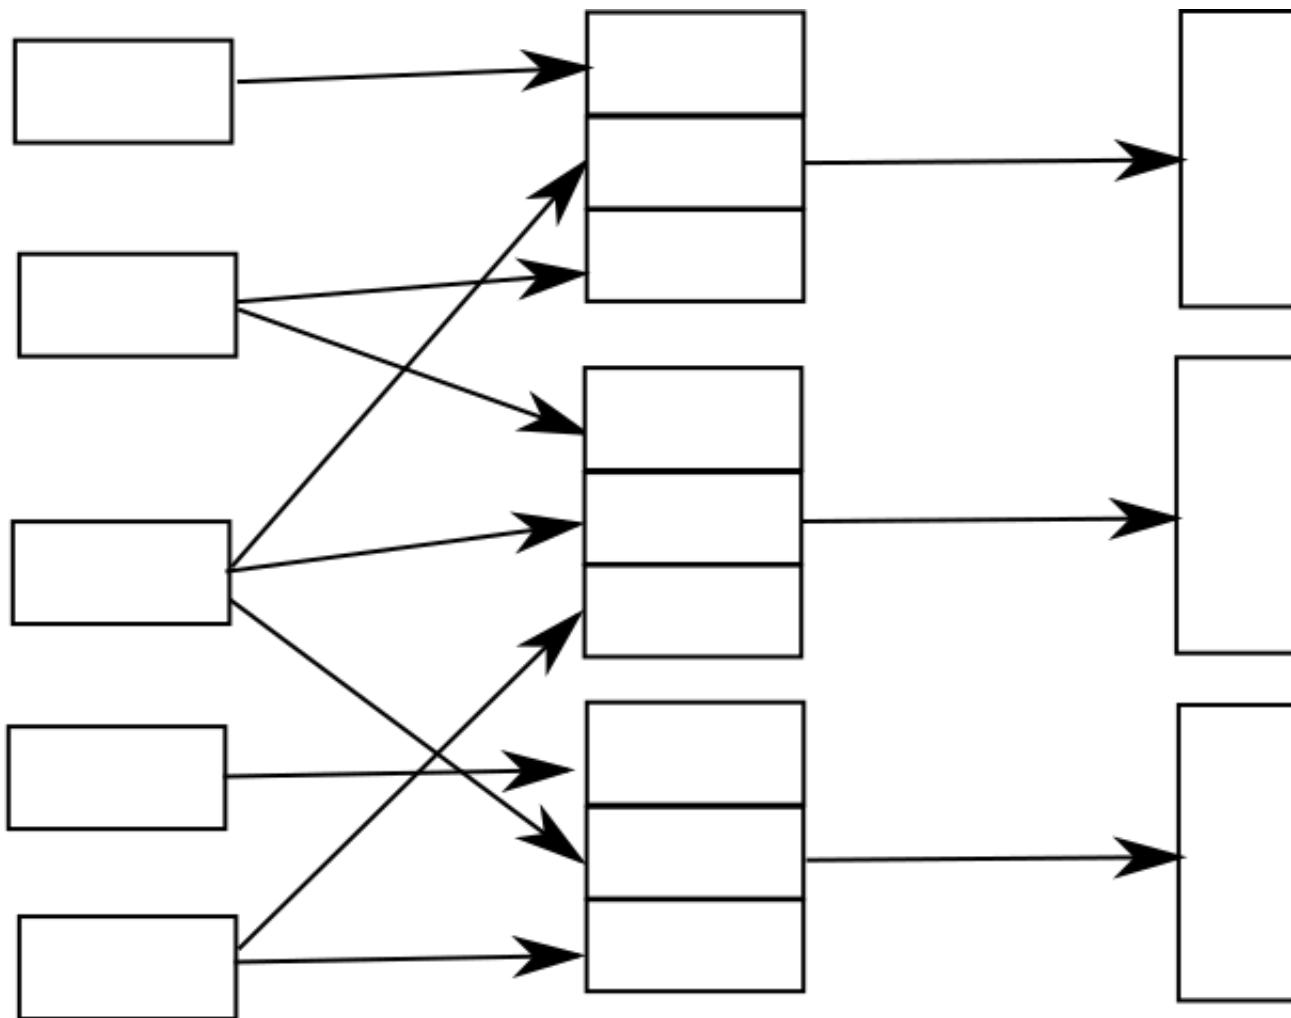
```
keys:
key:
    name: %main-features
    resource: from vehicles select features where model = $model
```

This retrieves a different set of features from the database depending on how the variable \$model is defined for the build. Launch the build with \$model = 'compact' and you get the feature set for the compact model. Launch the build with \$model = 'van' and you get the feature set for the van model.

Naturally, this is leaving out a whole lot of detail about how this query gets executed and how the results get structured into a document domain list structure. But these are implementation details.

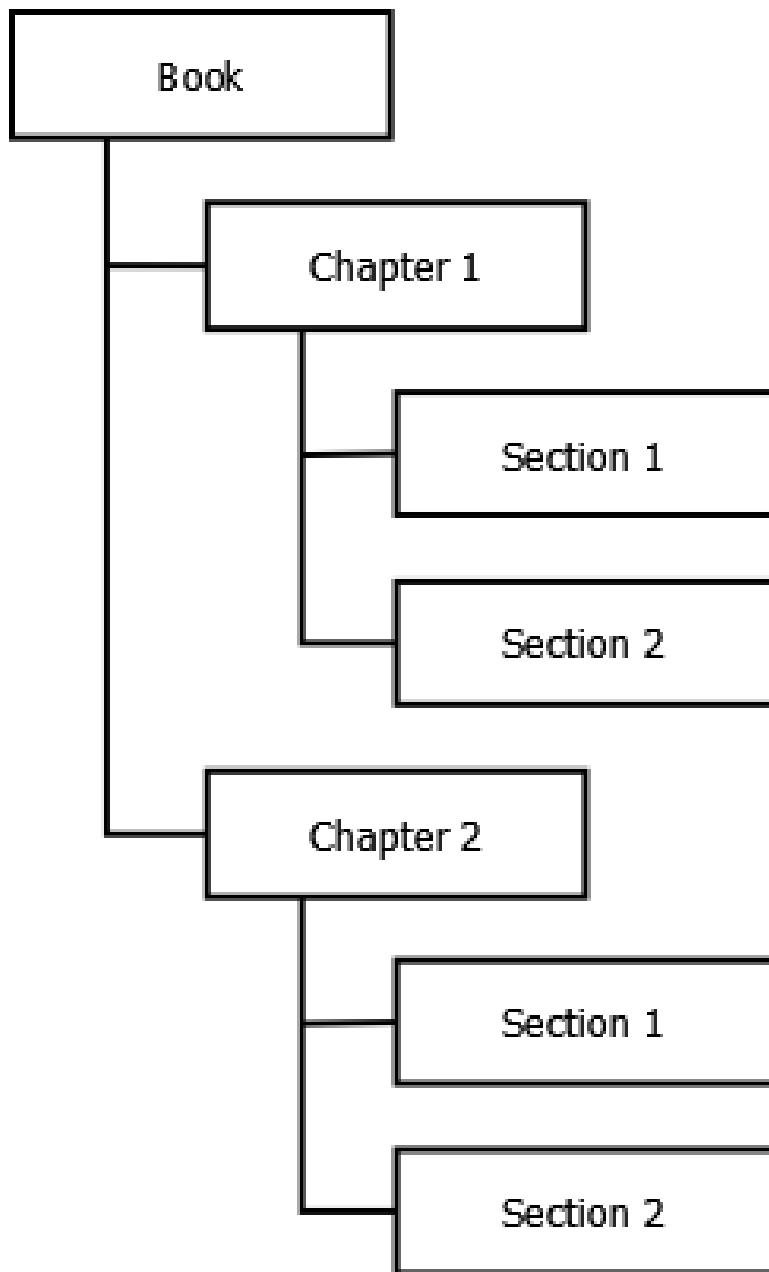
Assemble from pieces

In the assemble from pieces approach, there is no common vs. variable distinction and no single source document into which reused content is inserted or to which conditions are applied. Instead, there is a set of content units that are assembled to form a finished document.



For example, if you have a range of products with common features, you might assemble the documentation for those products using a common introduction with a piece representing each feature of each model.

This could be a flat list, or it could be a tree structure. For instance, you might assemble a chapter of a manual with a introductory piece and then several sections below it in the tree.



The assembly approach requires a structure to describe how the units are assembled. This structure is often called a map. (It is called a map in DITA, for instance.) Some applications may also refer to it as a table of contents.

```
map: Widget Wrangler Deluxe User Manual
  unit: units/ww/deluxe/intro
    unit: units/ww/shared/basic_features
    unit: units/ww/deluxe/deluxe_features
  unit: units/ww/shared/install/intro
    unit: units/ww/shared/requirements
```

```
unit: units/ww/deluxe/requirements  
unit: units/ww/shared/install  
unit: units/ww/deluxe/install_options
```

It is important to note here that a map does not always have to be written by hand by an author. In some cases the map may be created by an algorithm based on the metadata of the units themselves. Whether this is possible depends on what determines the desired order of units. If the assembled units are supposed to form a narrative flow, they may need to be ordered by hand as it is difficult to generate a narrative ordering from metadata. But if the units do not form a narrative flow, they can easily be ordered using metadata. For instance, if you are assembling a cookbook from recipe units you might order them by season or main ingredient without caring if the boiled egg recipe comes before or after the scrambled egg recipe.

Rather than using a map, you can allow the units themselves to pull in other units, which may in turn pull in other units. So the Widget Wrangler Deluxe install introduction unit might look like this:

```
unit: Installing the Widget Wrangler Deluxe
```

You should be very careful when installing the Widget Wrangler Deluxe. Follow these steps carefully:

```
>>>(unit units/ww/shared/requirements)  
>>>(unit units/ww/deluxe/requirements)  
>>>(unit units/ww/shared/install)  
>>>(unit units/ww/deluxe/install_options )
```

This avoids the need for a map, but the downside is that it can make the units less reusable. In the above example, for instance, you would need a separate introduction unit for the regular Widget Wrangler since the introduction file imports all the requirements and procedural units. By assembling units with the map, you can use a shared install intro, which increases the amount of reuse you can do.

When using the map approach, it is important to think about how your content will single source to paper-like media and hypertext-media. Some people will output the same map to both media. In hyper-text media that usually results in the map being turned into a table of contents, often displayed in a separate pane as in a help system. This may be fine if you are creating a help system, but it is not how Web content is usually displayed. Another approach to single sourcing content that uses the assemble from pieces approach is to create completely separate maps -- or even to use completely different assembly techniques that don't involve maps at all -- to produce paper-like and hypertext outputs. This can help you to work around some of the design limitations that we talked about in Chapter 12, *Single Sourcing*.

Combining multiple techniques

There is one problem with the idea of using a common install intro for both the regular and deluxe widget wrangler. The intro mentions the name of the product. To solve this problem without requiring two different units, we can use the variable into common or common with conditions reuse techniques within the intro unit. Here is an example using variable into common:

```
unit: Installing the >($product_name)
```

You should be very careful when installing the >(\$product_name). Follow these steps carefully:

```
>>>(unit unit/ww/shared/requirements)
>>>(unit unit/ww/deluxe/requirements)
>>>(unit unit/ww/shared/install)
>>>(unit unit/ww/deluxe/install_options)
```

There are a number of ways in which you can mix and match the basic reuse patterns to achieve an overall reuse strategies. Most systems designed to support reuse will allow you to do all these patterns and to combine them however you wish.

Content reuse is not a panacea

Content reuse can seem like an easy win, and in some cases it can return substantial benefits, but there are pitfalls to be aware of. You will need to plan carefully to make sure that you avoid the traps that await the unwary.

Quality traps

There are three main quality traps with content reuse.

- Making content too generic
- Losing the narrative flow
- Failure to address the audience appropriately

Many works on content reuse casually recommend making content more generic or more abstract as a means to making it more reusable, without saying anything about the potential downside. This is very dangerous and can do serious harm to the quality of your content. Statements that are specific and concrete are easier to understand and communicate better than statements that are generic and abstract. Replacing specific and

concrete statements with generic or abstract statements will reduce the effectiveness of your content significantly.

Unfortunately, human beings suffer from the curse of knowledge. The curse of knowledge is a cognitive bias that makes it very hard for people who understand an idea to appreciate the difficulties that idea presents to people who do not understand it. The curse of knowledge makes the generic or abstract statement of an idea appear equally communicative, and perhaps more succinct and precise, than the concrete and specific statement of it. This is a problem for writers at all times, pulling them away from the kind of specific and concrete statement that make ideas easier to comprehend. The desire to make content reusable reinforces this temptation.

Replacing the specific and concrete with the generic and abstract always reduces content quality and effectiveness. You may decide that the economic benefits of content reuse outweigh the economic costs of less effective content, but you should at least be aware that there are real economic consequences to this choice.

Another potential quality problem comes with the loss of narrative flow. Not all content has or needs a lengthy narrative flow, but if you start breaking your content into reusable units and putting them back together in different ways, the narrative flow can easily be lost. In some cases you can avoid this problem by making the topics you present to your audience more self contained using an Every Page is Page One information design. But don't assume that you have an effective Every Page is Page One design just because you have broken your content into reusable units. If that content was written in a way that assumed a narrative flow, it is not going to work when reused in a way that breaks that flow.

Finally, reuse can encourage us to come up with one way of telling our story that we present to all audiences. But not all audiences are alike, and the way we tell our story to one audience may not work for another audience. Good content tells a good story to a particular audience. Two different tellings of the same story do not constitute redundant content if they address different audiences.

Cost traps

It is easy to see content reuse as a big cost saving. Reusing content means you do not have to write the same content over and over again. It is easy to add up the cost of all that redundant writing and regard that number as pure cost savings from a content reuse strategy.

But all of the reuse techniques create multiple artifacts that need to be managed. This includes both content and processing code. You need mechanism to make sure that your content obeys the constraints required to make the pieces of content fit together reliably. You need a mechanism to make sure that way you have done reuse actually produces the documents you want. The cost of such management can be non-trivial and the consequences of the management breaking down can be significant.

Where costs can really mount, though, it when it comes to modifying the content when the subject matter changes. It is often not until the subject matter changes that you find out if the content we have treated as common is really common once the subject matter changes. If not, you may have a complex management task to sort out what is really common and what isn't. This can involve complex edits that then have to be tested and verified. If you get everything right, you can realize major savings when it comes time to modify your content, but if you get it wrong, it can multiply costs.

If your content collection and its web of reuse relationships is not audited and validated regularly, it can become chaotic over time and lose cohesion. This can make adding new content or changing existing content increasingly difficult and expensive.

Depending of the techniques you use, content reuse strategies can complicate the lives of authors, which may reduce the pool of authors you can use or reduce their productivity. As the size of the content set grows, it can take longer and longer to determine if reusable content exists and to find and reuse it. It is possible for this to end up costing more time than was saved by not rewriting content. (Reuse techniques that factor out the reuse from the author's work avoid this problem.)

Some content reuse techniques are easy to use in non-structured ways and early in a project it may seem like a non-structured approach to reuse speeds things up by allowing writers to reuse content wherever they find it. Over time, however, this approach can lead to a rat's nest of dependencies and relationships between bits of content that makes it hard to update or edit the content with any confidence.

Once the cohesion and discipline of a content set starts to break down, the decline tends to accelerate. As it becomes harder to find content to reuse, more duplication occurs, which further complicates the search for reusable content, creating a vicious circle. As links and other content relationships break down, people tend to form ad hoc links and relationships to get a job finished, further tangling the existing rat's nest. Under the gun, it is almost always easier to get the next document out by ignoring the structure and discipline of the content set structure, but the effects of this are corrosive. Without consistent discipline, even in the face of deadlines, a reuse system can fail over time.

All of these issues can be managed successfully with the right techniques and the right tools, but they all introduce costs as well, both up-front costs and ongoing costs. Those costs have to be reckoned up and subtracted from the projected cost savings before you can determine if a content reuse strategy is really going to save you money.

Chapter 14. Composition

The composition algorithm deals with the problem of composing a larger work out of smaller pieces. Many traditional writing tools produce files that are meant to encompass a whole work. If you take two Word files, for example, or two FrameMaker files and simply join them together, the result will not be a document that is the combination of the two files, it will be a corrupt file that will not open.

Both Word and FrameMaker have features that allow you to compose Word and FrameMaker documents out of smaller pieces. To a limited extent, they may even allow you to compose documents out of files other than their own. Some systems are specifically designed to allow you to compose document from many different source files. In many cases, though, some cleanup is required before the imported files can be used, particularly if they contain media domain formatting information. This cleanup requirement means that this is usually a one time import. You can't keep editing the original files and have the changes immediately reflected in the importing system.

Systems that require a high degree of composability have often turned to structured writing for solutions to the composability problem.

There are several parts to the composability problem, each of which structured writing helps address in different ways.

Fundamental composability

The first requirement of composability is that you must actually be able to combine the pieces. Most structured writing formats consist of a hierarchy of structures. Those structures tend to be self similar in form. For instance, all structures in an XML document are composed of XML elements. This means that you can take an XML document apart at any point in the structural hierarchy and insert, remove, or rearrange the structures at that level. To compose a larger structure out of smaller structures, you simply wrap new elements around them. Thus XML, and many other structured writing formats, provide a fundamental composability often lacking in other formats.

Structural composability

The second requirement of composability is that the result of combining markup structures must be a valid document (must conform to the appropriate constraints for that document). The simplest way to assure this is to plan all of your pieces to fit the constraints of the documents they will be inserted into. The most obvious way to do this is to make sure that all of your pieces come from the same tagging language. Thus DITA has good support for composability, but only for sources that are themselves DITA.

But belonging to the same tagging language is not enough. Tagging languages constrain where certain structures can occur and you must also make sure that the pieces go into a place where they are structurally allowed in that language. Just because all the pieces come from the same languages does not mean that every possible combination results in a valid document. This requires planning and careful management to make sure the combinations you create are valid.

However, it is not essential to composability that all the pieces you want to combine come from the same language as the document you are composing. You can also take content from different sources and with different structures, as long as you can transform their structures on input to match the structures of the destination document. This can be a very powerful technique. For instance, you can use it to compose documents from content in a database. (Indeed, all database reporting systems are exactly this: systems that compose documents in one format from tabular data in another format.) Structural composability depends on the semantic equivalence of structures, not common syntax.

For this approach to work, however, it is important that all of the sources you draw from have a high level of conformance to their own constraints. If you don't know, or cannot rely on, the structure of the pieces you are drawing in, you cannot reliably combine them with an algorithm. Since design for conformance is often more about being specific to the writer's task and knowledge, composability is often best served by creating file formats that do the best job of ensuring conformance from each type of contributor rather than in trying to get everyone to use the same format.

Stylistic composability

While structural composability is vital, it is not always sufficient. You could have pieces in a media domain language that are structurally composable but formatted differently. The resulting document would be valid and would publish successfully, but it would be a mess of competing styles and fonts.

For practical purposes, then, you should not try to create composable content in the media domain. You should at least use the document domain. The document domain separates content from formatting so you can compose a document in the document domain and then apply consistent formatting to the result.

This is true even if the pieces are in different document domain languages. All document domain languages essentially describe the same set of abstract document structures -- document are documents after all, they all have the same basic features which all document domain languages seek to represent. As long as you can recognize the same basic features in each of the source languages, you can compose a document from pieces in different document domain languages by converting to a common output language. (Embedded management-domain markup may spoil the party, however, since there is not the same level of semantic equivalence between management system. For this reason you

generally want to do composition after the resolution of management-domain structures, unless those structures are actually doing composition themselves.)

Narrative composability

Even if you can assemble pieces from different document domain sources and format them all with a single consistent look, that does not mean that the result will be a complete, correct, coherent narrative. To create a coherent narrative, you need to ensure narrative composability.

This is not necessarily a matter of making the document sound like it came from a single person. Many business documents are the result of several different writers, sometimes working together, sometimes inheriting and maintaining a document over time. Truly making such a document sound like it was all written by one person is a tall order and usually not necessary to achieve it business purpose.

What does matter is that the document be cohesive and coherent. The terminology it uses should be consistent from beginning to end. The end should flow logically from the middle and the middle from the beginning. There should not be obvious duplications of content nor omissions (obvious or not). This clearly requires a number of constraints on the content affecting both composition and style.

There are a couple of approaches to narrative composability. One is the information typing approach that you find in systems like DITA or Information Mapping. In this approach, content is broken down into certain broad types such as procedure, process, principle, concept, structure, and fact (Information Mapping) or task, concept, and reference (DITA -- though DITA allows you to define others through specialization.) The idea here is that if you keep reference information in a separate chunk from a concept, for instance, the chunks will compose more reliably, since there will not be duplicate concept information in a reference chunk.

The difficulty with this approach is that these abstract categories don't always make a lot of sense to writers when they are writing about concrete subjects, and different writers may interpret the chunk types or their boundaries differently, resulting in material that does not compose as well as you might hope.

Also, this approach, while it has been shown to improve the quality of writing in some cases, can also impose an artificial clunkiness and lack of flow on the content, leaving it feeling choppy or disjointed.

And, if one of your aims is to impose a specific rhetorical style or structure for particular kinds of content, any composition that you do needs to be subject to the same rhetorical constraints as if the piece had been written as a whole. This is often difficult to achieve, not least because it is often difficult for authors to write material that conforms to a rhetorical structure when they can't see the whole structure.

The other approach to literary composability is to move content to the subject domain. A subject domain structure for a particular subject does not have to be structured as a collection of abstract chunk types. The structure is highly specific to the subject matter at hand and is therefore much more concrete and less susceptible to varying interpretation by writers. Also, you can use the subject domain to factor out many of the style issue that might otherwise compromise composability. (This is just like factoring out formatting issues by moving from the media domain to the document domain.)

Of course, not all material fits into obvious strongly typed subject domain structures. Content that is more conceptual or theoretical in nature does not have a strong subject domain structure because it does not approach its subject matter in such a systematic or regular way. Then again, the ability to compose such content out of existing pieces is limited anyway. By its very nature such content requires a continuous flow of exposition that is very hard to assemble from pre-written chunks.

Chapter 15. Single source of truth

The third form of single sourcing is the single source of truth. This means that when a piece of information is needed anywhere in your organization, there is one and only one place where that information is stored and from which it should be accessed.

This is, of course, a case of content reuse, but whereas the content reuse algorithm deals with how reuse happens, the single source of truth algorithm is concerned with ensuring that only one copy of a given truth is created and maintained. The content reuse algorithm can work without any such assurance, and frequently does. The writer finds some content to reuse and reuses it, without asking if it represents the one and only source of the truth it expresses.

The task of finding some content that meets your needs and can be reused is very different from that ask of assuring that there is only one piece of content in existence that can meet a specific need. Finding content to reuse is a task for the writer of the content doing the reusing. Ensuring that there is only one source for that content is a task for the writer of the content being reused. They have to make sure when they create and store some piece of content that there really is no other piece of content serving the same purpose in the entire body of content.

The problem with the single source of truth algorithm is, as Pontius Pilate asked, “What is truth?”¹

For some types of content, this is an easy question to answer. What is the customer’s birthday? A person can only have one birthday, so there is no difficulty creating a clear policy that says that a customer’s birthday may only be recorded once across the organization and should be accessed from that single source whenever it is needed.²

The set of constraints that defines a piece of information as unique are not universal. The definition of what constitutes unique for different pieces of information is complex and specific to the subject matter at hand. Take a recipe for guacamole, for instance. Is guacamole a single dish for which there can only be one recipe? Clearly not. There are many different ways in which you can prepare guacamole, some differing only slightly from one another and some presenting welcome variations that different people might like to try. Clearly a recipe site would not want eight essentially identical guacamole recipes, but nor would they want to pick one variation to the exclusion of all others. So then the question becomes, how do you decide when a recipe text is an effective duplicate of another text and when the subject is a welcome variation on how to prepare guacamole. If you decide the variation is welcome, how do you differentiate it from other guacamole recipes in your collection?

¹John 18:38

²Stating the policy is straightforward; implementing and enforcing it may be more difficult, since it means every system or document that wants to include the customer’s birthday has to be capable of retrieving it dynamically from the central data store.

Clearly the answers to these questions are not universal. The way you decide these questions for recipes are not the same way you decide them for API reference topics, used car reviews, or conceptual discussions of ideas. Whatever constraint you decide upon, the business processes and systems that ensure that these constraints are followed are not universal, but specific to each function and organization.

Content, by its nature, deals with those subject that do not fit neatly into rows and columns, and thus cannot be formally normalized according to database rules. Where databases describe relationships formally by use of records and keys, content describes relationships informally in prose. In particular, content deals with complex, unique, and potentially ambiguous relationships that could not be reduced to rows and tables at any reasonable cost. The same fact may be mentioned in many different pieces of content for many different purposes. It may be elaborated on in one place, explained briefly in another, and merely mentioned in a third. What is the single source of this fact? Can it be effectively factored out and stored separately when it is described (legitimately) in different levels of detail in different documents?

Content is always designed for a particular audience, both to serve a particular need and to suit a particular background and level of knowledge. Everything we know about effective content tells us that we need to address different audiences and different tasks differently. Taking a piece of content designed for one audience and using it for all other audiences, or attempting to write generic content that takes no account of any audience's needs or tasks is certain to produce content that is significantly less effective. What then is the single source of truth for this information?

It comes down to this: there is a difference between a truth and an expression of a truth. You can have one source of a truth, but you often require different expressions of that truth for different purposes and different people.

In the database world, you can store the customer's birthday in a single field of a single database, but you will produce many different expressions of this truth for use by different functions in your organization and for presentation to this or other customers in different forms.

The way we achieve the storage of a single source of a truth is by abstracting it out of all of the expressions of that truth. Abstracting variants from invariants is, of course, what structured writing is about. But content is not a truth but an expression of a truth (or of several truths and their relationships).

A simple case of the single source of truth algorithm is the replacement of certain pieces of text by variables. In technical communication, it is a frequent practice to replace things like company names and product names by variables:

Thank you for buying >(\$company-name)'s >(\$product-name).
We hope you enjoy it for many years to come.

There are a couple of reasons for doing this. One is that company names and product names often have a precise formal variant that the marketing department wants everyone to use, and one or more informal variants that people actually use. There is constraint here: use the formal name for company and product not the common name. But because authors are more used to using the common name (like everyone else) they are likely to slip it in without thinking. And if they do remember that they are supposed to use the formal name, they may not remember it correctly and so use the wrong form. Using the variable instead ensures that the correct version of the name is used at all times.

The second reason is that products and companies are sometimes re-branded. The names change. If the names are written into the source in a hundred places, you will need to find those hundred instances and change them all. If a variable is used, you just need to change the definition of the variable and you don't have to touch the source content at all.

This is clearly a very tactical approach to creating a single source of truth. We have a single source of truth for the name of the company and the name of the product. We don't have a single source of truth for every name. We don't for example, have a variable for the name of every ingredient in a recipe:

```
ingredients:: ingredient, quantity, measure  
    >($carrots), >($5), >($lb)
```

We don't do this because there is no abstraction here. Carrots are carrots, 5 is 5, and pounds are pounds. Whereas, the company name is not company name, but "Acme Corporation".

When we talk about single source of truth in regard to content, therefore, we are not (or should not be) talking about designating a single piece of content as the sole truth on a particular subject. We are (or should be) talking about abstracting out key information from various pieces of content to store separately.

Single vs Findable

In many cases, the desire to create a single source of truth is to manage change in a content set. Since facts may be scattered across a content set in all kinds of content designed for all kinds of purposes for all kinds of audiences, it can be time consuming and expensive to find and change all of these instances when the subject matter changes. The biggest challenge, perhaps, is making sure that you have found all the instances in which the fact is mentioned, which can be particularly challenging when information is scattered across multiple formats and systems. We will look at managing change in more depth in Chapter 25, *Change management*.

Factoring out that information into a single source of truth should, in theory, make the whole update process much easier. Lots of organization use variables to insert product

names into documents, for instance, so that they can just update the value of the variable when the product name changes. However, there are a number of problems with this approach:

Historical usage

Sometimes there is historical usage in a document: a reference to past products or services that remain true after the contemporary changes are made. Not every instance of a product name, for instance, should change when the current product name changes. But it is often hard to make the distinction between historical and contemporary use at the time the variable is inserted into the content.

Validation

It is hard to make sure that the variable is always used by writers. There is generally no way to structurally remind or constrain the author to use the variable instead of the name. Automated checking systems can sometimes catch instances, but they can be confused by historical usage or by other facts that happen to be expressed with the same words. (Think of a company like Henry's Camera, for instance, who cannot guarantee that every use of "Henry's" is a reference to the company name.)

Anticipating change

Finally, there is the problem of which facts to factor out and which not to. Factoring out every fact would be prohibitively expensive and make authoring virtually impossible. You can only afford to factor out those facts that are likely to change. But none of us has a crystal ball, and some of the biggest changes we have to deal with often come from places we did not expect.

For all these reasons, factoring out facts into a single source of truth is a limited proposition. It can be highly useful if we apply it in the right way to the right facts, but it can be cumbersome and error prone if we try to take it too far. Thus we are certainly going to end up with all sorts of duplication of facts and names in our content. And from time to time those facts are going to change and we are going to have to find all the instances and make the appropriate changes.

If we structure our content in a way that makes search easier, we can potentially save ourselves a large amount of time and money. Making expressions of a truth more findable has many of the same benefits as creating a single source of truth and can be less expensive, easier to implement, and more resilient in the face of unanticipated changes.

This is all about the ability to recognize information in context, and the ability to recognize non-canonical references to facts (a person or thing referred to by a nick name, for instance, like calling John Wayne "The Duke" or IBM "Big Blue"). Having your

content in the subject domain make this much easier. Second best is having your content in the document or media domains but stored in a content management system with a cataloging scheme based in the subject domain.

Chapter 16. Publishing

All structured writing must eventually be published. Publishing structured content mean transforming it from the domain in which it was created (subject domain, document domain, or the abstract end of the media domain) to the most concrete end of the media domain spectrum: dots on paper or screen.

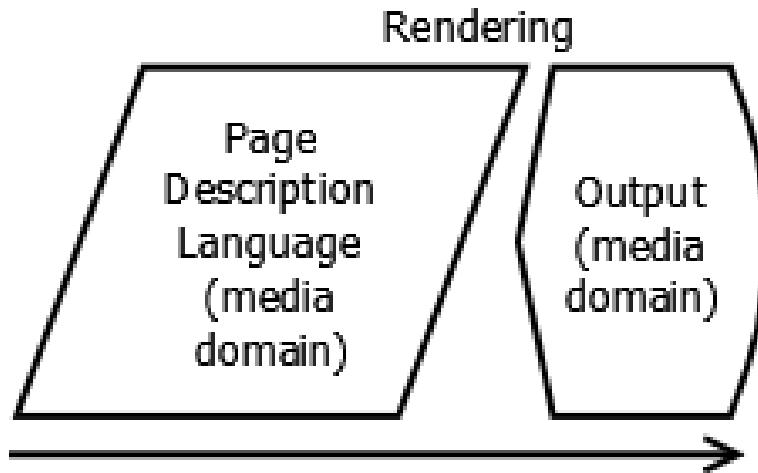
In almost all structured writing tools, this process is done in multiple steps. Using multiple steps makes it easier to write and maintain code and to reuse code for multiple purposes.

In this chapter, I am going to describe the publishing process as consisting of four basic algorithms which I have mentioned in passing in earlier chapter: the synthesis, presentation, formatting, and encoding algorithms. These four stages are formalized in the SPFE architecture, which I will talk about later, but I think they are a fair representation of what goes on in most publishing tool chains, even if those tool chains don't divide responsibilities exactly as I describe them here, or make such clear separation between them as I do here.

The Rendering Algorithm

There is actually a fifth algorithm in the publishing chain, which we can call the rendering algorithm. The rendering algorithm is the one responsible for actually placing the right dots on the right surface, be that paper, screen, or a printing plate. But this is a low-level device-specific algorithm and no one in the structured writing business is likely to be involved in writing rendering algorithms. The closest we ever get is the next step up, the encoding algorithm.

The rendering algorithm requires some form of input to tell it where to place the dots. In writing, this usually comes in the form of something called a page description language. Like it sounds, this is a language for describing what goes where on a page, but in higher level terms than describing where each dot of ink or pixel of light is placed. A page description language deals in things like lines, circles, gradients, margins, and fonts.



One example of a page description language is PostScript. Here is the PostScript code for drawing a circle:

```
100 100 50 0 360 arc closepath  
stroke
```

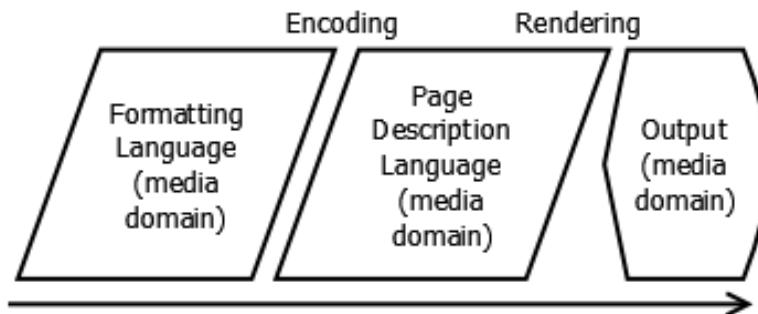
This code is basically moving a virtual pen over a virtual output device; the equivalent of a hand guiding a pen over paper. But it is a much lower-level operation than we need to worry about in structured writing.

The Encoding Algorithm

Since most writers are not going to write directly in a page description language, the page descriptions for your publication are almost certainly going to be created by an algorithm. I call this the encoding algorithm since it encodes your content in the page description language.

While it is possible that someone responsible for a highly specialized publishing tool chain may end up writing a specialized encoding algorithm, most encoding algorithms are going to be implemented by existing tools that translate formatting languages into page descriptions languages.

There are several formatting languages that are used in content processing. They are often called typesetting languages as well. XSL-FO (XSL - Formatting Objects) is one of the more commonly used in structured writing projects. TeX is another.



Here is an example of XSL-FO that we looked at in Chapter 12, *Single Sourcing*:

```
<fo:block space-after="4pt">
    <fo:wrapper font-size="14pt" font-weight="bold">
        Hard Boiled Eggs
    </fo:wrapper>
</fo:block>
```

You process XSL-FO using an XSL-FO processor such as Apache FOP. Thus the XSL-FO processor runs the encoding algorithm, producing a page description language such as PostScript or PDF as an output.

Writers are not likely to write in XSL-FO directly, though it is not entirely impossible to do so. In fact some boilerplate content such as front matter for a book does sometimes get written and recorded directly in XSL-FO. (I did this myself on one project.) But when you are constructing a publishing tool chain, you will need to select and integrate the appropriate encoding tools as part of your process.

The job of the encoding algorithm is to take a high level description of a page or a set of pages, their content and their formatting, and turn it into a page description language that lays out each page precisely. For publication on paper, or any other fixed-sized media, this involves a process called pagination: figuring out exactly what goes on each page, where each line breaks, and when lines should be bumped to the next page.

It is the pagination function, for instance, that figures out how to honor the keep-with-next formatting in an application like Word or FrameMaker. It also has to work out how to deal with complex figures such as tables: how to wrap text in each column, how to break a table across pages, and how to repeat the header rows when a table breaks to a new page. Finally, it has to figure out how to number each page and then fill in the right numbers for any references that include a particular page number.

This is all complex and exacting stuff and depending on your requirements you may have to pay some attention to make sure that you are using a formatting language that is capable of doing all this the way you want it done.

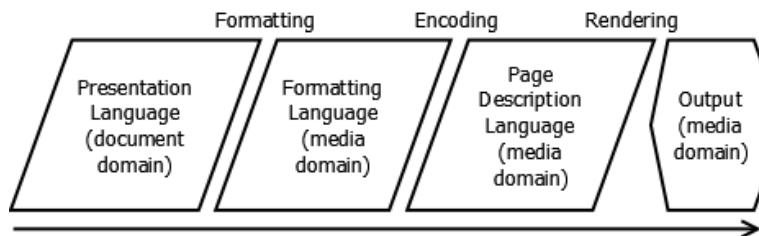
Also, you are going to have to think about just how automatic you want all of this to be. In a high-volume publication environment you want it to be fully automatic, but this could involve accepting some compromises. For example, it is not uncommon for writers and editors to make slight edits to the actual text of a document in order to make pagination work better. This is very easy to do when you are working in the media domain in an application like Word or FrameMaker. If you end up with the last two words of a chapter at the top of a page all by itself, for instance, it is usually possible to find a way to edit the final paragraph to reduce the word count just enough to pull the end of the chapter back to the preceding page. This sort of thing gets much harder to do when you are writing in the document domain or the subject domain, particularly if you are single sourcing content to more than one publication or reusing content in many places. An edit that fixes one pagination problem could cause another, and a major reason for writing in those domains is to take formatting concerns off the author's plate.

For Web browsers and similar dynamic media viewers, such as E-Book readers or help systems, the whole pagination process takes place dynamically when the content is loaded into the view port, and it can be redone on the fly if the reader resizes their browser or rotates their tablet. This means the publisher has very little opportunity to tweak the pagination process. They can guide it by providing rules such as keep-together instructions through things like CSS, but they obviously cannot hand tweak the text to make it fit better each time the view port is resized.

The formatting language for these kinds of media is typically Cascading Style Sheets (CSS).

The Formatting Algorithm

The job of the formatting algorithm is to generate the formatting language that drives the encoding and pagination process. The formatting algorithm produces the media domain representation of the content from content in the document domain.



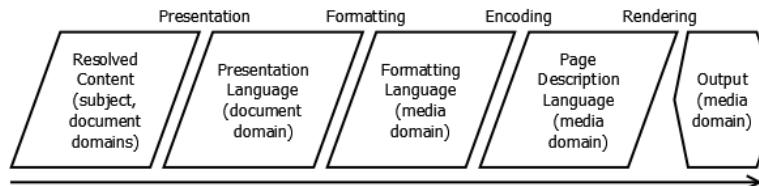
In the case of HTML output, the formatting algorithm generates HTML (with connections to the relevant CSS, JavaScript, and other formatting resources). This is the end of the publishing process for the Web, since the browser will perform the encoding algorithm internally. In the case of paper output, the formatting algorithm generates a formatting language such as TeX or XSL-FO which is then fed to the encoding algorithm as implemented by a TeX or XSL-FO processor.

In some cases, organizations use word processing or desktop publishing applications to tweak the formatting of the output by having the formatting algorithm generate the input format of those applications (typically RTF for Word and MIF for FrameMaker). This allows them to exercise manual control over pagination, but with an obvious loss in process efficiency. In particular, any tweaks made in these applications are not routed back to the source content, so they will have to be done again by hand the next time the content is published.

This algorithm is usually the province of publication designers. One of the most elementary structured writing algorithms is to separate content from formatting (Chapter 10, *Separating Content from Formatting*) which means removing formatting as one of the writer's concerns. Almost every structured writing implementation will involve writing formatting algorithms, however. Even if you use off-the-shelf languages like DITA or DocBook, you will have to write or modify formatting algorithms to get the specific formatting your want.

The Presentation Algorithm

The job of the presentation algorithm is to determine exactly how the content is going to be organized as a document. A pure document domain document is a representation of the presentation of the content. The job of the presentation algorithm produces a pure document domain version of the content. This may mean producing the entire presentation from purely subject-domain content, or simply handling the occasional subject-domain structure in a largely document domain file.



The organization of content involves several things:

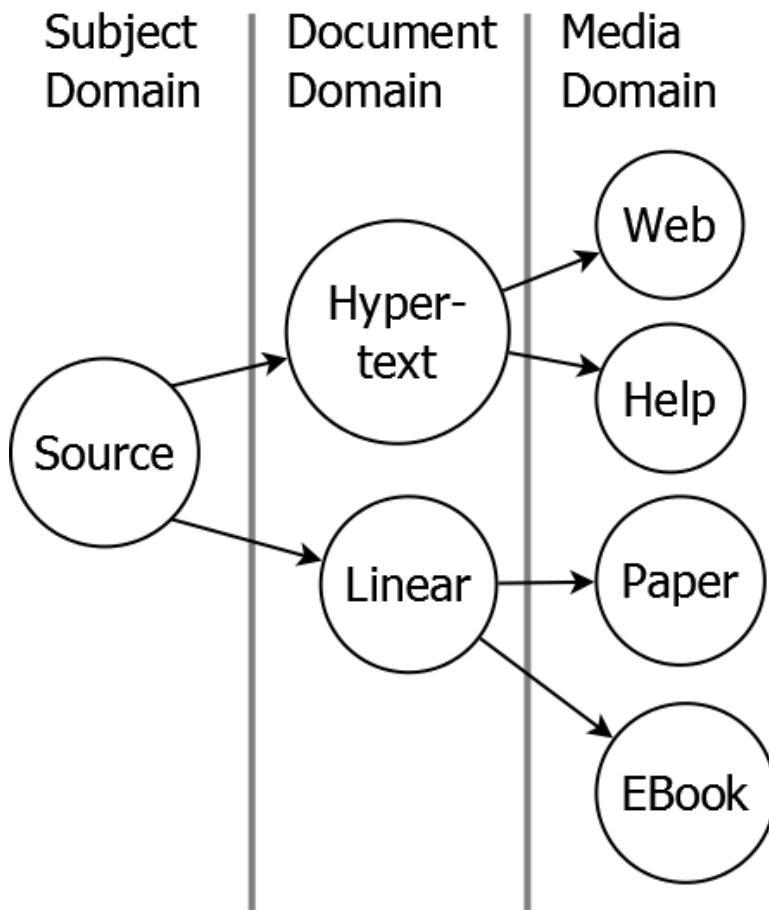
Ordering At some level, content forms a simple sequence in which one piece of information follows another. Authors writing in the document domain typically order content as they write, but if they are writing in the subject domain, they can choose how they order subject domain information in the document domain.

Grouping At a higher level, content is often organized into groups. This may be groups on a page or groups of pages. Grouping includes breaking content into sections or inserting subheads, inserting tables and graphics, and inserting information as labeled fields. Authors writing in the document domain typically create these groupings as they write, but if they are writing in the subject domain, you may have choices about how you group subject domain information in the document domain.

- | | |
|----------|--|
| Blocking | On a page, groups may be organized sequentially or laid out in some form of block pattern. Exactly how blocks are to be laid out on the displayed page is a media domain question, and something that may even be done dynamically. In order to enable the media domain to do this, however, the document domain must clearly delineate the types of blocks in a document in a way that the formatting algorithm can interpret and act on reliably. |
| Labeling | Any grouping of content requires labels to identify the groups. This includes things like titles and labels on data fields. Again, these are typically created by authors in the document domain, but are almost always factored out when authors write in the subject domain (most labels indicate the place of content in the subject domain, so inserting them is a necessary part of reversing the factoring out of labels that occurs when you move to the subject domain). |
| Relating | Ordering, grouping, blocking, and labeling cover organization on a two dimensional page or screen. But content can be organized in other dimensions by creating non-linear relationships between pieces of content. This includes hypertext links and cross references. |

Differential presentation algorithms

As we saw in Chapter 12, *Single Sourcing*, the organization of content is an area where the document domain cannot ignore the differences between different media. Although the fact that a relationship exists is a pure document domain issue, how that relationship is expressed, and even whether it is expressed or not, is affected by the media and its capabilities. Following links in online media is very cheap. Following references to other works in the paper world is expensive, so document design for paper tends to favor linear relationships where document design for the web favors hypertext relationships. This is an area, therefore, in which you should expect to implement differential single sourcing and use different presentation algorithms for different media.



Presentation sub-algorithms

The presentation algorithm may usefully broken down into several sub-algorithms, each dealing with a different aspect of the presentation process. How you subdivide your publishing algorithm is something you need to decide based on your particular business needs, but the following are some operations that it may well pay to treat as separate algorithms.

The linking algorithm

How content is linked or cross-referenced is a key part of how it is organized in different media, and a key part of differential single sourcing. We will look at the linking algorithm in detail in Chapter 17, *Linking*.

The navigation algorithm

Part of the presentation of a document or document set is creating the table of contents, index, and other navigation aids. Creating these is part of the presentation process.

Because these algorithms create new resources by extracting information from the rest of the content, it is often easier to run these algorithm in serial after the main presentation algorithm has run. This also makes it easier to change the way a TOC or index is generated without affecting your other algorithms.

The public metadata algorithm

Many formats today contain embedded metadata designed to be used by downstream processes to find or manage published content. One of the most prominent of these is HTML microformats which is used to identify content to other services on the web, including search engines. This is a case of subject domain information being included in the output. Just as subject domain metadata allows algorithms to process content in intelligent ways as part of the publishing process, subject domain metadata embedded in the published content allows downstream algorithms (such as search engines) to use the published content in more intelligent ways.

If content is written in document domain structures, public metadata is generally created by writers as annotations on document domain structures. But if content is created in the subject domain, the public metadata is usually based on the existing subject domain structures. In this case the public metadata algorithm may translate subject domain structures in the source to document domain structures with subject domain annotations in the output.

This does not necessarily mean that the public metadata you produce is a direct copy of subject domain metadata you use internally. Internally, subject domain structures and metadata are generally based on your internal terminology and structures that meet your internal needs. Public terminology and categories may differ from the ones that are optimal for your internal use. But because this is subject domain metadata, and thus rooted in the real world, there should be a semantic equivalence between your private metadata and the public metadata (the public usually being more generic and less precise than the private). The job of the public metadata algorithm, therefore, is not merely to insert the metadata but sometimes to translate it to the appropriate public categories and terminology.

The document structure normalization algorithm

In many cases, content written in the subject domain will also include many document domain structures. If those document domain structures match the structures in the document domain formats you are creating, all the presentation algorithm needs to do is to copy them to the document domain. In some cases, however, the document domain structures in the input content will not match those required in the output, in which case you will need to translate them to the desired output structures.

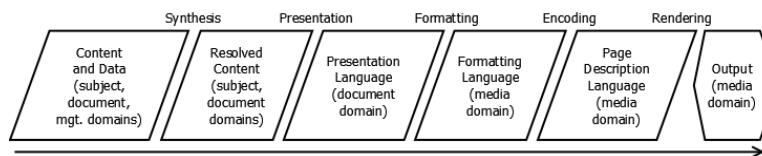
The Synthesis Algorithm

The job of the synthesis algorithm is to determine exactly what content will be part of a content set. It passes a complete set of content on to the presentation algorithm to be turned into one or more document presentations.

Among other things, the synthesis domain resolves all management domain structures in the content (unless some are to be retained for downstream post-publication algorithms to work with). This means that it processes all inclusions and evaluates all conditions. The result is document domain or subject domain content with all of the management structures removed and replaced with the appropriate document or subject domain structures and content.

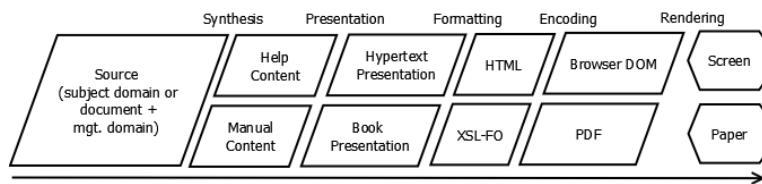
In the case of document domain content, processing the management domain structures yields a document domain structure which may then be a pass-through for the presentation algorithm (that is, the document domain markup may already express the desired presentation).

In the case of the subject domain content, processing management domain structures yields a definitive set of subject domain structures which can be passed to the presentation algorithm for processing to the document domain.



Differential synthesis

We noted above that you can use differential presentation to do differential single sourcing were two publications contain the same content but organized differently. If you want two publications in different media to have differences in their content, you can do this by doing differential synthesis and including different content in each publication.



Synthesis sub-algorithms

The synthesis algorithm can involve a number of sub-algorithms, depending on the kind of content you have and its original source.

The inclusion algorithm

If your content contains management domain include instructions, such as we identified in discussing the reuse algorithm these must be resolved and the indicated content found and included in the synthesis.

As we noted in the chapter on the reuse algorithm you can also include content based on subject domain structures, without any management domain include instructions. Such inclusions are purely algorithmic -- meaning that the author does not specify them in any way. It is the algorithm itself that determines if and when content will be included and where it will be included from. This too is the job of the inclusion algorithm.

The filtering algorithm

If your content contains management domain conditional structures (filtering) they must be resolved as part of the synthesis process. In most cases, you will be using the same set of management domain structures across your content set, so maintaining your filtering algorithm separately makes it easier to maintain.

Again note that you may be filtering on subject domain structures as well (or instead of) on explicit management domain filtering instructions. Such filtering is, again, purely algorithmic, meaning that the author has no input into it. The filtering algorithm is then wholly responsible for what gets filtered in and out and why.

Coordinating inclusion and filtering

It is important to determine the order in which inclusion and filtering are performed. The options are to filter first, to include first, or to include and filter iteratively.

Generally you want the filtering algorithm to run before other algorithms in the synthesis process so that other algorithms do not waste their time processing content that is going to be filtered out. However, if you run the filtering algorithm before you run the include algorithm, any filtering that needs to happen on the included content will not get executed.

Doing inclusion before you filter addresses this problem, but creates a new one. If you include before you filter, you may end up including content based on instructions or subject domain structures that are going to be filtered out, which could then leave you with a set of included content that was not supposed to be there.

The preferred option, therefore, is to run the two algorithms iteratively. Filter your initial source content. When you filter in an include instruction, immediately execute the include and run the filtering algorithm on the included content, processing any further include instructions as they are filtered in.

The rules based approach to content processing that we have looked at in this book makes this kind of iterative processing relatively easy. You simply include both the filtering and inclusion rules in one program file and make sure that you submit any included content for processing by the same rules the moment it is encountered.

```
match include
    process content at href
    continue
```

The extraction algorithm

In some cases you may wish to extract information from external sources to create content. This can include data created for other purposes, such as application code or data created and maintained as a canonical source of information, such as a database of features for different models of a car. We will look at the extraction and merge algorithms in Chapter 22, *Extract*.

The merge algorithm

In one sense, every structured writing is a merge algorithm. As we saw in Chapter 11, *Processing Structured Text*, most algorithms consist of factoring into content information or metadata that was factored out as we moved the point of recording from the media domain to the document domain and the subject domain. But it is also possible, and often useful, to combine information from different sources to create a new set of content. For instance, if you are creating a reference for any kind of computer code or data, there is almost always a way to extract a lot of the information you need from the existing definitions of that code or data. Extracting the information from the source is both more economical and more reliable than rewriting it from scratch. But you will almost certainly want to add additional written information as well. So, if you write the additional information in a structured way that identifies the piece of extracted information it corresponds to, you can use a merge algorithm to pull the two together for publication. More about the merge algorithm in Chapter 23, *Merge*.

The resolve algorithm

When we create authoring formats for content creation, we should do so with the principal goal in mind of making it as easy as possible for authors to create the content we require of them (see Chapter 9, *Authoring*). This means communicating with them in terms they understand. This may include various forms of expression that need to be clarified based on context before they can be synthesized with the rest of the content. This is the job of the resolve algorithm. Its output is essentially a set of content in which all names and identifications are in fully explicit form suitable for processing by the rest of the processing chain.

Content written in the subject domain is not always written in a fully realized form. When we create subject domain structures, we put as much emphasis as we can on ease of authoring and correctness of data collection. Both these aims are served by using labels and values that make intuitive sense to authors in their own domain. For example, a programmer writing about an API may mention and markup a method in that API using a simple annotation like this:

To write a Hello World program,
use the {hello}(method) method.

In your wider documentation set, there may be many APIs. To relate this content correctly in the larger scope you will need to know which API it belongs to. In other words, you need this markup:

To write a Hello World program,
use the {hello}(method (GreetingsAPI)) method.

The information in the nested parentheses is a namespace identifier. A namespace specifies the scope in which a name is unique. For instance, in most families, we only use first names to refer to each other because our first names are unique in our families. But when we go to school or to work, we often find that there are other people with the same first name. We then add people's last names when referring to them. The last names serves as a namespace identifier. We may need to add more specific namespace identifiers as we deal with larger groups of people (such as their addresses). But we don't use the full namespace identifiers of our family members at home. We add them as needed when we deal with the wider world.

In this case, the method name `hello` might occur in more than one API. The namespace specifies that this case of the name refers to the instance in the `GreetingsAPI`.

Deferred synthesis

For static presentation, all synthesis happens before the material is presented. But if you are presenting content on the web, you can defer parts of the synthesis algorithm to the browser, which can synthesize and present content by making calls to web services or other back-end data source, or by making a request to code running on the server to synthesize and present part of the page.

Combining algorithms

As we have seen, structured writing algorithms are usually implemented as sets of rules that operate on structures as they encounter them in the flow of the content. Since each

algorithm is implemented as a set of rules, it is possible to run two algorithms in parallel by adding the two sets of rules together to create a single combined set of rules that implements both algorithms at once.

Obviously, care must be taken to avoid clashes between the two sets of rules. If two set of rules act on the same structure, you have to do something to get the two rules that address that structure to work together. (Different tools may provide different ways of doing this.)

In other cases, though, one algorithm needs to work with the output of a previous algorithm, in which case, you need to run them in serial.

In most cases, the major algorithms (synthesis, presentation, formatting, encoding, and rendering) need to be run in serial, since they transform an entire content set from one domain to another (or from one part of a domain to another). In many cases the sub-algorithms of these major algorithms can be run in parallel by combining their rule sets since they operate on different content structures.

The consistency challenge

The biggest issue for every algorithm in the publishing chain is consistency. Each step in the publishing chain transforms content from one part of the content spectrum to another, generally in the direction of the media domain.

The more consistent the input content is, the easier it is for the next algorithm in the chain to apply a simple and reliable process to produce consistent output, which in turn makes the next algorithm in the chain simpler and more reliable.

Building in consistency at source is therefore highly desirable for the entire publishing algorithm. This is an interesting problem because good content by its nature tends to be diverse. Content is the way that we convey the information that is by its nature less consistent and more prone to exceptions. It is the stuff that does not fit easily into rows and columns.

One approach to this problem is to write all the content in a single document domain language such as DocBook. Since all the content is written in a single language it is theoretically completely consistent and therefore should be easy to process for the rest of the tool chain.

The problem with this is that any document domain language that is going to be useful for all the many kinds of documents and document designs that many different organizations may need is going to contain a lot of different structures, some of which will be very complex and most of which will have lots of optional parts. This means that there can be thousands of different permutations of DocBook structures. A single formatting algorithm that tried to cover all the possible permutations could be very large and complex and potentially very hard to maintain.

Another alternative is to have authors write in small simple subject domain structures that are specific to your business and your subject matter. You would then transform these to a document domain language using the presentation algorithm. This document domain language could still be DocBook, but now that you control the DocBook structures that are being created by the presentation algorithms, you don't have to deal with all the complexities and permutation that an author might create in DocBook, just the structures that you know your presentation algorithm creates.

These subject domain documents would have few structures and few options and therefore few permutations. This would mean that the presentation algorithms for each could be simple, robust, and reliable, as well as easy to write and to maintain. You would also be able to do differential single sourcing by writing different presentation algorithms for each media or audience.

The trade-off, of course, is that you have to create and maintain the various subject domain formats you would need and the presentation algorithms that go with them. It's a trade-off between several simple structures and algorithms and a few complex ones. This trade-off comes up time and time again in structured writing. We will see it again when we look at other algorithms.

Functional lucidity and a layered architecture

A layered publishing architecture can be very valuable in providing functional lucidity to a variety of authors. To build a publishing system that is capable of managing a wide variety of content, it is often necessary to create a lot of abstractions in your document markup. If you ask authors to write content directly in those abstract structures you may find that they struggle with the abstraction, and with the complexity and the range of options that go with a language that is designed to handle so many different source and publication types. Providing a set of separate authoring languages in the document or subject domain, can deliver functionally lucid authoring to different types of authors, greatly improving the authoring process.

At the same time, the simplicity of these special-purpose languages can allow you to be much clearer about what the names and labels mean in the context of that language. Providing and agreeing on precise meanings for terms is much easier the smaller the group of people you are dealing with and the smaller the range of subject matter you are addressing. We have relatively limited vocabularies and we reuse words and phrases between different domain of discourse all the time. Agreeing on what a term means across all domains of discourse is virtually impossible. Agreeing on what it means in a limited domain with a limited audience is much easier.

For purposes of rolling up content from many domains into one larger content set, you will definitely have situations in which the same names and labels are used to mean

different things. But this is not a problem as long as you know which domain each piece of content comes from. In other words, every tagging language defines a set of names and labels for content in the context of a particular namespace.

Placing names into namespaces does not magically resolve all disagreements about what to call things in a wider information space. But it does allow for an information architect to choose an definitive mapping of names from each namespace into the enterprise namespace. The results may still be disputable, but at least they will be consistent. And if they are disputed, and a different mapping is accepted, only the mapping has to change to put the new system into effect. As long as each pieces of content is tagged correctly according to the rules of its local namespace, it does not have to change just because the rules of the enterprise namespace change.

Achieving agreement (and, what is really more important, functional lucidity) within a local domain is easier in some domains than others. The media domain is highly concrete, so there really is not much room for disagreement there. Styles, though, are often given names from the document domain, as they are really a step into the document domain, and this can lead to disagreements because of the more abstract nature of the document domain.

The document domain is the most difficult place to achieve either agreement or functional lucidity. (By functional lucidity, I mean that the language seems easy to use and obvious when you are actually writing as opposed to when you are attempting to hammer out agreement in a committee room.) The document domain is inherently an abstract place, and there are always different ways to abstract from the concrete reality of web pages and books, particularly because of how different hypertext media domains are from paper media domains and the difficulty of truly abstracting beyond those differences. Thus you will frequently hear argument between the proponents of various document domain languages about the correctness and usability of their choices. For instance, partisans of Markdown may praise its functional lucidity (it is very easy to write in) which partisans of ReStructuredText may praise its greater abstraction and range of application.

The management domain is quite concrete, like the media domain, but more arbitrary. It consists either of commands or of management metadata, both of which are particular to a specific management system. Difficulties here are likely to be more about the management processes to be implemented rather than the correct naming of things.

Chapter 17. Linking

Few readers read content straight though. Unless the content is perfectly matched to their experience and their goals, there will be points where they need more information, points where they need less, or points where they decide that they need something else altogether. These are deflection points in the content, points in which the reader’s “next” may not be the thing that comes next in the linear order of the work. That may mean deflecting to other content or to a different way of finding information such as asking a friend or posting a question on a forum.

Deflections are a natural part of information seeking. They are part of what is called information foraging. The reader, in pursuit of their individual ends, will follow where the scent of information leads and that will not always be the next paragraph of the current text.

Writers know they can’t meet everyone’s needs perfectly every time so they use links and other devices to help readers deflect when they need to. If you have Model A, do this. If you have Model B do that. Supporting deflection helps readers achieve their goals and helps keep the reader in the writer’s own content, or content they approve of, rather than deflecting elsewhere. The Web and other electronic media have made the cost of deflecting much lower, meaning the readers deflect much more readily, and writers need to pay far more attention to deflection in the design of their content.

Deflection points can be handled in various ways. The writer may choose to do nothing, leaving it up to the reader to look up a word if they don’t understand it, for instance. They may use footnotes, cross references, sidebars, parenthetical material, or hypertext links to provide deflection choices. They may use tables or flowcharts to allow readers to choose different paths through content. They may even attempt to anticipate and forestall deflection by using information about the individual reader to dynamically reorder the content to suit the reader’s needs.

Deflection costs the reader more on paper than online. For paper, we may design content to minimize the need to deflect, or to keep deflections inside the local work. For the Web, we may organize content with deflection in mind, allowing different readers to choose their own course, rather than trying to optimize one course for all readers. This difference in the ease of deflection between paper and the Web and other hypertext media is one of the main reasons we want to practice differential single sourcing.

Deflection also enters into the discussion of content reuse. We reuse content in multiple documents so that readers don’t have to deflect from one document to another to find it.

If we reuse content in different media, we might want to have a different deflection strategies between paper and hypertext outputs. We may want to include the same chunk of content in multiple paper documents but link to a single copy of it when creating a hypertext. (Linking, in other words, is a kind of reuse: reuse by reference rather than copying.)

Thus we should not be thinking solely in terms of managing links in our content. We should be thinking about implementing the right deflection strategy in each of our outputs. To see how that works, we need to look at deflection and linking in each of the structured writing domains.

Deflection in the media domain

In the media domain, we simply record the various deflection devices as such: cross references, tables, links, etc. For example, in HTML a link simply specifies a page to load:

```
<p>In Rio Bravo,  
<a href="https://en.wikipedia.org/wiki/John_Wayne">the Duke</a>  
plays an ex-Union colonel out for revenge.</p>
```

The phrase “the Duke” is a deflection point. The reader may not know who “the Duke” is, or may want more information on him. The link supports the reader at the deflection point. The reader can either deflect by clicking the link or stay the course and read on.

But if the HTML page gets printed, the link is lost. The phrase “the Duke” is still a deflection point. The reader can still deflect, by doing a search for “the Duke”, perhaps, or asking a friend what it means. But the printed version lacks any support for that deflection.

If the content had been written for paper, the deflection point might be supported in a different way. For example, it might be supported by adding an explanation in parentheses:

In Rio Bravo, the Duke (John Wayne) plays an ex-Union colonel out for revenge.

Or it might be handled with a footnote:

In Rio Bravo, the Duke* plays
an ex-Union colonel out for revenge.

...

* "The Duke" is the nickname of the actor John Wayne.

Clearly this is a case in which we would like to do differential single sourcing and handle a deflection point differently in different media. To accomplish this, we need to move the content out of the media domain.

Deflection in the document domain

Moving to the document domain is about factoring out the formatting specific structures of the media domain. But a link is not really a piece of formatting, so conventional refactoring into abstract document structures is not going to apply. For this reason, people working in the document domain often enter hypertext links exactly the way they would in the media domain: by specifying a URL. Thus in DITA you might enter a link as:

```
<p>In Rio Bravo,  
<xref href="https://en.wikipedia.org/wiki/John_Wayne"  
      format="html">The Duke</xref>  
plays an ex-Union colonel out for revenge.</p>
```

The difference from HTML is slight here. The link element is called `xref` rather than `a`. But the meaning of `xref` is bit more general. The HTML `a` element is saying, create a hypertext link to this address. The DITA `xref` element is saying, create some sort of reference to this resource. (As we will see in a moment, it is capable of linking to things other than HTML pages, which is why it requires the `format` attribute to specify that in this case the target is an HTML page.) This generality gives us a little more leeway in processing. We can legitimately create print output from this markup that looks like this:

In Rio Bravo, the Duke (see: https://en.wikipedia.org/wiki/John_Wayne)
plays an ex-Union colonel out for revenge.

This is not the way we would handle the deflection point if we were designing for paper, but it is a definite improvement from a differential single sourcing point of view. At least the link is now visible to the reader.¹

Fundamentally, though, this is not a satisfactory differential single sourcing solution. Unless there were no alternative, you would not normally direct a reader of paper to the web for more information, nor vice versa. Linking to an already published file, such as an HTML page, means linking to the address where the published file resides. It commits us to a particular format for the link target. If instead of linking to the published address, we link to the address of the unpublished source file, or to an identifier for that file, we gain the freedom to link to any format of that content that we choose to publish.

In DITA, you can link to another DITA file (the default format, so we don't need the `format` attribute):

¹Technically we could do this from the HTML markup as well, but that would be cheating. The HTML markup is not really giving us permission to do this. It is telling us to create a hypertext link and nothing else. The problem with cheating is that you are basing your algorithm on constraints that are not being promised or enforced, and this can fail in ways you may not expect or catch. Some cheats are more reliable than others, but you probably don't want to get into the habit. Better to move your content creation to a format where cheating is not required to get the output you want.

```
<p>In Rio Bravo, <xref href="John_Wayne.dita">The Duke</xref>  
plays an ex-Union colonel out for revenge.</p>
```

We don't yet know if that DITA file will be published to paper or the Web, what the address of the published topic will be, or if that topic will stand alone or be assembled into a larger page or document for publication. This means that the publishing system is taking on responsibility for both ends of the link. It has to make sure that the target page is published in a way the source page can link to, and that the source page links to the right address. But transferring this responsibility to an algorithm gives us the leeway to publish this link as we see fit.

If we publish as a book on paper and the target resource ends up as part of a chapter in the same book, we can render the `xref` as a cross reference to the page that resource appears on. We could format that cross reference inline or as a footnote. These are all legitimate interpretations of the `xref`'s instruction to create a reference to a resource.

If we publish to a help system and the target resource ends up as a topic in the same help system, we could render the `xref` as a hypertext link to that topic.

This is a big step forward, but it still does not let us do this:

In Rio Bravo, the Duke (John Wayne) plays an ex-Union colonel out for revenge.

In other words, we can render the `xref` as a cross reference or a link or a footnote, but we can only handle the deflection point as a reference to the specified resource. We can't decide to link to a different resource or handle it by parenthetical clarification instead. To give ourselves the ability to link to different resources, we can turn to the management domain.

Deflection in the management domain

Linking to a source file rather than to an address gives us more latitude about how the link or cross reference is published, but we are still always linking to the same resource. If we are doing content reuse, this is a problem because you do not know if the same resource will be available everywhere we reuse our topic. We need to be able to link to different resources when our topic is used in different places.

To accommodate this, we can factor out the file name and replace it with an ID or a key. IDs and keys are a management domain structure that we looked at in Chapter 13, *Reuse*. They allow us to refer to resources indirectly. Using IDs lets us use an abstract identifier rather than a file name to identify a resource. Using keys lets us remap the resources we point to. This makes keys the more efficient way to address this problem. So instead of referring to a specific resource on John Wayne, we refer to the key `John_Wayne` using a reference to a key. In DITA this would look like this:

<p>In Rio Bravo, <xref keyref="John_Wayne">The Duke</xref> plays an ex-Union colonel out for revenge.</p>

Somewhere in the DITA map for each publication, the key John_Wayne points to a topic. Publications link the keyref to the resource pointed to by that key in each of their DITA maps. This allows you to link to different resources in each publication.

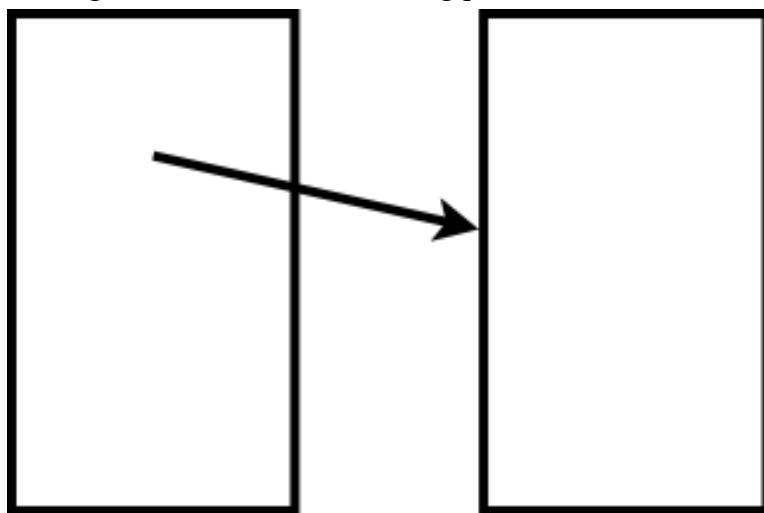
The problem with IDs and Keys

However, there is still a problem with linking based on IDs and keys. Keys will let you vary which resource a keyref resolves to, but what happens when there is no resource to which that key can reasonably be assigned?

The xref demands that a reference to a resource be created, but there is no resource to link to. You are going to have a broken link, and fixing it is not easy. You can't simply go in and remove the xref from the source for one publication, because it defeats the purpose of content reuse if you have to edit the content every time you reuse it. Removing the key reference would fix your broken link in one publication, but that would result in the link being removed from all the publications, even where the resource does exist and the link ought to be created.

Relationship tables

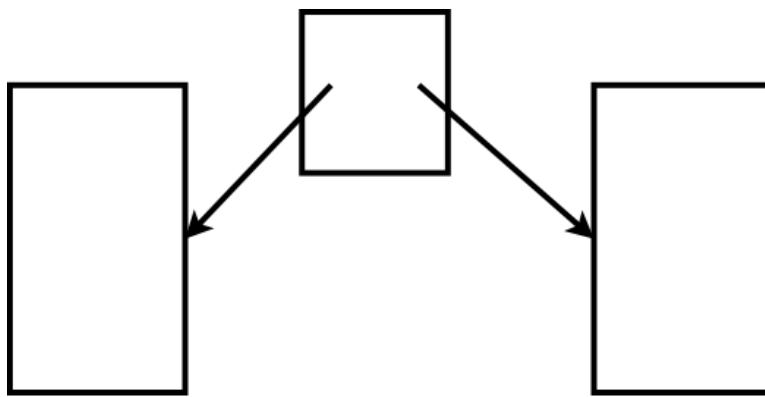
One approach to the link-only-when-resource-available problem is to use a relationship table. In a conventional linking approach, the source page contains an embedded link structure pointing to the target page. The source knows it is pointing to the target, but the target does not know it is being pointed to.



The idea that the target resource does not know it is being pointed to is important because it means it does not have to do anything in order for other resources to point to it. The

fact that only the source and not the target has to know about the link is fundamental to the rapid growth of the Web. If the target resource had to participate in the link process, every link would require negotiation between authors and the agreement of the author of the target resource to edit it to accept the link. It would be impossible for the Web to grow explosively and organically as it has under these conditions.

A relationship table takes this one step further. When you create a link using a relationship table, you factor the link out of the source document and place it in a separate table. The relationship table says resource A links to resource B, but neither resource A nor resource B knows anything about it. (Think of it like being introduced to a stranger by a third party because you share a common interest. I collect china ducklings. You make china ducklings. We don't know each other, but our mutual friend Dave introduces us. You and I are the source and destination resource; Dave is the relationship table.)



Notice that Dave has three choices about how he does the introduction. He can tell me about you as a seller of ducklings, or you about me as a buyer of ducklings, or he can introduce us both to each other. In the same way a relationship table can describe a link A to B, B to A, or both ways.

Once the links are factored out of a piece of content, you can reuse it anywhere you like. If there is a suitable resource available to link to, you enter it in a relationship table for that build and have the presentation algorithm create the link at build time. If no suitable resource is available for a different publication, no entry is made in the relationship table for that publication, and the presentation algorithm does not create a link.

The problem with link tables

The problem with link tables is that while they address the problem of link management in reused content and allow you to link to different targets in different publications, they separate the link from the deflection point it supports. The link that was created to serve the deflection point in the content has been factored out of the content into a separate file. The deflection point itself has not been marked up, so there is no way to put the link back where it belongs. Links generated by relationship tables end up in a block, usually at the end of the page.

The fundamental problem here is that links exist as a way to service deflection points in content. If we lose sight of this and think only in terms of managing links, it makes sense to pull links out of the content and manage them separately. But this is to ignore the reason the links exist in the first place, which is to provide the reader with options at a deflection point. It is a classic case of a focus on content management in structured writing leading to a neglect of content quality, and thus damaging content quality in the name of process efficiency.

Since the end of a page is a deflection point in a hypertext system, there is a legitimate case for creating and managing page-level links. But if the page is well designed to fulfill a discrete purpose for the reader, the end of the page is actually the point at which the writer knows least about what the reader might want to read next. The relationship table approach is fine as a way to manage deflection points that occur at the end of a page (or at the beginning). But they are not the way to support the full array of foreseeable reader deflection points which occur in the body of the page.

The other problem with the relationship table approach is that it is time consuming. You have to rewrite the links for each content set, and because the deflection points are not recorded in the content source, you have to figure out the appropriate links each time. This goes against the spirit of recording something once and using it many times. A mechanism intended to help you reuse content ends up forcing you to redo the work of linking for each publication you create.

Conditional linking

Before we leave the management domain, it is worth mentioning a management domain approach that we could use to address our differential single sourcing problem and get the appropriate deflection strategy for online and paper publishing. We could use conditional structures to define both options in the source file. With a little specialization to support `media` as a conditional attribute, you could do this in DITA:

```
<p>In Rio Bravo,  
<ph media="online"><xref keyref="John_Wayne">The  
Duke</xref></ph><ph media="paper">The  
Duke (John Wayne)</ph> plays an ex-Union  
colonel out for revenge.</p>
```

In DITA, the `ph` element is used to delineate an arbitrary phrase in the content that you want to apply management domain attributes to. Here we define two different versions of the phrase “the Duke”, each with different forms of deflection support (parenthetical expression for paper, link for online), and each with a corresponding media condition. The synthesis algorithm would then choose the appropriate version of the phrase for each publication based on the conditions set for the build.

There are some pretty obvious problems with this approach. It is twice the work for authors to create every link, and it doubles the maintenance cost of the content as well. It also flies in the face of the idea of creating formatting-independent content.

Unfortunately, in a general purpose document domain tagging language with management domain support, it is pretty much impossible to prevent writers from doing things like this in order to achieve the effects they want. And in practice writers do end up using conditional markup like this for all kinds of differential single sourcing and reuse problems that are not easy to solve in the document and management domains. In some cases this can lead to tangles of conditions that are hard to maintain and debug.

For an alternate approach to this problem, and the others we have discussed, we can turn to the subject domain.

Deflection in the subject domain

There are a number of problems with the management domain approach to links.

- Like all management domain structures, they are artificial. They don't correspond to things in the author's everyday world, which makes them harder to learn and use.
- You can't link to a key or an ID that does not exist. This means that as you are developing a set of content, the first pages you write have very few other pages to link to. Authors cannot enter links to content that has not been written yet.
- In reuse scenarios, the use of IDs and keys does not solve the whole problem because it cannot guarantee that the resource that an ID or key refers to will be present in the final publication. You can use relationship tables to address this problem, but they create additional complexity for authors and don't correctly handle deflection points in the middle of the content.
- Unless you resort to ugly conditional structures, you can't use media-appropriate deflection mechanism for differential single sourcing.

As we have seen before (Chapter 13, *Reuse*), we can often remove the need for management domain structures by moving content to the subject domain. The same is true with deflection points. In the document domain we handled a deflection point by specifying a resource to link to, specifying both that the deflection mechanism would be a link and that the link target would be a particular page. In the management domain we used keys to factor out the target resource but not the deflection mechanism (it was still an `xref`). In the subject domain, we can factor out the target resource as well. We do this by marking up the subject of the deflection point:

```
<p>In <movie>Rio Bravo</movie>,
```

```
<actor name="John Wayne">the Duke</actor>  
plays an ex-Union colonel out for revenge.</p>
```

This markup clarifies that the phrase “the Duke” refers to the actor named John Wayne. These are respectively the type of the subject (actor) and its value (John Wayne).

The subject domain markup is not link markup. Unlike the document domain markup, it does not insist that a reference should be created nor does it specify any resource to link to. This markup is a subject annotation. It clarifies that the phrase “the Duke” refers to the actor named John Wayne (and not the Duke of Wellington or the Duke of Earl) and that the phrase “Rio Bravo” refers to the movie (and not to the city in Texas or the nature reserve in Belize²).

Given this markup, we can easily create the paper-style deflection mechanisms we have been looking for. We simply have the presentation algorithm take the value of the name attribute and output it between parentheses:

```
<p>In Rio Bravo, The Duke (John Wayne) plays an  
ex-Union colonel out for revenge.</p>
```

It also allows us to create a link if we want to. We’ll look at how in a moment. But first we should look at the implications of subject annotation more deeply.

Subject annotation markup says, “this is an important subject that we care about in this context”. How is this an appropriate way to handle a deflection point? The intersection of subjects -- where a topic on one subject makes a significant mention or another significant subject -- are likely deflection points. When we create a link in the media domain or the document domain it is because our text is mentioning a significant related subject that the reader may want to know more about. The difference is that in the document domain we handle the mention of an important subject by creating markup that specifies a resource on that subject and in the subject domain we create markup that specifies the subject itself. This leaves us with more options about how to handle the deflection point, and that is what we have been looking for.

Marking it up a phrase as a significant subject does not oblige the publishing algorithm to create a link. If you decide to have the publishing algorithm create a link on the Web and a cross reference on paper, nothing in the markup obliges you to use any particular formatting or target any particular resource. There is no question of cheating here if you decide to create one kind of deflection device or another, or not to create one at all.

In all our previous examples, mentions of “Rio Bravo” were not marked up, even though it is clearly an important subject and a potential deflection point. This reflects the author’s decision not to create a link to support this deflection point. But what if we want to make

²[https://en.wikipedia.org/wiki/R%C3%A9gion_Bravo_\(disambiguation\)](https://en.wikipedia.org/wiki/R%C3%A9gion_Bravo_(disambiguation))

a different choice later? By marking up “Rio Bravo” as a significant subject, we keep our options open. Now we tell the presentation algorithm to create links on the names of movies if we want to, or not if we don’t want to.

But there are additional reasons to annotate Rio Bravo as a significant subject, because that annotation can be used for other purposes as well.

1. The subject annotation says that “Rio Bravo” is the title of a movie. In the media domain, the titles of movies are commonly printed in italics. We can use the subject domain `movie` tags to generate media domain italic styling.
2. We could use this subject annotation to generate document domain index markers so that we can automatically build an index of all mentions of movies in a work.

Subject annotation thus serve multiple purposes, and correspondingly reduces the amount of markup that is required to support all these different publishing functions. This is a common feature of subject domain markup. None of it is directly tied to specific document domain or media domain structures which will be required to publish the content. Each piece of subject-domain markup may be used to generate multiple document domain and media domain structures. For example, we could generate the following document domain markup from from the subject domain markup above (the example is in DocBook):

```
<para>
    In
    <indexterm>
        <primary>Rio Bravo</primary>
        <secondary>Movies</secondary>
    </indexterm>
    <citetitle pubwork="movie">Rio Bravo</citetitle>,
    <indexterm>
        <primary>John Wayne</primary>
        <secondary>Actors</secondary>
    </indexterm>
    <ulink url="https://en.wikipedia.org/wiki/John_Wayne">The Duke</ulink>
    plays an ex-Union colonel out for revenge.
</para>
```

This sample contains index markers, formatting of movie titles, and links on actor’s names, all generated based on the subject annotations in the source text. It should be clear how much less work it is for an author to create the subject domain version of this content than the DocBook version. Yet all the same publishing ability is maintained in both version.

Generating links from subject annotations has a number of other advantages:

- In a reuse scenario, you never have to worry about broken links or creating relationship tables. You generate whatever links are appropriate to whatever topics are available in the presentation algorithm.
- In a differential single sourcing scenario, you are never tied to one deflection mechanism. You can generate any mechanism you like in whatever media you like.
- You don't have to worry about maintaining the links in your content because you source content does not contain any links. The subject annotations in your content are objective statements about your subject matter, so they don't change. All the links in the published content are generated by the presentation algorithm, so no management is required.
- There is no issue with wanting to link to content that has not been written yet. The subject annotation refers to the subject matter, not a resource. Links to content that is written later will appear once that content becomes available to link to.
- It is much easier for authors to write because they do not have to find content to link to or manage complex link tables or keys. They just create subject annotations when the text mentions a significant subject. This requires no knowledge of the publishing or content management system. It does not even require knowledge of any other resources in the content set. It only requires knowledge of the subject matter, which the author already has.

Finding resources to link to

Of course, the question remains, what resources do we link to, since they are not specified in the text? If we choose to translate subject annotations into links, we need a way to find resource to link to. We do this by looking up resources based on the subject information (type and value) captured by the subject annotation. For this we need content that is indexed using those types and values (or their semantic equivalents). So naturally this means that we need to index our content. If you have a page on John Wayne, you can index it like this:

```
topic:  
    title: Biography of John Wayne  
    index:  
        type: actor  
        value: John Wayne  
body:  
  
John Wayne was an American actor known for westerns.
```

Now the linking algorithm looks like this:

```
match actor
    $target = find href of topic where type = actor and name = @name
    create xref
        attribute href = $target
    continue
```

However, content stored in the subject domain may already be indexed effectively enough by its inherent subject domain structures. Suppose your content collection includes thai subject-domain actor bio:

```
actor:
  name: John Wayne
  bio:
    John Wayne was an American actor known for westerns.
  filmography:
    film: Rio Bravo
    film: The Shootist
```

Here the topic type is actor, and the name field specifies the name of the actor in question. This is all the information we need to identify this topic as a source of information on the actor John Wayne.

Only very minor changes to the linking algorithm are required to use this:

```
match actor
    $target = find href of actor topic where name = @name
    create xref
        attribute href = $target
    continue
```

There is a lot more to how this mechanism works in practice, including what you do about imperfect matches and what happens when the query returns multiple resources. But that takes us into the specifics of individual systems and that is more detail than we need for present purposes.

Indexing of topics may also be done by a content management system, in which case the linking algorithm would query to CMS to find topics to link to.

A useful feature of this approach is that you can have the publishing algorithm fall back to creating a link to an external resource if an internal one is not available. If a search of the index of your own content fails, you can search indexes of external content. You can build such an index yourself, but some external sites may also provide indexes, APIs, or search facilities that you can use to locate appropriate pages to link to.

Deferred Deflection

Readers don't always deflect the moment they reach a deflection point. In some cases, they choose to set the alternate material aside for later reading. This is particularly easy to do on the Web, where you can simply open pages in new browser tabs for reading later.

The idea of the deferred deflection can also occur in document design. A document design that gathers a set of links together at the end of a document, rather than including them inline, is recommending deferred deflection to the reader. It attempts to keep the reader following the writer's default course to the end of the document before they go off to other things. The relationship table approach to link management that we mentioned earlier can only produce deferred links.

The merits of deferred links are debatable. Some argue that inline links are a distraction, that they actually encourage deflection. But the lack of links does not stop the reader from deflecting if they want to, and if they do deflect, the lack of a link means they may leave your content set and land on competitor's content or content that is of poor quality or that contradicts what you have been saying. The fact that the debate exists, however, suggests that we may want to factor this design choice out of our source content so that we can choose between inline and deferred links later.

To leave open the option of deferring or not deferring links, we have to records links at the deflection points they belong to. We can choose to defer them at publishing time if we wish, but if we defer at writing time, we can't put the links back inline at publishing time because we don't know where they belong.

But for this strategy to work, we need to be able to tell the difference between links that can be deferred and those that cannot. An simple example of a link that cannot be deferred is one that says "For more information, click here." Obviously this link has to remain on the words "click here".

But there is a more subtle issue as well. For a link to be deferred on publishing, it must be possible to contextualize the link in the deferred location. In other words, when the deflection point occurs inline in a paragraph the reader should be able to infer where the link will lead from the paragraph and from the text the link is applied to. But lifting the same link text out of the paragraph and putting is somewhere else is not guaranteed to provide the same context.

For example, a link marked up like this is hard to defer algorithmically:

```
<p>In Rio Bravo,  
<xref href="https://en.wikipedia.org/wiki/John_Wayne">The Duke</xref>  
plays an ex-Union colonel out for revenge.</p>
```

We could generate a list of links and insert it later in the document. It might look like this:

```
<p>For more information, see:</p>

<ul>
    <li><a href="https://en.wikipedia.org/wiki/John_Wayne">The Duke</a>
        ...
</ul>
```

But will it be clear out of the context of the original text what the words “the Duke” refer to? (The answer here is maybe, but it is not hard to image cases where it would be a definite no.)

On the other hand, if the deflection point is marked up in the subject domain like this:

```
<p>In <movie>Rio Bravo</movie>,
<actor name="John Wayne">The Duke</actor> plays
an ex-Union colonel out for revenge.</p>
```

Then, given that we know what the subject of the deflection point is, we could use it to create a list of links that are categorized by type and use the real names of actors even when the original text use a nickname:

```
<p>For more information, see:</p>

<ul>
    <li>Actors:
        <ul>
            <li><a href="https://en.wikipedia.org/wiki/John_Wayne">John Wayne</a>
                ...
        </ul>
    </li>
    <li>Movies:
        <ul>
            <li><a href="https://en.wikipedia.org/wiki/Rio_Bravo_(film)">Rio Bravo</a>
                ...
        </ul>
    </li>
</ul>
```

In short, algorithmically deferring document domain links is always tricky, but we can comfortably defer linking of subject annotations if we want to.

Different domain, different algorithm

What the linking algorithm illustrates perhaps better than any other is that the movement from one domain to another changes the structured writing algorithms in fundamental ways. While the algorithm has the same end in each domain, the way it achieves that end can be significantly different.

One of the points I have tried to make about structured writing algorithms is that they always start with the content structures. How you design the content structures -- the way the author records the content -- determines everything you can do with the content. You create content structures to support algorithms. You create algorithms to improve content quality or streamline content management and publishing. You then choose the content structures that support the algorithm you want to implement.

In the document domain, the data structures tend to have a one to one correspondence with their algorithms. As system designers determine they need a particular algorithm, they create structures to support that algorithm. Thus document domain languages that require support for linking, reuse, indexing, and single sourcing have separate data structures for linking, for reuse, for indexing, and for single sourcing. (Some of these may be management domain structures, of course.)

In the subject domain, though, the data structures reflect the subject matter. If you go looking for a one to one correspondence between a structure and the algorithm it supports, you won't find it. Thus you will not find link markup or reuse markup or index markup or single sourcing markup in the subject domain. You will find markup that clarifies and delineates the subject matter of the content it contains. Any algorithm we want to apply has to interpret that subject domain annotation and use it as the basis for creating whatever kind of document or media domain structure you want for publishing.

System designers do still have to think about what algorithms they want to apply, but that is to make sure that the aspects of the subject matter needed to drive the algorithms are captured. Since every subject structure can potentially drive many publishing algorithms, however, you will often find your subject domain content already supports any new algorithms you want to apply. This helps future proof your content.

Moving from the document domain to the subject domain is not a matter of asking what the subject domain equivalent of a document domain structure is, therefore, but a matter of asking what information in the subject domain drives the creation of document domain structures. Subject domain content can look very different from its document domain counterpart and will often be starkly simpler and easier to understand.

Chapter 18. Conformance

Structured writing is about constraints. Texts that meet appropriate constraints will be of higher quality and greater consistency and will be easier to manage. Texts that record the constraints they conform to can be reliably processed with algorithms. The key to any structured writing systems, therefore, is the ability to assure conformance with the desired constraints.

Constraints have always been part of the authoring process. Style guides and grammatical reference works express constraints that content is expected to follow. Editorial guidelines tell writers what kind of content a publisher is looking for, at what length, and in what format. These are constraints. If a publisher says that manuscripts must be delivered in DocBook or Word format, that is a constraint. When the government says that you must submit your online tax return in a particular file format, that is a constraint.

Some of these constraints are merely statements of requirements. Authors are not given any assistance in following them nor is there any verification mechanism to tell them if they have followed them or not (other than perhaps an email from an irate editor). Others are highly mechanical. Good tax preparation software will guide you all the way in filling out your tax forms and will run all kinds of checks to make sure that you did it correctly. It will also factor out many of the complexities of the tax code and ask you for information in a way you can understand.

This higher level of conformance checking and support helps make the process easier and the results more reliable. Structured writing is really all about improving conformance to enable automation and improve quality. The conformance algorithm is thus the lynchpin of structured writing. Without it, none of the other algorithms will work reliably.

How many constraints you need to place on your content depends on your quality and process goals. The larger your content set becomes, the more critical your quality constraints, the more frequent and dynamic your outputs, and the more of your processes rely on algorithms, the more constraints you need and the more pressing the issue of conformance becomes. Content reuse, for example, relies on conformance to the constraints for writing content that fits when reused, and on conformance to the constraints on what can be reused where. If you want to do any kind of real-time publishing of content, meaning there is no time to do quality assurance on the output of the algorithm, then reliable content is key, and conformance is how you ensure that content is reliable.

One of the ways in which a structured writing project can get into trouble is by introducing new constraints to meet management or publishing automation goals without considering how conformance to those constraints will be achieved. In some cases, this results in a highly imperative approach to conformance, in which writers are trained to implement the constraints, but where the structures they are creating provide no guidance

or validation of those constraints. The system constraints, in other words, are not reflected in the content structures. If you are creating complex structures and also creating complex constraints that are not reflected or implemented in those structures, you are going to have a twofold conformance problem: conformance will be expensive, and it will be inconsistent.

The first and best way to ensure conformance with a constraint is to factor out the constraint. If the constraint is that the titles of works should always be formatted in a certain way, you can factor out this constraint by moving to the document domain and using something like DocBook's `citetitle` element to mark up the names of works. Now it is the publishing algorithm that is responsible for conforming to the formatting constraint. The formatting constraint has been factored out of the content.

When we move content creation from the media domain to the document domain we are factoring out all of the formatting constraints of the document. When we move content from the document domain to the subject domain we are factoring out many of the document or management constraints.

But while this factors out one set of constraints, it also creates a new set of constraint in the new domain. When we factored out the formatting constraint for the titles of works, we introduced a new constraint, which is to markup the title of works using `citetitle`. Factoring one constraint into another is useful if it makes the constraint easier to conform to or easier to validate. A constraint may be easier to conform to if it is simpler, easier to remember, or does not require knowledge that is outside the writer's concerns. For instance, the `citetitle` tag is a single tag, not a set of formatting instructions, and the author knows when they are citing the title of a work. It may be easier to validate if it has fewer components or can be limited to a narrower scope. For instance, the set of things that are titles is smaller than the set of things that are formatted in italic, so validating the `citetitle` constraint requires looking at a smaller and more homogeneous set.

If the constraint you are introducing is not easier to conform to or easier to validate, you probably should not introduce it. There are certainly cases where moving content to a more formal document domain model introduces more constraints than it eliminates without making those constraints easier to comply with or validate. On the other hand, sometimes those additional constraints are required for algorithms. Inevitably, authors have to make some concessions to the needs of algorithms, but we don't want algorithms to impose such a burden on authors that we distract them from doing good research and quality writing. So if the structures you create for the sake of algorithms prove too complex for authors to conform to, or too difficult to validate, you should ask if you could refactor those constraints again, perhaps to the subject domain so that you get both the precision and detail that the algorithm needs and the ease of use and validation that the writer needs.

Completeness

Completeness is an obvious aspect of content quality. Unfortunately, lack of completeness is often hard for writers and reviewers to spot. The omission of information they already know is hard to see unless there is an obvious hole in a predefined and explicit document structure. Defining structures that highlight information requirements can significantly improve completeness.

In the subject domain chapter (Chapter 6, *Writing in the Subject Domain*) we saw how calling out the preparation time and number of servings for a recipe help ensure that the author always remembers to include that information, whether or not we decide to present it in fields or as part of a paragraph.

But this is not the only way that subject domain structured writing helps ensure completeness. Every subject-domain annotation highlights a subject that is important to your business. You can use an algorithm to scan those annotations and build a list of subjects that are important to your business. You can use this list to make sure that all the subjects you need to cover are actually covered.

For example, structured writing allows you to annotate your text to call out the types of certain phrases such as function names, feature names, or stock symbols.

When installing widgets, use a {left-handed widget wrench} (tool) to tighten them to the recommended torque for your device.

This sample annotates the phrase “left-handed widget wrench” and records that these words describe a tool. If all mentions of tools are annotated this way, you can then compile a list of all the tools you mention in all your topics and make sure that you have suitable documentation for each of them.

Consistency

Similarly to completeness, consistency can make a big difference to readers, but lack of consistency is hard to spot if the structure of the content is not explicit in all the ways we want it to be consistent.

Being consistent simply means abiding by constraints. As we have seen, this can take the form of either enforcing the constraint through required structure or, preferably, factoring out the constraint so that it is handled by algorithms rather than people. We have looked at how you can factor out constraints in both the document domain and the subject domain.

If you annotate the important things in your content set, such as tools in the example above, you can use the annotations to check for consistency of names. Suppose a writer accidentally writes “spanner” rather than “wrench” in the name of the tool:

When installing widgets, use a {left-handed widget spanner} (tool) to tighten them to the recommended torque for your device.

Since you can now generate a list of tools mentioned in the text, you can check each mention against a list of approved tool names. This can reveal both incorrect names (consistency) and tools that may be missing from the official list (completeness).

The same would apply to values in fields, such as the wine match field in the recipe example. With the wine match in a separate field you can compile a list of wines mentioned or check each mention against an approved list.

Accuracy

Accuracy problems are often hard to spot. Typos, using old names for things, or giving deprecated examples are all hard for writers and reviewers to see. But there are structured writing techniques than can catch many of these kinds of problems.

For example, if you were documenting an API, you could annotate each mention of a function.

Always check the return value of {rotateWidget()} (function) to ensure the correct orientation was achieved.

API function names can be quite tricky to remember sometimes and small typos can be difficult to spot. But with this annotation, you can validate all mentions of functions against the API reference or the code base. This technique not only catches misspellings. I have seen it catch the use of deprecated functions in examples, for instance.

Semantic constraints

We can usefully divide constraints into two types: structural constraints and semantic constraints. Structural constraints deal with the the relationship of various text structures. Semantic constraints deal with the meaning of the content.

For instance, consider this structure:

```
<person>
  <name>John Smith</name>
  <age>middle</age>
  <date-of-birth>Christmas Day</date-of-birth>
```

```
</person>
```

Some people certainly describe themselves as middle aged, and Christmas Day is certainly a date of birth, if an incomplete one. The author has complied with the structure of the document. But the creator of this markup language was probably looking for more precise information, probably in a format that an algorithm could read. What they wanted was:

```
<person>
  <name>John Smith</name>
  <age>46</age>
  <date-of-birth>1970-12-25</date-of-birth>
</person>
```

Some schema languages, such as XML Schema, let you specify the data type¹ of an element. You could specify that the type of the age field must be whole number between 0 and 150 and that the date-of-birth field must be a recognizable date format. Here's what the XML schema for these constraints might look like:

```
<xss:schema
  xmlns:xss="http://www.w3.org/2001/XMLSchema"
  elementFormDefault="qualified">

  <xss:element name="person">
    <xss:complexType>
      <xss:sequence>
        <xss:element name="name" type="xss:string"/>
        <xss:element name="age" type="age-range"/>
        <xss:element name="date-of-birth" type="xss:date"/>
      </xss:sequence>
    </xss:complexType>
  </xss:element>

  <xss:simpleType name="age-range">
    <xss:restriction base="xss:int">
      <xss:minInclusive value="0"/>
      <xss:maxInclusive value="150"/>
    </xss:restriction>
  </xss:simpleType>
</xss:schema>
```

¹The data types referred to in the example above are not data types as they are commonly understood in programming terms (which refers to how they are stored in memory). In XML, as in all major markup languages, the data is all strings. What a data type really is in a schema is a pattern. There is a language for describing patterns in text that is called regular expressions. Regular expressions are a bit cryptic and take some getting used to but they are incredibly powerful at describing patterns in text. XML schema lets you define types for elements using regular expressions, so there is a huge amount you can do to constrain the content of elements in your documents.

```
</xs:simpleType>
</xs:schema>
```

This schema uses the built in types `xs:string` and `xs:date` for the name and date-of-birth elements and defines a new type called `age-range` for the age element. Using this schema, the example above would now fail to validate with type errors reported for the age and date-of-birth fields.

Applying these kinds of semantic constraints to content is not going to work if most of your text is in free form paragraphs. It is hard to define useful patterns for long passages of text. If you want to exercise fine-grained control over your content, therefore, you must first break information down into individual fields and then apply type constraints to those fields.

In some cases, we create text structures purely for the purpose of being able to apply semantic constraints to their content. We use structural constraints to isolating semantic constraints so that they are testable and enforceable.

This can be particularly effective when you are creating content in the subject domain since you don't have to specify information in sentences, even if you intend to publish it that way. You can break the content out into separate structures and define the data type of those structures to ensure you get complete and accurate information, and to ensure that you can operate on that information using algorithms.

The recipe text that we have used before is a good example of how content that could be expressed entirely in free-form paragraphs can be broken down in a fine-grained way that allows us to impose a variety of structural and semantic constraints.

```
recipe: Hard Boiled Egg
introduction:
    A hard boiled egg is simple and nutritious.
ingredients:: ingredient, quantity
    eggs, 12
    water, 2qt
preparation:
    1. Place eggs in pan and cover with water.
    2. Bring water to a boil.
    3. Remove from heat and cover for 12 minutes.
    4. Place eggs in cold water to stop cooking.
    5. Peel and serve.
prep-time: 15 minutes
serves: 6
wine-match: champagne and orange juice
beverage-match: orange juice
nutrition:
```

```
serving: 1 large (50 g)
calories: 78
total-fat: 5 g
saturated-fat: 0.7 g
polyunsaturated-fat: 0.7 g
monounsaturated-fat: 2 g
cholesterol: 186.5 mg
sodium: 62 mg
potassium: 63 mg
total-carbohydrate: 0.6 g
dietary-fiber: 0 g
sugar: 0.6 g
protein: 6 g
```

This entire recipe could be presented free form. But when structured like this we can enforce detailed constraints like ensuring that there is always a wine match listed or calories are always given as a whole number. The publishing algorithms could actually stitch all this content into paragraphs again if that was how you wanted to publish it. But when it is created in this format the author is provided with a huge amount of guidance about the information you want, and you are able to manipulate and publish the content in many different ways. Reader's will benefit because all the recipes will conform to what you know reader's need and want in a recipe.

Entry validation constraints

The ability of algorithms to read the data in your structures can have another conformance benefit because it allows you to check one piece of information against another. For instance, if you have date of birth and age you can calculate current age from the date of birth and compare it against the value of the age field. If the values don't match, you know the author made an error and you can report it. Here is a Schematron assertion that tests this constraint (in a slightly imprecise fashion: date arithmetic is surprisingly hard):

```
<schema xmlns="http://purl.oclc.org/dsdl/schematron"
queryBinding="xslt2">
  <pattern>
    <title>Age constraint</title>
    <rule context="person">
      <assert test="age =
        xs:int(days-from-duration(current-date() -
        xs:date('1970-12-25')) div 365.25)">
        Age does not match the given date-of-birth.
      </assert>
    </rule>
  </pattern>
```

</schema>

Referential integrity constraints

In the management domain, there are a set of constraints that we could call referential integrity constraints. Referential integrity simply means that if you make a reference to something, that something should exist. In the management domain, we often give IDs to structures and use those IDs to refer to those structures for purposes such as content reuse.

If you are going to reuse a piece of content by referring to its ID, there is an obvious constraint that a piece of content with that ID must exist. This constraint is important enough that the XML specification actually builds in direct support for it. XML itself requires that if an attribute is defined as having the type IDREF, then there must be an element with an attribute of type ID with the same value in the same document. This can be useful for checking things like a footnote reference in a document actually corresponds to a footnote somewhere in the document.

Many management domain algorithms, however, require referential integrity not only within a single document but between documents. The conformance of a document to these referential integrity constraints can sometimes only be judged by the publishing algorithm when it is published in a particular combination. In fact, it is possible for a document to have referential integrity when published in one collection and to lack it when published in another.

Since it is always better to validate a constraint as early as possible, a content management system that is aware of the referential integrity constraints of a system (such as a DITA CMS, for example) may validate the referential integrity of content in all its potential combinations without the need to actually publish it.

Needless to say, however, referential integrity constraints of this complexity still present a management and authoring headache, even with content management system support. It is worth considering if there is a way to factor them out. One example of how this can be done is found in the various approaches to the linking algorithm (see Chapter 17, *Linking*).

Conformance to external sources

Referential integrity constraints can span multiple documents. So can semantic constraints. For example, you may want values in your document to match values in databases or in other documents. A technical writer documenting an API may produce an API reference, much of which may be extracted from the program source code, and also a programmers guide, which they will write from scratch. The programmer's guide will obviously mention the functions in the API many times. There is the possibility that the writer may misspell one of the names, or that the API may be changed after parts of the document are written, or that a function the writer has mentioned no longer exists.

It is clearly a semantic constraint on the programmer's guide that all the API calls it mentions should actually be present in the API. Since the API reference is generated from the source code, we can express this constraint as: functions mentioned in the programmers guide must be listed in the API reference.

This is an important constraint. When we implemented this constraint on one project I worked on, it revealed a number of errors:

- Misspelled function names in the programmer's guide.
- The inclusion in the programmer's guide of material related to a private API that was never released to the public.
- The failure of the API guide to include an important section of the API due to incorrect markup in the source code.
- A section of the programmer's guide that discussed how to do things using a deprecated API and failed to discuss how to do them with the new API.

All these errors were present in the programmer's guide despite several thorough reviews by multiple people over multiple software release. These are all the kinds of errors that human beings have a hard time spotting in review. But they all have significant impact on users who are trying to actually use the API.

As part of the conformance algorithm for the programmer's guide, we added a check that looked up each reference to an API call, including those in code blocks, in the source files for the API reference and reported an error if they did not match. None of the errors listed above would have been detected without this check.

Of course, for this check to be possible, the algorithm that did the checking had to be able to identify every time the programmer's guide mentioned an API call, and it had to be able to find all the API call names in the API reference. For this to be possible, both documents had to be written in a specific structured format that made the function names accessible to the algorithm. Here is a simplified example. First, a code sample from a programmer's guide:

code-sample: Hello World

The Hello World sample uses `{print}()` to output the text "Hello World"

```
```(python)
print("Hello World.")
```

Next a function reference listing from the API reference:

```
function:
 name: print
 return-value: none
parameters:
 parameter: string
 required: yes
 description:
 The string to print.
 parameter: end
 required: no
 default: '\n'
 description:
 The characters to output after the {string}(parameter)
```

Because the API reference labels “print” as a function name and the code-sample annotates “print” as the name of a function, we can look up “print” in the API reference to validate the annotated text in the programmer’s guide.

By adding these structures and annotations to the content, we isolated the semantics of the function call names so that we could apply semantic conformance checks to them.

## Conformance and change

Requiring conformance to outside sources means that a document’s conformance is neither static nor absolute. A document that was conforming may stop being conforming because of outside events. But this reflects reality. One of the most difficult aspects of content management, in fact, is detecting when a document ceases to be conforming because of a change in the reality that it describes. Using structured writing techniques to validate the conformance of a document against an external source can go a long way to addressing this class of problem.

## Design for conformance

So much for the mechanical aspects of conformance. But conformance is fundamentally a human activity and we may need humans to conform to constraints that we cannot easily express or validate in purely mechanical terms.

Schemas and downstream algorithms can verify that a writer created the required mechanical structures, and, in some cases, that they created them with content in the required form. But in most cases there is little or nothing that they can do to verify that the content of those structures is actually what the structure says it is. Unless the author understands what the structure is for, they are not likely to create content that actually obeys the constraints that the structure expresses.

If the author is not on-board with creating content according to the structure, there is nothing the mechanical conformance can do to ensure that they did not just write it the way they wanted to while wrapping the required structures tags around it to make it pass the validation tests.

If there are optional elements to a structure, there is nothing that mechanical validation can do to determine if the options parts of the structure were included when they should have been or omitted when they should have been. Only the author can know this, and only if they understand the purpose of those structures and are on board with them as the appropriate structure for what they are writing.

If the structures are difficult for the writer to understand (if they are complex management domain structures that require complex abstract IDs, for instance) then the author has to think about issues outside of their domain expertise. They have to stop thinking about writing and their subject and start thinking about the system. Since most writers don't really understand how complex structured writing systems work internally, they tend to do the simplest thing they can think of that makes the system stop throwing errors and accept their work. This often produces dumb compliance to mechanical system constraints rather than intelligent conformance to content constraints.

The real key to achieving conformance is to create structures that are easy to conform to. For most content, conformance is not about trying to catch evil doers. The authors are on side and trying to produce good content. Authors who understand structured content may impose constraints as an aid to their own work, just as a carpenter, for instance, might design a jig to guide their saw. Constraints are a tool for writers, not a defense against them. Constraints may force a lazy writer to pull up their socks and do some more research. They may force an inattentive writer to recast their first draft into a more consistent format. But they should never prevent a good and diligent writer from doing good work.

The real core of compliance in structured writing, therefore, is not enforcement, but creating structures that:

- Are clearly and specifically appropriate to the subject matter and the audience being addressed.
- Clearly and specifically supply the information that the reader needs to accomplish a clear and specific goal.
- Clearly communicate to authors what is expected of them in terms they understand.
- Either remind authors of what is required or (preferably) factor it out.

Writers may disagree, of course, about what goals should be addressed, what information readers need to achieve those goals, or how best to express information to that audience. Where we have many authors contributing to a common information set, it is important

that these differences be addressed and resolved professionally and that all authors involved are on board with the plan going forward. It is at that point that well-thought-out content structures can capture the decisions made and make sure that everyone stays on track and is consistent.

For example, suppose we have a set of cooks contributing to a common recipe information set. Everyone gets on board with the principle that every recipe should include the preparation time and the number of servings. At that point, creating a recipe structure that explicitly calls out those fields helps all the cooks to remember, and therefore comply with, the agreement they have made. This will significantly improve compliance compared to merely stating the requirement in the style guide. (And, of course, it will make that information available to other algorithms as well.)

Auditing and enforcement still have a role to play, not because authors are hostile to the system, but because they are human. But auditing and enforcement are secondary to the main aim of conformance-friendly design. And in that spirit, auditing and conformance should be seen as part of a feedback loop that is constantly seeking to improve the design. If you keep finding the same mistakes over and over again, that is not a training problem or a human resources problem, it is a design problem.

Finally, be aware of how your authoring interface affects conformance. How do you know if you are meeting constraints -- in any activity? Feedback. With any activity, we need a way to know when we are done and when the work is correct. In the media domain, there is one form of feedback: how the document looks. With a true WYSIWYG display, if it looks right on the screen, it will render correctly on paper, or whatever media you are targeting as you write. That is the writer's signal that they are done, and done correctly. That's fine in the media domain. But if you are trying to create content in the document domain but giving your author's a media domain display (thanks to a structured editor that is trying to look as much as possible like Microsoft Word) then the feedback they receive will be media domain feedback, and the conformance you will get will be media domain conformance. However, there is a very real risk that content created this way will not work as intended when output to other media, or reused, or processed with any of the other structured writing domains. Don't expect to get real conformance to any structure that is hidden from the author.

---

# Chapter 19. Auditing

If the conformance algorithm is about making sure that an individual item meets its constraints, the audit algorithm is about making sure that the content set as a whole meets its constraints.<sup>1</sup>

As such, auditing is fundamentally a content management function. It is about making sure that:

- The definition of the content set is correct (we know what types of content it should contain, and which instances of each type)
- The content set is complete (it contains all the items of each type that it should)
- The content set is uncontaminated (it does not contain any items or types it should not)
- The content set is integrated (it contains all of the relationships between items that it should)
- Each item in the content set conforms to its constraints

Auditing a large content set is difficult and many CMS solutions are deficient in audit capabilities. The main reason for this is that with the way most content is recorded and stored (media domain or generic document domain formats), it is very difficult to mechanically assess what content you have and what state it is in. It is hard to know if you have all the pieces you should have if you can't tell exactly what the pieces you have are.

Auditing is a content management function and this book is about structured writing, not content management. However, one of the biggest, and least appreciated, benefits of structured writing is that it makes content more auditable. When content management systems fail or become unmanageable, (which they do with disappointing frequency), the root cause is often either lack of attention to regular audits, or the lack of ability to audit effectively. Without the ability to audit effectively, content sets often end up incomplete, corrupt, and poorly integrated, which reduces quality and increases costs at every stage of the process. And a vicious cycle can develop in which writers, frustrated with the difficulties of the system, create workarounds that further corrupt the information set. Whatever expenses you may incur to implement a more structured structured writing approach could well be offset by the savings associated with more effective auditing of the content set alone.

---

<sup>1</sup>Content strategists often use the term “content audit” to mean a current state analysis performed at the beginning of a website redevelopment project. A content strategy content audit is about cataloging, and possibly categorizing, the content you already have. I am using the word audit to refer to an ongoing and or recurring activity in which you ensure that a content set is meeting or continuing to meet its goals.

# Correctness of the definition of the content set

Content strategists will spend a great deal of time and effort developing a content plan (usually this is for a website, but the same principle applies to any content set). How they do this is beyond the scope of this book, but the result should be a definition of the content set: which types of information it is supposed to contain and what instances of those types. (This definition is based, of course on the goals it is designed to achieve, which is the business of the content strategist to define.)

The definition of a content set is not necessarily static. It is not necessarily a fixed list of topic types or of specific topics to be developed. For one thing, the subject matter may change during the course of content development, which would change the content pieces needed, and perhaps require new content types or modifications to existing types. Second, the exact set of pieces or types may not be knowable at the outset. Content development explores a complex set of relationships between subject matter and the needs and background of the reader that cannot be fully known without traveling the ground in detail. An agile approach to content planning is essential in most content projects. But remember that agile is a disciplined approach to evolving a plan. It does not mean hacking away regardless; it means constantly refining your model in a disciplined and deliberate way as you learn more about your subject matter, your readers, and your business needs.

But it is hard to be disciplined and deliberate in evolving the big picture model of the content set if you are not disciplined and deliberate in how you create the pieces. If the author is given no guidance and is then asked to assign CMS metadata to the article after the fact, they will tag it using the terms that seem like the closest fit, but they will not revise the content to fit the labels they are applying to it. You won't really know what type of content you have.

If you don't really know what type of content you have, you can't really tell if the definitions for your content are correct. Some content may perform poorly because it does not fit the type definition properly. But you can't tell whether it failed because it didn't fit the type definition or because the type definition is wrong. Thus you don't know what to fix.

Similarly, an author may write something that is better than the current type definition. But if it is tagged as that type in the CMS, the type definition will get the credit for the success. But then the problem is that you don't detect the problem with the type definition, and so you don't change it. But then the content of other authors who follow the type definition does not work as well, giving you inconsistent results that are hard to interpret. Unless your content types are codified and auditable, their success won't carry over to other content. And if the content fails, you won't know whether it failed because it was true to a badly defined type or false to a well defined one.

Having strong well defined topic types makes it easier to audit your topic types to make sure they are doing the job they were designed to do. Similarly, having strong well defined topic types means that you can have greater assurance that each topic is doing the job it is supposed to do, which helps you make sure you have covered all the subjects you should have.

But structured writing can do more than this to help you audit the definition of the content set. If you create content in the subject domain, including extensively annotating the subjects that you mention in the text, you can use algorithms to extract a list of the types and subjects that your content is actually talking about. In your initial top-down plan, you may not have thought about the need for content on a certain subject or to support a certain activity, but if that subject or that activity start showing up in the body of your content, that is a strong indication that those subjects and activities are related to the purpose of your content set and should probably be included in the definition of the content set.

Subject domain markup is how you know what your content is actually talking about, what every author is discovering or thinks needs saying. Without this information it is very difficult to establish that the content set is meeting its coverage goals.

This actually attacks two audit problems. If writers are writing about things outside your current coverage definition, either your coverage definition needs updating, or writers are polluting the content set with irrelevant material.

## Ensuring the content set remains uncontaminated

Subject domain content structures and annotations can help you prevent contamination of the content set by irrelevant material. But more important than catching writers in the act is catching the flaws in content types that allow for contamination to creep in.

A major form of contamination in any content set is redundant content. As we noted in Chapter 15, *Single source of truth*, we have to be careful in how we define redundancy, because it is not simply a matter of only addressing a subject once, it is a matter of addressing an audience need only once, and that may require several topics on the same subject addressed to different readers. But it is all too easy for duplicate content to sneak into a content set. Some of it comes in because the same functionality is repeated in many products or in content delivered to different media. Some comes in through authors simply not knowing that suitable content already exists.

Content reuse is a major motivator for structured writing for exactly this reason. But the content reuse algorithm only addresses the problem of how to reuse content. It provides a method to reuse content you are aware of. It does not prevent you from duplicating content because you did not look for or did not find existing reusable content. You need

to audit your content regularly to make sure that content is findable by writers who may want to reuse it, and to make sure that duplication is not creeping in.

There are natural language processing algorithms that will attempt to identify redundant content in a content set, but such algorithms focus on similar texts. This is not enough. The same or similar sequences of words may occur in different places without being redundant. They may mean different things, or perform different roles, in context. On the other hand, redundant pieces of information may be expressed in very different words. In other words, it is redundant information, not redundant phrases, the we care about.

Even when redundancies are found, they may be very difficult to consolidate if they don't have similar boundaries within their respective documents (the composability problem -- see Chapter 14, *Composition*). Strongly typed content, meaning content that conforms to a model that breaks down and enforces the various components that make up a document, makes it possible to detect duplication in a much more formal and precise way.

A person who consults a repository to see if there is a piece of content they can use relies on the ability to query the repository in a sensible way for the type of content they are looking for. They also rely on their ability to recognize the content when they see it, and on it actually being strongly conformant so that they can use it with confidence. Strong topic typing helps with all of these things. The easier it is to correctly identify reusable content and use it, the less corruption of the repository will occur.

## Ensuring that the content set is well integrated

A content set is never a collection of wholly independent pieces. The items in the set have relationships to each other that matter to the reader. Whether you express those relationships through links or cross references, or whether you rely entirely on tables on contents and indexes, it is still important to understand and manage the relationships between items.

Relationships between items may matter for management reasons as well. If you have documentation for multiple releases of a product, the relationship between the documentation for feature X in version 3 and that for feature X in version 2 matters to you. It may matter because the feature has not changed and you can reuse the item. It may matter because an error was found in version 2 and you want to fix it in version 3 as well. And if you put this content online, the relationship may matter for the reader as well, if they search for feature X and get the result for version 2 when they are using version 3.

You can describe the relationship between items externally. Items are related whenever they have the same value in any one of their metadata fields. Which field it is tells you what the relationship is. Finding the relevant metadata field to look at allows you to manage the relationship. But the same problem exists here as it always does with all external metadata (see Chapter 26, *Content Management*) -- the content may not conform

to the metadata, and without structured writing in the content itself, it is hard to audit the conformance of the content to its metadata.

But the bigger problem is that external metadata does not map the important relationship that can exist a part of one item and the whole of another item. Are function names appearing in the programming topics all listed in the API reference? Are the utensils mentioned in a recipe all covered in the appendix of kitchen tools? These are important content relationship questions, but their importance cannot be mapped with external metadata. You need the markup inside the piece that identifies function calls and the names of kitchen tools.

Structured writing, particularly in the subject domain, helps you discover and manage these relationships by making clear the subject on which these relationships are based.

## Making content auditable

I have talked all through this chapter about how using strong topic types makes content easier to audit. What is a strong topic type? Fundamentally, a strong topic type is one that makes explicit what the content is supposed to say and how it is supposed to say it. Or, to put it another way, a strong topic type is one that captures, enforces, or factors out the major constraints of the content, including its major rhetorical constraints.

It is possible for content to conform to all of its rhetorical constraints without the use of structured writing techniques. But strong topic types provide explicit guidance to the author and facilitate the use of conformance algorithms. They are created to meet your conformance goals. Similarly with auditing, you specify the content structures you need in order to meet your auditing goals.

Auditing is sometimes not as straightforward as conformance, even with structured writing techniques in place. Auditing often requires human review, not only to make sure that all subjects have been covered, but to discover new issues or subjects that need to be addressed. Human review of a large content set is difficult, though, due to the sheer amount of content. An audit algorithm can simplify this work by creating different views of the content set that humans can review more easily.

Suppose, for instance, that an organization is using subject-domain annotations to drive linking as described in Chapter 17, *Linking*. Every topic in the collection is supposed to be indexed to state the type and names of the subjects it covers. Every mention of a significant subject is supposed to be annotated with its type. The linking algorithm can certainly use these annotations and index entries to link the content without any need for authors to create or manage links in the source text. But that does not guarantee that all the right links get made. There could be errors in indexing or annotation that are impossible to detect when conformance testing individual topics.

We can use those same index entries and annotations to create audit reports for several purposes. These are some of the things we can do:

- We can create a sorted list of all the phrases that have been annotated and see if they are being annotated consistently. (Everyone is using the same annotation type, for instance.) This will tell us a lot about the types we are using, how well they are understood, and what instances of each type we should be covered.
- We can create a list of all the phrases that have been indexed and check it against our content plan (perhaps against a taxonomy if we have one). This will tell us a lot about if our coverage is complete, if writers are getting off track, or if our content plan or our taxonomy is off base with reality.
- We can create a sorted list of all the index terms and check it against the list of annotated phrases to find phrases are that being indexed but not annotated or annotated but not indexed. This can tell us is there are subject we are not covering, if writers are discussing subjects they should not be, or if some topics are not being indexed or annotated properly.
- During the content development phase, the list of things that are annotated but not indexed will inevitable grow, as subjects are being referred to before the content that describes them is written. The trend line of the growth of new subjects being annotated vs subjects being indexed will allow you to track how close a content set is to completion, even in cases were defining the boundaries in a advance is difficult.

## Performance auditing

Auditing marketing content can be easier than auditing other types of content because you can measure content against behavioral goals. Are readers taking the actions you want, and does a content change produce more of the desired action? All content aims at changing reader behavior in one way or another, but not all behavior changes are easy to measure. Unless the behavior in question is an interaction with the web page that contains the content, behavioral changes are both hard to observe and hard to attribute.

Whether auditing the performance of your content is easy of difficult, however, when you do have evidence that a particular piece of content is performing well, you want to be able to reproduce that success.

If a piece of content follows a well defined set of constraints, it is much easier to figure out why it was successful. Rather than trying A/B testing of individual pieces, for instance, you can try A/B testing of content that follows different structures. If one structured consistently outperforms the other, you can adopt that structure for all your content.

Of course, to ensure continues success (or to ensure that the same structure is continuing to perform well) you need to be sure that the new content you are producing does actually conform to the proven successful structure. This means that you need to apply the conformance algorithms to your new content creation, as well as auditing it for overall compliance and continued success.

---

# Chapter 20. Relevance

Establishing the relevance of content is essential for readers, and many of the structured writing algorithms.

- Every reader who looks for content has to decide if a given document is relevant to their needs.
- A search engine needs to determine if a document is relevant to a reader's search.
- The reuse algorithm needs to determine if content relevant to a certain point exists or not.
- The composition algorithm needs to find all the pieces of content relevant to a certain purpose.
- The linking algorithm needs to find any documents relevant to a subject mentioned in the current document.
- The conformance algorithm needs to determine if there is information in one document that should conform to information in another document.
- The quality algorithm needs to ensure that a document contains all the information relevant to the reader's needs.

There are two fundamental parts to the relevance algorithm: being relevant, and demonstrating your relevance.

## Being relevant

All too often, making sure that content is relevant is neglected. CMS systems will attach metadata to unstructured documents asserting their relevance to certain topic or requirements without verifying that they actually meet those needs. In many cases, the CMS has a cataloging system for content that requires certain metadata fields to be completed, often from a predetermined list of values. If developed correctly, this metadata may have the potential to define the relevance of content, but these requirements are not always communicated to the author before they write. Instead, writers are left to compose content without any guidance as to what would make it relevant, and then asked to tag it with the CMS metadata after the fact. In many cases, this tagging is done hastily and lazily, simply picking whatever terms seem appropriate at first glance without really thinking about whether the piece fulfills the promise the metadata is making.

The result is that the content returned by the CMS is often not relevant to the reader's needs because it simply does not possess the relevance that the CMS metadata claims. In many cases, bypassing the CMS's navigation mechanisms and doing a simple site

search will produce better results, since the search engine looks at the content itself to determine what it is relevant to. But while this can rescue the reader's quest for relevant content where such content exists, it does nothing to ensure that relevant content actually gets created. Without this assurance, the CMS may be telling lies to its owners as well as to its readers.

The plain fact is that if you are going to implement strict standards of relevance, and a strict vocabulary for proclaiming relevance, this needs to be communicated to the writer long before they have finished writing and are trying to submit their content to the CMS. One way to do that is through structured writing.

## Showing relevance

Being relevant is important, it is equally important to show that you are relevant. Readers and algorithms can't find you if you don't show your relevance.

Actually, that is not entirely true. Search engines can determine the relevance of content with an amazing degree of accuracy simply by reading the text itself and looking at how the content is used and linked to. This attests to the first importance of being relevant.

But if the search engine decides a topic may be relevant and adds that topic to search results, but the reader then looks at it and can't see the relevance, you still don't reach the reader. (And the search engine may take note as well and downgrade that topic's relevance ranking.)

Showing your relevance to the reader is all about establishing your subject matter and context and doing it clearly and unambiguously in the few seconds that the reader is likely to look before dismissing you and moving on to the next item in their search results. When developing a tagging language (or other structured content container) for your content, therefore, one of the most important things to consider is how the reader will recognize what your documents are relevant to.

Relevance is not established the same way for all readers or for all subjects. For a recipe, for instance, a picture may do a lot to establish relevance. For an API reference, a version number and the description of a return value may be key relevance indicators.

But relevance is not just about the subject matter. It is also about the reader's purpose. For a business page, the inclusion of a stock price chart may tell you that the page is of interest to an investor rather than a potential customer. Placing that chart at the top of the page helps establish the relevance quickly. For an article about a place, pictures of beaches or nightlife show that the page is relevant to potential tourists, not to residents trying to decide what schools to send their kids to.

Creating content in the subject domain allows you to make sure that writers produce all the pieces of information that make a page relevant to the intended audience. Because the

subject domain also factors out the order of presentation of content, it helps you to make sure that every page is organized in a way that best shows its relevance.

Better still, because subject domain content is organized by algorithms, you can experiment to see if one organization of content works better than another and adapt your presentation algorithm accordingly without having to edit any of the content.

Also, because storing content in the subject domain allows us to use the extract and merge algorithm to pull in content from other sources, we don't have to include the beach pictures or the stock chart with our subject domain content. Instead, we can have the synthesis algorithm query other content sets or feeds to find the best current tourist shots or to generate the stock chart in real time.

But that, of course, depends on the mechanical part of the relevance algorithm, because in order for the synthesis algorithm to do this it has to be able to find relevant pictures or a relevant stock chart. It is not enough to demonstrate relevance to readers, you also need to demonstrate relevance to algorithms.

## Showing relevance to algorithms

Showing relevance to algorithms comes down to breaking information up into clearly labeled fields containing clearly unambiguous values. For example, to query a web service that generates stock charts, you would need to provide it with a clearly unambiguous identifier for the company whose chart you wanted. The best way to do that would not be the company name, since there can be duplicate or similar names in different industries (Apple Computer vs Apple Music, for example), but the company's stock symbol, which is guaranteed to be unique by the exchange on which it is listed (though you do have to provide the exchange code as well to be globally unique: NASDAQ:AAPL, not just AAPL).

This means that the web service needs to index its information on stock prices according to stock symbols (which is exactly what it does, of course). Actually, it is more than likely that the stock chart drawing service does not hold this information at all, but requests it as needed from yet another web service. The stock ticker symbol is the unambiguous key that identifies the company in each of these transactions.

While humans and algorithms assess relevance in different ways, the foundations of relevance are the same for both: clear identification of the type and subject matter of the information. This means that you need to maintain metadata that clearly identifies the type and subject, and also present the content in a way that makes its type and subject evident at a glance.

It is possible to maintain metadata on type and subject of content in a content management system that hosts content written in the document or subject domains. In this scenario, ensuring that the organization of the content makes the type and subject evident at a

glance is a separate problem from correctly labeling the content in the CMS. While it is certainly possible, with the right authoring discipline, to make sure that the content demonstrates what the metadata claims about it, it is also possible for the content not to match what the metadata claims, or not to demonstrate its type and subject effectively.

Moving content to the subject domain allows you to show relevance to both humans and machines with the same structures, assuring that the two do not get out of sync.

---

# Chapter 21. Active content

When you publish to electronic media, you can create active content, that is, content that has behavior as well as formatting. Some examples:

Dynamic arrangement

Part of the presentation algorithm is arranging content on the page or screen, but with online media you can allow the reader to arrange the content. For instance, you can publish tables that readers can sort for themselves.

Adaptive content

Similarly, you can create content that adapts itself dynamically to the view port in which it is displayed. For instance, displaying in multiple columns on a wide view port, and in a single column on a narrow one.

Progressive disclosure

You can present content in a way that only shows part of the content on the screen initially but reveals more when the user clicks on a link or takes another action. For instance, you might show the high-level of a procedure and provide a link that opens detailed steps for those who need them. This is a way to cater to audiences with different levels of preparedness.

Transclusion

You can pull in content from another source.

Feeds and dynamic sources

You can include content that comes from an external source which updates independently of your content such as a feed or a web service.

Interactive media

You can include graphics and other media that the user can interact with.

Active content is simply the postponement of one or more of the structured writing algorithms to the time of reading. Changing the layout of the page when you resize your browser or rotate your phone is executing the presentation algorithm in the browser. Allowing the reader to change the font size is allowing the reader to configure and execute the the formatting algorithm in the browser. Pulling a live stock quote into a page that mentions a company is having the web server or the web browser executes the synthesis algorithm when the page is served/displayed.

In principle, therefore, you support active content in exactly the same way you support those algorithms in any form of structured writing: by creating your content in the appropriate domain and with the appropriate structures to reliably support the algorithms

you want to run. Reliability is key here, of course. If you are creating static books or web pages you may be able to get away with reviewing the results before they are published and even hand tweaking the output to fix any issues, but you can't do that with active content. You need to be able to trust the algorithm 100%, which means you need to trust the content it is working on 100%. This makes the conformance algorithm crucial to an active content project, and, as we have seen, different domain support the conformance algorithm with varying degrees of reliability.

Another factor to consider, though, is that while active content is essentially just postponing the point at which you execute various structured writing algorithms until the content is in the user's hands, you don't necessarily want to postpone running those algorithms on the whole page. You usually only postpone them for the parts of the page that you want to make active. For instance, if you want real-time stock quotes embedded in your business stories whenever a company is mentioned, there is no reason to postpone the synthesis, presentation, and formatting of the entire page; you only need to postpone the synthesis of the stock quote itself. And if you want that quote to continue to update in real time while the page is displayed, you certainly don't want to be redrawing the whole page every time the price changes.

The display of stock quotes depends on subject domain metadata that unambiguously identifies the company mentioned in the content. For example:

{Microsoft} (company "NASDAQ:MSFT") is a large software company.

In a normal static publishing flow, this subject domain information would be resolved into a static piece of content by the presentation algorithm:

\*Microsoft\* US\$60.29 +0.34 (+0.57%) is a large software company.

But for active content the main presentation algorithm needs to pass through the subject-domain information to the browser so that that bit of synthesis can be executed in real time when the content is displayed.

This means that only those parts of the content that you want displayed as active content need to be captured in the domain, and specific structures, required to reliably run the applicable algorithms. The rest of the content can be in a different domain. In practice, this usually means that you have subject domain structures to support the active content embedded in document domain structures. Of course, having all your content in the subject domain is also an option, in which case you would perform the presentation algorithm on the content you want to deliver statically, while passing on the subject-domain representation of the content you want to be active to the system that will perform the active content algorithms.

However, this is not necessarily, or even likely, to be a straightforward pass through. The downstream system, such as the web browser, that is going to execute the active content

algorithm, may expect or require the subject domain information it uses in a particular format. For instance, you might have to embed the ticker symbol in a call to a function that returns stock quotes:

```
<p>Microsoft ()
is a large software company.</p>
```

Which domain should you work in if you want to deliver active content?

## Media Domain

There is relatively little that you can do in the way of active content in the media domain. Active content involves the action of algorithms on content and the media domain represents the final or almost-final form of content, leaving little scope for algorithms to act, and little in the way of metadata for them to act on. In the media domain, there are really no constraints on the interpretation of the content, and therefore not much that an algorithm can know about content that it could use as a basis for any action. In practice, therefore, active content almost always works on content written in the document or subject domains.

## Document Domain

The document domain represents the desired presentation of a piece of content. Therefore the only kinds of active content algorithms that you can execute of ordinary document domain structures are those that leave the presentation alone and only affect the formatting. For example, you can allow readers to select the font size of a web page, which is a purely formatting concern.

If you want to do active content that dynamically changes the presentation of the content, or that changes the text of the content (as in our stock price example) then you need either subject domain structures or document domain structures that are explicitly designed to support active content algorithms.

For example, a generic table structure does not support the action of allowing the reader to sort on any column. Sorting by column only makes sense if the content inherently sortable.

item	legs	price
table	4	\$400
stool	3	\$20
shooting stick	1	\$75
chair	4	\$60

For example, in the table above the columns are sortable. The reader could choose to sort the table on the item name, the number of legs, or the price, all of which might be useful arrangements, based on their needs.

But consider what might happen if you sorted this table by on the second column:

- |    |                                 |
|----|---------------------------------|
| 1. | Don protective clothing.        |
| 2. | Clear the area.                 |
| 3. | Block all entrances.            |
| 4. | Activate the destruct sequence. |

Thus you can't just decide to make all the tables on a page sortable by column. Unless the document domain structure explicitly states which columns of a table are sortable and which are not, you can't implement this kind of active content.

To implement column sorting at the document domain level, therefore, you need some sort of sortable table structure which assures that the sorting behavior is only applied to columns where it makes sense in tables where it makes sense.

## Subject Domain

While the document domain typically requires different markup for different algorithms, the subject domain typically does not. In the subject domain you capture the semantics of the subject matter, which are the same no matter what algorithms you are applying to the content.

In the subject domain, the destruct sequence would be structured as a procedure, not a table, and therefore there would be no ambiguity about whether the step could be sorted or not.

```
procedure: Self Destruct
 is-it-dangerous: yes
 steps:
 step: Don protective clothing.
 step: Clear the area.
 step: Block all entrances.
 step: Activate the destruct sequence.
```

In the subject domain, the product list would be a structured data set with known semantics (probably maintained as a separate database).

```
products:: item, legs, price
```

```
table, 4, $400
stool, 3, $20
shooting stick, 1, $75
chair , 4, $60
```

Knowing what the semantics are, you would know whether that data set is sortable and therefore whether it can be presented as a sortable table.

Creating your content in the subject domain gives you the greatest flexibility to generate active content in ways that are appropriate to the subject matter and the device. And because the subject domain does not require any different structures to support active content, your authors don't have to understand or even think about how the active content might work.

This does not mean that active content is always going to be a free gift of the subject domain. Apart from the fact that you still have to design and implement the content behavior, you are going to have to make sure that you are capturing the subject domain semantics that are needed to drive the active content behaviors you want. Those semantics may not be different in kind from any other subject domain semantics, but you may need more of them, or to break them down in greater detail, than you might have to for other algorithms. You will also need to make sure that you get a high degree of conformance to these structured from your writers, as it is difficult to validate the correct operation of every active content algorithm on every content set at run time. The success of your active content strategy is going to depend heavily on the quality and consistency of your input data.

---

# Chapter 22. Extract

A great deal of the content we produce, particularly technical and business content, is essentially a report in human language on the specific features of a product, process, or data set. If those features are described in any kind of formal data set, such as a database or software source code, we can use structured writing techniques to extract information from those sources to create and/or validate content.

## Tapping external sources of content

We have talked throughout this book about moving content from the media domain to the document domain and from the document domain to the subject domain. We have seen the advantages that come from creating content in the subject domain, and we have looked at the processing algorithms that can use subject domain content to produce various kinds of publications in different media.

Subject domain content is simply content that is created and annotated in structures that are based on the subject matter rather than on the structure of documents or media. Subject domain structures tell you what the content is about, rather than how it should be published. You can therefore write algorithms that processes it based on what it is about rather than how it is presented.

Any data source that is contained and annotated in subject domain structures is therefore a source of subject domain content, whether it was intended to produce content from or not. This includes virtually all databases and quite a bit of software code. It also includes all authored content anywhere available (under an appropriate license) that contains any usable subject domain structures or annotations. All of this is potential material for generating content as part of your overall publishing process. As such, the extract algorithm can work effectively with many other structured writing algorithms.

Perhaps most obviously, extract and merge works with the composition algorithm. Extract provides new sources of content for the composition algorithm to work with.

As a source of subject domain content, the extract algorithm also naturally separates content from formatting and contributes to the differential single sourcing algorithm.

By tapping existing information to build content, the extract algorithm also works hand in hand with the content reuse algorithm. In fact, it is really the highest expression of reuse, since it not only reuses content in the content system, but information from the organization at large, or even beyond the organization -- a process that further reduces duplication within the organization.

Because the extract and merge algorithm taps directly into external sources of information, it is also a great source of information for the conformance algorithm. At one

level, it provides a canonical source of information to validate existing content against. At another level, it factors out part of the conformance problem from the authoring function.

If the authoring function is required to conform to information in these sources, the best way to do this is to generate content directly from these sources. Then responsibility for the correctness of the content is shifted to the people who maintain the source you are extracting from. Since they were already responsible for the correctness of the information, this does not create any additional work for them, which means that there is a significant net gain in efficiency for the organization.

## Information created for other purposes

There is nothing new, of course, about generating content from database records. Database reporting is a highly important and sophisticated field in its own right and it would be entirely correct to characterize it as a type of structured writing. What sets it apart, largely, from other structured writing practices, is that the databases it reports on serve other business purposes besides being sources of content. An insurance company policy database, for instance, may be used to publish custom benefit booklets for plan participants but it is also used for processing claims. The design of the structures and data entry interfaces of these systems has tended to fall outside the realm (and the notice) of writers and authoring system designers.

This is a pity, because it has often resulted in organizations developing separate processes, tools, and repositories for content creation in which the information already contained in databases is researched, validated, recorded, and managed entirely independently on the content side of the house. Rather than treating code and databases as sources of content, writers treat them as research sources. They look information up in these sources and then go away and write content (in a separate repository) to relate the information from those sources.

The essence of the problem is that many content organizations choose to work in the media domain or the document domain and have neither the tools nor the expertise to bridge the gap to all this material already available in the subject domain. But even when content organizations do extend their efforts into the subject domain, they are often blind to the fact that the subject domain content they are proposing to create already exists in the systems of another department.

Another way in which this is a pity is that when content is produced from these existing systems, for instance by a database reporting process, it exists in isolation from the rest of the content produced by the organization. Such content can often be quite sophisticated and beautifully formatted and published. But it is the product of an entire structured publishing chain that has to be separately developed and maintained.

In the field of software documentation we see the same pattern in regard to programming language API documentation. Much of the material of an API reference guide is a

description of each function, what information is required as input (its parameters or arguments), the information it produces as output, and the errors or exceptions that it can generate. All of this information already exists in the code that implements the function.

API documentation tools such as JavaDoc or Sphinx extract this information from code and comments and turn it into publishable content. This is an application of subject-domain structured writing and the API documentation tools that do this implement an entire structured publishing system internally, producing final output, often in multiple formats.

And here we see all the same problems again:

- An entire publishing chain is maintained separately from the main content publishing chain.
- The content produced from this publishing chain is isolated from all the other content produced by the organization.
- Much of the same information is often created and maintained separately in a different repository and tool chain in the form programming guides and/or knowledge base articles.

There are other cases of entirely separate publishing chains producing information that is isolated from the rest of an organization's content. Technical support organizations create knowledge bases to answer commonly asked questions. The material in these knowledge bases is technical communication, plain and simple, yet it usually exists in isolation from the product documentation set, even to the extent that users may not be able to search both the documentation and the knowledge base from the same search box. Most users, however, have no way of guessing whether the answer they are looking for is going to be in the docs or the knowledge base (or in the users forum, often yet another independent publishing system.)

There are a couple of ways to address these redundancies and the isolation that goes with them. One is to attempt to unify all content authoring and production in a single enterprise-wide system, often with a single set of content structures intended for use across all enterprise departments. However, this is a highly expensive and disruptive approach and tends to create interfaces and structures that are less usable and less specific to various business functions than the ones they replace. It also ignores the fact that many of the systems from which we wish to extract content exist for other purposes besides the content that is generated from them. Their subject-domain structures are specific and necessary to the database functions or software code generation they were built for.

Another approach is to leave the subject domain systems in place (and perhaps create more of them) and feed their output into a common document-domain publishing tool chain. It is a normal part of the publishing algorithm for subject-domain content to pass through the document domain on its way to media-domain publication. Subject domain

content, by its nature, is not strongly tied to a particular document domain structure, so integrating many sources of subject domain content into a single publishing chain is not particularly onerous. (Specific management domain features of certain tool chains make things more complicated, but since the subject domain tends to factor out a lot of the management domain, this is not an insurmountable problem.)

Most enterprise-wide content systems are based on document-domain languages. (There is, after all, no way to create a single enterprise-wide subject-domain system, since an enterprise creates content on many subjects.) In principle, a document-domain system should be capable of integrating content from domain-specific subject-domain systems. Unfortunately, it is not common for either the subject domain systems or the enterprise content systems to be designed with this kind of integration in mind.

Because of this, we sometimes have to find ways to extract content from these sources and feed them into a unified publishing chain. This creates the need for extraction algorithm (which are part of the synthesis algorithm within the overall publishing algorithm).

## The Extraction Algorithm

A common example of the extraction algorithm is found in API documentation tools such as JavaDoc. These tools parse application source code to pull out things like the names of functions, parameters, and return values, which it then uses to create the outline, at least, of reference documentation. Essentially, it generates a human language translation of what the computer language code is saying.

How the extraction algorithm works depends entirely on how the source data is structured, but it should usually create output in the subject domain that clearly labels the pieces of information it has culled from the source. For instance, a Java function definition is a piece of structured content in which the role and meaning of each element is known from the pattern and syntax of the Java language itself (its grammar):

```
boolean isValidMove(int theFromFile,
 int theFromRank,
 int theToFile,
 int theToRank) {
 // ...body
}
```

This same information can be extracted by an algorithm that knows the grammar of Java to produce something that looks more like subject domain content:

```
java-function:
 name: isValidMove
```

```
return-type: boolean
parameters:: type, name
 int, theFromFile
 int, theFromRank
 int, theToFile
 int, theToRank
```

This is the same information, but in a different structure. In this structure, however, it is easily accessible by content processes and can then be processed through the rest of the publishing tool chain just like any other content.

Any structured data source that expresses the semantics of its data in a consistent way is a source of subject domain content. We just need to find a way to get at it.

## The diversity of sources

When it comes to drawing content from diverse sources, the term single sourcing can mislead us. Single source can lead us to think that it means all source content is kept in a single place. Some vendors of content management systems would like to encourage this interpretation. But a better definition is that each piece of information comes from a single source. That is to say, that each piece of information comes from only one source, but there may be many sources, each maintaining different information (this is an aspect of the single source of truth algorithm).

This has nothing to do with keeping it all in the same place. Nor is keeping all content in the same place a particularly useful approach to ensuring that it is only stored once. Ensuring that content is only stored once, a process formally called “normalization” is actually about making sure that information, and the repository in which it is stored, meets an appropriate set of constraints.

This is not to say that there are never trivial differences between the ways in which different bodies within an organization store and manage information that should not be rationalized. There are all kind of isolated and ad hoc information stores in most organizations that could potentially be much more efficient and (vitally) much more accessible, with a degree of rationalization and centralization. But it does not follow at all that absolute centralization into a single system or a single data model is appropriate, useful, or even possible.

The best way to ensure that information is stored once is to have it stored in a system with the right constraints and the right processes for the people who create and manage that information. This may mean that an integrated publishing system will draw from diverse sources of information and content. The ability to extract content from these sources and to merge it with other content for publication is therefore central to an effective strategy.

---

# Chapter 23. Merge

If the extract algorithm gives us access to new sources of subject domain content, it does not always give us everything we need for a complete document. Sometimes we need to combine the content we have extracted with content we have written in order to present a complete set of information. This is the job for the merge algorithm.

The information that we can extract about an API function from the code that implements it is useful content, there is not enough detail to build an API reference with this information alone. A good reference entry also requires some explanation of the purpose of the function, a little more detail on its parameters, and possibly a code sample illustrating its use.

To address this, we can merge authored content covering these topics with the content we have extracted from the source.

In the case of API documentation tools, the authored content for merging is often written in the source code files. It is contained in code comments and is often written in small subject domain tagging languages that are specific to that tool. (Though as with all subject domain structures, any other tool can read them if it wants to.)

Here is an example of authored content combined with source code in JavaDoc<sup>1</sup>:

```
/**
 * Validates a chess move.
 *
 * Use {@link #doMove(int theFromFile,
 * int theFromRank,
 * int theToFile,
 * int theToRank)} to move a piece.
 *
 * @param theFromFile file from which a piece is being moved
 * @param theFromRank rank from which a piece is being moved
 * @param theToFile file to which a piece is being moved
 * @param theToRank rank to which a piece is being moved
 * @return true if the move is valid, otherwise false
 */
boolean isValidMove(int theFromFile,
 int theFromRank,
 int theToFile,
 int theToRank) {
 // ...
}
```

---

<sup>1</sup><https://en.wikipedia.org/wiki/Javadoc#Example>

}

In this example, everything between the opening `/*` and the closing `*/` is a comment (as far as Java itself is concerned), and the rest is a function definition in Java. However, JavaDoc sees the comment block as a block of structured text using a style of markup specific to JavaDoc.

The JavaDoc processor will extract information from the function definition itself (the extract algorithm) and then merge it with information from the authored structured content (the merge algorithm). In doing so, it has the chance to validate the authored content (the conformance algorithm), for instance by making sure that the names of parameters in the authored content match those in the function definition itself. This ability to validate authored content against extracted data is an important part of the conformance algorithm.

However, the merge algorithm does not require that the authored content be part of the same file as the data you will be extracting other information from. You can just as easily place the authored content in a separate file. All you need to be able to merge the two is an unambiguous key that you can find in the source data. You then enter that key as a field in the authored content where it can be used to match the authored content to the relevant extracted data.

One of the downsides of API documentation tools like JavaDoc is that they tend to be tightly coupled systems that produce media domain output such as formatted HTML directly, often providing little or no control over presentation or formatting. This is a problem because it means that your API reference content does not look like the rest of your content. And worse, it is not integrated with or linked to the rest of your content. This has obvious consequences like mentions of API routines in your programmer's guide not being linked to the documentation of that routine in the API reference. It would be much better to generate subject domain content from the API documentation tool and then process it with the rest of your content. For many tools this is actually possible because many of them offer an XML output which may be either subject domain or document domain. Even if it is document domain, it may be regular enough that you can extract the subject domain structures reasonably easily.

---

# Chapter 24. Information Architecture

Information architecture is the arrangement of content so that it can be found and navigated. We have always had information architecture. The term is new, but not the need or the concept. In the past, though, information architecture was divided into two pieces. The basic unit of information was the book and the “architecture” of the book was an integral part of the responsibility of the author and editor. Larger sets of information were created by collecting and organizing books and that was the responsibility of the librarian or bookseller. (Libraries and book stores have different information architectures to serve different purposes.)

With the advent of online media, first in the form of large capacity electronic media such as CD-ROMs and then the Internet and the Web, this division of architectural responsibilities was overthrown. The basic unit of information in electronic media is not the book but the page. Thanks to hypertext linking, the relationships between pages in electronic media are much more complex than on paper. Also, the architecture of online media has to account for the ability to add, modify, and delete individual bits of content at any time. It is possible to think of book or library architectures in largely static terms. It is a serious mistake to think of Web architectures in this way.

This leads to the development of architectures of much smaller units with much more complex relationships to a much larger, more diverse, and more rapidly changing set of resources. These architectures include not only text and static graphics but active media: videos, animation, and dynamic feeds and information widgets. Given these factors, the old separation of roles between writer and librarian no longer works. Authors have to be much more conscious of how their pages interact with other pages in the collection, including those created by others. The scale at which these small pieces of content relate with each other is much greater than the scale at which the pieces of a book related to each other. This calls for a whole new approach to information architecture, and for the appearance of a function and a role that had no equivalent in the paper world. Thus the term “information architecture” was born, not to name something entirely new, but something transformed by our technology.

Information architecture as a discipline in its own right, as opposed to being an aspect of authorship or librarianship, has arisen to combat the chaos that emerged in many websites as they began to grow, lacking an overall organizing principle or influence. But we should recognize that information architecture is as much part of the book world as it is part of the web world, even if it was not traditionally a job title in the book world. And if you are producing both web-like content and book-like content, your information architecture has to comprehend both.

Because information architecture involves the organization of large bodies of content it can benefit greatly from structured writing techniques. Structured writing can give works of content a more definite character and identity which makes them easier to organize. By providing guidance and validation to authors it allows information architects to

better communicate and validate requirements. By making the content more accessible to algorithms, it allows the use of algorithms to do information architecture tasks, such as the automated organization and linking of content.

## Top-down vs. bottom-up information architecture

How can structured writing structures and algorithms support information architecture in structured writing?

I'm going to start with making a basic distinction between two types of information architecture: top-down and bottom-up. Top down information architecture deals with navigational aids and organizing systems that stand apart from the content and point to it. A table of contents or a website menu system is a piece of top-down information architecture. Bottom-up information architecture deals with navigation and organization that exists within the content itself. A web site with a consistent approach to hypertext links within its pages is an example of a bottom-up information architecture.

But bottom-up information architecture is not just about linking, it is about the way content is written. A topic in a bottom-up information architecture is designed to be entered via search or links from almost anywhere (as opposed to being designed to be entered exclusively from a previous chapter). It is designed not as a stand-alone entity but as a hub of its local subject space, offering readers many onward vectors according to their needs and interests. I call this approach to information design Every Page is Page One, and it is described in my book, *Every Page is Page One: Topic-based Writing for Technical Communication and the Web* (???). One of the key principles of Every Page is Page One is that a topic should follow a well defined pattern or type. Structured writing, particularly subject domain structured writing, is very useful in developing Every Page is Page One content.

Bottom-up and top-down information architectures are not incompatible with each other. In fact almost every information architecture has both top-down and bottom-up elements. (Books, for instance, which are principally top-down, based on a table of contents, may also have internal cross references, which are a bottom up mechanism.)

Structured writing can be used to drive both the top-down and bottom-up aspects of information architecture.

## Categorization

One of the key elements of top-down information architecture is categorization. An information architect develops categories of content and develops an organizational

schema (such as a table of contents) based on those categories. This may include levels of subcategories forming a hierarchical categorization scheme.

Not all categorization is hierarchical, though. In some cases content can be classified on several independent axes, allowing for the development of what is called faceted navigation. The easiest place to see faceted navigation in action is on a used-car site where you can narrow down your selection using any set of criteria that matter to you, such as selecting blue convertibles or all-wheel drive vehicles with manual transmissions.

Categorization may be implemented as part of the content management algorithm, with categories implemented as part of the external metadata that a CMS applies to a content object. This is common practice when dealing with content in the media domain or the document domain.

For content in the subject domain, however, the metadata required to assign a piece of content to a category may be inherent in its subject domain markup. It is the nature of the subject domain to describe the subject matter and any markup that describes the subject matter may already contain the fields that you need for categorization. This is one of the attractions of the subject domain: the markup can serve many purposes, which simplifies both markup design and content authoring and often means that you don't need to create additional structure to support a new algorithm.

If the subject domain information you need is not already there, you may have to add it to the design. But be careful not to reinvent the wheel. Some organizations create complex taxonomies as a basis of categorization, and the terms in the taxonomy may not always match the terms used for the same subjects in the content or the subject domain markup of that content. This does not mean that you need to manually relabel your content. You may well be able to map the two sets of terminology to each other so that your algorithms can find the same subjects under different names. Quite simply, categorization schemes are based on some aspect of reality, and if that corresponds to the aspects of reality captured by the subject domain markup of a piece of content, then categorization can be done algorithmically, even if terms don't match perfectly.

Relying on the subject domain metadata already in the content, rather than creating a separate metadata record, can be a tremendous advantage, because it makes submission of content to a repository so much easier for authors. But in some cases it can also avoid the need for a costly CMS altogether, since it allows the publishing algorithm to categorize content at build time without the need of a separate metadata store or a separate system to manage categorization.

## Linking

We have covered the linking algorithm already (Chapter 17, *Linking*), but linking is at the heart of a bottom up information architecture. In a bottom-up architecture, a page is not simply a leaf on a tree: the prize you find at the end of the search. It is a junction

point in the exploration of an information space and the quest to understand a subject. In reading a page, a reader may discover new subjects that they need to understand and new options that they need to consider. They may discover that what they thought they knew is wrong, or what they thought they wanted to do was not the right choice. They may find that their search or their traversal of the categorization system has led them to the wrong place, or they may discover whole new worlds they wish to explore. At a more mundane level, they may discover that they need additional information to complete their task, such as reference data.

These are all pointers to some next topic that the reader needs. Even the most prescient writer cannot have chosen all of them as the linear next topic in a linear narrative. To serve the reader they need to pave all of these possible paths for them, and the way you do that is with hypertext links.

This means that linking is not something that happens at arbitrary points where the author feels like adding a link. It is something that is planned for as part of the information architecture. Whether you specify hard links in the media domain or the document domain, manage them with keys in the management domain, or generate them from subject annotations in the subject domain, they should be created in a disciplined and consistent manner according to a deliberate plan.

## Tables of Contents

Tables of contents can serve various purposes depending on the nature of the work. Some describe a linear reading order for a work, some provide a classification scheme for random access to the content, some are simply a list of chapters that does not necessarily imply an intended reading order.

A table of contents may seem like a document domain structure, but it is really more of a media domain structure, for two reasons. First, it contains specific links to specific resources at specific addresses, or specific page numbers in a paper or a virtual paper format such as PDF. Secondly, it is virtually always factored out in document domain markup languages. Tables of contents are not written, they are generated.

From an information architecture point of view, what matters is how they are generated. In DocBook, for instance, it is typical to write each chapter of a book in a separate chapter file and then pull them together into a book using a book file. The order of the table of contents is then determined by the order in which the chapters are listed in the book file. The TOC itself is generated by extracting chapter and subsection headings from the chapter files in the order they appear in the book file.

In DITA, the normal process is to assemble a book using a map file. A map file may assemble a book out of DITA topics or other maps, and this may include assembling the chapters from topics as well. In the end, though, the TOC is generated in the same way, by traversing the document assembled by the map.

In both these cases, the order of the TOC is specified by hand by the person who creates the book or map file. But there are other ways to determine the order of content in a TOC.

For instance, a reference work such as an API reference may be organized by listing each library in order by name, and each function in alphabetical order by name within its library, creating a table of content with two levels. There is no need to write a map file to create this table of contents. There is an algorithm for creating this table of contents. In fact, it is the algorithm stated in the first sentence of this paragraph, “listing each library in order by name, and each function in alphabetical order by name within its library”. Here is that algorithm expressed in pseudo code:

```
create toc
 for each library sorted alphabetically
 create toc-entry library name
 for each function in library sorted alphabetically
 create toc-entry function name
```

Tables of contents serve different purposes. Some describe a curriculum, a designed reading order. Others are simply a means of navigation, a way to select one topic out of a collection of many. If your content is written in the subject domain, the chances are that it already contains the structures on which such a classification could be based, and again the TOC can be generated based on the metadata already in the content.

One advantage of this approach is that if a TOC is assembled based on metadata in the content, that means that when new content is added, it is automatically included in the TOC the next time it is generated. This simplifies the task of adding new content to a collection by avoiding the need to update multiple files or systems when a update occurs. This makes life easier for authors as they do not need to know how the TOC is constructed. They only have to create an individual piece of conforming content and submit it to the right location. This also constitutes yet another example of the single source of truth algorithm, since the basis for the content’s inclusion at a particular point in the TOC is stored only in one place.

## Lists

A major feature of a bottom-up information architecture is the list. Like tables of contents, lists are a catalog of resources. But while a TOC is a list of resources defined by their container (contents = things in a container) a list may have any principle of organization or inclusion.

For instance, you might want to have a list of all the movies starring each actor in a collection of movie reviews. Such list are not only a useful piece of information, they are also an important aid for navigating around a site. Maintaining such list by hand would

be laborious and error prone, especially with new movies being added to the collection all the time.

If you have your movie reviews in a structured format that lists the actors in the movie in a format accessible to algorithms, like this:

```
movie: Rio Bravo
 starring:: actor
 John Wayne
 Dean Martin
 Ricky Nelson
 Angie Dickinson
 Walter Brennan
```

you can generate the filmographies for all your actors, like this:

```
create-filmographies
 for each unique actor in movie/starring/actor
 create filmography actor with link to actor
 for each movie where starring/actor = actor
 create entry movie with link to movie
```

Tables of contents are a top-down information architecture device. You expect to find them at the top of the information set. Lists are a bottom-up device. You expect to find them as independent pages or as elements within a page. Thus if our collection includes the biographies of actors, and we want each biography to include the filmography, we can omit the filmography from the subject domain version of the biography and add it to the output in the presentation algorithm.

```
match actor-bio
 create html
 create h1 "Biography: " + actor-name
 continue
 create h2 "Filmography"
 for each movie-review where starring/actor = actor-name
 create li
 create a with attribute href = address of movie-review
 output movie-name
```

It is worth noting that, besides being part of the information architecture algorithm, this kind of thing is also a sophisticated example of the reuse algorithm. It takes a set of movies, each with a list of actors, and uses it to generate a list of movies for each actor, reusing existing information to create new content. This happens without any explicit reuse related markup.

# Dynamic content

Another key feature of modern web architecture is dynamic content, which means content that is generated in response either to what the site already knows about you (from your account information, or a transaction token such as a cookie, or the selections or entries that you make on the page). For example, when you log into Amazon, the first page you see is crafted for you based on everything Amazon knows about your browsing and purchasing history. As you make selections, such as adding an item to your shopping cart or wish list, that information is used to shape the next page you see.

If you browse a used car site like Autotrader.com, you can select those features of a car that you are interested in (red convertibles with manual transmission under \$20000, for instance) and the next page will be built based on that input.

The ability of a site to generate pages dynamically depends on its ability to identify content that is relevant, based on everything they know about the reader, and to assemble those pieces to form a page. For this to work, the content has to be easy to identify unambiguously, and it needs to be highly composable.

As we have seen, these properties are maximized when content is stored in the subject domain, both because the subject domain makes the relevant metadata available, and because working in the subject domain helps authors produce more consistent content that works better with these algorithms.

The consistency of the content is most important in any dynamic content application. There is no possibility for an author or editor to inspect the output of a dynamic content publication before the reader sees it, since it is assembled in real time based on the unique things we know about each reader. This requires total confidence:

- that the content conforms to its constraints
- that those constraints are completely and correctly expressed by its markup
- that the algorithm correctly processes and delivers the content

All three of these requirements depend on the soundness and simplicity of the markup design. They require precise content structures with few alternatives, clear guidance for authors, and good audit capability. Without these properties, content and its markup will be inconsistent and reliable algorithms will be hard to write and test because of the wide variety of markup combination they may encounter.

Most dynamic content applications model their content in relational database tables for these very reasons. However, with the correct markup design, almost certainly in the subject domain, there is no reason why you cannot use markup-based tools and solutions to achieve the same thing. (Dynamic content is an example of active content, which we discussed in Chapter 21, *Active content*.)

---

# Chapter 25. Change management

One important motivation for structured writing is what is often called “future proofing”. Future proofing means building a system or product with a view to making it able to survive future changes in environments or requirements. Future proofing is difficult because you cannot know with certainty what changes will occur, how likely they are, or what they will cost.

Building a future proof platform can increase up-front costs, delaying the time it takes to get to market and possibly missing a window of opportunity. Nor can you be sure that your investment will ever pay off, since the future you prepared for may not be the future you get.

But not building a future proof platform can result in you not being able to keep up with developments in a market and losing your early lead. It may require massive and expensive changes when future events render your current system obsolete. Instances of both problems abounded when traditional publication systems were confronted with the rapid rise of the Web.

The safest approach to future proofing is not to try to anticipate the particular way in which the future will develop, but to create features that will be of value no matter what happens in the future. Creating content in the subject domain is the best way to practice this kind of future proofing for content, because writing in the subject domain creates metadata that contains only true statements about the subject matter itself. Those statements are going to remain true as long as the subject matter itself remains unchanged. That is as future proof as you can make your content.

For example, suppose you write your ingredient list in reStructuredText as a table:

=====	=====
Item	Quantity
=====	=====
eggs	12
water	2qt
=====	=====

Later you decide that you want to present ingredients as a list instead. To do this, you will have to go back to your content and change the markup. Doing this across a whole collection of recipes will be expensive.

Suppose instead that you use subject domain markup:

```
ingredients:: ingredient, quantity
eggs, 12
```

water, 2qt

Now you don't have to change the content to make the change in presentation. You just change the presentation algorithm. Thus the subject domain markup has future proofed your content against this change of layout. The document domain reStructuredText markup specified the use of a table, which is not a truth about the subject matter but a decision about layout that can change independent of the subject matter. The subject domain markup simply specifies that "eggs" is an ingredient and "12" is a quantity. These are truths about the subject matter that will not change. Thus they are invulnerable to future changes outside of the subject matter itself.

Moving your content from the media domain to the document domain provides a degree of future proofing. By factoring out the formatting details, it protects your content against changes in formatting rules. Moving your content from the document domain to the subject domain provides additional future proofing. By factoring out the content and organization of documents, it allows you to target different publications and to create different document designs for different media.

Structured writing imposes specific structures on content for specific purposes. It does not make content magically immune to change nor does it guarantee you will not have to rewrite the content or change the structure to accommodate future changes in your subject matter or your business requirements.

You can, however, design your content structures to help you manage specific and foreseeable changes. If you are lucky, the structures you create may also allow you to adapt content for unforeseen circumstances, particularly if your content is stored in the subject domain. But this is a bonus. You cannot guarantee any content or structure will work for things you have not foreseen.

But changes in content happen all the time. Many of them are entirely predictable and you can use structured writing to support the management of those changes. For instance, companies re-brand from time to time. If the content is in media domain structures, the effort to change to a new appearance could be significant. If the content is in the document domain, however, changing how it is formatted is simply a matter of changing the formatting algorithm to produce different-looking output.

It is worth noting, though, that while changing the formatting algorithm is less work than changing the formatting of a large body of content, it is also more complex work. It requires a skill set that is not as widely available as the skill of changing fonts in a word processor, for instance. It also cannot be done incrementally. Once the entire new algorithm is written, all the content can be converted to the new look almost instantly. But until it is finished, none of the content can be converted. A structured writing system is not the kind of thing you can set up once and walk away from. You need to maintain an ongoing capability for making these kinds of changes efficiently and effectively.

A general move to the document domain (or the subject domain, or even a disciplined use of styles in a word processor) will allow you to handle font and layout changes. But

what if the re-branding goes further? Suppose it involves changing the names of products or even the company. Should your structured writing approach explicitly support that change? Some organization like to mandate that writers insert a variable rather than the actual name for the company name and all product names. That way, when a product name or the company names changes all you have to do is redefine the variables.

I have always been skeptical of the value of this practice. It forces the writers to remember to use the variable every time. This interrupts their chain of thought, which slows their writing down and uses up some of their precious attention, thus impacting content quality. And it is virtually impossible to ensure compliance. Writers will sometimes simply forget and write the names out normally, which means you always have to search for these instances anyway when a change happens. Then there are issues with historical usage of the names, where you don't want the change to happen, and with inflections if the new or old names end in 's' (in English; other languages may have different inflection problems).

Company and product names are distinct strings that are easy to search for when you need to make a change. The overhead of creating and maintaining the variables is greater than the overhead of doing a search and replace through the content when a change occurs. And doing a search and replace allows you to make intelligent choices about historical usage, inflections, and even changes in line breaks. If your content is held in text form (in a markup language) in a repository (file system or content management system) that allows you to do a search and replace across multiple files, this is probably easier and more reliable than using variables. (And it is what you are going to have to do anyway if there is a name change that you did not anticipate and therefore did not use variables for.)

You may well need some markup for company and product names, however. You may want to format them differently or link from them to more information about the product or company.

Rather than use a variable like this:

```
We here at >($company-name) do not recommend using
our product to catch roadrunners.
```

I would rather use an annotation like this:

```
We here at {Acme Corporation}(company) do not recommend
using our product to catch roadrunners.
```

This second approach identifies the words Acme Corporation as a company name. Creating this markup requires no extra thought from the writer. They do not have to remember what the appropriate variable name is. (They do have to remember company as an annotation type, but that is a type, not an individual name, and if your markup

is well designed your annotation types should be few and memorable.) And this same markup can be used to format the company name appropriately and to generate links to information on the company.

This does not guarantee that the writer will always remember to add the annotation, or that they will always spell the company name correctly. (There is no way to guarantee that a free-floating annotation will always be remembered. The best you can do is make them easy to do.) But you can use the company annotation to find all the phrases marked as company name, sort them, and look for variants. This then allows you to go back and fix incorrect spellings. But it also allows you to identify the ways in which writers are misspelling the company names and search the whole text for those misspellings. This improves your success rate catching both misspellings and the failure to annotate. This kind of content hygiene operation should be performed regularly on any content set, and subject domain annotation makes it easier to do while removing a distraction for writers.

At another level, re-branding can involve the organization deciding to change its tone or voice. It may wish to go from professional and reserved to friendly and jocular. There is no way, of course, for any structured writing process to recast content from formal to funny. You can't make your content proof against every kind of change.

One form of change that is so common that it may be overlooked is simply the ongoing creation of new content and the editing of old content. In the age of the Web, this is a particular concern because we can now add a new piece of content whenever it is ready, edit an existing piece whenever it needs it, and delete an old piece whenever it becomes obsolete. We don't have to wait for a major publication release for all these changes to roll out. Each rolls out when it is ready. And each time one rolls out, it impacts the information architecture or the entire content set.

Adding, editing, or deleting one topic does not mean that all the other topics are unaffected:

- There may be topics that link to the deleted topic.
- There may be topics that should link to the newly added topic.
- There may be topics that should no longer link to a changed topic, and topics that should now start linking to it.
- Topics in a category may now have a new neighbor or may have lost one.
- Any top-down navigation tools need to be updated for the topic changes.
- Deleted topics may leave holes in the information set that need to be filled.
- New topics or edited topics may mention subjects that are not adequately covered by existing topics, revealing the need for yet more topics.

- Deleted or edited topics may leave other topics orphaned, needing to be removed or edited to serve a current purpose.
- Events in the world can change the status of a whole set of topics, for instance, those relating to the current version of a product suddenly become “previous version” when a new topic is released.

When there is the potential for adding, editing, or deleting a topic to have ripple effects through the whole content set, and when such additions, deletions, and edits happen on a daily basis, it is vital to have algorithmic support for change management. Managing all the effects by hand is doomed to failure.

Content management systems often have change management features that can be helpful. For instance, many of them will inform you if a change or deletion of an existing topic will break any existing links. They will also help you find topics on related subjects or manage the membership of categories and the navigation aids that are based on them. What they won’t do is tell you things like which pieces of existing content should be linking to the new content you just added. The only way to discover that is if the existing content contains subject domain annotations that relate to the subject for the new topic. With these in place, a linking algorithm (as described in Chapter 17, *Linking*) can discover these relationships automatically.

As should be coming clear by this point, the change management algorithm is really an aspect of all the other algorithms. For any of the algorithms to keep working over time and as content changes, the structures that support those algorithms have to stay conformant. Part of designing those structures, therefore, should be thinking seriously about what it will take to maintain them in a conformant state when changes happen. Change management, therefore, relies heavily on the conformance algorithm. And this is a reciprocal relationship. To maintain conformance of your content, you have to be able to manage change successfully. But to manage change successfully, you need content to be conformant so that you can reliably identify the changes that are required and, as far as possible, execute those changes algorithmically.

As we noted in Chapter 18, *Conformance*, the heart of conformance is designing structures that are easy to conform to. Change management is not something you can tack on as an afterthought, but nor is something that necessarily requires a separate set of structures. Content that is highly conformant and highly auditable is easy to change consistently, which in turn helps maintain conformance. The heart of the problem lies in designing content for conformance.

---

# Chapter 26. Content Management

To many people, structured writing exists mainly as an aid to content management. In this book, I have taken a very different view, focusing on structured writing principally as an aid to content quality. Nonetheless, any content project of significant size requires some form of content management, and failure to manage content well will have a deleterious effect on content quality. So, any structured writing project has to take into account how the content it creates will be managed.

Content management principally deals with two problems.

- Orchestrating content for the publishing process. That is, making sure all the right bits get into the right documents at the right time. In many cases, there is a workflow aspect to this, making sure that every piece of content is seen, acted on, and or approved by every interested party in the organization prior to release. It may also include rights management issues, if there are different people with a financial stake in the content who may be entitled to compensation based on its use. These functions are well covered by other works, and I don't intend to discuss them in any detail here.
- Orchestrating content for writers during the writing process. For instance, if you are practicing content reuse, writers need to find content to reuse, and a content management system can help. Change management is also a huge part of content management: when a change happens in subject matter or in business requirements, how to you make sure all the necessary content changes are made, and made efficiently.

## Metadata is the foundation of management

Content management systems do their job largely through the collection and management of metadata. Metadata is the record of the identity and status of content. Management actions are actions on metadata: either creating and updating metadata or performing actions (running algorithms) based on metadata.

As we have seen, metadata is pervasive in structured writing. We defined structured writing as writing that not only obeys constraints, but also records the constraints that it follows. The record of those constraints is metadata. In fact, most of the metadata that a content management system manages is simply the record of the constraints that content obeys. This includes much of the metadata related to workflow, since the workflow requirements of a system are also constraints on the content.)

There is, therefore, a continuity between content management and structured writing. Both collect and manage information about the constraints that a piece of content follows. And this means that the boundaries between structured writing systems and

content management systems are fluid. Constraint metadata that is held in a CMS in one organization may be embedded in the content in another.

You might expect that the principle type of metadata contained in a CMS would be management domain metadata. After all, we described the management domain as an intrusion into the structured writing world, since it does not actually describe the structure of content. The reason for the intrusion of the management domain into content is to allow for the management of the content below the level of whatever file or chunk size you store in the CMS.<sup>1</sup>

But while you will rarely find much in the way of media domain or document domain metadata stored at the CMS level, CMS's often contain a great deal of subject domain metadata.

If you are managing a large volume of content, you will need some way to find content on a particular subject. If you are doing content reuse, for example, you will constantly be asking if content already exists on the subject you are preparing to write about. If your CMS is managing the delivery of content dynamically to the Web, it will need to respond to queries based on subject matter. And if you are optimizing your content for search you will need to provide the search engine with subject metadata in the form of keywords or microformats. All of this depends on subject domain metadata. Subject domain metadata is therefore central to CMS operations.

## The location of metadata

It is a very common pattern for a CMS to store document domain content and attach subject domain metadata to it. For instance, a CMS might store recipes written in MarkDown and attach separate metadata records to each recipe listing the key recipe metadata needed for retrieval and sorting of recipes. One of the things that writers often complain about with CMS systems is that they are not allowed to submit content to the system without filling out complicated metadata records.

An alternative approach would be to write recipes in a subject domain format in which all the recipe metadata is included in the content from the beginning. The CMS then requires no external metadata label, though it does obviously require a way to access and query the metadata embedded in the content. (CMSs based on XML databases often have this capability as a natural consequence of the XML database architecture.)

Which approach is preferable? The conventional CMS approach arises because most CMS's are based on relational databases, which are good at storing metadata records and

---

<sup>1</sup>In some CMSs, this distinction between the chunk stored in the CMS and the structures expressed inside that chunk is moot. A CMS based on a native XML database, for instance, makes no distinction between the chunk and the structure of the chunk, but treats the entire repository as a single XML resource that it can query and manage down to any level of granularity. Even with such a system, however, this distinction remains for the author, who had to deal with the structure of whatever sized chunk of content they are being asked to author.

attaching them to blobs, but are not good at storing or querying the hierarchical structure of content. It has several disadvantages.

1. It gives no support for subject-domain validation of the content. It does nothing to help improve content quality. By requiring document-domain content as the storage format, it precludes the use of the subject domain for authoring and cuts you off from all the advantages it provides.
2. The system has no way of telling if the content meets its constraints. It records the content constraints in a separate record without ever validating that the content meets them.
3. It separates the metadata from the content it describes. This allows for drift between the content and the metadata.
4. It can only record the characteristics of a chunk of content as a whole. It cannot look down into the content to find more fine grained metadata. One of the advantages of writing a recipe is the subject domain is that it allows you to do things like querying the collection of recipes for all those with a calorie count below 100. But unless the metadata record for the recipe includes that level of detail, the CMS cannot respond to that query. And if the CMS does store that level of detail, it is effectively asking the author to write the entire content twice, once in the document domain and once in the subject domain. Not only is this more work, it is quite likely that the two versions will fall out of sync with each other.

But storing metadata in the content presents some challenges as well. Having each piece of content stored in the subject domain makes a lot of sense from a semantic point and makes it easy to submit content, since not additional metadata forms have to be filled out. The problem is how to retrieve it. A CMS is essentially a database, and the way you retrieve information from a database is to write a query. A query is different from a search. A search is fuzzy. A search engine takes a plain text question of search phrase and tries to figure out which documents are the best match. Search engines may be powerful and sophisticated, but their results are essentially a sophisticated mechanical guess, and sometimes they get it wrong. Ask a search engine for a list of recipes with less than 100 calories, and it will give you a bunch of guesses based mostly on the plain text of those documents. Chances are it will catch some, miss others, and give you some false hits.

A query, on the other hand, is a precise request for items whose metadata precisely matches specified criteria. If you write a query to return recipes for which the value of the field recipe/nutrition/calories is less than 100, it will return all the results, miss none, and give you no false hits. However, it will work only for content that is stored that way. To write that query, you will need to know exactly how recipes are stored in the system.

If you have many different content structures in your repository, you will need to know how each of them is structured in order to create the queries to return them. This is not the end of the world. You can save writers from having to remember how to do all

of the queries by creating saved queries that they can run at any time. But it is still a complicating factor.

To avoid this complexity, and make querying easier to do, some systems may provide a generic metadata label that is attached to all pieces of content in the repository. This will still contain subject domain metadata, but in a less precise format. This has two downsides. First, while it is easier to write queries, because you are always querying the same structure, it can be harder to make those queries return the right content because the subject domain metadata you are querying is not held in such precise containers. Secondly, it forces writers to record the subject domain metadata in a format that does not make it obvious or intuitive what metadata to include or how to express it. This means that the metadata records will be inconsistent, which means that queries will return inconsistent results.

In the end there is no way around this. Accurate reliable queries depend on precise consistent metadata. Precise consistent metadata is specific to the object it belongs to. There is no such thing as a generic metadata record. They are always specific to the things they describe. Subject domain metadata is specific to its subject. If you want to be able to find all recipes with calorie counts less than 100, you need recipe specific metadata that specifically records the number of calories in the recipe. If you want to find a used car listing for a blue convertible, you need metadata that specifically records the car color and body style. There is no generic metadata format that supports both these queries.

## Managing the process

When we create an individual piece of content for one-time publication in a single media, there is really not much of a role for management in the process. Content management becomes a concern when you want to manage the production and publishing of many pieces of content, to manage the relationship between them, ensure consistency and quality, or to publish them many times in different ways.

Of course, many of the reasons we have looked at for moving content from the media domain to the document domain or the subject domain have to do with managing the production and publishing process. But managing a body of content and the processes and tools that create and process that content, requires a whole set of metadata of its own.

Structured writing is about imposing constraints on content. Content management is about imposing constraints on the content process. But it is also about managing the constraints we impose on content.

Doing structured writing requires recording content in document domain or subject domain structures, factoring out invariants into separate files, expressing constraints, and creating algorithms to translate the content to the media domain for publishing. All of this creates a lot of artifacts to keep track of, and requires a process both for keeping

track of them and for running the structured authoring and publishing tool chain. Thus there is a need to manage both the artifacts and the process.

Of course, all content is managed. Sometimes by organizing folders on your hard drive. Sometimes by using an elaborate content management system. But management is not part of the content itself. It is a process that exists around the content, that helps the content come into being. But it is not content.

This is not a book about content management. That would be a subject for a book in itself -- and there are several such books already. It is worth noting, however, that content management systems, particularly component content management systems, often use structured writing as a means to enable content management practices. This is absolutely a legitimate use of structured writing, though it is not one that I will focus on.

You should be aware, however, that some discussions of structured writing are based on the premise that it is being used to enable content management. In many cases, the driver for such systems is the content management methodology being used, rather than the writing process of individual pieces, which is the focus of this book. When you are looking at a particular structured writing system, particularly one tied to a content management system, it may be much easier to understand if you look at it as an enabler of content management functionality.

One prominent system that is widely used as an enabler of content management is DITA. This is not to say that DITA is only a content management enabler, but this is the motivation for a great deal of DITA adoption.

The fact that structured writing can be used both as a content management enabler and a means of improving content quality means you need to think carefully about what your goals are when you start a structured writing project. Are you looking strictly for quality improvement, strictly for content management improvements, or a mix of both?

If you are looking at both, or even if you are looking purely for content management improvements, it is important to understand that some of the techniques used to facilitate content management can have a negative effect on content quality. For instance, if you make your content more generic in order to make it more reusable, that has a negative effect on quality. If you practice extensive assemble-from-pieces reuse, this can compromise the narrative flow of the resulting narrative unit. Even though these things can be managed, managing them creates an overhead which not only costs money, it also divides the writer's attention, which can also reduce quality. Finally, the large scale use of management structures to enable content management seriously impacts the functional lucidity of your structured writing system, compromising quality and restricting your range of contributors.

As we have seen, it is possible to factor out a lot of management domain structures by moving content to the subject domain. This can help to maintain quality while making your management algorithms more reliable. There are limits to this, as this approach

does not let you do ad hoc management operations, only those operations that are the logical consequence of subject domain markup. That may not necessarily be a bad thing, since any ad hoc structures in your content present a maintenance headache that may not be worth the management effects that they implement. Alternatively, you can use some management domain structures in your subject domain markup in order to support those management functions that you can neither factor out nor live without.

## Conflicting constraints

As we have said all along, structured writing is about imposing constraints on content, and about recording the constraints applied to content so that algorithms can detect them. The first reason to impose constraints on content is to improve content quality. But as content collections grow, and as online content becomes more integrated, we also have a growing need to manage content, and to use algorithms to help manage content.

For algorithms to manage content they need to know what constraints it meets, so structured writing is a natural place to turn when you want to implement content management.

But the constraints you use to implement content management may not be the same as those you implement to improve content quality. It is easier to manage content (or anything else) if it is more uniform. The constraints that you will naturally wish to impose to make content more manageable are those that make it more uniform. Thus a system like DITA which, as a starting point, proposes that there are just three types of content (concept, task, and reference) has an obvious appeal from a management point of view.

The constraints that you impose to improve content quality, on the other hand, are those which make sure that a piece of content does just the job it is supposed to do. They are the kind of constraints that make sure that a recipe contains everything a recipe needs and is presented in the way a recipe should be presented. They are highly specific to the subject matter and to the audience. Three generic content types are not going to provide all the constraints we need to effectively manage content quality. Indeed, some of the constraints that are designed to facilitate content management may be positively damaging to content quality.

We have something of the same issue with the publishing algorithm. To manage publishing effectively is to good to have a single highly constrained presentation file format (a pure document domain description of how the content should be presented). Such a format obviously does nothing to constrain how subject matter is expressed, so it does very little to address content quality issues beyond those related to consistent presentation and formatting.

But as we saw when we looked at the publishing algorithm, you don't need to write in the presentation file format. You can write content in the subject domain and translate

it to the presentation file format as part of the publishing process. And if you are doing differential single sourcing, you may be translating it into two different presentation file formats for presentation in different media. (In fact, the desire to do differential single sourcing may be one of the motives for writing in the subject domain.)

This works fine for the publishing algorithm. The algorithm gets the consistent presentation format it needs to work reliably and efficiently, but authors still write in the subject domain and thus have the benefit of subject domain constraints to manage content quality. The consistent document domain format is just a temporary artifact created as part of the publishing process.

But for the content management algorithm, this is no solution. The content management algorithm needs to manage the original source files in which the content is created. It does it no good to manage temporary artifacts created by the publishing process. If it is to manage the source files in a system where content is created in the subject domain, it is going to have to deal with many more source formats.

But while this may to be the approach you would naturally choose if you set out to create implement a content management system -- if improved content quality were not one of your business goals -- I would argue that managing subject domain content may actually lead to better content management in the long run.

In any system that relies on constraints, on data that is known to meet certain constraints, it is necessary to make sure that the constraints are actually being met. This is the role of the authoring algorithm and the conformance algorithm. And as we have seen it is often much easier to provide effective guidance and perform effective validation in the subject domain. Also, the subject domain allows you to factor out many constraints, which is the most effective way of making sure they are obeyed. The document domain provides far fewer opportunities for factoring out constraints and providing effective guidance and is much more difficult to audit correctly.

Thus while a simple document domain system of concept, task, and reference topics meets the content management algorithm's desire for uniformity, it provides little opportunity for ensuring that the full range of constraints necessary to make content management and reuse work are actually followed. The result can be deterioration of the quality of the content set over time, a process that tends to be self perpetuating, as disorder in current content makes it harder to impose order on new content. (Just as you cannot put things away neatly in a messy drawer.)

The variety of constraints and formats found in a subject-domain system may present a greater content management challenge initially, but it can go a very long way to ensuring that the necessary constraints are met. And, as we have seen, you can often use subject domain structures to factor out management domain concerns (Chapter 7, *The Management Domain: an Intrusion*), which can go along way to removing the conflict between quality structures and management structures in content. This not only leads to more effective management, but also to a simple authoring experience.

# Content management and reuse

The reuse algorithm requires two things:

- that you create content that will fit together with other content
- that you be able to find reusable content when you need it.

Finding reusable content is a content management function.

Find reusable content requires answering the question: does this unit of content already exist? This is a complex question. It depends on what is required to definitively identify a piece of content as the canonical content unit for a particular purpose.

- Is it about the same subject?
- Does it have the same scope?
- Is it addressed to the same audience?
- Does it meet my quality constraints?
- Is it in a compatible form?
- Is it blackboxed across all uses? Could one instance change independent of the others.

It is also important to remember that answer to the question “does this unit of content already exist” will often be no. Every query costs time, and divides the writer’s attention, but only those where the answer is yes save the organization any time or money. Thus the search for reusable content adds an overhead to the writing process even when there is no content to reuse. This can have two undesirable effects:

1. It increases the cost of all authoring activities, which can create greater costs than any successful reuse is saving. It is quite easy for a reuse system to end up making an content organization less productive rather than more so.
2. It creates a disincentive for authors to look for reusable content. This can result in content being duplicated. If this happens the duplicate content might itself be reused, resulting in your systems having the same content stored twice and each stored instance being reused in other content. This not only increases costs and makes searches more difficult, it creates a problem for the change management algorithm because it raises the possibility that only one of the instances will be changed when a change occurs, resulting in the change being missed in the other reuse chain. If enough inconsistencies like this accumulate, people can start to lose confidence in the whole system.

You can greatly reduce the impact of this problem if you can give authors a clear idea of when they should expect to find reusable content and when they should not.

## Creating manageable content

We saw with the conformance algorithm that the key to conformance was to create structures that are easy to conform to. The content management algorithm relies heavily on the conformance algorithm. Successful management of content depends on knowing exactly what assets you are managing. The more you know about each asset, and the more reliable what you know about that asset is, the more confidently you can manage it, and the less likely the management system is to slowly descend into chaos.

What this means is that the content management algorithm depends on content that is written being easy to manage. It may seem like simple generic units of content would be the easiest to manage, but the problem with such generic units is that you know very little about them and what you do know is often unreliable. Generic units may be easy to create and easy to store, but they are not easy to manage. It may require more initial thought and effort to plan for the management of highly specific well constrained content units, but such units will in fact prove to be the easiest to manage over time, especially as your content set grows.

---

# Chapter 27. Collaboration

Many organization produce large information sets that cannot be produced and maintained by a single person. Thus collaboration is required in producing and maintaining all that content. Collaboration is a difficult and complicated process, and the more people you have trying to collaborate, the more complex it becomes. How does each collaborator know what others are doing? How do they know which parts of the wider work they are responsible for? How do they integrate their work with the work of others? How do you manage the overhead created when collaborators have to be aware of other people's work as well as their own? As always, any complexity that is not handled in the organization gets dumped on the customer in the form of quality and findability problems.

The fundamental constrain on collaboration is the amount of time that collaborators have to spend on collaborative activities -- orienting themselves to the work of others -- as opposed to creating new work of their own. Unless you do something to ease this burden you quickly reach the point where adding more collaborators actually slows the project down because every new person you add to the project increases the collaboration overhead for everyone already working on the project, which means they have less time left for new work. At a certain point, therefore, the next person added takes more new-work time away from every other contributor than the new-work time they add to the project.

The essence of managing collaboration, therefore, is to minimize the amount of collaborative overhead in the system. But unless you do this in a way that allows all of the necessary coordination of work to get done, the unmanaged complexity gets dumped on the customer. Before the Web, companies handled the overhead of collaboration largely by assigning different books, pamphlets, and other publications to different writers and issuing them with some basic style guidance for language and physical appearance and layout. All the other aspects of collaborative complexity, like making sure that everyone was saying the same thing at the same time in the same way, or that people were not creating the same content over and over again, or that at least one person was saying everything that needed to be said, were ignored and the results were dumped on the customer in the form of inconsistent information that was incomplete, incorrect, contradictory, and hard to find.

With the advent of the Web, all the same content, produced by the same non-collaborative approach to collaboration, got dumped onto the company website. It wasn't necessarily any better or worse than it was before, but because it was now all searchable in one place, the quality problems became much more obvious. Early attempts to allow employees to post content to the Web, however, arguably made things worse, since they combine a technically difficult process with a complete lack of support for collaboration on coverage, consistency, or style. Out of this mess, the discipline of content strategy was born.

As always, we can address these problem by moving the complexity around to make sure that each part of the burden is placed on a person or process best able to handle it. Structured writing can help by enabling your to more effectively redistribute the complexity while reducing the amount that slips through the cracks.

DITA, as a language, moves a lot of the complexity of content management, and therefore of collaboration, into the source file in the form of management domain structures. If the writer is working directly in DITA, therefore, that complexity is transferred to the writer. In a collaborative situation, this means that that part of the burden is being transferred onto the collaborators, which means that it takes more time, discipline, skill, and knowledge to be an effective collaborator. Not all DITA-based systems do this, however. Some systems -- off the shelf or customized -- may put an interface over DITA that makes it simpler for collaborators to contribute. This means they are transferring the complexity from the writer to the interface, and therefore the person who designs the interface. However, if the interface does not capture the full sophistication of the DITA language underneath, resulting in certain functions not being supported, it is actually distributing complexity to the reader again.

Don't think that all authors in a collaborative system have to be treated the same way. Some of your contributors may be full time professional writers while others are engineers or marketers, field personnel or support people for whom communication is an important, but not central, part of their jobs. It makes perfect sense in these situations to design a collaborative system that distributes authoring complexity from your occasional contributors to your full time writers. People can handle far more complexity in their core task than then can in any of their peripheral tasks.

This kind of approach can be very fruitful, but it is not well supported by most off the shelf tools which essentially present the same interface to everyone. A structured writing approach in which different contributors use different structured writing languages, each best suited to their contributions, can be very effective. Of course, any use of structured writing distributes some degree of complexity towards the writer, since they are not obliged to know and follow the structure. On the other hand, this can also distribute a lot of complexity away from the writer. While the blank page may seem like the simplest possible interface, it actually give no task guidance at all. It is an interface without any affordances. It leaves it to every author not only to write, but do design the information they are creating, to decide what needs to be said and how to say it. Structured writing can provide this design information to the writer, thereby distributing the design complexity to the information architect who designs the structured writing language. What is vital here, though, is that the structured writing language not distribute any other complexity to the writer. Any language that requires them to master publishing or content management concepts, for instance, is not going to work well for this purpose. What works is a simple subject-domain language the addresses the reader in terms they already understand and ask for annotation using concepts and ideas that the writer already knows.

One of the more common reasons for introducing content management and/or structured writing to an organization is to improve collaboration. It is certainly not the only way

to facilitate collaboration. In fact, the more common approach is to create dead simple and largely unstructured tools such as message boards and wikis. One of the most widely collaborative projects in the world -- Wikipedia -- runs on a wiki using a fairly simple document domain markup language which is often hidden behind a simple WYSIWYG editor.

Why then might one turn to a more complex structured writing system for collaboration. In a word: integration. The kind of collaboration that is supported by message boards and wikis is one in which collaboration chiefly means everybody being able to see what everyone else is doing. Any connections between the pieces created by different people are loose and non-critical, mostly taking the form of hypertext links. And such connections are there are are manged by large scale community efforts. Wikipedia is full of links betteen articles largely because anyone can go in an edit an article to add a link to an article on a related subject.

But not all collaboration can rely on such loose and uncoordinated activity. In many cases you need to bring the pieces created by collaborators together to form an integrated and cohesive whole. You can do that by hand, of course, but that can be cumbersome and time consuming. It may be hard for any one coordinator to keep up with all the content that is being created, especially if new content is continually being written, without any freeze period to allow the integration to take place.

Equally important, the collaborators on a message board or a lager wiki like Wikipedia are largely ignorant of each other and each others activity. Duplication of effort and even outright contradiction may be frequent. If you have huge amounts of volunteers constantly reading the site and finding duplication and contradiction, as Wikipedia does, you can live with this (though there will always be parts of the system that are in error at any give time). But an organization that is paying its writers and editors may not be able to afford this slow and effort-intensive approach. It may need a way to coordinate the activity of its collaborators to avoid duplication and error.

## Bridging silos

There is a lot of talk in content management circles about breaking down content silos. The naive way to do this it to have everyone use a single system and a single markup language. But as we have seen, this means either adopting a simple document domain language that everyone can learn, like Markdown, which does not have enough structures to meet everyone's needs, or adopting a large complex document domain language like DITA or DocBook that meets a lot of needs but has poor functional lucidity, especially for part-time contributors.

The most pernicious myth about collaboration is that it requires everyone to use the same tools and to understand each other's work. In fact, this is the worst way to collaborate, because it create a huge amount of overhead which can swamp the ability to actually get anything done. Efficient collaboration is actually achieve by limiting the amount that

collaborators have to know about each others work and each others tools. This allows each group or individual to work efficiently while still creating a product that can be integrated successfully with the work of other.

This approach to collaboration is seen through the worlds of engineering and computer programming. The secret ingredient that allows workers to collaborate with minimal knowledge of each other's work is the interface.

## **Expanding your pool of authors**

It is often desirable to bring other authors into the fold. However, as soon as you get outside your pool of professional authors, the ability to teach them to do complex publishing or content management tasks is greatly reduced.

Groups who want to involve software developers in the authoring of documentation, for instance, often turn to simple document domain markup languages such as MarkDown or WikiMarkup (using a Wiki as a collaboration platform).

Functional lucidity is the key to expanding your pool of authors.

---

# Chapter 28. Timeliness

Information tends to change quickly these days, and readers no longer have any patience with outdated content. But there are many difficulties in ensuring that content is always timely. How do you detect when content is out of date? How do you push updated content to the reader quickly? How do make sure updates in one place don't break other content? Structured writing provides ways to address all of these questions.

There is no separate structured writing algorithm for timeliness. Timeliness rests on executing the entire writing process, from ideas in the head to dots on a page in the hands of a reader, in the shortest time possible, while maintaining all of your quality goals. Thus all of the structured writing algorithms can contribute to timeliness.

## Quality

The maintenance of quality is key to timely delivery. Any way in which you can use algorithms or structures to assist the writer to produce quality content therefore helps your deliver in a timely manner. In particular, if your structured writing formats are specific about what pieces of information are required and factor out not only formatting but presentation, your can accept new content and publish it very quickly with minimal concerns about maintaining quality.

## Authoring

Functional lucidity in structured writing design helps authors work faster and with more confidence. Factoring out any non-content concerns, particularly those related to content management and publishing, will help maintain quality whey improving delivery speed.

## Separating content form formatting

Formatting content by hand takes time and is subject to unintended variation through human error. If you want publish quickly and with consistent formatting, you need to factor out the formatting from the content and hand the formatting task over to a formatting algorithm. This will allow you to shorten the time it takes to publish by making formatting and output virtually instantaneous once the writing is finished. It also helps speed up the writing process and improve content quality by avoiding the division of the writer's attention between writing and formatting. However, it is important to make sure that the structures you create to factor out formatting are

not more complicated for authors to deal with than the formatting they replaced.

### Single sourcing

If you are delivering content to multiple media, you need to avoid having to separately prepare the content for publication in different media. The less manipulation of the content you have to do to output to each target media, the better. If you want to maintain equal quality across different media, taking a differential single sourcing approach will ensure high quality in each media without slowing down the publishing process.

### Content reuse

Content reuse can improve timeliness by avoiding the need to create content twice. If a new piece of content can be created by pulling in pieces of existing content, they may allow you to get it out faster. Be careful not to assume that this is an automatic win, however. The reuse process itself takes time. A lot of its payoff comes not from reducing end to end authoring time but from avoiding re-translation of content. Complex reuse systems also require the work to be done by authors who are conversant with the tools and with the content set (so that they can find reused content efficiently). The availability of those authors may impact your ability to move quickly. Reuse can be a timeliness win under the right circumstances, but don't lose track of the fact that simplicity can often be a key virtue when you need to act quickly.

### Composition

If you are going to use content reuse as part of your timeliness strategy, ensuring the composability of your content becomes doubly important. If you want to deliver reused content quickly, you need the pieces to go together seamlessly and reliably. But don't forget that it is also important to maintain a cogent narrative flow across a piece of assembled content. Avoid composability tactics that may compromise the quality of the finished piece.

### Single source of truth

Where it can be achieved, a single source of truth can be a huge win for timeliness. If you can assure

that there is only one source for a particular truth (however you define truth for this purpose) then you simply deliver it where required.

### Publishing

Automating publishing can allow you to release content much more quickly. This is particularly important if you are going to release content to multiple media or if the changes you need to make affect multiple pieces of content.

### Linking

One of the most challenging aspects of adding new content or removing old content in a timely fashion is managing the links. The new content should link to all relevant content in the current content set, and the current content should link to the new content where appropriate. When content is removed, all the content that links to it should be updated, either to remove the link or to link to something else. Any structured writing approach that manages links will help with this, but by far the most efficient way to deal with it is to use the soft linking approach detailed in Chapter 17, *Linking*.

### Conformance

Conformance is fundamental to publishing new content quickly while maintaining quality. The faster and more reliably you can ensure the conformance of content to its governing constraints, the faster you can publish it. Structured writing support for the conformance algorithm thus provides a direct benefit to timeliness.

### Auditing

A huge part of timeliness is knowing when new content, or edits to old content, are required. The world changes all the time, but unless you have an efficient way to determine the impact of those changes on your content set, you can't respond to those changes in a timely manner. Taking a structured writing approach to auditing, as outlined in Chapter 19, *Auditing*, can go along way to ensuring that you always know when changes need to be made, and that you can find all of the content that needs to be changed.

### Extract

The extract algorithm is not just a great shortcut for generating content, it is also a great way to keep up with changes in the real world. If the data source that you are drawing content from is updated, simply repeating the extraction and republishing the content will bring it up to date. If you want to be really up to date, you can also run the extract algorithm dynamically, only pulling content from source when the reader requests it. If you have a fully automated publishing process that does not require any user intervention, you can effectively publish live data right out of the source file.

### Merge

In many cases the extract algorithm works with the merge algorithm to combine content extracted from another source with content written by authors. In this case, you can use changes in the source as a flag to tell you when changes in the written content are required.

### Information architecture

Changes to subject matter don't just affect individual pieces of content. They affect the overall information architecture. If you know that your content is going to change on a regular basis, it makes sense to create an information architecture that is highly adaptable to change. The more of your architecture that can be generated by algorithms, the easier it will be to make architectural changes in response to change. But having your information architecture implemented by algorithms requires structured writing techniques that enable the architectural routines to access the aspects of the content they need to organize it. A bottom-up information architecture is particularly effective at handling rapid and constant change. Implementing a bottom-up architecture, which relies heavily on well structured topics connected by rich linking, benefits enormously from structured writing techniques. By changing the way content is organized and linked, these techniques can allow you to add and remove individual pages from a content set without fear of breaking things.

### Change management

While the audit algorithm will alert you to the need for changes in a timely manner, the execution of those changes still needs to be managed in order to respond in a timely manner without compromising quality. Using structured writing techniques to facilitate change management can help ensure that the required changes get made quickly and reliably.

### Content management

Every change in a content set is a content management action. Unless you design your system very carefully, the cost of each content management action will increase as the size for the content set grows, as each action has to take account of the possible impacts on a greater number of resources. Using structured writing techniques to automate aspects of content management can help avoid the cost of content management escalating as the content set grows and can make each individual content management action easier and more reliable.

### Collaboration

Collaboration can do a great deal to improve the timeliness of content by making more authors available to contribute content and by moving content creation close to the source of information changes. Structured writing can be a double edged sword when it comes to collaboration. By automating publishing and information architecture, it can make it much easier to coordinate the efforts of multiple contributors. By supporting composability, it can make sure that the contributions of different collaborators work well together. And by enforcing content constraints, it can help assure that all contributors create content that does the job it is supposed to do. At the same time, however, many structured writing systems are difficult to learn and use, making it difficult for collaborators to contribute. They can also be expensive, making it uneconomical to include occasional contributors as authors. Both of these problems can be avoided by focusing heavily on the functional lucidity of your structures and markup, which not only avoids the heavy learning curve of structured

writing, but removes the need for complex editing tools.

While there is no separate algorithm for timeliness, the way in which you select and integrate all of the other structured writing algorithms that you use has a big impact on your ability to deliver in a timely fashion. Optimizing individual pieces may not gain you the time you are looking for, or may not maintain your quality standards while delivering quickly. You need to think careful about where bottlenecks can occur in your overall process and about any points in the chain where quality could be lost.

While much of your content structure design activities will be done in terms of the specific algorithms you want to support, a concern with timeliness demands that you think about how those structures affect the timeliness and quality of the entire process. Complex document domain/management domain reuse structures, for instance, may save time by avoiding duplicate work, but you need to think about the time it takes authors to use these features, and any quality traps that lurk in their use. You also need to think about whether the use of such features could leave you overly dependent on particular writers with particular training to make changes when they are needed. The most technically efficient system is of little value if the only person who can execute it is on vacation. All structured writing systems require special skills in certain key roles, but it pays to avoid setting up your system in a way that creates skill bottlenecks for executing short term content changes.

One of the key features of the subject domain is that subject domain structures support multiple algorithms with the same markup. This can be an enormous benefit in ensuring timeliness. Not only does it simplify the content changes that need to be made to respond in a timely fashion, it hands more of the management and production phases over to algorithms, who are always faster than people. It also means that content updates require no knowledge of the management systems, which helps avoid a skills bottleneck for content changes.

---

# Chapter 29. Repeatability

Content is one of the hardest products to test. One of the biggest barriers to consistently producing quality content is being able to test that quality has actually been achieved -- that is, that the content is serving and/or influencing the reader as you want it to do. The problem, quite simply, is that it is hard to observe the effect of content. What you can't observe, you can't measure, and what you can't measure, you can't manage.

A second problem is that even when you do get to observe the effect of content on the reader, how do you know what aspects of the content achieved the created the effect you observed. Making these observations on individual pieces of content is of limited value. Yes, you can do continual A/B testing, taking down content that does not perform well and continually trying other things until you finally create a piece of content that performs well. But then how do you reproduce that success with the next piece of content you write. How do you determine which aspects of the successful content contributed to its success and should be emulated in future content?

The fact is that when you ask this question about a single piece of successful content your answer is little more than an educated guess. You are abstracting from a single data point. To be certain which aspects of the content are doing the most to achieve it aims, you need to observe multiple samples that exhibit the same basic features. It is when you observe that multiple pieces of content exhibiting the same features are all successful that you know with some degree of reliability that it is those features that make the content work. Once you know this, you are in a position to reliably produce new content that will be similarly successful. You have achieved repeatability.

To achieve repeatability, you need structured writing. More specifically, you need structured writing that is focused on content quality. In order to generalize the measurements you are taking, so that you can draw conclusions that apply to more than one piece of content, you need to make sure that each piece of content you are measuring has the same structures and features. That is the only way you can measure if it is the structures and features, rather than the individual texts, that are performing well. You need consistent conformant content structures that express those features of a text that impact is quality. Until you have that, you have no reliable basis on which to extrapolate any measurements you do of individual pieces of content.

Once you have confirmed that these structures contribute to content quality, you need to insist on them in all the new content you create in order to ensure that it conforms to the constraints and structures you have established through testing. In order to achieve repeatability in content creation, therefore, you need structured writing.

The most fundamental part of content quality is to give the reader the information they need in a form they can use. A confounding aspect of content quality is that every reader is different. They are doing different thing, they have different experiences, and they have different vocabularies.

## Repeatability

---

For many writers in corporate environments, lack of knowledge about the readers, their tasks, and their backgrounds, is one of the biggest problems in determining what content is required and what form it should take. While direct contact with the customer is undoubtedly the best way to address this, you can actually learn what content to create and how without every meeting the reader, as long as you can measure the performance of your content, generalize your results, and repeat the structures that perform best.

---

# Chapter 30. Conversion

There is one final algorithm we should look at, not because it is part of the structured writing process, but because it is a big part of how many organizations make the move to structured writing: the conversion algorithm. The conversion algorithm is how you get content from an unstructured system to a structured system.

There are lots of products that advertise that you can easily and cleanly convert all your existing content to a new structured writing system like DITA. Such claims should be taken with a large grain of salt.

Structured writing is about constraints. A piece of structured content, therefore, is one that follows a set of constraints and records the constraints it follows.

Of course, it is possible that a piece of content may follow a set of constraints without recording the fact that it does so explicitly. As we shall see in Chapter 39, *Patterns* it is also possible for a structure to be discernible in the pattern of the text used to express the content. For instance, a date expressed in the format December 13, 2016 follows a pattern that is easy to describe and to detect mechanically in text. You can fairly easily write an algorithm that detects dates in this format in a text and replaces them with an explicitly marked-up date format such as <date><month>12</month><day>13</day><year>2016</year></date>. Other ways of detecting implicit structure could be formatting of the source text. For instance, you might be able to detect that the use of a heading style such as H2 was the start of a section, even if there is no explicit section structure in the source text.

If you can detect these implicit structures reliably, you can transform them into explicit structures. Thus it is possible to write an algorithm to convert a non-structured document into a structured document by recognizing implicit structure and turning it into explicit structure. However, for this to be successful, certain conditions must be met:

1. The source document must already obey all the relevant constraints of the format you are converting it to. You can't convert a document to an explicitly structure if it does not already follow that structure implicitly.
2. The structures in the source document that you want to make explicit must be detectable as patterns in the text of formatting of the source document. In some cases that governing constraint may have been followed but there may not be any way to detect that fact automatically from the text. In many cases, subject domain annotations clarify the subject of an otherwise ambiguous word, and no algorithms short of a sophisticated AI is going to be able to detect which of the meanings of the word was intended.

Of course, a conversion does not have to be complete to be useful. It is common practice to have an algorithm convert those structures it can detect and then to have human authors

do the rest of the work. This may include not only finishing the markup of a piece, but also editing it to make it meet its governing constraints.

But don't let anyone deceive you about the quality for such conversions. There are three issues to be aware of in particular:

1. If there is a structure in the source that the conversion algorithm cannot detect, it will treat it as a generic structure. The resulting file will validate according to the schema of the target language, but the markup will be incorrect with respect to the structures and the text they contain. The fact that the file validates can hide these errors from you.
2. There may be cases where the conversion algorithm detects a structure but cannot place it in a valid place in the current document (because the source does not follow the applicable constraints). In these cases, the conversion process should produce an error so that you are made aware of the aberrant structure. However, if you conversion vendor boasts of their conversion success rate, they may be tempted to ignore the error and mark the structure up as something generic, thus hiding the problem.
3. If there are any structures that are optional in your target language (most annotations are options by their nature) then the conversion does not have to detect them to create a valid output file. Catching all those that could be detected is extra work for the person writing the conversions algorithm. Be careful that you are getting as many of these detected as possible.
4. Many conversion services are merely offering generic format to format conversions such as Word to DocBook or Flare to DITA. These are necessarily lowest common denominator conversions that almost certainly add no real value to the text. They may get you through a change of tools, but they are unlikely to help you effect a change of process.

The limits of such conversions should be fairly obvious based on our review of the structure writing domains and the structured writing algorithms. Much of the practice of structured writing is about factoring out invariants in text. While an algorithm can sometimes detect implicit structure and turn it into explicit structure, it is much harder to write an algorithm to reliably factor out any but the most trivial invariants (the bullets on list items, for instance). You certainly should not expect that a conversion algorithm can move content successfully from the media domain to the subject domain.

Any given markup language requires certain constraints to be followed. Content cannot validly fit into that markup language if it does not follow those constraints. So, there is no mechanical process -- no algorithm -- that converts content that does not follow the constraints of a markup language and make it follow them.

At best, such a conversion is going to produce a file that validates against the schema for that language, but that validation is going to be a lie. And, of course, conformance that is a lie is of no use to any structured writing algorithm. The only way to bring content

that does not conform to the target constraints into conformance is to rewrite it. Anyone who tells you different is selling snake oil.

Finally, we should note that the move to structured writing is often prompted by a change of outputs. As more of an organization's content moves from paper to the Web, there is a entirely appropriate desire to move to a different tool chain for producing content. But in this case is is vital to remember that the difference between paper and the Web is not merely one of format. If it were, dumping PDF's of your books on a Web server would be the only change you would need to make. But the move from paper to the Web involves a fundamental change of information architecture, one that goes deep into the content itself and how it is designed and written. In many cases, structured writing can be a huge help in executing a suitable information architecture for the web, including such important consideration as automating or managing linking. But you cannot automatically convert content from one information architecture to another. Preparing content for a bottom-up information architecture where content is located primarily by search and navigated primarily by links, requires that content be written in a very different way. Automated content conversion simply cannot achieve this for you.

There are therefore just two significant use cases for automated conversion tools:

1. Where you are transferring content from one format to another at more or less the same place on the structured writing spectrum, and the structures in the source format are more or less the same as those in the destination format (though they may be less explicitly recorded).
2. As a first step in what you know will be a largely human-executed conversion project. The chief value of such a conversion is often to get the content into the tool set that you are going to use for structured writing going forward, as a first step to actually applying structure to it.

---

# Chapter 31. Exchange

One of the earliest motivations for structured writing was portability. This was in the days when there were many different computing platforms all using different data formats. Moving a document written on one platform to another platform while retaining any kind of formatting or structure was a major challenge.

To enable the algorithm of porting a document from one system to another, structured writing techniques were used to move the representation of the document away from a series of embedded typesetting codes towards more abstract structures to which style information could be attached (separating content from formatting). The interpretation of this structure and style information into platform-specific typesetting codes or screen display instructions (the media domain) could be done separately on each platform.

System portability is now much easier because the number of platforms has decreased and much of the low level representation of data has been standardized (thought ASCII and Unicode, for instance). Major applications, or compatible equivalents, now run on all major platforms. The reason they can do this is because the files are structured to factor out specific display instructions for different platforms: structured writing at work.

Portability remains a concern, however, but now it has more to do with portability between organizations. How does one organization send content to another organization in a format that it can use?

The problem here is how exactly the receiving organization wants to use the content. Making content exchangeable is not difficult in itself. The most common formats for exchange between organizations are PDF, HTML, and Microsoft Word. There are PDF viewers for every platform. Every platform has browsers that can read HTML. Word files can be read by a variety of applications on virtually every platform while retaining all or most of their source formatting.

If what the receiving organization wants to do is print and read the content, the PDF format fits the bill. If they want to be able to edit it while making minor modifications to formatting, Word or does the job. If they want to edit it and put it on an internal or external Web server, HTML is fine.

But if the receiving organization want to take the received content, manipulate it, query it, reorganize it, and include it in their own publication system, these format are not structured enough to meet the need. Or at least, it will involve a lot of human execution of algorithms because there is not enough structure in the content to hand the job over to machines. This may be fine for occasional or ad hoc transfers, but for large or regular exchanges of content it can become inefficient and error prone.

A commonly suggested approach is to exchange content in a common or standardized document domain format. Some common formats suggested are DocBook and DITA.

But this is not as easy as it sounds. DocBook and DITA are both large and complex formats that different organizations may use in different ways. Different organizations may specialize either format, and while specialized version should still be formally compatible, they may not process as intended in the recipient's system, especially if one organization has chosen to interpret the standard in different ways -- or to violate it in order to achieve a specific business goal. It is impossible for a standard to express all the constraints that individual businesses or individual writers want to use, so standards are inevitably used in non-standard ways that compromise portability.

Exchange, in other words, is a composition problem and all the levels of composability that we discussing in relation to the composition algorithm (Chapter 14, *Composition*) apply to exchange. Choosing an exchange format merely because it is common or even a standard does not fully solve the exchange problem unless that format is also highly composable.

As we noted, composability requires a high degree of constraints, particularly if you want to achieve full literary composability. This is not to say that you cannot use one of these standardized languages to exchange content with a particular set of partners, but you will probably need to agree on an additional set of constraints beyond those provided by the standard format, and figure out how to enforce and validate them.

Of course, the best way to enforce and validate content constraints is with a structured writing format that either enforces them or (ideally) factors them out. This, of course, requires that both parties agree to create and accept such formats. As we noted, in the case of composability this does not have to mean using the same format, just using formats that are semantically equivalent or can be used to generate a semantically equivalent output.

Some organizations choose formats for their internal work on the basis of those formats being exchangeable with others. As we have seen, different structures support different algorithms. The ideal format for a given organization is the one the supports all the algorithms they want to run on their content. If exchange is the primary goal for adopting structured writing, or if the form that support exchange also supports all the other algorithms you are interested in, choosing a common public format makes sense, since it saves you the cost of any custom development.

But such formats tend to be large and loose and therefore have limited validation and audit capability. This can compromise composability at several levels. It is perfectly possible for one organization to send a file in a public format to another organization that uses the same format and have the exchanged file not be composable at all with the receiving company's content.

A common format, in other words, does not guarantee composability and therefore does not guarantee successful exchange. Different organizations may use a common format like DocBook, for instance, in very different way, rendering their files non-composable with other company's DocBook files. And, of course, other companies may use a different format altogether. To successfully exchange content with another

company, you will need to deliver content in the format they use, and in the way they use it.

The capacity to exchange content, then, does not lie so much in having the same format as everyone else, as it does in being able to deliver content in whatever format your customer wants to receive it. And, of course, in having a supplier that will deliver content to you in whatever format you want to receive it in.

We have seen that moving content from the media domain to the document domain helps you deliver the same content to different media or to format it differently. We have also seen that moving your content to the subject domain makes it easier to create different forms of document organization and behavior. The same is true for exchange. An exchange format, in the end, is just another output. The ability to create multiple exchange formats comes from exactly the same techniques that let you create different document designs and output formats.

Portability is often advanced as a principal argument against developing your own structured writing structures. But in fact portability is often best served by custom formats that enforce or factor out the constraints you need to port content reliably to other organizations.

---

# Chapter 32. Translation

Content often needs to be translated into multiple languages. There are two ways to approach translation when you practice structured writing. One is to use your entire structured writing process to produce a finished first-language output, and then sent that output to a translation company to be translated. This has the benefit of simplicity, but in every other way is unsatisfactory. The result is not timely, and may involve translating the same material multiple times, even if the content has not changed between versions or different uses of the content.

In practice, therefore, it makes sense to translate your structured content at the source level and then apply all of your structured writing algorithms to the translated versions of the content to produce all your second-language versions. In other words, your translated content is structured content in the same way as the original content.

There are two main consequences of this.

All of the structured writing algorithms are available to govern and assist in the translation process.

Functional lucidity for translators is just as important, if not more important, than for authors. If translations are performed by people outside of your organization they probably don't spend all of their time working in your structured writing system, and they almost certainly don't know anything about the specifics of your content management practices. ...

## XLIFF

There is a standard XML vocabulary for dealing with localization data. It is called XLIFF (XML Localisation Interchange File Format). Its purpose is to improve functional lucidity for translators by hiding the complexities of source formats from them. A description of XLIFF, or how it might work with your structured writing system is beyond the scope of this book, but if translation is an important consideration for you, it would make sense to talk to someone with a background in localization as you design your system.

---

## **Part III. Structures**

---

# Table of Contents

33. Rhetorical Structure .....	234
Irreducible rhetorical models .....	236
Rhetorical metamodels .....	238
Meta-models vs generic models .....	240
Making the rhetorical model explicit .....	242
Structure and randomness .....	242
34. Mechanical Structure .....	244
Flat vs. nested structures .....	245
The role of syntax .....	249
Agreeing on boxes, names, and labels .....	250
Constraint and annotation .....	252
Secondary structures of interpretation .....	254
Elements vs. attributes .....	256
35. Metadata .....	261
The recursive nature of metadata .....	262
Where should metadata live? .....	263
Taxonomy and controlled vocabulary .....	266
Top-down vs. bottom-up terminology control .....	267
Local and global names .....	271
Ontology .....	272
36. Blocks, fragments, paragraphs, and phrases .....	274
Semantic blocks .....	274
Functional blocks .....	277
Narrative blocks .....	278
Fragments .....	279
Paragraphs and phrases .....	280
37. Wide Structures .....	283
Tables .....	283
Alternatives to tables .....	287
Alternate presentation .....	288
Semantic structure .....	288
Record data as data .....	289
Don't use tables for layout .....	289
Code .....	290
Pictures and graphics .....	291
Inline graphics .....	297
38. Subject structures .....	299
39. Patterns .....	305

---

# Chapter 33. Rhetorical Structure

Structured writing is an attempt to improve the quality of written work. Its intended outcome is a piece of writing that works well for its intended purpose. All writing has structure, from the basic grammatical structures that make sentences comprehensible, to the larger rhetorical structures that make information accessible or frame an argument cogently and persuasively.

While structured writing is not an ontology and does not attempt to express the actual information in a piece of writing, it does aim to support the creation of an effective rhetorical structure. There is, therefore, a relationship between the rhetorical structure of a document and mechanical structure (Chapter 34, *Mechanical Structure*) created by markup.

The rhetorical structure of a piece of content is how it tells its story. For many types of stories, the optimal rhetorical structure is quite consistent and often well known. In other cases, the best rhetorical structure can be determined both by a careful consideration of what needs to be said and by experience and testing with readers. Content quality is greatly enhanced when the rhetorical structure is well defined and followed consistently. Also, a well defined rhetorical structure provides an effective baseline against which to compare and measure proposed improvements. Using an explicit predefined rhetorical structure helps enhance and maintain content quality.

In Chapter 6, *Writing in the Subject Domain* we looked at how structured writing in the subject domain can capture the rhetorical structure of a recipe in various ways to serve various business purposes. A recipe is a fairly well known type of information. The various parts can be organized and presented in different ways, but most people recognize a recipe when they see one. This makes a recipe an example of a rhetorical pattern. Rhetorical patterns are likely to be common in texts from many different sources. In many cases, readers will recognize content that conforms to a rhetorical pattern more easily.

Individual organizations, however, may require more specific structures within instances of the rhetorical pattern. As we saw in Chapter 6, *Writing in the Subject Domain*, a wine magazine may require every recipe to have a wine match. A health-oriented magazine may require every recipe to contain a complete set of nutritional information. Other organizations may have specific requirements about how recipes are to be presented, such as requiring ingredients to be presented in a table rather than a list. The specific set of requirements of an organization -- their unique constraints -- constitute a formal content structure.

In some cases, rhetorical patterns are immediately obvious because they have a visual shape. The various components of a recipe just happen to look physically different on a page (which is why recipes are the most popular structured writing example). There is

the picture; the introduction, which is a block; the ingredients, which are a list; and the steps, which are a numbered list. The recipe pattern is visually distinctive even without looking at a word of the text.

However, rhetorical patterns are not about elements that are visually distinct. They are about the different types of information that are required, the way they are expressed, and the order they are presented in. There may be considerable variation in the second two properties. Whether we would count these variations as options within one rhetorical pattern or as defining different rhetorical patterns should probably depend on their effect. Any organization and means of expression that has the same rhetorical affect we can reasonable count as variations on a single rhetorical pattern.

When you look at a page that appears to be just a sequence of paragraphs with perhaps some subheadings thrown it, it is easy to assume that there is no particular rhetorical pattern present. But this is not necessarily true at all. If a consistent set of information is being presented for a particular purpose, and we can find (or reasonably imagine) that same set of information being assembled for the same purpose to describe another object of the same type, then we have a rhetorical pattern. And where we have a rhetorical pattern, we can define a structured content type for a specific business purpose.

This is not to say that once you define a structured content type, you will be able to take the markup that you define and wrap it around every existing example of the rhetorical pattern without changing a word. (Again, the point of structured writing is to improve content, not to faithfully represent its current state.) Structured content types are more precise than rhetorical patterns and existing unstructured topics that follow a rhetorical pattern will almost always have to be edited to fit a more precise structured type.

What you will find when you start to move content that meets a rhetorical pattern into a defined structured type is that a lot of the content does not fit the rhetorical pattern particularly well. You will find some instances that only partially fit, but which omit information commonly found in the pattern (and perhaps required in the structured type). You will find that some instance contain information not found in most instance of the pattern, and not supported by the structured type. You will find information not expressed in the way that the structured type expects.

These discoveries mean one of four things:

- The discernment of rhetorical patterns in incorrect and you are trying to make content with a different pattern fit your structured type. You need to define a new structured types for this new rhetorical pattern.
- The definition of the structured type is incorrect. You need to modify the structured type to more correctly reflect the rhetorical pattern.
- The content is a variation of the rhetorical pattern that is deliberately not supported by the structured type. The content needs to be edited to fit.

- The content is deficient. It needs to be upgraded so that it fulfills its purpose correctly, as defined by the structured type.

Interpreting the mismatch between existing content and the structured type can make or break your entire structured writing project. There is a huge temptation to treat existing text as canonical and try to shape the model to fit it. But as I have stressed several times, the purpose of structured writing is not to represent existing texts, but to make content better. If your current content processes are so good that all your existing content fits your new structures perfectly, then you are not realizing any gain in content quality and you are wasting your time by adding additional mechanical structure. Finding content that does not fit the models is not a sign that the models are broken, but that the process is working.

This does not mean that the models never need to be changed. But it does mean that you change the models to match the things you discover about the best rhetorical structure for your content to achieve your business goals, not to make your existing content, or even the new content that authors want to write, fit the model.

This means that applying structure to your existing content is not a trivial or mechanical task. The purpose, after all, is to improve the quality of the existing content, and that is going to mean additional research and writing work to bring the content up to standard.

(Let's make this distinction clear: people often convert content from one file format to another, including for binary formats to markup formats. This is a mechanical process, though one that may require some cleanup. It does not, in itself, impose any additional constraints on the content. It merely changes the syntax that expresses existing structures. This kind of conversion is often possible to document domain formats like DITA and DocBook. This does not mean that the resulting DocBook or DITA output will correctly express the full range of constraints or structures that these formats are capable of. You can also do a reliable transformation from one subject domain format to another (say from a relational database to XML markup). But you cannot do a reliable mechanical transformation of media domain content to the document domain or of document domain content to the subject domain. The subject domain imposes constraints that may be expressed rhetorically in the document domain, but are not expressed mechanically. These conversions are writing tasks, not something than can be done mechanically.)

## Irreducible rhetorical models

In some cases, as we have seen, it is possible to factor out the rhetorical structure of a item and move the content entirely into the subject domain. This is possible with a recipe, for instance (see Chapter 6, *Writing in the Subject Domain*). In these cases, any reasonable rhetorical design can be created from the subject domain content by the presentation algorithm because the rhetorical design consists of a particular arrangement of facts.

But not all rhetorical models are reducible to an arrangement of facts. This is clearly true of philosophical essays and even of books like this one. In works of this sort,

the rhetorical structure -- the course of the argument -- cannot easily be reduced to a repeatable structure. But there are certainly cases in which rhetorical structures can be highly repeatable and yet do not consist merely of an arrangement of facts.

One such is a rhetorical pattern useful in technical communication (and perhaps in other fields) which I call the think-plan-do pattern. Many technical communication tasks simply involve telling a user how to perform specific functions on specific pieces of machinery. But there are cases in which the user's task has highly complex input conditions and potentially far-reaching consequences. In this case, the technical communication task goes well beyond telling the user how to operate the machine. It is about helping them to correctly plan their actions to achieve the desired business outcome.

You can approach this problem by simply collecting all the relevant facts that the user would need to make a correct decision. But a mere listing of relevant facts is not helpful to a user who does not fully understand the complexity of the task or the seriousness of its potential consequences. For example, a user may well not understand the security implication of a particular configuration option of a computer system. The safety of that option may depend of a variety of factors, such as who has access to the system, what software is running on the system, what data it contains, and how other settings are configured.

If the user does not appreciate the seriousness or complexity of the issues involved, they may skip all of the additional information and go straight to the beginning of the procedure. If they do, a mere listing of relevant facts may scare them out of trying to change the setting (which may be just as unsafe as changing it, and may have other negative consequences).

A potential rhetorical approach to addressing this problem is to walk the user through each of the decisions that need to be made in order to plan their changes correctly. This can consist of a number of carefully designed discrete questions designed to help the user figure out which issues apply to their situation and, if they do apply, how to deal with them.

In other words, the model presents a formal planning methodology in the form of a set of questions which break down the planning of the change into manageable pieces that the user can successfully comprehend and act on.

Depending on the material, it may be possible to find a common patterns in the subject matter of these questions. (The exact same set of questions need to be considered for each configuration setting, for example.) But in many cases, the questions that need to be asked are particular to the individual case. It is the rhetorical device of breaking the planning process into a set of discrete questions that is most important to improving the quality of the content and ensuring that the reader is successful.

Most document domain models used in structured writing are used simply to separate content from formatting or to facilitate content reuse. A focus on quality often takes us

into the subject domain. But this example shows that the document domain can also be used to enforce or encourage rhetorical structures that can be highly beneficial to readers. While these structures are not in the subject domain themselves, they are often highly specific to the subject matter. It is certain particular subjects, and certain particular audiences that create the need for these kinds of rhetorical structure. In many cases, such structures will occur within a larger subject-domain structure that establishes the appropriateness of the rhetorical tool that is used.

## Rhetorical metamodels

There are different ways of thinking about the rhetorical structure of content. Above, I describe the topic pattern of a recipe as consisting of a picture, an introduction, ingredients, and a list of preparation steps.

However, we could notice that there are a great many other type of information with a similar pattern. For instance, a knitting pattern usually has a picture of the garment, an introduction describing the project, a list of the yarns and needles required, and a list of steps for knitting and assembling the pieces. Lots of other things look similar. Instructions for assembling flat pack furniture, for example, or planting flowers in your garden.

These are not the same rhetorical pattern. You would not confuse a recipe with a knitting pattern. And each of them can have specific information fields that would make no sense for the others. A pot roast will never have washing instructions. A flat pack bookcase will never have a wine match. Nonetheless, they all have the basic pattern of picture, description, list of stuff you need, steps to complete. We might call this the make-thing-out-of-stuff-with-tools pattern.

The make-thing-out-of-stuff-with-tools pattern is what we might call a meta-pattern. It is not based on seeing similarities between texts, but on seeing similarities in the rhetorical patterns of texts. The meta pattern is not intended for creating content directly, but it can potentially provide hints that help us develop individual topic patterns.

Not only are there meta-patterns for topics, like the make-thing-out-of-stuff-with-tools meta pattern, there are also meta patterns for the different types of information that go into a meta pattern, such as the picture, description, list of stuff you need, and steps to complete. These are sometimes called “information types” (a confusing term, since text at any scale expresses information, and therefore the structure of information at any scale is an information type).

Two notable examples of these information type meta patterns are found in Information Mapping and DITA. Information Mapping proposes that documents are composed of just six information types: procedure, process, principle, concept, structure, and fact.<sup>1</sup> Documents are then constructed of some arrangement of information blocks of one of these six types, which it calls a map.

---

<sup>1</sup><http://www.informationmapping.com/fspro2013-tutorial/infotypes/infotype1.html>

In other words, Information Mapping proposes that every rhetorical pattern is always composed of some combination of these six information types.

DITA proposes something similar, but it proposes just three types: concept, task, and reference,<sup>2</sup> which, confusingly, it calls topic types. Like information mapping, DITA assembles documents out of these topic (information) types using a map.

In the concept/task/reference metamodel, our recipe topic pattern would consist of one concept topic (the introduction), one reference topic (the list of ingredients), and one task topic (the preparation steps). And our make-thing-out-of-stuff-with-tools meta-pattern would similarly consist of one concept topic (description), one reference topic (list of stuff you need), and one task topic (steps to complete). (DITA's information model does not include pictures. It just provides a mechanism for including them in textual topic types.)

What neither DITA nor Information Mapping provide is any way to model the larger recipe pattern. DITA will let you write a map to combine a concept topic containing an introduction (which presumably is where you would include the picture), a reference topic containing a list of ingredients, and a task topic containing preparation instructions. But it does not give you a way to specify that a recipe topic consists of one concept topic, one reference topic, and one task topic. In other words, DITA does not provide any way to define larger types or the overall rhetorical structure of documents.

What DITA does do is provide a way to specify a concrete instance type of any of its meta-types. The list of ingredients in a recipe is an instance of the meta-pattern list-of-stuff-you-need in the make-thing-out-of-stuff-with-tools meta-pattern. But a list of ingredients has a specific structure that is not the same as the list of pieces in a flat-pack furniture box. It consists of an ingredient name, a quantity, and a unit of measure. The unit of measure is vital in an ingredient listing because not all ingredients are quantified in the same way. You don't measure eggs the same way you measure flour, for instance.

To express this constraint, DITA will let you specialize the reference topic type to create a list-of-ingredients topic type that imposes (and records) this constraint. You could then construct a recipe using an introduction-to-recipe topic (a specialization of concept), a list-of-ingredients topic (a specialization of reference), and a preparation-steps topic (a specialization of task). However, it still would not give you a way to specify that a recipe consists of these three topic types in this order.

Actually, it is possible to define a recipe topic types in DITA, but this involves having a different idea about how atomic the basic DITA topic types are. Some DITA practitioners might say that a recipe is not a map made up of three information types, but a single task topic. In this view, a task topic is much more than what Information Mapping would call a procedure. It allows for the introduction of a task, a list of requirements, and the procedure steps all within the definition of a single topic. (I have asked a number of

---

<sup>2</sup>Or, at least, it originally proposed these three types. The DITA specification now includes other topic types, some of which are much more concrete than these original three.

DITA practitioners how a recipe should be modeled in DITA and have received both answers from multiple people.)

One of the reasons for this uncertainty about what an atomic topic is in DITA is DITA's focus on content reuse. DITA topics are not only units of information typing, they are units of reuse. The approach in which a recipe is a single topic leaves you with fewer larger units of content, which makes individual topics harder to reuse. The atomic unit of content that is small enough to maximize potential reuse is much smaller than the atomic unit of content that contains a complete rhetorical pattern. The atomic unit of reuse is smaller than the atomic unit of use.

Because DITA has no mechanism for describing model larger than a topic, a DITA practitioner is left with a choice between modeling for maximum reuse and modeling to constrain a topic type to rhetorical pattern. In practice, it seems that different DITA users make different decisions about how atomic their topic types should be, based on their business needs.

## Meta-models vs generic models

Ideally, a meta-model should just be a model of models. You should not be about to use it for anything other than to create concrete models. It should not only suggest those things that each specific model should have in common, but also the unique things that are specific to particular instance of the pattern. For instance, it should in some way suggest that a recipe instance of the make-thing-out-of-stuff-with-tools meta pattern might want to include a wine match. (A thing-goes-with-other-thing relationship, perhaps.)

In practice, meta-models often turn out to be simply a list of those things that all instances of the meta-model have in common. In many cases, instead of inventing an entirely new notation for describing meta-models, people just create a model with only the common properties. Thus the expression of the meta-model takes the form of a generic model, which means that it is perfectly possible to write content using that generic model. Thus while DITA's concept, task, and reference topic types are intended as meta-models to be specialized into concrete models, they are implemented as generic models which can be used directly.

A great many DITA users don't specialize at all. They write all of their content in the base task, reference, and concept topic types (or the even more basic "topic" topic type, of which task, reference, and concept are actually specializations). This means that the topic type imposes no specific rhetorical pattern. But at the same time, the generic pattern can be confining. For instance, it does not allow you to have two procedures in a single task topic.

Are meta-models useful for defining topic patterns? If a concrete topic pattern describes the kinds of information that are needed to help a particular audience perform a particular task, do we arrive at that pattern more easily by derivation from a meta-model or from observation of multiple concrete examples.

The obvious problem with the current generation of content meta--models is that none of them alert us that a recipe might need a field for a wine match. It is not impossible to imagine that a meta-model could do this. A meta-model could observe that objects are commonly used with other objects and lead us to ask what other objects is a steak dinner used with. There are obviously multiple aspects of this question. A steak dinner is used with a knife and fork. A steak dinner is used with a table and chair. A steak dinner is used with family and friends. A steak dinner is used with a glass of wine. How do we characterize each of these thing-used-with-thing relationships in a meta-model, and how do we decide which of these types of thing-used-with-thing types is relevant to a recipe?

Perhaps, for instance, we might decide that because a recipe describes a foodstuff, thing-used-with-thing relationships are relevant when the other thing is also a foodstuff. In other words, we might decide that a thing-used-with-like-thing relationship is part of the meta-model.<sup>3</sup> (I am not, by the way, suggesting that this is a useful part of a meta-model, I merely wanted to illustrate the problem of defining a meta-model that would comprehend all the specific models we might care about in the real world.)

This is getting complicated enough for me to conclude that, while the {ontologists} (concept “ontology”) may one day come up with a such a model and a reliable way to derive concrete content models from it, for most writers, information architects, and content strategist, building a concrete topic model from the observation of instances is probably the preferable method.

Creating a good subject-domain structured model for content is actually pretty simple when you get down to it. You simply ask yourself five questions:

1. What information does the reader need to fulfill whatever purpose this unit of content is meant to serve?
2. What is the best way to express each of these elements?
3. What is the best way to organize these elements so that the reader can quickly recognize that this is the content they need and can effectively find the information then need in the content?
4. What constraints do I need to apply to the content to make sure that what the author creates conforms with the answers to the first three questions, and how will I test their conformance to those constraints?
5. What level of detail and precisions do I need in the content structures I create to make sure that the structured writing algorithms I want to apply to the content can run reliably.

---

<sup>3</sup>Rob Hanna’s Enterprise Content Metamodel[<https://www.oasis-open.org/committees/download.php/41040/Enterprise%20Content%20Metamodel.pptx>] does attempt to do something like this for business information, attempting to describe the relationships between pieces of business content based on the business functions they serve as a basis for deriving specific information types.

Answering those question is not necessarily easy, of course. But the process itself is simple and you should not make it any more complex than it needs to be.

## Making the rhetorical model explicit

I noted above that there can be a rhetorical model in a piece of text that is just a sequence of paragraphs. You can discern the topic pattern in those paragraph and model that pattern in a topic type, and still present the output as a sequence of paragraphs. Presumably, in each instance of the topic type, those paragraphs would now be more consistently expressed with fewer errors and omissions than before, but the presentation itself would be the same.

Alternatively, you may choose to make the structure of the rhetorical model more explicit to the reader as well as to the writer. In this case, the sequence of paragraphs might be replaced with a distinct combination of headings, graphics, tables, lists, pictures, and text sections that would repeat in every topic of that type.

The question, of course, is whether making the rhetorical type explicit in this way improves the content or not. In its favor, the more explicit rhetorical type makes it easier for the reader to recognize the type. (As we noted above, you can recognize a recipe by its shape, without reading a word.) This makes it easier to identify relevant content, which is particularly important on the Web. It can also make it easier to scan the content to pick out the parts you need. (This is a property that the Information Mapping company tends to focus on in their promotional material.) The argument against this treatment is that it can lead to a noisier page that is harder to read straight through.

Whether you want to make the rhetorical structure of your pages explicit in these ways, therefore, is a matter to be decided on a case-by-case basis. But don't fall into the trap of supposing the because you have chosen a plain presentation, that means there is no rhetorical structure, and therefore no structured type. The rhetorical structure of the content is a separate thing from the presentation of the content, and the aim of structured writing is to improve the rhetorical structure, not just to make the presentation more uniform.

## Structure and randomness

However much success we may have in defining common rhetorical structures, most content does not surrender to the analytical knife entirely. There is usually an element of randomness in every piece of content: material that does not fit the model perfectly but is vital to an individual instance, relationships that do not fit the pattern of relationships in the type but are crucial to the individual instance.

These arise because the world that the content describes has irregularities. If the world were perfectly regular, we would not need content. All of its features and

relationships of a perfectly regular world could be described in database tables. We turn to written descriptions in human languages because the features and relationships we want to describe are, to a greater or lesser extent, unique, unusual, or irregular. The orderliness that we create when we create a structured type for a particular subject is, in part, a reflection of the orderliness of the human designed things, and partly an imposition of structure on the world by humans in an effort to make it easier to understand, communicate, or reason about. Creating that order is essential to effective understanding and reasoning, but dealing with, and adequately expressing the significant irregularities that remain after that order and structure is imposed, is also vital to effective understanding and reasoning.

Discerning when you are dealing with randomness within the scope of the model and when you are trying to stuff things into the model that do not belong is much more art than science. But it is important to make good decisions when these cases arrive.

---

# Chapter 34. Mechanical Structure

Structured writing is writing that obeys constraints and records the constraints it obeys. If the constraints it obeys are its rhetorical structure, the way it records those constraints is its mechanical structure. The mechanical structure of structured writing is a compute data structure. However, conventional computer data structures like relational database tables do not work well for content because they are too regular to fit the shape of content. Creating structures that are regular enough for algorithms to deal with yet irregular enough to fit written language is an interesting problem to which more than one solution has been proposed.

Raw text is just a stream of characters. Inside that stream of characters, there are rhetorical structures like headings, bibliographical entries, bold text, chapters, allusions, references, links, tables, and labeled lists. The question is, how do you express these various structures on top of the raw stream of characters?

One way is to indicate with some kind of markup the place where a rhetorical structure begins or ends.

This text is [B>bold<B].

What happens if a different rhetorical structure overlaps with this one?

This text is [B>bold, [I>this is bold and italic,<B] and this is just italic<I].

For some markup systems, such as that used in the word processor WordPerfect, this works perfectly fine. But this approach really puts a damper on the kinds of algorithms you can run. Many algorithms depend on their ability to isolate certain elements of a text for processing. Suppose your source text contains markup like this:

Use the [B>[function-name>print( )<B]<function-name] to print.

If a conformance algorithm is trying to validate function names in a text, it is not going to see `print()` as the value of the function name, but `print()<B]`. That is messy to deal with, to say the least. It is messy to write and maintain as well, which is why WordPerfect had a reveal codes mode that allowed the writer to go in and fix up broken markup when things got messy. Thus this approach to markup is not used much anymore.

To avoid this mess, most modern markup systems treat the text stream as consisting of a set of non-overlapping blocks. Blocks can nest inside other blocks, and even inside the content of other blocks, but two blocks can never overlap each other. This makes it

much easier for algorithms to process markup but it means that in cases where rhetorical structures do overlap, you have to resort to convoluted markup like this:

```
<p>This text is bold, <i>this is bold
and italic,</i> <i>and this is just italic</i>.</p>
```

There is more to this question than trivial examples like this. In the academic study of text markup, where markup is used not for the preparation of documents but to mark them up for academic study, there is a debate about whether the hierarchical structure of XML-based languages actually reflects the real structure of the text.

From a structured writing perspective, we don't have to worry so much about whether the markup is objectively true to the text. Our concern is to create structures that improve the quality of content and enable the structured writing algorithms we want to use.

A hierarchical model is almost universally preferred. Whether they truly represent the structure of text or not, hierarchical models are easier to define and process and provide an easier way to express constraints.

## Flat vs. nested structures

But even when we choose the hierarchical approach to markup design, we are still left with some fundamental choices about structure. The first is flat vs. nested structure.

We noted in Chapter 5, *Writing in the Document Domain* that in HTML, you have six levels of heading (`h1` through `h6`) whereas in DocBook you have only `title`. In DocBook, you can divide a document up into sections and nest sections inside sections. You can then print the titles of sections inside sections in a smaller font than the titles of first level sections. You get to have differences in heading size without having six different heading tags.

But the DocBook model assumes that the real structure of a document is a hierarchy of nested sections and that the size of titles announces the steps up and down that hierarchical tree. HTML makes no such assumption. It will let you put a `<h4>` immediately after an `<h1>` if you want to. It treats documents as essentially flat structures punctuated by headings of various sizes as and where appropriate.

Which model of a document is correct? You can think of a document as being organized hierarchically, with major ideas expressed in sections, sub-ideas supporting the major ideas in subsections, etc. There are doubtless documents that fit that model. But you can also think of documents as being more like a journey in which headings function more like road signs. A city gets a big sign, a hamlet a small sign, and a town a medium sign. But the town is not inside the city, not the hamlet inside the town, and there is no guarantee that on leaving the city you will come to the town before you come to a hamlet.

Studies by Peter Flynn indicate that most authors think of the documents they are writing much more in terms of a punctuated linear model than a hierarchical model.

The classical theory, derived from computer science and graph theory, is that the document is a hierarchical tree (actually inverted: a root-system) and that all necessary actions can be seen in terms of navigation around the tree, and of insertion into and withdrawal from the nodes which form the branches and leaves.

The conventional writer, however — and we expressly exclude the markup expert, as well as the experienced technical authors who responded to the survey — is by repute probably only marginally aware of this tree; but we have been unable to measure this at present. In this view, the document is seen as a continuous linear narrative, broken into successive divisions along semantic lines, and interspersed with explanatory material in the form of figures, tables, lists, and their derivatives.

—Flynn2009

The return from subjection problem, which comes up again and again in structured writing, both at the rhetorical and mechanical level, is an illustration of this point. Within a section, you can create one or more nested subsections. If one applies a hierarchical logic to the design of the document, the role of the subsection is to treat a secondary or supporting point for the main section. In this logic, one ought to be able to return to the section level after a set of subsections in order to sum up the overall point of the section. Logically you might also wish to return to the section level of discourse between subsections so as to locate the matter of the subsection in the narrative of the section.

You could certainly create structured writing section and subjection structures that would let you follow this logic. The problem is, how would you publish this structure. While it is not impossible to imagine a way to format something like this, it would clearly be a violation of all the norms of document layout and organization. In a normal document, if you want to sum up a section after several subsections, you introduce a subheading “Conclusion” after the end of the last subsection. The role of the conclusion is to sum up the section as a whole, so it is not a subsection hierarchical logic of the document, it is a return to the section level.

But no one, to the best of my knowledge, has ever tried to actually mark it up that way. The conclusion will be marked up as another subsection. And any material that relates a subsection to the matter of the section will be contained within the subsection itself. Thus the section/subsection structure of a document does not reflect any hierarchical logic in the rhetorical structure of the argument, it actually just groups content under different levels of headings.

And as Flynn’s research illustrates, there is a good chance that the author is not thinking of their argument as consisting of a strict hierarchy of points either, but as a sequence of

points with the occasional insertion of heading to break up the text or perhaps signal a change in emphasis or subject matter.

To think of it yet another way, a text is a flattening of the representation of a complex reality -- the world -- a reality whose complexity can no more be fairly represented by as a hierarchy than it can by a simple sequence. The building up of the picture of the real world, in all its complexity, is not the job of the document outline, but the job of the whole of the text, of every verb, adjective, and noun, or every sentence, and simile and symbol and story that the writer can assemble to guide the reader to the full grasp of the complex world they are describing. To suggest, therefore, that the mechanical structure of a markup language can capture the full rhetorical structure of a piece of writing is absurd. It can, at best, capture certain aspects of that rhetorical structure, and outside of that it works best when it stays out of the writer's way as much as possible.

But if the constraints that we want to express in structured writing demand hierarchy, while functional lucidity demands more of a punctuated linear model, how do we reconcile these two opposing requirements in markup language design?

This is of greatest concern in the design of document domain languages. The structure of media domain languages is largely dictated by the shape and relationship of the media-domain object they are modeling. In the subject domain, we have abstracted content out of strict document order. Hierarchy in the subject domain tend to match the hierarchy of relationship in the subject matter itself.<sup>1</sup> In the document domain, however, it is a real concern. The document domain consists of abstractions of document structures and the nature of their relationship to the structure of thought in the text is not obvious.

The options available are:

- Create a really flat document domain language. Examples are HTML and Markdown. The problem here is that they impose few constraints, and the lack of context-setting hierarchy makes it hard to model different types of document structures without creating hundred of tags -- which would negate any functional lucidity that you gained by keeping the language flat.
- Create a hierarchical language that has a really permissive structure, so that you can put boxes inside boxes in lots of different ways. An example of this is DocBook. The problem here is that the possible permutations make writing algorithms difficult and you often need to impose additional constraints on your authors that are not expressed or enforced by the markup itself. This again diminishes functional lucidity, and compromises conformance. (An interesting property of this approach is that the flexibility of the language means that authors can choose to create documents that are deeply nested or very flat. This is not really a virtue, however, as it is not clear how this choice contributes to improved content quality.)

---

<sup>1</sup>Though this is not universal. Addresses, for instance, which are based on hierachal locations, are modeled as flat ordered lists. The order reflect the hierarchy, but the nesting of city inside country and of street within city is not reflected in the structure of an address record.

- Define a smaller, stricter document domain language that is appropriate to the particular types of documents you want to write, possibly as restricted subset of an existing language like DocBook. The main difficulty with this approach is that it involves you in having to do your own language design, which many organization try to avoid. Once you have decided to go this route, going to the subject domain instead may be no more expensive while providing better functional lucidity and conformance.
- Define a strict hierarchical document domain language that expresses the constraints you need and make people learn it. This works if you are able to recoup the expense of training your authors. It does not work if you want to include occasional authors in your pool of contributors.
- Move content creation to the subject domain.

In the document domain, however, you have to make a choice between enforcing a hierarchical model of a document or allowing a more linear one. You will almost certainly do this in the context of hierarchical syntax. But you do need to decide if you are going to allow arbitrary heading levels. Or, to put it another way, to model a document as a hierarchy of sections with each section having a title, or a single flow of text randomly interspersed with headings, or something in between.

The something-in-between option can seem appealing but you can end up with something really unconstrained, meaning that it is difficult to write reliable algorithms to process it. Here is the kind of issues you can run into:

Is a list part of the paragraph in which it occurs? Some markup languages, such as Markdown and SAM, take the view that lists are separate blocks that come after a preceding paragraph. But DITA, DocBook, and HTML will all let you place a list either inside or outside a paragraph:

```
<p> The primary colors are:

 Red
 Blue
 Yellow

</p>

<p> Their complementary colors are:</p>

 Green
 Orange
 Purple

```

The interesting question here is whether these two structures should be formatted differently for output. Should the list that is inside the paragraph be indented, which the one that comes afterwards be printed flush left?

The default for HTML, at least in the browsers I tried it on (Edge, Internet Explorer, and Chrome) is that both are indented by the same amount, but of course you could change this with CSS and DocBook and DITA processors can do what they like.

But would it be a good idea to make a distinction between the two? Only if authors clearly understood the difference between the two and clearly know when they were choosing one or the other. (If they are writing in a visual editing view of an XML editor, then they almost certainly won't know the difference.)

If you are designing your own markup language, do you want to allow both these forms? Probably not. The more permutations of structure your markup language allows, the more work it will be to write the algorithms to process it. Some allowance may be made for ease of authoring, if you wish, but it may be better to wait for authors to complain about a restriction like this than to build in all the possible variations in advance.

## The role of syntax

The syntax of your markup language also plays a role in striking the right balance between hierarchy and functional lucidity. Fully explicit markup syntax like that of XML forces every element of the hierarchical structure of the document to the fore in the syntax. The syntax has explicit hierarchical constraints of its own (what XML calls the well-formedness constraint). This pushes the hierarchical structure in the author's face, which detract from functional lucidity.

Languages like Markdown with implicit syntax are just as hierarchical as their equivalents with fully explicit syntax but they feel flatter. Thus where HTML make you write:

```
<html>

 <p>The first item.</p>
 <p>A second paragraph in the first item.</p>

 <p>The second item.</p>

 </html>
```

In MarkDown you write:

\* The first item.

A second paragraph in the first item.

\* The second item.

The latter has just as much structural hierarchy as the first, but feels much flatter. It uses indentation to indicate that the second paragraph belongs to the first list item, but that feels natural and obvious rather than contrived or imposed.

If a structured writing language divides all content up into blocks and nested blocks, XML-based languages make all such blocks completely explicit by markup up the beginning and end of every block. Lightweight languages like Markdown use a more implicit syntax better reflects the implicit structure of formatted text itself.

One of my motivations for creating SAM, which is designed for structured writing, and therefore, unavoidably, for hierarchical structures, was to express hierarchy implicit where possible and as naturally as possible where it cannot be implicit, for the sake of improve functional lucidity for strictly constrained structured writing languages. (First and foremost, I created it for myself, to improve functional lucidity for me in the structured writing I do, such as the writing of this book.)

## Agreeing on boxes, names, and labels

The names and labels on blocks tell us what kind of thing is in the box. What kind of things they tell us depends on the domain we are dealing with.

In the media domain they tell us what the words in the box look like, either directly (an actual format description) or by reference (the name of a style). In the document domain they tell us what part of a document the words are. In the management domain, they tell us what to do with the words under different circumstances. In the subject domain they tell us what the subject matter of the words is.

The block does not only have to have information from one domain on it. It is not uncommon to have a block with a name in the document domain and a label in the management domain or the subject domain. We saw and example of this in Chapter 6, *Writing in the Subject Domain*:

```
<section publication="Wine Weenie">
 <title>Wine match</title>
 <p>Pinot Noir</p>
</section>
<section publication="The Teetotaler's Trumpet">
```

```
<title>Suggested beverage</title>
<p>Lemonade</p>
</section>
```

Here the section box is in the document domain, but the publication label is in the management domain.

The names and labels on the blocks tell us what constraints the structures inside the block obey, or, as we are writing, what constraints the words we write are expected to obey. Algorithms use the names and the labels to process the content, assuming that they accurately reflect the constraints.

For structured writing to work, it is essential that everyone involved understands and agrees on what the names and labels mean. If we don't, the names and labels will not accurately reflect the constraints we expect, and the whole content quality project falls apart and the algorithms stop working.

Confusion and disagreement about what the names and labels of a particular language mean are not uncommon. Large document domain languages like DocBook and DITA have large vocabularies, and many of the names they offer are quite abstract. Questions about the right way to tag certain passages are common in the communities around these languages, and opinions can vary considerably in some cases. These disagreements don't only affect low-level structures. In DITA, for example, it is common to debate if a topic that is operational but not procedural is a concept or a task. Some writers choose to use only generic topics because they don't feel the models of the task, concept, and reference topics fit the content they are creating.

Having precise definition of terms is important, therefore, in developing a structured writing language. But it is equally important that the language be functionally lucid. The authoring algorithm requires that creating structure should not come at the detriment of the writing itself. This requirement is not met if writers are constantly having to puzzle out or debate the right way to mark something up.

In the software world, meta models and abstraction are powerful tools for modeling systems. They provide clear high-level rules for the design of specific structures and create opportunities to reuse code for objects with a shared base model. But these tools can also lead to very abstract naming schemes and even to abstract structures. In the content world, such abstract names and structures can be formally correct but lack the kind of functional lucidity required for effective authoring.

The problem of defining a (mechanical structure) to express a (rhetorical structure) is not only one of defining a correct representation of the content. We also have to design and name structures that can be written by our intended set of authors without imposing a heavy burden on their attention. In other words, designing content structures, at least those intended for use by authors, is as much about interface design as it is about data structure design.

Clear concrete and specific names, and an organization of boxes that intuitively fits the subject matter, all make for easier authoring. There is no reason that such structures cannot be derived from abstract models, or that they cannot be mapped to abstract models after the fact, but it is important not to let the abstractions intrude too much into the world of the author.

Of course, functional lucidity only matters for the formats that authors actually write in. As we have seen, the publishing algorithm typically consists of multiple steps, and each one of those steps can create a format that is closer to the media domain than the one before it. It is perfectly possible to design a document domain structure for the sole purpose of serving as a step in the publishing chain. Separate authoring formats are created for authors to actually write in (perhaps subject domain formats or simplified highly constrained ad-hoc document domain formats). Content is transformed from these formats to the document domain format by the presentation algorithm and then the document domain format is translated into various different media domain languages by the formatting routine. An arrangement like this eliminates the need to compromise between different demands in designing a single language, generally making each language in the chain simpler and more constrained, which in turn makes it each one easier to validate and to process.

## Constraint and annotation

Broadly speaking, structured markup provides two things: constraint and annotation. Constraint governs the relationship of pieces. It is constraint that says that a list consists of list items, that an API reference must begin with a function signature followed by a list of parameter values, that an ingredient listing consists of the ingredient name followed by a quantity followed by a unit of measure, or that a section must begin with a title.

Annotation adds information to the constrained structures. The name of the structure is its first and most important annotation. It is the annotation that expresses the core constraint that has been imposed. This may be supplemented by other annotations on the structure that further clarify the constraint that has been imposed, or contain other metadata, such as management domain annotation.

But we can also have unconstrained annotation, or at least, annotation whose usage is very lightly constrained. In some XML-based languages, for instance, there are sets of attributes that are allowed to appear on just about any element in the entire language. (ID attributes are an obvious example.) But the main source of unconstrained annotation are those that can float in the middle of a piece of text.

As we noted, we cannot capture the full rhetorical semantics of any piece of writing using simple linear or hierarchical structures. At some point we have to leave the writer free to express themselves in paragraphs of text on which we can impose no meaningful restriction. But we may still be very much interested in constraining the interpretation of many of the words and phrases used in those paragraphs.

Such annotations are constrained in the sense that there is a limited set of them available for use, but they are almost always unconstrained as to where and in what order they appear in the text -- for the very good reason the their text that they annotate obeys no regular rule or pattern.

Every structure has a name (though the name is implicit in some forms of markup. The name of the structure is an annotation of the constraint that the structure imposes. (Remember that we defined structured writing as both imposing a constraint and recording that the constraint has been applied.)

The different structured writing domains are defined by what these annotation of structure are saying about the content they contain.

Media Domain	The annotation says what the content should look like.
Document Domain	The annotation says what role the content plays in the document.
Subject Domain	The annotation relates the content to its subject matter.
Management Domain	The annotation says how the content should be managed.

But the name of a structure only allow you to say one thing about the content. In some cases you may want to add more than one annotation to a structure. For instance, you might want to add a condition to a paragraph. In this case, the name of the structure annotates that it is a paragraph and you need an additional management-domain annotation to hold the condition.

This use of additional annotations is quite common. For instance:

- In DITA, the keyref attribute is used to add management domain metadata to a number of document structure elements.
- In HTML, the style attribute can be used to add media domain metadata to the document structure.

However, not all annotations are attached to structures. At a certain point in defining the structure of content we get down to what is essentially free-form narrative content -- that is, paragraphs of text. But even within a paragraph, there are pieces of text that we may want to annotate. Examples from each of the domains include:

Media Domain	A <b>bold</b> or <b>strong</b> annotation on a piece of text you want to emphasize.
Document Domain	An <code>xref</code> annotation to create a reference to another part of the content.

Subject Domain A function annotation to identify that a piece of text is the name of a function.

Management Domain A conditional annotation to identify a piece of text that may be conditionally included or excluded from output.

Different markup systems support structure and annotation in different ways.

## Secondary structures of interpretation

Under normal circumstance, the structures that constrain the content also constrain its interpretation. But there are cases in which annotations are added to structures to constrain a different interpretation of the content from the one specified by the structure that contains it, and not only of individual fields, but of a structure as a whole.

Consider these examples of HTML Microformats from Wikipedia<sup>2</sup>. The first example shows an address formatted as a list.

```

 Joe Doe
 The Example Company
 604-555-1234
 http://example.com/

```

Here the phrase The Example Company is contained in li tags. This is part of a list structure delineated by ul tags, so the markup is largely structural in the document domain. The li does not really tell you anything useful about what the content itself is about. It does not tell you anything useful beyond what document structure it belongs to.

The second example adds hCard microformat markup:

```
<ul class="vcard">
 <li class="fn">Joe Doe
 <li class="org">The Example Company
 <li class="tel">604-555-1234
 http://example.com/

```

This example adds subject domain metadata in the form of the class attributes. For example, it says that the phrase The Example Company is a reference to an

---

<sup>2</sup><https://en.wikipedia.org/wiki/Microformat>

organization (`org`). This annotation is not modifying or refining the constraint expressed by the `li` tag. It is saying something else entirely.

But this is not just about suggesting a different way of interpreting The Example Company (as an organization name, as well as as a list item). There is actually a complex subject domain structure being expressed. Not only is the list item The Example Company annotated as `org`, the list that contains it is annotated as `vcard`. The constraint of interpretation provided by `org` is actually dependent on it being part of a `vcard` structure.

In other words, the secondary structure created by the annotations in the sample above is equivalent to pure subject domain markup like this:

```
vcard:
 fn: Joe Doe
 org: The Example Company
 tel: 604-555-1234
 url: http://example.com/
```

In other words again, the microformats are overlaying a second structure on the list structure. In the world of HTML, this makes sense. HTML needs to be a standardized document domain language so that browsers can display it for human reading. Humans don't need the vCard annotations to recognize that the content is an address, but algorithms do. So the microformat adds a second, hidden, subject domain structure to the document for readers who are algorithms rather than people.

Structured writing constrains both the creation and the interpretation of content. In the normal case we expect that the creation of the content would be just as constrained as the interpretation. After all, it is hard to rely on the interpretation of structure if the creation of the structure is not constrained. However in this case the interpretation of the data is more constrained than the creation.

Authoring our content this way would obviously be inefficient and error prone. But this is only a problem if the content is actually authored in this format. If the content is authored in a format in which its creation is constrained to the same extent as we wish the output to be constrained, it does not actually matter that the resulting output constrains interpretation more than it constrains creation. Our concern as content creators is simply to make sure that any content we produce that promises to abide by a constraint actually does so, whether the format we deliver it in actually imposes that constraint or merely annotates it.

So, we can confidently produce this information using subject domain markup and then deliver it as HTML with hCard annotation markup using a presentation algorithm something like this (as with all example algorithms in this book, this is pseudocode):

```
match vcard
 create ul
 attribute class = "vcard"
 continue

match fn
 create li
 attribute class = "fn"
 continue

match org
 create li
 attribute class = "org"
 continue

match tel
 create li
 attribute class = "tel"
 continue

match url
 create li
 create a
 attribute class = "url"
 attribute href = contents
 continue
```

## Elements vs. attributes

Another issue in the design of the mechanical structure of a markup language is that of elements vs attributes. Elements and attributes are the core markup structures of XML, and many other markup systems make the same distinctions in one way or another. The simplest way to state it is that attributes add additional annotation to an element. So, for instance, in:

```
<author-name first="Mark" last="Baker">
```

the element is named `author-name` and it has two attributes named `first` and `last` that contain my first and last name respectively.

Why is this marked up like this and not like this:

```
<author-name>
```

```
<first>Mark</first>
<last>Baker</first>
```

Both these constructs express the same information in a way that clearly constrains how author names are to be written and how the markup of author names is to be interpreted by algorithms.

Is one of these options correct and the other incorrect? When should you use attributes and when should you use elements?

Consider our vcard example. It could be written this way, using just elements:

```
<vcard>
 <fn>Joe Doe</fn>
 <org>The Example Company</org>
 <tel>604-555-1234</tel>
 <url>http://example.com/</url>
</vcard>
```

Or it could be written this way, using attributes:

```
<vcard
 fn="Joe Doe"
 org="The Example Company"
 tel="604-555-1234"
 url="http://example.com/"
/>
```

The first says that `fn`, `org`, `tel`, and `url` are independent structures that belong as members to a `vcard` structure. The second says that the `vcard` structure has a number of data fields -- annotations -- that complete its meaning.

Does this distinction matter terribly? Both allow you to get at the information you want. Both constrain both the creation and the interpretation of data. There are limits to the version that uses attributes. You can't have more than one attribute with the same name, whereas you could have more than one member elements of the same name (multiple `tel` elements for someone with more than one telephone number for instance). Also, XML specifies that attributes are unordered, so can't restrain either the order in which writers create them or the order in which the parser reports them to a processing application.

Given this, you may be wondering why people bother with attributes, since you can do the same things with elements and have both more flexibility and more capacity to impose constraints. Yet people continue to use attributes extensively when designing markup languages in XML. When people create XML document types for representing data, rather than for writing document, they almost always use the attribute format, perhaps

because it is slightly less verbose and slightly easier to read, or perhaps because as programmers they are accustomed to representing data as key/value pairs linked with = signs.

But for documents it is more complex. To understand why XML even has attributes, and why other languages, such as SAM or reStructuredText, also have similar mechanisms for adding annotations to blocks, we need to go back to the original concept of markup as something written onto a manuscript after the fact. Markup is an addition to the text, not part of it. The content of an element is part of the underlying text. Anything you want to add, therefore, cannot be element content, since that would be adding to the text. Everything else has to be added to element definitions as attributes.

This view is reinforced by the academic interest in markup as a way to prepare texts for study. Again, here, the text is preexisting and canonical. The markup is external to it and so everything that is external to the original text must be contained in the markup itself (as attributes) and noting that is internal to the original text must be removed or replaced by markup. Thus in this fragment of a Shakespearian play marked up by John Bozek we see that the original text is kept perfectly intact:

```
<ACT>
<TITLE>ACT I</TITLE>
<SCENE>
<TITLE>SCENE I. Rousillon. The COUNT's palace.</TITLE>
<STAGEDIR>
Enter BERTRAM, the COUNTESS of Rousillon, HELENA, and LAFEU, all in
</STAGEDIR>
<SPEECH>
<SPEAKER>COUNTESS</SPEAKER>
<LINE>
In delivering my son from me, I bury a second husband.
</LINE>
</SPEECH>
<SPEECH>
<SPEAKER>BERTRAM</SPEAKER>
<LINE>And I am going, madam, weep o'er my father's death</LINE>
<LINE>anew: but I must attend his majesty's command, to</LINE>
<LINE>whom I am now in ward, evermore in subjection.</LINE>
</SPEECH>
```

Had this markup employed the normal refactoring of text into markup that we have seen in our discussions of the document domains, then the number and title of scenes would have been factored out so that instead of:

```
<SCENE>
```

<**TITLE**>SCENE I. Rousillon. The COUNT's palace.</**TITLE**>

we might have factored out the scene number and the word ‘SCENE’ like this:

<**SCENE**>

<**TITLE**>Rousillon. The COUNT's palace.</**TITLE**>

Going further, we might have noted that the introduction of a scene is invariably the name of its location, so we might have done this:

<**SCENE**>

<**LOCATION**>Rousillon. The COUNT's palace.</**LOCATION**>

or even this:

<**SCENE location=Rousillon. The COUNT's palace."**>

And similarly, we might have replaced:

<**SPEECH**>

<**SPEAKER**>BERTRAM</**SPEAKER**>

<**SPEECH SPEAKER="BERTRAM"**>

Making a few changes like this in the markup would leave us with only the words actually spoken by the actors as the “text” of the play and everything else expressed as elements or attributes.

This actually makes quite a lot of sense, because all of the stage directions and attribution of speakers in a play is actually metadata annotating the speeches of the actors which are the only things the audience is actually supposed to hear.

So is the right way to markup a play to preserve the original printed text -- which includes all of the playwright’s metadata -- or is it better to separate the playwrights metadata from the speeches which are the ‘real’ play?

The answers to questions like this need not concern us greatly. The fact that such question exist, however, help us to understand both why a markup language like XML is structured the way it is, and why so many texts are marked up the way they are -- and why so many tagging languages are designed the way they are.

The way we, as practitioners of structured writing for content quality, settle these matters is by asking ourselves which approach best supports the structured writing

algorithms that we want to implement, and always remembering that the reliability of every other algorithm depends on how well the conformance algorithms works, and that the conformance algorithm depends to a large extent on how the authoring algorithm designs for conformance and for functional lucidity.

In none of this are we in the least concerned about preserving the canonical nature of a preexisting text. There is no preexisting text. We are all about creating new texts. Because of this we have every reason to prefer to use sub-structures rather than annotations on existing structures, to express things like our vCard example. In fact, SAM, which is designed specifically for structured authoring, only supports this format:

```
vcard:
fn: Joe Doe
org: The Example Company
tel: 604-555-1234
url: http://example.com/
```

SAM only supports a very limited set of annotations on blocks, all of which have predefined meetings. You could, in fact, eliminate annotations on blocks altogether, and use child blocks for everything, but I have supported a limited set of common management-domain block annotations in SAM, mostly to improve functional lucidity.

Support for free-floating annotation in text is another matter. Those you will always need for document-oriented markup and structured writing.

In summary, when defining the mechanical structure of your structured writing, don't get hung up on what is text and what is markup. In each domain, text and markup together form a body of constrained content which can be successfully created by an author and successfully processes by one or more algorithms. Only when you resolve the content all the way to the media domain do you finally have to sort out exactly which characters appear in which order and decoration to represent that content to a particular audience. When we choose to create content in the other domains it is precisely because we want to exercise more control over these things, and to use algorithms to help us create and manage them. Whether some idea or constraint is expressed by text or markup in those domains should be based solely on what works best in those domains. (Which is precisely why attempting to do structured writing using a WYSIWYG editor is so counter productive, and why it is so important to ensure the functional lucidity of your markup as markup.)

---

# Chapter 35. Metadata

We live in the age of metadata, so much so that the word metadata has almost come to replace the word data itself and has come to be applied to almost any form of data that describes a resource. For example, we hear a lot about law enforcement getting access to metadata related to phone calls, which simply means the data about which number called which number and for how long.

The standard definition of metadata is data that describes data, but that definition misses the central point. Metadata does not merely describe data, metadata creates data. Metadata turns an undifferentiated set of values into useful data by constraining the interpretation of those values.

In the case of our recipe example, we have seen that the ingredients of a recipe in a document domain markup are just items in a list.

```
section: Ingredients
 * 12 eggs
 * 2 qt water
```

Adding subject domain markup allows us to clarify to algorithms exactly what these string.

```
ingredients:: ingredient, quantity, unit
eggs, 12, each
water, 2, qt
```

This subject domain markup is metadata, and it turns what were just list items into a set of ingredient data. It is not the the data existed and the metadata came along afterward to describe it. It is that the data exists as data only because the metadata is there to describe it. Take away the metadata and data becomes noise. We (as designers of algorithms) simply do not know what it means.

In the definition of structured writing (Chapter 3, *How ideas become content*) I said that structured writing is writing that obeys constraints and that records the constraints it obeys. In other words, structured writing is the application of metadata to content which turns content into data by constraining the interpretation of the content. If structured writing is writing that obeys constraints and records the constraints it obeys, Metadata is the means by which those constraints are recorded. Without it, you have no formal assurance that the information is what it seems to be.

This is not to say that a list structure is not metadata also. It is. It is document domain metadata. It formally identifies a piece of text as a list item. This makes it easy to write algorithms that recognize lists. This allows you to reliably format list items in whatever

media you choose to publish in. That metadata still transforms content into data, but not into ingredient data, just list item data. The operations we can perform on list item data are obviously far less sophisticated than those we can perform on ingredient data.

There is a very important point here: The same set of values, the same string of letter, words, and numbers, can be turned into different kinds of data by applying different kinds of metadata to them. This is a very clear illustration of the fact that it is metadata that creates data. And what it tells us is that we can choose what kind of data we want to turn our content into by choosing what type of metadata we want to apply to it.

Basic publishing algorithms can be performed using document domain metadata alone. To publish a list of ingredients it is enough to know that the string “6 eggs” is a list item. If we want to present ingredients different in different media, we need to know that “eggs” is an ingredient and “6” is quantity of that ingredient. But as we have seen, publishing-related algorithms such as differential single sourcing, content reuse, and content generation also require subject domain metadata. Establishing the relevance of content requires subject domain metadata because we search for content based largely on its subject matter. Validating, which deals with whether content meets its defined constraints, depends on the metadata of the domain in which the content is recorded. Auditing, which deals with whether content meets business requirements, generally requires subject-domain metadata as most business requirements are in the subject domain.

## The recursive nature of metadata

Metadata is confusing because it is recursive. If metadata is the data that describes (and thereby creates) data, it is also the data itself. And since data is created by metadata, this means that metadata is itself created by metadata. In other words, if the line `ingredients:: ingredient, quantity, unit` turns the ingredient lines into metadata in this markup:

```
ingredients:: ingredient, quantity, unit
 eggs, 12, each
 water, 2, qt
```

Then what makes `ingredients:: ingredient, quantity, unit` a piece of metadata and not just another string of characters? Whatever it is, it is also metadata, and the content depends on that metadata as well, since it cannot exist without its metadata and that metadata cannot exist without the metadata that defines it.

In structured writing, we add structure to content to replace the things we have factored out. That structure is metadata to the data that is the text of the file. But if we store that file in any kind of repository, the information that identifies the file in that repository is metadata to the file as a whole. But that is not the only kind of metadata for that file. If the structure of the file is described by a schema, the schema is also metadata for the file.

But we're not done yet because the specification of the schema language is the metadata that tells you what the schema means. And then of course, there is the specification of the markup to consider. The XML specification is part of the metadata for every XML document in existence. And we are still not be done, because the XML specification uses a formal grammar description language, called BNF. The BNF specification is metadata for the schema language description.

How do we break out of this infinite recursion of metadata. Data is information that has been formalized for interpretation by algorithms. Fortunately, human beings can understand natural language without that degree of formalization. Eventually, then, we reach a point where the last piece of metadata is not described by metadata but by human language. That human language document essentially bootstraps the whole metadata cascade that eventually yields pieces of data that can be unambiguously interpreted by algorithms.

So, every piece of data has a spreading tree of metadata supporting it, which, if traced to its roots, eventually leads to plain language documents that explain things in human terms.

## Where should metadata live?

One of the great questions about metadata is where it should live: with the data it describes or separate from it? For much metadata the answer is obvious. It lives separately because its scope is wider than one resource. An XML document does not include the XML spec, for instance. But for metadata that is unique to an individual resource, the question is an important one.

Most early graphic file formats only stored the image. Most modern format also store extensive metadata about the image. The pictures you take with your digital camera include lots of information about the camera and the settings that were used to take the shot, and, if the camera has a GPS receiver, even the geographic coordinates of where the shot was taken. Having that metadata embedded in the file ensure that the picture and its metadata stay together.

In part the argument is about who is responsible for creating the metadata. In the case of the photo, the metadata is in the file because the camera is the best placed instrument to record it. Sometimes it is about whether the metadata is intrinsic or extrinsic to the content. Sometimes it is about control.

For example, should the history of a file be stored in the file or in the repository? Storing it in the file lessens the file's dependence on the repository and makes it more portable. But a repository vendor may prefer to sell you a system in which to uninstall their repository would be to lose all your file history. If file status information is only stored in a workflow system, for instance, it is very hard to move away from that system. If it is stored in the file, it is easy to move away, and also to edit when not connected to the system, which can save you on licenses.

Writing your content in the subject domain means that more of your metadata is stored in the same file as the content, increasing its independence and portability. Also, as we have seen, the use of subject domain structures can lessen the need for management domain structures for algorithms like single sourcing and content reuse, which reduces the need for external management domain metadata. All of this contributes to improved functional lucidity, referential integrity, and change management.

There is another way to apply subject domain metadata to content besides recording the content in the subject domain. This is to apply subject domain metadata externally to the content or an overlay to the document domain structures in which the content is recorded.

For example, if you wanted to have access to the data on which ingredients were mentioned in each recipe (perhaps as an aid to finding recipes), you could record the recipe itself in the document domain and then place it in a relational database table with a many-to-many relationship to a table that listed all known ingredients.

Systems designed primarily for document domain content creation, such as DITA, can include formalized structures for attaching subject domain metadata to document domain structures. In DITA, the primary such mechanism is the subject schema.

But the two approaches are not equivalent. They vary substantially in capability, implementation requirements and ease of use.

For example, when we talked about differential single sourcing we saw that subject domain content could be transformed into radically different document domain structures for publication in different media. (A table for paper vs. an interactive widget online, for instance.) Attaching subject domain metadata to a piece of content that contains a document domain structure (such as a table) does not give you the capability to differentially single source that table to different structures for other media.

One area in which external subject-domain metadata excels is finding and filtering. You can certainly find and filter on subject domain content, but because there are many different subject domain structures, it can be non-trivial to figure out how to write the query expression that does the finding and filtering you require. By contrast, external subject domain metadata can be stored in normalized data structures that are consistent and therefore easier to write queries for. This consistency also allows tools and systems to put high-level interfaces over these query mechanisms for users that don't know how to write queries themselves. As such, the approach is widely favored in content management systems.

On the other hand, there is a limit to how fine grained a find and filter operation can be with external metadata. If you want the advantage of external metadata stored in normalized structures for ease of querying, it can only discriminate down to the level of the pieces of content it is applied to. It can retrieve that whole content unit, but not pieces of it.

For validation, it comes down to a matter of what constraints you want to validate against. As we have seen, moving content to the subject domain allows us to factor out many of the constraints that we want to apply to the document domain and to express and enforce the constraints we want to apply to how the content treats its subject matter. If content quality is your primary reason for adopting structured writing, these advantages are hard to ignore.

For auditing, the situation is much the same. If you are looking to audit the consistency of your content, document domain metadata will not suffice, as we noted in Chapter 18, *Conformance*. If you are looking to audit your coverage of subject matter, external subject domain metadata attached to document domain content will give you a high level look at your coverage, but tell you little about quality or how well subjects are being covered.

When it comes to ease of use, there are things in favor of each approach. If you have standardized document structures and standardized metadata structures, then users only have to learn those structures.

On the other hand, separating subject metadata from the writing of the content itself means that writers have to do two separate jobs. Rather than creating the metadata as they write in the subject domain, they have to write in the document domain and then create separate subject domain metadata to describe the document they have just written. This creates more work for the writer, and has frequently proven error prone.

This can significantly impact the quality of both the content and the metadata. Because the document domain does not impose any constraints on how a subject is treated, it is quite easy to apply subject domain metadata to content that simply does not merit it. Since metadata is supposed to express a promise that certain constraints have been met, this creates a broken promise.

The fact is, promising that a constraint has been met without validating that it has actually been met is a recipe for broken promises. Attaching subject domain metadata externally to document domain structures always runs this risk. Subject-related constraints have not been factored out or enforced. There has not even been effective subject-domain guidance provided to the writer while writing the content.

I noted above that external subject domain metadata excels in the area of finding because it makes it easy to write queries on standardized metadata structures. But there is a large caveat to this, which is data quality. If the promises the metadata makes are not kept by the content itself, the value and reliability of searches is severely diminished.

At the end of the day, the value of any system is limited by the quality of its data. Content management systems often fail or underperform because of poor quality data caused by a poor fit between content and its metadata. In particular, there is a tendency for data quality to decline over time. Systems based on external subject-domain metadata need to be aware of the problem and vigilant to prevent it.

We should also note that the two approaches are not mutually exclusive. The subject metadata used in both cases is essentially the same, since it describes the same subject. There is no reason, therefore, why you cannot take subject domain content, with its integrated metadata, and extract the metadata to an external metadata store. This way you can use the same query interface to find the content without giving up the validation and auditing capability of subject domain content or the finer-grained application of metadata that the subject domain enables. This approach can significantly improve the quality of the metadata in the system because all the subject domain guidance and constraints have been applied to the content before it is stored. Thus, as long as the content itself is correct, there cannot be a mismatch between what the metadata promises and what the content delivers.

There are ease-of-use advantages to this approach as well. Since the CMS metadata is extracted from the content metadata, the writer does not have to enter the metadata separately from the content after it is written. If you take such an approach, however, it is vital to be very clear about which is the canonical source of the metadata. The subject-domain content is the canonical source and the only one that should ever be edited. The exported metadata in the CMS is a cache of the content metadata created for the sake of optimizing searches. Of course, if your CMS was based on an XML database, you would not need to cache the metadata in relational tables; you could query it directly.

## Taxonomy and controlled vocabulary

With subject domain metadata, whether it is stored internally in the content or externally in a CMS, it is important to be consistent in how subjects are named. Of course, it is also important in the content itself that key subjects be named consistently across the content set to avoid confusion for readers. If you are supplying a large-scale top-down navigation scheme for your content, that scheme is also going to be concerned with the names of things, and with choosing the right terms to name the subjects so that readers can find them.

For all of these reasons, large scale content projects need to exercise some control over terminology. (Control of terminology is important for translation as well, but that is outside the scope of this book.) There are two parts to the terminology control problem: establishing terminology, and enforcing its use.

The principal difficulties in establishing terminology are:

- Human beings have a fairly small “use vocabulary” and we reuse words all the time. In safety critical functions like air traffic control or the operating room, we train everyone to use a special unambiguous vocabulary (which is generally undecipherable to the layperson) but for most uses, the words we want to control are likely to be used in multiple ways that are not easy to disambiguate formally.
- People in different fields (even within the same organization) often use different terms for the same concept. What the chef calls “pork”, the farmer calls “pig”. What the

English call “boot”, North Americans call “trunk”. Trying to force everyone to use the same term means forcing them to say things that don’t make sense in their own field.

- People in different fields (even within the same organization) often use the same words to mean different things. What the conference organizer calls a function is not what the programmer calls a function. What a programmer calls a function is (in more subtle ways) not what a mathematician calls a function.
- People in one field may have ten terms that make fine-grained distinctions among things that people in another field lump together under a single term. For instance, programmers make a distinction between subroutines, functions, methods, and procedures. I know of one documentation project on which it was mandated that all of these should be called “routines”. This was OK most of the time, but because a problem when function pointers were introduced into the product. (You can’t use any of the other terms with “pointer”.) Enforcing a single term obscured a distinction that turned out to matter.

In short, very little of our terminology is truly universal. The meaning of words changes depending on the context you use them in and the audience you are addressing. This is, in fact, at the core of how content works and why it is different from other forms of data. Content tells stories, and what words mean depends on the context in which they are used in the story. Stories are not as precise as formal data, but the only way we have to define formal data is with stories. This is why the spreading tree of metadata that supports and explains any point of data always ends with a human-language document. All of structured writing is an attempt to bring some small part of the orderliness and manageability of data to stories to improve the quality and consistency of stories so that we can communicate more effectively.

Thus terminology control is important, because it is key to making our structures and metadata work, but at the same time it must be done with care and sensitivity and a real sense of the limits within which it is possible to control the terms we use to tell stories. Good taxonomies always confine themselves to specific domain and define their terminology within the confines of those domains, but even within all but the most strictly controlled domain (such as the operating theater or the control tower) there are still commonly many shades of meaning on individual words which are only fully disambiguated by the story.

## Top-down vs. bottom-up terminology control

You can manage terminology top-down or bottom up. The principle tools of top-down management are controlled vocabularies and taxonomies. A controlled vocabulary is essentially a list of terms and their proper usage within a specific domain.

A taxonomy is a more elaborate scheme for controlling and categorizing the name of things. Taxonomies are frequently hierarchical in nature, defining not only the terms for individual things but the names for the classes of things. Thus a taxonomy does not just list sparrows and blue jays and robins, it also classifies them as birds, and birds as animals, and animals as living things. A good taxonomy should be specific to the domain for which it is intended. (Blue Jays and Cardinals occupy a very different place in a baseball taxonomy than in an ornithological taxonomy.) As a classification scheme, a taxonomy may be used not only as the basis for controlling vocabulary, but also as a basis for top-down navigation of a content set.

Alternatively, you can control terminology from the bottom up. To do this, you use subject domain annotations in your content to highlight key terms and place the usage in the appropriate domain.

In {Rio Bravo}(movie), {the Duke}(actor "John Wayne") plays an ex-Union colonel.

In the passage above, the annotations call out the fact that "Rio Bravo" is the name of a movie and that "the Duke" is the name of an actor called "John Wayne".

This is taxonomic information. It places "Rio Bravo" in the class "movie" and "John Wayne" in the class "actor", with the added information that "the Duke" is an alternate term for the actor John Wayne. By placing these terms in these classes, it makes it clear that "Rio Bravo" in this context is not the Mexican name for what Americans call the Rio Grande or any of the several American towns named Rio Bravo, and that "the Duke" refers to John Wayne and not to The Duke of Wellington (who was often called by that nickname) or any of the possible meanings of "Duke". In other words, the taxonomy is embedded in the content. The subject domain internalized taxonomic metadata.

To appreciate why this might be useful, we can look at some of the difficulties of maintaining and enforcing vocabulary constraints.

As mentioned above, constraining vocabulary is hard because the same term can mean different things in different contexts, and different terms can mean the same thing in different contexts. If you attempt to build a taxonomy from the top down it can be very difficult to anticipate the various meanings a term may have in different contexts. Even if you study existing texts, there is no guarantee that you will exhaust all the possibilities, and such searches are tedious and time consuming.

Secondly, even once you have defined your taxonomy from the top down, there is the question of how it is going to be enforced. You can require authors to use terms from the taxonomy, but how is that going to work? Do you expect them to carry the entire taxonomy around in their heads? And as they are writing, do you expect that they will recognize that they are using a word not on the taxonomy every time they refer to a subject

covered by the taxonomy? Any attempt to comply with such requirements is going to create a lot of mental overhead for writers, and, as we have noted (Chapter 9, *Authoring*) dividing author's attention has a negative impact on content quality.

There are mechanical solutions that attempt to catch vocabulary problems, but the fact that words can mean so many things in different contexts means no such process can get it right all the time.

An alternative to first trying to define and then trying to enforce a taxonomy from the top down is to let it emerge, in a disciplined way, from the content itself. The key to this is that in structured writing, particularly in the subject domain, you annotate those things that are significant to your content. If the vocabulary you are trying to enforce is not vocabulary that is significant to your content, you are probably wasting your time, so the vocabulary you want to use and the subjects you want to annotate should overlap.

If you have writers annotate the significant subjects in their content as they write, they will be annotation those very terms whose vocabulary you want to control. So when they mention the name of a bird like Blue Jay or Robin, they annotate it as {blue jay} (bird) or {robin} (bird).

Of course, you will sometimes need to revisit the list of all birds mentioned in your content to check them for accuracy and consistency. To get the current list, you simply have an algorithm scan your content and create a list of all the birds named in the content.

If an author mentions a new bird that you would not have thought to include in your taxonomy, it will get added to the list. You can have an algorithm alert you every time a new bird is mentioned by an author. Effectively, now, your taxonomy is bubbling up from your content. Your authors are not having to worry about whether the terms they are using are in the taxonomy or not, as long as they mark up what type of thing they are naming. But you will still get an alert every time a new term is added, and will be able to evaluate if it is being used correctly.

It is possible that some authors will forget to annotate some birds, or will annotate them incorrectly (as something other than a bird). But we can easily catch most of these mistakes as well. Incorrect annotations will tend to show up as anomalous entries in other annotation categories. In the cases where there are genuine name conflicts between two different domains, you can make a list of such conflicting names and use an algorithm to compile a list of all tagged instances of those names to review for incorrect tagging. For failure to tag at all, you can use the generated list of tagged terms to compile a list of terms to check for and periodically scan the content set to pick up significant cases.

Another useful conformance tool that can come out of the bottom-up approach is a stop list. A stop list is a list of terms that should not be used, but frequently are. It can be used by an algorithm to scan content for inappropriate vocabulary. Stop lists can only really be created bottom up. You can't anticipate or ban every term anyone might ever

come up with. You should only ban terms that are both problematic and which occur frequently. (The chance of false hits -- of banning terms that are perfectly legitimate in other contexts -- rises with every word you add to the stop list.) With a bottom up approach to terminology control, you get an accurate measure of which terms are being misused, and the frequency and nature of the misuse. This is an excellent basis for compiling a useful stop-list.

Also, because subject annotation can specify the type of a term (that is distinguish between {Blue Jays} (baseball-team) and {Blue Jays} (bird) you can make your stop list type specific: banning words used in one sense but not another. This can greatly reduce false hits, which is important because people tend to abandon the use of conformance tools if they produce very many false hits, as we can see from the very infrequent use of the grammar checkers in word processors.

As your taxonomy develops from the bottom up, it may be tempting to start to impose it from the top down, but this is probably not a good idea (though there may be exceptional circumstances). Having to comply with an top-down taxonomy while writing divides the writer's attention and compromises functional lucidity, both of which are bad for content quality and the author's efficiency. Continuing to rely on feedback to catch errors after initial authoring allows writers to stay in the moment as they write. If the feedback is given to the authors and they fix the errors themselves, they will learn the correct terminology and tagging over time and will make fewer mistakes. However, the door will remain open for new terms to be allowed into the content as appropriate, and to bubble up to the taxonomy, rather than banning them from the content until the taxonomy can be updated from the top down.

This approach does not make all vocabulary constraint problems go away, but it does have a number of advantages.

- It turns an up-front taxonomy development effort into an permanent part of the content development process. This not only reduces the spike in effort, it means that the taxonomy is based on real experience writing real content and that it is continually maintained as subjects and business objectives change. Indeed, if you are already annotating your content to support any of the other structured writing algorithms, you essentially get taxonomy development, maintenance, and control almost for free. (Someone does have to review the reports, make the edits, and add to the canonical taxonomy, of course.)
- It improves functional lucidity by not forcing authors to refer to the taxonomy while writing. If you are already annotating for other reasons, you are imposing no additional burden on authors at all.
- By improving the consistency of the annotations, it makes all the other algorithms that rely on them more reliable as well. (This is one of the greatest virtues of the subject domain. Subject domain markup can serve multiple algorithms, meaning you get the benefits of multiple algorithms with less cost.)

# Local and global names

A big part of functional lucidity is the ability to call things by the names that are familiar to you. We noted in the discussion of the relevance algorithm (Chapter 20, *Relevance*) that the use of clearly labeled unambiguous identifiers can help algorithms determine the relevance of a piece of content. The problem is, not everything has a globally unambiguous identifier (such as a company stock ticker) and even for things that do, the author may know what those identifies are off the top of their heads (for instance, many authors may know the names of companies like Apple, Google, and Microsoft but not know their full ticker symbol). Forcing them to look them up every time you want to unambiguously identify a company will add a lot of overhead to the authoring process, and it will also make the markup of the information more complex, again impacting the efficiency of authoring.

Fortunately, we don't have to unambiguously identify everything we write about at a global scale. We only have to ambiguously identify it within the context of the content we are writing. There is actually a universal truth about language in this. Very few words have only one meaning or identify only one thing. "Sun" is a big ball of burning gas and a computer company that was bought by Oracle. We distinguish these meaning by context, and we can distinguish content identifiers by context in the same way.

In fact, doing the identification in context is actually more accurate, since there is less possibility of accidentally introducing a confusion with a usage you are not aware of. The more highly contextualized an identifier is, the less ambiguous it is (as long as you specify both the context and the identifier you are looking for).

Rather than forcing the author to add this additional namespace information when they write (the equivalent of calling everyone by their full name and address around the dinner table), we can have the synthesis algorithm add it based on what it knows about the source of the content. This simplifies the author's task, which means they are more likely to provide the markup we want. (It is also another example of factoring out invariants, since we know that all method names in this particular piece of content will belong to the same API.)

The namespace of a person's name is implicit in where they live. It is not that people lack a last name or address when they are at home. It is that we don't use those additional identifiers because they are cumbersome and the first name is sufficiently unique to establish who we are talking about. Similarly, content created within a domain has its namespaces implicit within the domain in which it is created.

When we pull content from multiple domains, however, as we might do for purposes of content reuse or simply because we have divided authoring of a large content set into multiple domains to make it easier to create and manage, we need to make the namespaces of all the local names we use explicit in the content as we pull it out of the local domain.

This can be done by algorithm because the algorithm can know which domain each pieces of content is coming from.

## Ontology

Finally it is worth saying a word about ontology. Ontology (in the information processing sense) is an attempt to create a formal mapping of the relationships between entities in the real world such that algorithms can draw inferences and reach conclusions about them.

In many way, therefore, an ontology is an attempt to do for algorithms what content does for humans. After all, one of the main reasons that we read is so that we can understand the world better, understand what various objects and institutions are and how they relate to each other, statically and in action, so that we can decide what to do.

In some sense, therefore, ontology is the ultimate in subject domain markup. And just as subject domain markup can be used to drive content management functions, so there are attempts to use ontologies to drive content management functions, especially at the enterprise level. One should also be able to generate human-readable content from an ontology, given a sufficiently sophisticated algorithm and a sufficiently sophisticated ontology.

All of this is very much outside our scope in this book. Subject domain markup is an attempt to capture certain aspects of the subject matter of a work. But it is not an attempt to model the argument of a work. Consider the passage:

In {Rio Bravo}(movie), {the Duke}(actor "John Wayne") plays an ex-Union colonel.

Here the subject domain markup formalizes the fact that Rio Bravo is a movie and that "the Duke" is a reference to the actor John Wayne. It does not model the relationship between the two. An ontology would want to model the "starred in" relationship between John Wayne and Rio Bravo, whereas subject domain structured writing is normally content to leave this to the text.

Similarly, this subject domain markup does not bother to denote that Union is a reference to both a country and its armed forces, and that colonel is a rank in those armed forces. It does not denote these things because this particular markup language is concerned with movies and these facts are entirely incidental to the movie business. Actors, directors, and movies are significant subjects in the movie review domain. The names of nations and their armies that figure in the plot of individual movies are incidental in that domain. A full ontological treatment of the passage above, however, would need to model those relationships.

Structured writing does make certain aspects of content clear to algorithms, but not with the intention of making it possible for the algorithms to make real-world inferences and

decisions based on the information in that content. It only does what is necessary to allow human authors to use algorithms as tools to improve the quality of the content they prepare for human readers.

---

# Chapter 36. Blocks, fragments, paragraphs, and phrases

A common theme in many structured writing systems is the idea that documents are made up of various types of blocks. This idea affects both how you look at information design and how you look at the creation of data structures and algorithms.

Every hierarchical markup system, of course, is made up of blocks, and of blocks within blocks. XML, for example, is simply a set of elements inside elements. But when we design markup languages and structured writing systems, we use these basic building blocks to create more complex blocks that have a wholeness of their own. We recognize them not simply as one step in the hierarchy of a document, but as a distinct type of object that we can give a name to and assign processing logic to.

## Semantic blocks

At the risk of adding further burden to an already overloaded term, I am going to call these recognizable objects semantic blocks because they are blocks that mean something in whichever domain they belong to. Higher level markup design is essentially a matter of defining semantic blocks and the ways they go together.

An easy example of a semantic block is a list. Note that I am not using “semantic” to mean subject domain; blocks have semantics in all domains. A list is a semantic block because “list” is an idea with meaning in the document domain independent of its exact internal structure. A writer can say to themselves, “I want a list here”, independent of any specifics of markup. If a structure has a name like this in the real world, the block that implements it (in the terminology I am coining for the purpose) is a semantic block.

Semantic blocks generally contain other blocks that we might not talk about independently if we did not need to describe the detailed construction of a semantic block. We can call these “structural blocks”. This is not to say that the building blocks of a semantic block may not also be semantic blocks. In fact, this happens frequently. Nor is it to say that the distinction between semantic and mere structural blocks is fixed or inflexible. A semantic block is a block that implements a structure that has meaning to you independent of its function as markup, and what you regard as meaningful, I may regard as an implementation detail.

The point is not that we must agree on exactly what is a semantic block and what is not, but that you should think of markup design in terms of semantic blocks. Blocks, with whatever internal structure you require, that will capture the structure of something that is real and meaningful to you.

DITA, DocBook, and HTML all define lists, and each of them defines the internals of a list differently. Nonetheless we recognize that each of them is an implementation of the idea of a list.

A list is made up of structural blocks that build the shape of a list. I'll illustrate this with XML since it makes the blocks explicit:

```


 <p>This is the first item.</p>

 <p>This is the second item.</p>


```

Here the semantic block is the ordered list bounded by the `ol` tags. The `li` and `p` tags inside are structural blocks that together implement the structure of an ordered list.

Other document domain examples of semantic blocks include tables and procedures. (Again you will find that DocBook, DITA, and HTML, not to mention S1000D, and reStructuredtext, all have tables, all with different internal structures, and that both DocBook and DITA have procedures, again internally different. It is possible to disagree greatly about how to structure a semantic block while still recognizing different implementations as examples of the same semantic block.

In the subject domain, examples of semantic blocks would include the ingredients list from the recipe example we have been using:

```
ingredients:: ingredient, quantity, unit
 eggs, 3, each
 salt, 1, tsp
 butter, .5, cup
```

and the parameter description from an API reference:

```
parameter: string
 required: yes
 description:
 The string to print.
```

One characteristic of semantic blocks is that they often tend to repeat as a unit, as this example does in an API reference entry:

```
function: print
 return-value: none
 parameters:
 parameter: string
 required: yes
 description:
 The string to print.
 parameter: end
 required: no
 default: '\n'
 description:
 The characters to output after the {string}(parameter).
```

They may also be used as a unit in different places in a markup language, or in different markup languages. For instance, the ordered list semantic unit may be allowed in more than one place in a document domain language, such as in a section or in a table cell.

Designing in terms of semantic blocks not only helps keep markup design and processing simpler, it also improves functional lucidity. Present the markup language to the writer as a set of familiar objects like lists or tables, or logical structures like ingredient list or parameter description, rather than a sea of tags, and the task becomes easier to understand (and the tags easier to remember).

An XML editor that implements a WYSIWYG interface to XML authoring may provide tool bar buttons for inserting semantic blocks such as list or tables as whole structures.

The structure of a semantic block can be strict or loose. A strict semantic block has one basic structure with few options. A loose one allows a much wider variety of structure inside, sometimes to the point that it acts more as a semantic wrapper than a defined semantic block.

DocBook is an example of a language with very loose semantic blocks. DocBook has the same high-level semantic blocks as any other generalized document domain tagging language, but so many tags are allowed in so many places that none of these objects are simple and easy to understand. This supports DocBook's goal of being able to describe almost any document structure you might want to create, but at the expense of simplicity and constraint.

How do you balance flexibility with functional lucidity and constraint in creating semantic blocks? Sometimes it is best to have more than one implementation of a particular semantic block. For instance, both DITA and DocBook have two tables model, as simple model and a more complex one based on the CALS table model.

# Functional blocks

We have looked at examples of semantic blocks whose semantics are in the document domain (lists and tables) and in the subject domain (ingredients list and parameter description). There is another way in which some structured writing systems divide content into blocks, which is according to the function they perform for the reader. I'm going to call these functional blocks. (To be clear, a functional block is a type of semantic block.)

Information Mapping is a structured writing system which views all content as being made up of just six types of information block: Procedure, Process, Principle, Concept, Structure, and Fact. These are functional blocks. They don't directly describe a physical or logical element of document structure (except for procedure), nor are they specific to any one subject. They describe the kind of idea that the content conveys -- they are actually based on a theory about how humans receive information.

Which structured writing domain do functional blocks belong to? Clearly they are not media, subject, or management domain structures. Are they a kind of document domain structure or something else again? I believe it is more useful to regard them as document domain structures than to invent another domain. Information mapping is a theory about the construction of documents to make them more effective. It regards a document as a mapping of information blocks, so IM's blocks are components of documents, and therefore in the document domain.

DITA also adopted this idea of documents being made up of functional blocks. In DITA's case, these blocks are named topics, which leads to some confusion since the word topic can be used to refer to both functional block, and also to a complete document (as in a "help topic" for instance).

DITA has popularized the idea that all content (or all technical content, at least) is made up of just three functional blocks: concept, task, and reference.<sup>1</sup> (DITA actually defines

---

<sup>1</sup>There is evidence that DITA is moving away from this vision of information typing. In DITA 1.3, the technical committee puts the emphasis on topic and map as the core types, rather than concept, task and reference.

The DITA Technical Committee wants to emphasize that topic and map are the base document types in the architecture.

Because DITA was originally developed within IBM as a solution for technical documentation, early information about DITA stressed the importance of the concept, task, and reference topics.

Many regarded the topic document type as nothing more than a specialization base for concept, task, and reference.

While this perspective might still be valid for technical content, times have changed. DITA now is used in many other contexts, and people developing content for these other contexts need new specializations. For example, nurses who develop evidence-based care sheets might need a topic specialization that has sections for evidence, impact on current practices, and bibliographic references.

more topic types than this today.) This idea is appealing because it is easy to see a correspondence between these three types and the reader activities of learning (concept), doing (task), and looking stuff up (reference).

This simple triptych is very appealing because it promises (though it does not necessarily deliver) easy composability for content reuse. Some people also maintain that it makes content easier to access for readers, though other (myself included) criticizes it on the ground that it tends to break content down too finely for be useful and robs content of its narrative thread.

It can fairly be said that this depends very much on how you use it. But from a structured writing point of view, the purpose of structured writing is to use constraints to improve content quality and this approach, by itself, lacks some of those necessary constraints. As we have noted, some approaches to content reuse reject significant constraints in favor of easy composability. DITA does, however, provide mechanisms for creating more constrained functional blocks, though not a means to constrain how blocks are combined.

All of this is quite distinct from the subject-based information typing of the subject domain. The instructions in a recipe and a knitting pattern are both tasks in DITA terms and procedures in IM's terms. But in the subject domain they are distinct types because there is a distinct way in you which write knitting instructions.

DITA and Information Mapping's approaches are broad and analytical, trying to find commonalities across many different kinds of information. The subject domain is very much specific and synthetic, concerned which how specific pieces go together to successfully describe a particular subject. All three approaches break content up into blocks, and the subject-specific blocks of the subject domain can probably be categorized according to the information typology of either DITA or Information Mapping. However, the specific structured writing structures that you would create in each case are different. Both the Information Mapping and DITA approaches, when reduced to specific markup structures, create abstract document domain structures. The subject domain, of course, creates subject domain structures.

## Narrative blocks

What is the next unit up from the semantic block? It is the unit that combines semantic blocks to form a complete coherent content item. We might perhaps call it the narrative block.

---

—<http://docs.oasis-open.org/dita/dita-1.3-why-three-editions/v1.0/cn01/dita-1.3-why-three-editions-v1.0-cn01.html#focus-of-dita>

The fact that the example of evidence-based care sheets clearly would include information from more than one of the abstract types, and that it is proposed as a specialization of topic rather than of concept, task, or reference, suggests a significant shift in thinking on this point, and that may indicate a shift away from abstract information typing towards a more concrete subject-domain approach.

The narrative block obviously comes in different sizes. It could be a help topic, a web page, a chapter, an article in an encyclopedia, or a recipe in a cook book.

It is probably fair to say that the narrative block is the largest unit of content to which we can usefully assign significant constraints. Or rather, the constraints that we impose on larger collections of content are of a different type and implemented in different ways from the constraints we impose on writing and through structured writing techniques. (As we have noted, DocBook a format primarily designed for whole books, is largely unconstrained, opting almost always to give the author multiple choices as opposed to strict guidance.)

The narrative block differs also between the paper world and the hypertext world of the Web. On paper a book consists of a basically linear series of chapters. On the Web, a site consists of a linked set of pages. But on the Web, pages are not restricted to linking to other pages on the same site, and search engines, which generate dynamic ad hoc links to pages rather than sites so that it is fair to say that the Web consists directly of linked pages, meaning that a narrative unit of hypertext relates itself to other content in a very different way from a narrative unit on paper.

Increasingly, though by no means universally, the term is used for the narrative block in structured writing terms “topic”. There is some confusion caused by DITA’s use of the word topic for its functional blocks, and some confusion in the DITA world too about whether Concept, Task, and Reference units should be presented as narrative blocks or not.

## Fragments

Another division of content that can occur, mostly in relationship to the management domain, is the fragment. By fragment I mean a chunk of text that is not either a semantic block, a functional block, or a narrative block, but is a block that you want to manage independently of the surrounding text.

For example, in a content reuse scenario, you might want to make a items in a list conditional based on which of the list items applies to different versions of a product.

Individual list items are not really semantic blocks. They are just structural blocks of a list. When you make list items conditional, what you are actually doing is creating multiple separate lists with some items in common, and recording them as a single list. You might be able to attach reasonably informative metadata to any one of those lists as a whole, but there is not a lot you can say about list items individually. They are fragments of a list. When you apply conditions to them, then, you are applying those conditions to fragments.

In some reuse systems, including DITA and DocBook, it is possible to apply conditions to arbitrary bits of text -- three words in sentence for instance. The block that sets off those three words in a fragment.

Some reuse systems also allow you to reuse arbitrary bits of text from other parts of the content set, simply because the text is the same in each case. Those bits of text would be fragments.

In some cases, you turn an existing structural block into a fragment by attaching management domain metadata to it. In other cases, you have to introduce additional markup into the document to delineate the fragment.

Fragments definitely solve some problems. They are also inherently unstructured and unconstrained. It is very easy to get into trouble with fragments. It is easy to create relationships and dependencies that are hard to manage because they don't follow any structural logic. You should approach their use with great caution and restraint.

## Paragraphs and phrases

Paragraphs are the thing that make structured writing different from most other computable data sets. This is not really because of the paragraph structure per se, but because of the phrases within the paragraphs that we want to annotate. It is rare in any other data set to see a structure floating within the value of another structure. But that is exactly what happens when we annotate phrases in a paragraph.

In {Rio Bravo}(movie), {the Duke}(actor "John Wayne") plays an ex-Union colonel.

In this examples, the annotation on the phrases "Rio Bravo" and "the Duke" float in the middle of the paragraph block. Here is the same thing in XML:

```
<p>In <movie>Rio Bravo</movie>,
<actor name="John Wayne">the Duke</actor>
plays an ex-Union colonel.</p>
```

Here the movie and actor elements float in the content of the p element. In XML parlance, this is called mixed content. If fact, XML breaks the structure of elements down into three kinds:

element content Elements that contain only other elements.

data content Elements that contain only text data.

Mixed content Elements that contain both text data and elements.

Mixed content is the reason that most traditional data format are not a good fit for content. They may be able to model element content and data content, but they lack an elegant way to model mixed content.

Even conventional programming languages have trouble with mixed content. In fact most libraries for XML processing invent an additional wrapper around each string of characters in a mixed content element, effectively representing it as if it were written like this (without mixed content):

```
<p><text>In </text><movie>Rio Bravo</movie><text>,
</text><actor name="John Wayne">the Duke</actor><text>
plays an ex-Union colonel.</text></p>
```

But while this makes the content palatable to conventional languages, it is clearly false to the actual structure of the document. Structured writing is essentially about reflecting the structure of thought or presentation in a narrative, and narratives have a structure that is not shared with other data. Indeed, we might say that all other data formats exist as an attempt to extract information from the narrative format to make it easier to process.

Thus we are taught in school that if we are presented a problem in this format:

John had 4 apples and Mary had 5 apples. They place their apples in a basket. Bill eats 2 apples. How many apples are left in the basket?

You solve it by first extracting the data from the narrative:

$$4 + 5 - 2 =$$

But in content processing, we cannot extract the data from the narrative because narrative is the output we are creating. Thus we have to call out the data (to make it processable by structured writing algorithms) while leaving the narrative intact.

When you move content to the subject domain, you will, in some cases, break down paragraphs and isolate the data. This may be done with the intention of recreating paragraphs algorithmically on output, or of switching from a narrative to a data-oriented reporting of the subject matter. Either way, it makes the data easier for algorithms to handle, and thus makes most of the structure writing algorithms work better. (You may have noticed that the subject domain provides the most constrained and elegant solution to many structured writing algorithms.)

Even so, it is rarely possible to do a complete breakdown of all paragraphs in refactoring content to the subject domain. Most subject domain markup languages still make considerable use of paragraphs and other basic text structures, and annotate phrases within the paragraphs were necessary. Only narrative is capable of expressing the full variety and subtlety of the real world relationships between things, and only narrative is capable of conveying these things effectively to most human readers. Even things that can be fully described to algorithms with fielded data must be described to most audiences with narrative, and even though companies like Narrative Science are working on how to turn data into narrative, they are far from producing a general solution.

Subject-domain structured writing extends the reach of more conventional algorithms into the world of narrative to enable specific structured writing algorithms and to provide constraints to improve the quality of the writing. Unlike ontologies, subject domain structured writing does not attempt to capture the whole semantics of a narrative, just to discipline and structure narrative to achieve specific content creation objectives.

Just as media domain structured writing needs to annotate phrases in the middle of a paragraph to describe formatting, so the document domain needs to annotate phrases to describe their role in the document, the management domain to annotate them to assign conditions or extract content for reuse, and the subject domain needs to annotate them to describe the subject the phrase refers to.

In planning your markup structures, therefore, it is important to think about which structures in your language need to be mixed content and which do not. Finding ways to avoid mixed content without violating the spirit of the essentially narrative nature of writing can pay dividends in an improved ability to express constraints and to execute virtually all the structured writing algorithms.

On the other hand, some of the most important subject matter that you need to model and make available to algorithms cannot effectively be factored out of paragraphs, particularly while maintaining functional lucidity. Be prepared, therefore, to think seriously about the types of phrases that you will need to annotate below the paragraph level and exactly which domain those annotations should be in.

---

# Chapter 37. Wide Structures

The notion of separating content from formatting works quite well when the content is a string of words. A string of words has only one dimension: length. A printed string, of course, has two dimensions: length and height, since each letter has a height and a width. But the height and the width of letters is a pure media-domain concern. Fitting a one-dimensional string of characters into a two dimensional font on a two dimensional page is the job of the rendering algorithm. It is one of the first things that gets factored out as we begin to structure content.<sup>1</sup> When we separate content from formatting, we separate the font from the character and are left with a string of characters whose length is measured not in inches but in character count.

Once these formatting dimensions are factored out, it becomes easy to create and manage text in the document and subject domains without thinking about how it will eventually flow onto a page or screen.

But when it comes to content that has dimensions that cannot be factored out, things get more difficult. The main problem cases are:

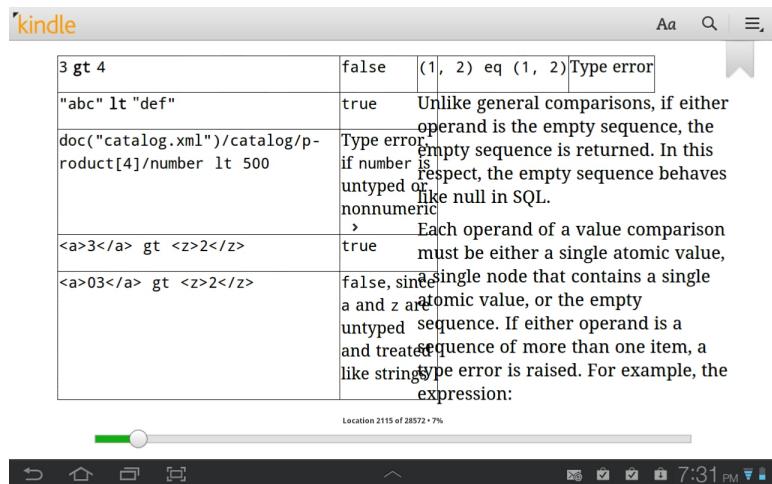
- tables
- graphics and other media
- preformatted text, such as program listings, that have meaningful line breaks

## Tables

Tables are one of the more complex problems in structured writing, particularly in the document domain. A table laid out for presentation in one publication can easily get messed up when an algorithm tries to fit it into another, as in this example from a commercially published book on my Kindle:

---

<sup>1</sup>There can, indeed, be some issues with rendering algorithms when it comes to hyphenation, widows and orphans, and the location of titles relative to the text, but with the appropriate document domain markup to delineate basic document structures, algorithms can handle these issues reasonably well. There are some fine points of typesetting aesthetics that may be difficult to automate, and algorithms can't edit the text itself to make a line break work better, as some human book designers will do. But that kind of manipulation is rarely needed. And if it is, you have the options of translating the content into the media domain and manipulating it by hand at that point. Of course, you won't do this for high-volume content that is delivered on a frequent schedule, but then you would not be able to do that kind of hand typesetting manipulation on that class of content anyway. Daily newspapers, for instance, are not typeset with the same level of attention to aesthetic detail as hardcover books.

**Figure 37.1. Broken table formatting**

This particular table is a particularly difficult case as it is not only one wide thing (a table), but it contains another wide thing (preformatted program code). I don't know exactly how this table was marked up, or which domain the content was written in, or how the formatting algorithm failed, resulting in the mess above, but including preformatted text in a table cell creates a no win situation for a rendering algorithm when it tries to shrink a table to a narrower view port. Does it:

- violate the formatting of the program code by introducing extra line breaks
- give the code the space it needs by squeezing all the other columns impossibly narrow
- resize the columns proportionally and let the preformatted text overlap the next column, but truncate it at the edge of the table
- resize the columns proportionally and truncate the preformatted text at the column boundary
- shrink the entire table so everything still formats correctly, even if it is shown in three point type.
- let the table expand outside the viewport so that it is either cut off or the reader has to scroll horizontally to read (Web Browsers tend to take this approach, but will it work on an e-reader? It certainly won't work on paper.)
- make the table into a graphic so that the reader can pan and zoom on it like they do with a large picture. (Some e-books seem to take this approach.)

If you are thinking that there is not one good option in the bunch, you are appreciating the extent of the problem. For books being transferred to e-readers, of course, there is not much that can be done to salvage the situation. Those books were probably prepared in a word processor on the more abstract edge of the media domain and the tables were prepared for a known page width in the printed book.

The tendency of readers to user small devices for reading, such as tablets, e-readers, and phones for reading means the wide tables are problematic for new content. On a phone, the amount of such a table that is visible on screen at any one time may be so small as to make the table essentially unnavigable, and to make it useless for such common table tasks as looking up values or presenting an overview of a subject at a glance.

Tables can cause problems with height as well as width. While most authors would instinctively know not to import a graphic that was six feet tall, we sometimes create tables that are that long or longer. On a web browser, the reader could simply scroll the table. But as soon as you start scrolling, you loose sight of the column headers and it becomes harder to read data across the table. On paper, it is common to repeat the headings at the top of each page when the table flows over several pages. This works, and it is possible to imitate the effect in a web browser by placing the body of the table in a scrollable frame under a fixed set of headings. But what happens on the page if the height of a table row is larger than the height of the page? Then a single row has to be broken over the page break, leading to questions about how you treat the break in the text of each cell in the row. In traditional typesetting, these things can be massaged by hand on a case by case basis. Getting a rendering algorithm to do it gracefully in every case is a very challenging task.

Creating tables in the document domain creates problems even when the intended output is paper and a sufficiently wide viewport is assumed. Since a table divides content up into multiple columns, there is always a question of how wide each column should be relative to the others, and whether or not the table should occupy the full width of the viewport or not. A table with just a few numeric value, for instance, probably should not be full page width because that would spread the numbers out too far and make comparisons difficult. One the other hand, a table with a lot of text in each cell needs to be full width, and needs to have column widths roughly proportional to the among of text in the each column. But this is tricky because some columns have side heads which means far fewer words in the first column than in the others, but you don't want to compress that column proportional to its word count because then the side headings will be unreadable.

In a media domain editor, which shows the formatting of the content as it will appear on paper, writers can create the table at a fixed width of their choice and then drag the column boundaries around to get the aesthetics of column boundaries right by eye. But tables created like this are not likely to format correctly on other devices, as Figure 37.1, “Broken table formatting” shows. And if you move the content creation out of the media domain and into the document domain, it is no longer possible to present the writer with a WYSIWYG page width for them to adjust column widths by eye. At this point you have to leave column width calculation to the rendering algorithm. The best you can do it to give it some hints about how to do its job.

This need to give the rendering algorithm hints about how to fit tables to pages has resulted in the creation of some very complicated table markup languages. Here is a simple example using the CALS table model, courtesy of Wikipedia:

```
<table>
 <title>Table title</title>

 <tgroup cols="3">
 <colspec colname="_1" colwidth="1*"/>
 <colspec colname="_2" colwidth="3*"/>
 <colspec colname="_3" colwidth="2*"/>

 <thead>
 <row>
 <entry>1st cell in table heading</entry>
 <entry>2nd cell in table heading</entry>
 <entry>3rd cell in table heading</entry>
 </row>
 <row>
 <entry>1st cell in table heading</entry>
 <entry>2nd cell in table heading</entry>
 <entry>3rd cell in table heading</entry>
 </row>
 </thead>

 <tbody>
 <row>
 <entry>1st cell in row 1 of table body</entry>
 <entry>2nd cell in row 1 of table body</entry>
 <entry>3rd cell in row 1 of table body</entry>
 </row>

 <row>
 <entry nameend="_2" namest="_1">cell spanning two columns</entry>
 <entry morerows="1">cell spanning two rows</entry>
 </row>

 <row>
 <entry>1st cell in row 3 of table body</entry>
 <entry>2nd cell in row 3 of table body</entry>
 </row>
 </tbody>
 </tgroup>
</table>
```

This sample is for a table with one case of a cell spanning two columns and one of a cell spanning two rows. As you can tell, this is not exactly obvious from the markup. In

practice, no one is going to create an CALS table by writing the markup by hand. They are going to use the table drawing tools in a graphical XML editor.

The problem with this is that while the view of the table in the editor looks just like the view of a table in a word processor like Microsoft Word, Word's graphical display is based on the actual page currently set up in printer settings and on the actual font that the document will be printed in. It can therefore show how things will fit in the table on an actual page (allowing the author to make media domain adjustments to the table). An XML editor cannot know what page size will be chosen or what font will be used when a document is printed. So while the display looks like it allows the same media domain adjustments to be made, this is an illusion and the table will not print as shown on screen.

Other markup languages take a different approach to tables. For instance, reStructuredText allows you to create a table like this:

Header 1	Header 2	Header 3
body row 1	column 2	column 3
body row 2	Cells may span columns.	
body row 3	Cells may	- Cells
	span rows.	- contain
body row 4		- blocks.

Like the DocBook CALS example, it allows you to span rows and columns, and in this case the effect is obvious from the markup. Equally obvious is that editing the content of this table, or creating a table in this style with any significant amount of text in the cells is going to be very difficult. Nor does this form provide a solution to any of the table rendering challenges described above.

## Alternatives to tables

One of the fundamental principles of structured writing is to factor out constraints wherever possible. Also, as I mentioned in the foreword, structured writing is not about creating markup that has fidelity to existing texts, it is about improving content quality and enabling effective content processing by algorithms.

One of the most fundamental of those algorithms is single sourcing, and tables are a media domain construct that are hard to create in the document domain and that do not single source well to different media.

Another fundamental algorithms is authoring, where the goal it to achieve functional lucidity by ridding the author of distractions. One of the basic ways of doing this is by moving the content to the document or subject domains to remove formatting distractions. Yet tables force the writer either to work in the media domain or to create complex and convoluted markup in the document domain to try and hint appropriate table rendering to the rendering algorithm.

Tables, therefore, as something we should be actively trying to factor out of the authored version of our content. What can you factor the content into? There are a number of alternatives, depending on what the table was being used for.

## Alternate presentation

In many cases the use of a table simply isn't necessary. There are other way to present the content with no loss of comprehensibility or quality. Some tables are just ways of formatting lists, particularly lists with two levels of nesting. If lists are an equally effective way of presenting content, choose lists rather than tables when writing in the document domain. In a structured writing environment, prefer the non-tabular version when available.

## Semantic structure

One way to present the list of ingredients in a recipe is to create a table with the ingredient name aligned left and the quantity aligned right. But as we have seen in our recipe examples, you can create a subject-specific ingredient list structure to capture your ingredient information, which you can then format any way you like for output.

```
ingredients:: ingredient, quantity, unit
eggs, 3, each
salt, 1, tsp
butter, .5, cup
```

A structure like this is a table in a different sense of the word: it is a database table and the `ingredients` structure creates a mini database table inside the body of the content. The difference between this table and a media domain table is that we know exactly what type of information each of the columns contains. This allows the presentation algorithm to make intelligent choices about column widths and all the other rendering issues that arise with tables and pass on appropriate hints to the rendering algorithm for rendering ingredient list tables in particular.

Another example where tables are frequently used in the media domain is procedures. Tables are sometimes used to create side heads for step numbers or high level descriptions of a step, which is then detailed in the right column. Instead of this, use explicit procedure

markup which can then be formatted different ways for output. Again, if a table is chosen as the output format, knowing that the contents are a procedure allows the formatting algorithm to provide appropriate layout hinting to the rendering algorithm.

## Record data as data

Many reference works have traditionally been presented as tables on paper. But most such works are really databases. They are not designed to be read but to be queried. That is, they are used to look up individual pieces of data in a large set. For a database of this sort, differential single sourcing requires that you provide the best method of querying the data that is available on each media (which included the method whose interface fits best in the available viewport). In these cases, the data should not be recorded in tables, at least, not in media domain tables. It should be recorded in whatever database format is must suited to the data and to the kinds of queries that the reader wants to make.

If one of the query mechanisms you want to support for this is printed tables on paper, then the content for those tables should be extracted from the database to create the printed table. Again, the additional semantic information available from the database structure allows the formatting algorithm to supply the appropriate rendering hints to the rendering algorithm.

## Don't use tables for layout

In the early days of the web it was common to use HTML tables as a way to lay out elements on a page. This practice is now heavily discouraged. CSS positioning is the preferred way to position elements on a page (in particular because it supports the use of responsive design to make pages display well on different sized display ports. This is old news. However, thinking about information in terms of page layouts is still an easy habit to get into.

We have talked about topic patterns as the starting point for the development of topic types. But topic patterns are often associated in our minds with a particular layout of page elements. When we create a topic pattern, we often do it by organizing page elements in certain ways. This is all perfectly legitimate. All content is formatted and displayed for consumption. The point of structured content is not to divorce content from presentation but to facilitate a happier marriage, especially when the same content must be married to different media. Thinking about your topic patterns in layout terms, therefore, is perfectly reasonable.

But when it comes to devising topic types based on those topic patterns, the specifics of page layout need to be factored out. This is essential to moving the content to the document or subject domains and to realizing all the benefits that come from that. As long as the semantics of each of the content elements is maintained, so that the presentation algorithm can tell which is which, the page layout can be recreated successfully and

consistently, and, if needed, different layouts can be created in different media. The trick is to learn to look at your page and not see tables, but objects with clearly defined types and names. Learn to see data as data, lists as lists, procedures as procedure, and prose as prose.

When you have done all of that, you will probably be left with two kinds of tables that you still have to deal with: Small ad-hoc grid layouts, and table which are database tables, but which are one of a kind, rather than something like ingredients, where the same table structure occurs in every recipe. For these, you will need some form of document-domain table markup. Which you choose will come down to how much fancy formatting of tables you want to be able to do, and how willing you are to let the rendering algorithm format your tables without extensive hinting from you.

## Code

There are some texts, particularly computer code and data, in which line endings are meaningful. Code is a form of structured writing and, in many languages, whitespace -- meaning line breaks, spaces, and indentation -- are part of the markup that defines the structure of the program. When you present code in a document, therefore, you have to respect line endings.

Furthermore, programmers usually work in a fixed-width font, meaning that all the letters are the same width. They tend to line up similar structures with whitespace to make them easier to read, so using a proportional width font for code in documentation will not only look weird to programmers, it will mess up that formatting. It will also make the code less recognizable as code, which could reduce findability.

All of which is to say that computer code, data, and other similar formats where line ends are meaningful have to be presented in a fixed width font and with line breaks where they are supposed to be. That makes code samples wide objects, just like tables, with many of the same issues when it comes to rendering them on small devices. One saving grace is that there are not usually any height issues with code samples.

Fortunately, or unfortunately, there is not much you can do to help the rendering algorithm when it comes to code. The options for fitting wide code on a narrow display are to shrink to fit, scroll to view, or truncate. It is not particularly likely that you are going to want your rendering algorithm to make a different choice for different kinds of code.

What is essential is that your document domain or subject domain markup clearly indicates when a piece of text is code. Preferably it should also indicate what kind of code it is, since knowing this can allow the formatting algorithm to do syntax highlighting for code in a known language, and can allow the linking algorithm to detect and link API calls to the API reference. In some cases it might even allow the conformance algorithm to validate the code to make sure it runs or uses the current version of the API.

# Pictures and graphics

Pictures and graphics are naturally wide objects. There are two basic formats for graphics, vector and raster. Raster graphics are made up of pixels, like a photograph, and have a fixed resolution. Vector graphics are stored as a set of lines and curves and can be scaled to meet any output requirement.

Sometimes the publishing algorithm needs to know how big the graphic is and how large it is supposed to be on the page. With raster files, the resolution of the file is set. However, its size may be in question. Is a graphic that is 600 pixels by 600 pixels a 1x1 inch picture at 600 dpi, a 2x2 inch picture at 300 dpi, or a 6x6 inch picture at 100 dpi? This is important if you are inserting a headshot into a document which will be published on both paper and the Web. You want a 1x1 photo at 600dpi for print, but you don't want that blowing up to a 6x6 photo when you add it to a web page which will display it at a typical 96dpi unless something intervenes to scale it appropriately.

Some raster file formats include metadata which may include the resolution (from which you can calculate the intended size). But there is no guarantee that that information is present in all files, or that the tools typically used to implement publishing tool chains can read that information.

Then there is the question of the intended size of the image, which is a design consideration independent of the resolution of the raster file. The intention of the person who created the picture and the intention of the person using it both play a role here. Photographs of small objects probably should not be blown up to 10 times life size (unless that is a specific intention, to show detail not visible to the naked eye). Diagrams showing complex relationships should not be shrunk down to where the relationships are unreadable. Simple diagrams should not be blown up to the size of a full page. Diagrams containing text should not be reduced or expanded so that the text becomes invisible or disproportionate to the text on the page. The person using the graphic may have some discretion, based on the role they wish the graphic to play, but their choices should stay within the range prescribed by the creator's intention.

If the rendering algorithm does not know how big a graphic is supposed to be, it has limited choices:

- Show a raster graphic at 100% of its resolution, regardless of whether it fits in the viewport or not (which means either cropping it or forcing the reader to scroll if it goes outside the viewport).

---

• Scale the graphic to the viewport (which may be stretching it as well as shrinking it).<sup>2</sup> Many web designers take an opposite approach, preparing a graphic to the exact size they intend it to be displayed at on a specific web page layout. This is a completely media domain approach, of course. In structured writing, we generally want a more flexible solution. We have all seen what happens to meticulously designed desktop website when they are displayed on a phone screen.

Since neither of these options will produce consistently good results, we generally need to provide the rendering engine with some information to help it render the graphic appropriately.<sup>2</sup>

The simplest way to supply this information is to include it in the markup that inserts the graphic. Thus HTML lets you specify the height and width of a graphic.

```

```

But do these values represent the size of the graphic or the size at which is it to be displayed in a particular media? In other words, do they define the size of the content (how big the image itself is) or do they define the size of the box that the image should fit in (the viewport for displaying the image)?<sup>3</sup>

DocBook allows you to make this distinction. Its `imagedata` tag supports attributes for specifying the size of the viewport (`height` and `width`) and for specifying the size of the image (`contentheight` and `contentwidth`). The specification also contains additional attributes related to scaling and alignment and complex rules about how the rendering algorithm is supposed to behave based on which combination of these attributes is specified.<sup>4</sup> In other words, it contains a sophisticated language to describe the sizing and scaling of graphics. It not only deals with media domain properties, it actually gives media domain instructions.

Working in the media domain is a problem, of course. It interferes with functional lucidity and it is problematic for differential single sourcing. But there is another issue to consider as well. Sometimes the best approach to differential single sourcing is to use the vector version of a graphic for one media and the raster format for another. For instance, you may want to use the vector version of a graphic for print and a raster version for online media. (Many online media cannot render common vector graphics formats.)

For all these reasons, we ought to make a distinction between the source of an image and the rendering of that image. For raster images, the source is the original high-resolution file recorded by the camera, the original screen shot, or the original raster file produced by an image editing program. For vector graphics, it is the original vector drawing file. From these source images, various image renderings may be made.

For this case we may go even further back and talk about the idea of the image. Rather than creating a single source file and generating various rendering from that, that artist may create several original renderings of the same image idea, optimizing each for

---

<sup>3</sup>This is actually quite a complicated question, and the meaning has changed between various versions of HTML. For some hints of the complexities involved, see <http://www.w3.org/TR/html5/single-page.html#attr-dim-width>.

<sup>4</sup><http://www.docbook.org/tdg/en/html/imagedata.html>

different uses. For instance, the black and white and color versions of a company logo will often be created separately, because the automatic gray-scale rendering of a color logo may not look good at all. You may also want to use different resolutions of the same raster graphic for different media or for different purposes. This may include manually redrawing a graphic to reduce fine detail, rather than simply scaling it mechanically.

Thus we could have a single ideal image with multiple source renderings and multiple derived rendering from each source. We want to use the correct version in each different rendering of our content. How do we go about it?

In DocBook we can use conditional processing to include a different image file under different conditions:

```
<mediaobject>
 <imageobject condition="epub">
 <imagedata
 fileref="../graphics/assemble.png"/>
 </imageobject>
 <imageobject condition="fo">
 <imagedata
 fileref="../graphics/assemble.svg"
 contentwidth="4in"
 align="left"/>
 </imageobject>
</mediaobject>
```

Here the `condition` attribute on the `imageobject` element specifies a different file to be used for two version of a book (this book, actually). The `epub` version is for eReaders, most of which cannot render SVG drawings, and so require a raster format (PNG in this case), while the `fo` version is for print publication using the XSL-FO page description language and uses the vector format SVG for high resolution rendering in print.

But this approach not only involves the use of media domain markup, it combines it with management domain markup. Is it possible to factor all of this out of the authored format?

To do so we have to go back to the distinction between the idea of a graphic and the rendering of a graphic. All of the approaches above have the author include a specific rendering of a graphic. To factor it out, we instead have the author include the idea of the graphic.

There are several ways to do this. In fact, this is really the same idea as we saw in the reuse algorithm where we factored out the filename of the content to be included and replace it with a semantic representation of the reason for the content. There we factored out an explicit filename from this example:

```
procedure: Blow stuff up
 >>>(files/shared/admonitions/danger)
 step: Plant dynamite.
 step: Insert detonator.
 step: Run away.
 step: Press the big red button.
```

and replaced it with a management domain key in this example:

```
procedure: Blow stuff up
 >>>(%warn_danger)
 step: Plant dynamite.
 step: Insert detonator.
 step: Run away.
 step: Press the big red button.
```

Then we refactored that into a subject-domain assertion of fact in this example:

```
procedure: Blow stuff up
 is-it-dangerous: yes
 step: Plant dynamite.
 step: Insert detonator.
 step: Run away.
 step: Press the big red button.
```

Any of these techniques can be applied to the insertion of graphics just as well as the insertion of text.

For instance, suppose you have a constraint that whenever a procedure mentions a dialog box you should have a picture of that dialog box, but your software is delivered on three different platforms and the dialog boxes look different on each. You can factor out the graphic simply by annotating the mentions of dialog boxes within procedures:

```
procedure: Save a file
 step: From the {File} (menu) menu, choose {Save} (menu-item). The
```

Now you can insert the correct screen-shot graphic for the current platform with an algorithm:

```
match procedure/step/dialog-box
 $dialog-box-name = contents
```

```
$graphic = find graphic where type = dialog box
 and name = $dialog-box-name
 and platform = $current-build-platform
insert graphic
```

There are several benefits to this approach:

1. Authors do not have to worry about finding or inserting graphics or remembering the rules about when they are supposed to use screenshots. They just have to remember to mark up the names of dialog boxes when they mention them.
2. Updating graphics for UI changes is simpler because you just update the catalog of images. You don't have to go searching through docs to find images that are affected by the change.
3. You don't need any conditional logic in the text to include the right graphic for the platform.
4. In media where the screen shot would not fit in the viewport, you can suppress the image or handle it a different way.
5. If you decide that most reader won't need to see the screenshot, you can use the markup to create a link to a topic describing the dialog box instead of putting a screenshot inline. In a reuse scenario, you might make different choices for content aimed at different levels of users (only including screenshots in material intended for novices, for instance).

In principle, there isn't any difference between factoring out text this way and factoring out graphics. Indeed, because you are not only factoring out text or a graphic, you are factoring out the decision about whether to express a particular idea with text or a graphics, and can potential make different choices about this for different audiences or different media.

However, when factoring out graphics, there is the added issue of the metadata that describes the various properties of the image and its renderings. One way to handle this is to create a metadata file for each image which provides the needed data for multiple renderings of the image and provides the path to each of the renderings.

The simplest way to implement this is to use an image include instruction that points to the metadata file instead of to an image file. This is what I did in writing this book. I noted above that the DocBook example, which conditionally includes two different versions of the graphics for epub and print was from this book. But this book is not written in DocBook, it is written in SAM. In the SAM source file, the image insertion looks like this:

```
>>>(image ../graphics/assemble.xml)
```

This is not the full factoring out of the graphics as described in the previous example. But note that the file that is being included is not a graphics file. Instead it is an XML file. That file looks like this:

```
<?xml version="1.0" encoding="UTF-8"?>
<image>
 <source>assemble.svg</source>
 <fo>
 <href>assemble.svg</href>
 <contentwidth>4in</contentwidth>
 <align>left</align>
 </fo>
 <epub>
 <href>assemble.png</href>
 </epub>
 <alt>
 <p>A diagram showing multiple pieces being
 combined in different ways to produce different
 outputs.</p>
 </alt>
</image>
```

This XML file describes the idea of the image, listing not only its source file and both of its renderings, but even a text description for use when the graphic cannot be displayed. By including this file instead of an image file, I was able to include the idea of the graphic in my content.

When the content was processed, the presentation algorithm loaded and read the `assemble.xml` file and used the information in it to generate the conditionalized DocBook file which became the source file for the formatting algorithm (which is implemented by the publisher's existing tool chain).

Could I have factored out the filename `assemble.xml` as well? Certainly. There are a number of other ways that I could have chosen to represent the idea of the graphic in the content. There are times when it makes a lot of sense to do that. If you are including screen shots in a procedure, for instance, the name of a dialog box is a good way of representing the idea of a graphic that is semantically relevant to the procedure itself.

But in the case of the images in this book, their relationship to the text is a little more arbitrary than the relationship of a screen shot to a step in a procedure, so factoring out the filename would have created an abstraction that was actually more difficult to remember as an author. The point is not to be as abstract as possible, but to combine the highest degree of functional lucidity with the constraints that improve content quality, and that will differ for different kinds of material and for different circumstances.

## Inline graphics

One further wrinkle with graphics is that authors sometimes want to place small graphics in the flow of a sentence, rather than as a separate block object. For instance, if giving instructions that involve the use of a keypad or keyboard, some authors may want to use graphics of the keys rather than simply print the character names. Under certain circumstances, this may make the content easier for a reader to follow.

Inline graphics can cause rendering problems. For instance, they may cause line spacing to be thrown off if the height of the graphics is greater than that of the font used. Inline graphics are something a writer can control and make judgments about when writing in the media domain but which may have unexpected and unwelcome consequences when formatted by algorithms from content created in the document or subject domains.

There are two techniques you can use to minimize problems with inline graphics. The first is to avoid their use altogether, where that is practical. If there is another way to present the same material just as effectively, it is better to choose that option.

The other is to factor out the graphic by using a structure to record its semantics. For instance, instead of including an Enter key graphic like this:

3. Press >(image enter\_key.png) to confirm the selection.

Do this:

3. Press {Enter} (key) to confirm the selection.

This leaves open the choice of how to represent the key in the output, and allows for differential single sourcing. For example, on a display that did not support graphics, or where graphics would be too fussy, the presentation algorithm could render this as:

3. Press [Enter] to confirm the selection.

But in for media where the use of a graphic is appropriate, the presentation routine could use a lookup table of key names and graphics to select the graphic file to represent the **Enter** key.

3. Press  to confirm the selection.

This approach allows the document designer to switch out the graphics used for keys to find ones that work best on different displays or at different scales. It is also much easier for the author who does not have to stop to think about which graphic to use. The same approach could be used in another common case, which is describing tool bar icons in a GUI application.

4. Press {Save} (button) to save your changes.

This has all the same advantages as mentioned for keys, with the additional benefit that if the interface designer decides to change an icon or to redefine the whole set of icons, you only have to update the lookup table used by the presentation algorithm. This could also be used to sub in different icons for different platforms if your application is run on more than one operating system. This is much more efficient than using conditional text to import different graphics for different configurations.

This is a good example of using the idea of the graphic rather than the graphic. The idea of the graphic is to represent the Enter key or Save button. This could be done in a number of ways, including a photograph of the key, by the use of a special font that creates the look of a key, or by a textual representation of the key such as [Enter]. The idea of a graphic is to represent a subject. So while you can insert the idea of a graphic in the form of a key or a reference to file that records the idea of a graphic and its implementations, you can also simply identify the subject itself.

As always, a common principle is at work here: better to capture the subject than a resource that represents the subject. Resources may change more often than subjects, and you may want different resources to represent a subject under different circumstance. But as long as the content remains current with its subject matter, the identification of the subject will not change.

---

# Chapter 38. Subject structures

When you create content in the subject domain, or add subject domain structures to document domain languages, the structures should reflect the subject matter rather than the presentation of the content. How exactly do you make this distinction? It is not always obvious, because even in the subject domain you are still often writing in document-domain structures like paragraphs and lists. As we have noted, a subject-domain structured writing language is not an ontology. It does not attempt to capture the entire semantics of a piece of content as formal data. It merely captures enough to enable the set of structured writing algorithms you are interested in, or to factor out constraints in other domains, or to increase functional lucidity.

This means that border between the document domain and the subject domain is somewhat fuzzy. There are subject domain structures in public document domain languages like DocBook, and there are document-domain structures in examples of subject domain languages that we have looked at, like recipes, movie reviews, and API references.

The key indication that you are in the subject domain is whether or not you could reasonably use the language to write about other subjects. DocBook has a number of subject domain structures related to software and computers, but you can ignore those structures (none of them are mandatory) and use DocBook to write recipes or movie reviews if you want to. Thus DocBook is a document domain language. Can you use our recipe markup language to write about software or movies? No. Can you use our API reference markup to write about preparing food? No. Therefore these languages are in the subject domain.

There are document domain languages that are suitable for only certain document types. DocBook's `article` document type and DITA's task topic are not specific to one subject, but they are specific to one type of document. Similarly, MarkDown is designed for simple web pages. eLesson Markup Language (eLML) is a format for creating electronic lessons, but it is not restricted to lessons on a particular subject. It is therefore in the document domain. Journal Article Tag Suite (JATS) is a language for scientific journals online. It is document domain because it is specific to a document type, not a subject.

The development of a subject-domain language generally begins by formalizing the rhetorical structure of a document on that subject . As we have shown, you can write a recipe in a document domain language. It will still follow the template of a recipe: introduction, picture, ingredients, preparation steps. But that template will not be formalized or recorded in the content. So, the first step into the subject domain is to formalize the rhetorical structure:

`recipe: Hard Boiled Egg`

introduction:

A hard boiled egg is simple and nutritious.

Prep time, 15 minutes. Serves 6.

ingredients:

- \* 12 eggs
- \* 2qt water

preparation:

1. Place eggs in pan and cover with water.
2. Bring water to a boil.
3. Remove from heat and cover for 12 minutes.
4. Place eggs in cold water to stop cooking.
5. Peel and serve.

The next step is to formalize the data structure of any data that is already presented in fields, lists, or tables. (In other words, not in paragraphs.)

recipe: Hard Boiled Egg

introduction:

A hard boiled egg is simple and nutritious.

Prep time, 15 minutes. Serves 6.

ingredients:: ingredient, quantity, unit

eggs, 12, each  
water, 2, qt

preparation:

- step: Place eggs in pan and cover with water.
- step: Bring water to a boil.
- step: Remove from heat and cover for 12 minutes.
- step: Place eggs in cold water to stop cooking.
- step: Peel and serve.

The next step is to look for data that is consistent across instances of the type and worthy of being constrained and made available to algorithms, and to pull it out of paragraphs and into structured fields, lists, or tables.

recipe: Hard Boiled Egg

introduction:

A hard boiled egg is simple and nutritious.

ingredients:: ingredient, quantity, unit

eggs, 12, each

water, 2, qt

preparation:

1. Place eggs in pan and cover with water.
2. Bring water to a boil.
3. Remove from heat and cover for 12 minutes.

4. Place eggs in cold water to stop cooking.
5. Peel and serve.

prep-time: 15 minutes

serves: 6

The next step it to think about other information that you might want to capture. Usually during this process, people start to say things like “Wait, shouldn’t we tell them X?” The process of formalizing the information you are proposing to offer often produces the realization that there is valuable information that you have not been providing.

recipe: Hard Boiled Egg

introduction:

A hard boiled egg is simple and nutritious.

ingredients:: ingredient, quantity, unit

eggs, 12, each

water, 2, qt

preparation:

1. Place eggs in pan and cover with water.

2. Bring water to a boil.

3. Remove from heat and cover for 12 minutes.

4. Place eggs in cold water to stop cooking.

5. Peel and serve.

prep-time: 15 minutes

serves: 6

wine-match: champagne and orange juice

beverage-match: orange juice

nutrition:

serving: 1 large (50 g)

calories: 78

total-fat: 5 g

saturated-fat: 0.7 g

polyunsaturated-fat: 0.7 g

monounsaturated-fat: 2 g

cholesterol: 186.5 mg

sodium: 62 mg

potassium: 63 mg

total-carbohydrate: 0.6 g

dietary-fiber: 0 g

sugar: 0.6 g

protein: 6 g

The next step is to look through the remaining paragraph data for the mention of significant subjects that should be annotated, and to define the types of those annotations. This could include annotating things that you decided not to pull out of paragraphs into fields. For instance, if you decided not to pull out all mentions of cooking utensils from

your recipe to make a list of required tools, you could annotate the mention of tools in recipe steps. (Or you could decide to do both, so you can validate one against the other.)

```
recipe: Hard Boiled Egg
introduction:
 A hard boiled {egg}(food) is simple and nutritious.
ingredients:: ingredient, quantity, unit
 eggs, 12, each
 water, 2, qt
preparation:
 1. Place eggs in (pan){utensil} and cover with water.
 2. Bring water to a boil.
 3. Remove from heat and cover for 12 minutes.
 4. Place eggs in cold water to stop cooking.
 5. Peel and serve.
prep-time: 15 minutes
serves: 6
wine-match: champagne and orange juice
beverage-match: orange juice
nutrition:
 serving: 1 large (50 g)
 calories: 78
 total-fat: 5 g
 saturated-fat: 0.7 g
 polyunsaturated-fat: 0.7 g
 monounsaturated-fat: 2 g
 cholesterol: 186.5 mg
 sodium: 62 mg
 potassium: 63 mg
 total-carbohydrate: 0.6 g
 dietary-fiber: 0 g
 sugar: 0.6 g
 protein: 6 g
```

The final step is to think about whether there is any management domain metadata that you need to add to the model. Usually this will be for management or tracking purposes, and may not be part of any current content examples you are looking at. This should be the last step because once you have formalized the rest of the model you may find that that data is usable for management and tracking purposes as well, so you won't need to add additional fields just for management purposes. For instance, lots of the subject domain information that is already formalized in the model can be used for managing which recipes you include in particular publications, such as quick and easy meals (less than 20 minutes prep time, fewer than 6 ingredients, fewer than 5 steps) or a low-cal vegetarian cookbook (fewer than 160 calories, no meat ingredient).

```
recipe: Hard Boiled Egg
author: bcrocker
rights: full
season: winter, spring, summer, fall
introduction:
 A hard boiled {egg}(food) is simple and nutritious.
ingredients:: ingredient, quantity, unit
 eggs, 12, each
 water, 2, qt
preparation:
 1. Place eggs in (pan){utensil} and cover with water.
 2. Bring water to a boil.
 3. Remove from heat and cover for 12 minutes.
 4. Place eggs in cold water to stop cooking.
 5. Peel and serve.
prep-time: 15 minutes
serves: 6
wine-match: champagne and orange juice
beverage-match: orange juice
nutrition:
 serving: 1 large (50 g)
 calories: 78
 total-fat: 5 g
 saturated-fat: 0.7 g
 polyunsaturated-fat: 0.7 g
 monounsaturated-fat: 2 g
 cholesterol: 186.5 mg
 sodium: 62 mg
 potassium: 63 mg
 total-carbohydrate: 0.6 g
 dietary-fiber: 0 g
 sugar: 0.6 g
 protein: 6 g
```

Most subject domain languages are small, simple, and fairly strict in their constraints. This is as it should be. Since you have to design them, and the algorithms that translate them into the document domain for publishing, you don't want them to be elaborate or full of different permutations of structure. If you find yourself needing a similar language for a related subject, it is usually better to create a new equally small, equally strict language for that subject rather than trying to make one language cover both.

A subject domain language should communicate with the author in terms that they understand. This means that the names of structures should make sense to them, but it also means that the how the formal structures break things up should make intuitive sense as

well. For an author with experience in the field, a subject domain language should be such a good fit that they don't really feel like they have to learn anything to use it. The vastly increases the functional lucidity of the language leaving more of the author's attention free to focus on content, while at the same time providing constraints and guidance that make sure that things are complete and consistent.

It is easy to get carried away with breaking things down into finer and finer pieces to formally describe the subject matter in finer and finer detail. Remember that all of this is wasted unless it enables one of the structured writing algorithms that matter to you. Remember too that authoring and conformance are among the algorithms. Making your markup mysterious, difficult, or tedious to create in the name of more precisely modeling the subject matter does more harm than good if reduces functional lucidity without creating a compensating increase in quality or efficiency.

Moving your content to the subject domain, in other words, is about doing the simplest thing that works to achieve a given outcome of a given quality. Sometimes MarkDown is the simplest thing that works, but there are lots of quality, management, and production goals that are hard to meet with MarkDown. Going to elaborate document domain languages is often the next step people take in pursuit of those goals, but these are not simple. In many cases, the subject domain may be the simplest thing that works for these larger goals. But if make it so, you need to make sure that your subject domain languages are as simple as they can be to get the job done.

---

# Chapter 39. Patterns

We noted that lightweight markup languages like Markdown rely on patterns in the text rather than explicit markup characters to delineate basic structures like paragraphs and lists. XML, on the other hand, makes a strict distinction between what is markup and what is text. Markup is always recognized explicitly by markup sequences that cannot occur as normal text in any part of the document. (If you want to enter those sequences as text, you either have to escape the markup characters with character entities or use CDATA marked sections.)

All the same, even XML markup is defined by patterns. If you look at the XML spec you will see that XML is defined by a series of patterns which are described using a pattern notation called EBNF. The difference between XML syntax and Markdown syntax is that in XML, patterns are defined absolutely, so that a markup start character is a markup start character no matter where it occurs, whereas in Markdown, the patterns are defined in context, so that a \* character starts a list item only at the beginning of a line, not the end.

Since markup is all about patterns, we are not confined only to the patterns defined by the markup language itself. There are lots of patterns that occur naturally in content and we can recognize those patterns in our content and act on them even if they are not formally marked up in our markup language.

A common example of this is the pattern of a URL. In most Web-based editors, if you enter a URL as plain text (no markup identifying it as a URL), the editor will recognize the pattern and will turn the URL into a clickable link in the HTML output.

Recognizing patterns rather than forcing the author to mark them up explicitly can increase functional lucidity significantly. Consider something as simple as a date. A date is a complex piece of data consisting of at least three elements, the year, month, and day. (It gets more complex if you include the day of the week or the time.) If you need those individual components of a date for processing purposes, you might be tempted to require dates to be marked up like this.

```
date: {2016}(year)-{08}(month)-{25}(day)
```

or

```
date:
 year: 2016
 month: 08
 day: 25
```

This is fully explicit markup. Whether markup should be fully explicit in this way is a debate in data representation circles. For structured writing, however, it is a practical question of whether such markup has functional lucidity, which clearly it does not.

A date in the format 2016-08-25 is already a piece of structured markup. In fact, it is a standardized piece of structured markup defined by the international standard ISO 8601. The only problem in identifying a date as such is that in another context 2016-08-25 could be a mathematical expression that resolves to 1983. But we can eliminate this potential confusion if use our markup language to isolate the context so that we can recognize when an ISO 8601 date is being used, and then use the ISO 8601 format specification to isolate the year month and day components of that data. In other words, all we need to do is this:

```
{1941-12-07}(date) is a day that will live in infamy.
```

It also means that if the author wants to enter the date in another format, we can annotate it with the ISO 8601 format to make the meaning clear to algorithms.

```
{December 7, 1941}(date "1941-12-07") is a day that will live in infamy
```

Of course, a sufficiently clever algorithm could recognize December 7, 1941 as a date as well, because that too is a fully predictable pattern. One of the defining features of patterns is that you can recognize the semantic equivalence between two patterns that express the same information. You can increase the functional lucidity of your markup by recognizing a wider variety of semantically equivalent patterns in content without having to impose uniform syntax on authors.

XML recognizes the value of patterns. The XML data types standard define a number of common patterns, and XSD schemas allow you to define patterns for use in your own markup. It calls these patterns “simple types” (as opposed to “complex types”, which are composed of multiple nested elements). Thus you can specify in XSD that a date field is of the simple type `xs:date`. The `xs:date` data type is a pattern that is based on ISO 8601, so you can do following in XML (as long as the schema define the type of the date element as `xs:date`):

```
<date>1941-12-07</date> is a day that will live in infamy.
```

It is important to understand what this does. If the type of element `date` in your XML tagging language is defined as `xs:date` then an XSD schema validator will validate 1941-12-07 as a valid date but will reject '`<date>last Thursday</date>`' as invalid. It will not break the date down into its separate year, month, and day components. If you need that, you will have to get your algorithm to do that for itself. But your algorithm can do that reliably if you know for certain that you are dealing with date in ISO 8601 format.

Most programming languages will have library functions that know how to manipulate ISO 8601 dates, so you probably don't have to do any work yourself to get the year, month, and day components of a date.

For other patterns, most programming languages have a library called "regular expressions" which you can use to rip a pattern apart and get at the pieces. (XSD actually lets you define new patterns by specifying a regular expression for the pattern you want to match.) Regular expressions will let your algorithms decode any sufficiently well defined and contextualized pattern in your content, removing the need to ask authors to explicitly break down many types of information into separate fields.

For example, in our most of our recipe examples, we have been explicitly breaking down the items in an ingredient list into ingredient, quantity, and unit of measure like this:

```
ingredients:: ingredient, quantity, unit
eggs, 12, each
water, 2, qt
```

That format is not terribly onerous to write, but there is still a small functional lucidity penalty here. An author might more naturally write:

```
ingredients:
 ingredient: eggs 12
 ingredient: water 2qt
```

Ingredient name, quantity, and unit of measure are still clearly distinct here. Any human reader can see them instantly. The trick is to get an algorithm to read them the same way. This trick can be accomplished fairly easily using a regular expression:[\*2};

```
(?P<ingredient>.+?) (?P<quantity>\d+) (?P<unit>qt|tsp|tbsp) ?
```

(The regular expression I am using here is deliberately simple and does not include a full list of all the possible names for units of measures. A more precise expression would likely be required for production quality code, but it would be distractingly verbose for our present purpose.)

To show you how this works, here is a set of ingredient values:

```
eggs 12
water 4qt
salt 1tsp
butter 2tbsp
pork chops 4
```

Here is the result of applying the regular expression above to those values:

```
Match 1
Full match 0-7 `eggs 12`
Group `ingredient` 0-5 `eggs `
Group `quantity` 5-7 `12`

Match 2
Full match 8-17 `water 4qt`
Group `ingredient` 8-14 `water `
Group `quantity` 14-15 `4`
Group `unit` 15-17 `qt`

Match 3
Full match 18-27 `salt 1tsp`
Group `ingredient` 18-23 `salt `
Group `quantity` 23-24 `1`
Group `unit` 24-27 `tsp`

Match 4
Full match 28-40 `butter 2tbsp`
Group `ingredient` 28-35 `butter `
Group `quantity` 35-36 `2`
Group `unit` 36-40 `tbsp`

Match 5
Full match 41-53 `pork chops 4`
Group `ingredient` 41-52 `pork chops `
Group `quantity` 52-53 `4`
```

As you can see, the regular expression has broken each ingredient down into ingredient, quantity, and unit fields, just like our explicit markup did. The conventional way of writing ingredients, in other words, is just as structured as or formal way. It just takes a line of code to pull them out of the string.

Essentially, any expression that we commonly use in text which follows a well defined pattern is already structured text, and we don't need to structure it again to treat it as structured. All we need to do is to use structure to place it into context so that we can recognize it reliably. For example, our regular expression for breaking down the components of an ingredient listing will only work when applied to text that we know for sure is supposed to be an ingredient listing. Thus we use normal structured text structures to delineate the ingredients list and each list item. This provides the necessary context to use of ingredient list pattern to pull out the ingredient, quantity, and unit fields that we want without forcing the author to make them explicit in the text.

Avoiding unnecessary markup of already recognizable patterns both simplifies our markup design and increase the functional lucidity of our markup languages.

If we were creating this recipe structure in XML and using XML Schema (XSD) as our schema language, we could create a data type for the ingredient field that would verify that the text entered met that pattern we are expecting:

```
<xsd:element name="ingredient">
 <xsd:simpleType>
 <xsd:restriction base="xsd:string">
 <xsd:pattern value="(.+?)(\d+)(qt|tsp|tbsp)?></xsd:pattern>
 </xsd:restriction>
 </xsd:simpleType>
</xsd:element>
```

This type definition would ensure that if the writer entered ingredients in the wrong form or using the wrong names for units, an error would be raised. For instance, these ingredient list items would trigger an error:

```
ingredients:
 ingredient: 12 eggs
 ingredient: water 2 cups
 ingredient: butter 4 lumps
```

The first would fail because the pattern would not find the quantity where it expected to. The next two would fail because cups and lumps are not on the list of strings recognized by the pattern (which probably should be extended to allow cups, but not lumps). In the XML schema example, the names of the fields that were present in the original regular expression are omitted because XML schema is not going to break the values apart and put them in separate fields for us. As far as XML schema is concerned, this is just the definition of a single element. It would be up to a processing application to break the pieces apart and make separate fields of them. It is design intent for SAM that its schema language will let you do just this: define multiple fields with a regular expression that tells you how to break a string apart and names the individual fields to extract from the pattern.

---

## **Part IV. Languages**

---

## Table of Contents

40. Markup .....	312
Why markup? .....	312
Markup vs. regular text .....	313
Markup languages .....	315
Concrete markup languages .....	315
Abstract markup languages .....	316
Instances of abstract markup languages .....	317
Concrete languages in abstract clothing .....	318
The ability to extend .....	319
The ability to constrain .....	319
Showing and hiding structure .....	320
Hybrid languages .....	321
Instances of hybrid markup languages .....	323
41. Lightweight Languages .....	325
Markdown .....	325
Wiki markup .....	326
reStructuredText .....	327
ASCIIDoc .....	328
LaTex .....	329
Subject Domain Languages .....	332
SAM .....	333
42. Heavyweight markup languages .....	334
DITA .....	335
DocBook .....	339
S1000D .....	340
HTML .....	340
Subject domain languages .....	341
43. Extensible and Constrainable Languages .....	345
XML .....	345
DITA .....	346
Specializing between domains .....	350
DocBook .....	351
RestructuredText .....	352
TeX .....	353
Sam .....	353
SPFE .....	355
44. Constraint Languages .....	357

---

# Chapter 40. Markup

Through much of this book I have talked about markup without really explaining it. Markup is the way much, though not all, structured writing is done.

Structured writing, in the broadest sense, is writing that applies any form of structure (any set of constraints) to writing. As such, most forms of writing on a computer, other than perhaps in a straight text editor without the use of markup, are structured writing since they all contain text in structures.

All writing programs have to store the writing in files. There are two possible file types they can use: binary and text.

For all intents and purposes, a binary file is one that can only be read or written by a computer program, usually the program that created it. Open up a binary file in a text editor and you won't be able to make heads or tails of it. And even if parts of it look like plain text, editing those sections and saving the file is likely to result in a corrupt file that the original application can no longer open.

A text file, by contrast, is one that you can open in a text editor and actually be able to read and write without breaking it. But to express structure in a text file, you need a way to interpolate information about structure into the text. The way we interpolate structure is with markup -- special sequences of text characters that are recognized as defining structure rather than expressing text.

## Why markup?

Structured writing is writing to which explicit structure has been applied. But binary formats can express such structures just as well as markup can. Why do we so often use markup rather than binary file formats for structured writing?

There are three main reasons:

Application independence

With binary files, it is possible to obscure how the file is interpreted, making it more difficult to write other applications that can edit the file. With markup, it is generally much easier to create other applications that can read and write the same file. This supports the development of tool chains, often based on open-source tools, to implement whatever set of structured writing algorithms are important to a business.

Ad hoc structure definition

If you need to define a structure to serve a particular purpose, perhaps one that will be used

for relatively few documents, then you need a simple and inexpensive way to create that structure. Markup can provide such a format. It is much easier, for instance, to define a subject-domain markup language for a particular purpose than it is to create a binary file format and a program to read and write it.

Human reading and writing

Markup makes it possible for humans to read and write files that are processable by computers. This makes it much easier to build a tool chain for different components and to integrate them to meet your needs. All markup formats make it possible for humans to read and write the files, but not all formats make it easy. As we shall see, some formats are much easier to write in than others.

# Markup vs. regular text

Some markup languages make the distinction between markup and regular text completely explicit. An example of explicit markup is an HTML tag. Tags are set off by opening and closing angle brackets:

```
<h1>Moby Dick</h1>
<p>Herman Melville's <i>Moby Dick</i> is a long book about a big whale.
```

HTML uses open angle brackets `<` to indicate the start of markup and closing angle brackets `>` to indicate the end of markup and a return to regular text. Actually the recognition of markup in HTML is a little more complicated than that, but that is more detail than we need to get into here. What matters is that there are certain sequences in the text which trigger a processing program (generally called a “parser”) to recognize when markup starts and when it ends.

What if you want to enter these “markup start” characters into the text of your document? You can't just type them in because the parser will think they are markup. To fix this, markup languages either define “escape” characters, that signal the parser to treat the following character as content, or they include markup for inserting individual characters in a way that won't be confused with markup characters. HTML takes the second approach. To include a `<` character in HTML, you use another type of markup called a “character entity.” A character entity is a code for a character. It begins with `&` (another markup start character), followed by a character code and ending with a semicolon. The character entity for `<` in HTML and XML is `&lt;`. (“lt” is short for “less than”, the name of the `<` character.)

```
<p>In HTML, tags start with the < character.</p>
```

This will display as:

In HTML, tags start with the < character.

Since & is also a markup start character, we need to replace it with a character entity as well if we want to include it literally. To include a literal & you use the character entity &amp;.

```
<p>In HTML, character entities start with the & character.</p>
```

This will display as:

In HTML, character entities start with the & character.

To include the literal sting &amp; ; therefore, you would write &amp; ;amp; ;.

```
<p>The character entity for an ampersand is & ;amp; ;.</p>
```

This will display as:

The character entity for an ampersand is &.

Other markup languages do not make such an explicit distinction between text and markup. For example, in Markdown a numbered list is created by putting numbers in front of list items:

1. First
2. Second
3. Third

Here the numbers are markup. That is, the Markdown processor recognizes them as indicating a list and will translate them into a structure in HTML like this:

```

 First
 Second
 Third

```

But at the same time, the numbers look like text to the writer or reader of the Markdown file, and there is no need to escape numbers followed by periods when they occur elsewhere in the text. Thus the following markdown file:

1. First comes 1.
2. Second comes 2.
3. Third comes 3.

will translate to HTML as:

```

 First comes 1.
 Second comes 2.
 Third comes 3.

```

Rather than thinking of markup as being something entirely distinct from text, therefore, it is better to think of markup as being a pattern within a piece of text that delineates its structure. In some cases those patterns may be absolute, meaning the same thing everywhere, and sometimes they may be contextual, meaning one thing in one location and something else in another location. Sometimes the markup characters may be entirely distinct from the text characters, and sometimes a pattern in the text may serve as markup as well.

## Markup languages

A set of markup conventions taken together constitutes a markup language. Markdown, DocBook, and JavaDoc are all markup languages. However, each of these languages recognizes markup in a different way. & may be a markup start character in HTML and XML, but it is just plain text in reStructuredText.

We can usefully divide markup languages into three types which I will call concrete, abstract, and hybrid.

## Concrete markup languages

A concrete markup language has a fixed set of markup that describes a fixed set of content structures. For example, Markdown is a concrete markup language that uses a markup that is designed to mimic the way people write plain text emails. Here is the passage about *Moby Dick* written in Markdown:

```
Moby Dick
=====
```

Herman Melville's Moby Dick is a long book about a big whale.

In Markdown, a line of text underlined with equal signs (=) is a level one heading. A paragraph is a block of text set off by blank lines. Emphasized text is surrounded with underscores or asterisks.

In Markdown, these patterns correspond directly to specific document structures. You cannot invent new structures without inventing a new version of Markdown.

## Abstract markup languages

An abstract markup language does not describe specific concrete document structures directly. It describes abstract structures which can be named to represent structures in any domain.

XML is an example of an abstract markup language.<sup>1</sup> The markup in an XML file does not directly indicate things like headings or paragraphs. Instead, it indicates a set of abstract structures called elements, attributes, entities, processing instructions, marked sections, and comments.

None of these abstract structures describes document structures in any of the structured writing domains. Instead, specific markup languages based on XML (or its cousin, SGML) indicate subject, document, management, or media domain structures as named instances of elements and attributes.

Here is the *Moby Dick* passage again, this time in XML (and more specifically, in Docbook):

```
<section>
 <title>Moby Dick</title>
 <para>Herman Melville's <citetitle>Moby Dick</citetitle>
 is a long book about a big whale.</para>
</section>
```

The structure described by the XML syntax here is that of an element which contains two other elements, one of which contains text, and the one of which contains a mix of text and another element. There is no separate syntax for titles or paragraphs as there is in Markdown. Everything is an element. To define specific document domain structures, we give names to the elements.

---

<sup>1</sup>The formal term for a language like XML is “meta language”, a language for describing other languages. In calling XML an “abstract” language, I am focusing on a different property, its use of structures that are not parts of a document but generic containers. A meta language needs such abstract containers. But I find that the term “meta language” is not helpful to most readers, so I have chosen instead to focus on this property of using abstract structures as opposed to the concrete structures of a language like Markdown.

Here, the document-domain structures “section”, “title”, and “para” are defined by XML elements named, “section”, “title”, and “para”. The generic `<i>` tag from the HTML example has been replaced with them more semantically specific `<citetitle>` tag, which in DocBook means the name of a literary work. Both “i” and “citetitle” are modeled as XML elements in DocBook.

XML elements are generic abstract structures. Named elements can be used to represent any media, document, subject, or management domain structure that you want, thus allowing you to create a markup language that suits a specific purpose.

Unlike a Markdown parser, an XML parser does not see paragraphs or titles. It sees elements. It passes the elements it finds, along with their names, down to a processing application which is responsible for knowing what “section”, “title”, and “para” elements mean in a particular markup language like DocBook. The parser is common to all XML-based languages, but the processing application is specific to DocBook. Thus while processing a concrete language like Markdown is generally a one step operation, processing an abstract language like XML is a two step operation, with the first step being to parse the file to discover the structures defined by elements and the second step to process those structures according to a specific set of rules applicable to a particular markup language.

## Instances of abstract markup languages

This means that DocBook is a instance of the abstract language XML. XML defines abstract structures. DocBook defines concrete structures by giving names to XML’s abstract structures. Many common markup languages are instances of XML<sup>2</sup>. XML is virtually the only abstract language used for content these days, so it is the only abstract language I am going to talk about.

Since specific markup languages like DocBook are instances of XML, I need to revise my earlier statement: We can usefully divide markup languages into **four** types: concrete, abstract, instances of abstract, and hybrid. In fact (spoiler alert), lets revise it again: We can usefully divide markup languages into **five** types: concrete, abstract, instances of abstract, hybrid, and instances of hybrid.

By these definitions, abstract and hybrid are not language types you can actually write content in, they are languages that you can use to define other languages that are instances of them. Despite how we often use the term, you don’t actually write in XML, you write in DocBook or DITA or some other language like BeerML, which happen to be instances of XML.<sup>3</sup> Thus in terms of actually writing content, we really do have just three types

---

<sup>2</sup>Sometimes also referred to as “applications” of XML, though this usage was far more common in the days of SGML.

<sup>3</sup>Or to put it another way, you write in DocBook semantics using XML syntax. Alternatively, since DocBook originated in the days of XML’s predecessor abstract language, SGML, you can write DocBook semantics in SGML syntax.

of markup languages: concrete languages, instances of abstract languages, and instances of hybrid languages.

Or to put it another way, as a designer of markup languages you can either:

- Design a concrete language from scratch (or modify an existing one)
- Use an abstract language (probably XML) to design a concrete language.
- Use a hybrid language to design a concrete language.

As a writer, you will either use:

- A concrete language with ad-hoc syntax (like Markdown)
- A concrete language (like DocBook) based on an abstract language (probably XML)
- A concrete language based on a hybrid language

## Concrete languages in abstract clothing

The key defining characteristic of an abstract language is the use of abstract named structures like XML elements. All XML elements share a common markup start sequence followed by the element name. This creates a named block of content. But concrete languages can use named blocks too. For example, JavaDoc, a concrete language for describing Java APIs, uses named blocks using @ as a markup start character:

```
/**
 * Validates a chess move.
 *
 * Use {@link #doMove(int theFromFile, int theFromRank,
 * int theToFile, int theToRank)} to move a piece.
 *
 * @param theFromFile file from which a piece is being moved
 * @param theFromRank rank from which a piece is being moved
 * @param theToFile file to which a piece is being moved
 * @param theToRank rank to which a piece is being moved
 * @return true if the move is valid, otherwise false
 */
boolean isValidMove(int theFromFile, int theFromRank,
 int theToFile, int theToRank) {
 // ...body
```

}

In this sample, @param and @return are named blocks. But in JavaDoc, there is a fixed set of named blocks that are defined as part of the language. You can't create a new language by defining your own block names. By contrast, XML itself defines absolutely no element names. Only instances of XML, like Docbook, define element names.

A particularly notable example of a concrete language in abstract clothing is HTML. HTML looks a lot like an instance of XML, but it is not. An XML parser cannot parse most HTML. HTML is nominally an instance of SGML but never did quite conform to it. XHTML is a version of HTML that is an instance of XML. HTML5 actually supports two different syntaxes, one of which is an instance of XML and one of which is not, meaning that it has both a concrete syntax and a syntax which is a instance of an abstract language. (Sigh. This is consequence of having made a mess early on and having to live with it for evermore. A lesson for all markup language designers.)

## The ability to extend

The downside of concrete languages is that their concrete syntax defines a fixed set of structures. If you want other structures, there is no way to create them short of inventing your own concrete language, or a variant on an existing one, and coding the parser and all the other tools to interpret that language. And designing new concrete languages is non-trivial because you need to make sure that any combination of characters that the writer may type is interpreted in an unambiguous way. Many versions of Markdown, including the original, contain ambiguities about how certain sequences of characters should be interpreted, which obviously detracts from its reliability and functional lucidity.

If you want to define your own structures to express the constraints that matter to your business, you need an easier way to do it. Abstract languages like XML make this much easier. You just write a schema describing the structures you want, and any algorithms you need to process those structures.

## The ability to constrain

Extensibility allows you to add structures to a language but does not place restrictions on where they can occur.

Extensibility allows you to have elements called ingredients and ingredient and wine-match. Constraints allow you to require that ingredient only occurs inside an ingredients structure and that the content of the ingredients structure must be a sequence of ingredient elements and nothing else. Constraints lets you say that writers can't put wine-match in the introduction or as a step in the preparation, they can only put it as a child of recipe after the servings field and before the prep-time field. Constraints allow, you to require that every recipe have the full list of nutritional information.

Constraints are what bring discipline to structured writing. They drive content quality and enable efficient reliable processing with algorithms. Constraints make it easier to write good algorithms because they limit the number of permutation of structures that you have to deal with.

All markup languages have constraints. With a concrete language, you get the constraints that are built into the language. Abstract languages allow you to define your own structures, and your own constraints. However, as we shall see in Chapter 43, *Extensible and Constrainable Languages*, not all languages that are extensible are also constrainable.

## Showing and hiding structure

The whole point of structured writing is to create content that meets constraints and that records the constraints it meets so that it can be reliably validated, audited, and processed. For this to happen, authors need to see the structures they are creating as formatted structures on screen. In the media domain a WYSIWYG interface shows you the media domain structures you are creating by visually rendering them on screen. But what about in the other domains? The document domain creates abstract document structures that are deliberately separated from their formatting. The subject domain creates subject-based structures that don't have a one-to-one relationship with any organization or formatting of a document. The management domain creates structures that have nothing to do with the representation of content at all. How does the author get to see these structures when writing in these domains?

This is a big problem with XML, the only abstract language in widespread use today. XML tends to hide structure. As an abstract language, an XML document is a hierarchy of elements and attributes -- not the concrete subject, document, management, or media domain structures the author is supposed to be creating. Those concrete structures are present in the markup because their names are there, but they are not visually distinguished the way the basic document structures are in a concrete language like Markdown. And XML syntax is verbose, meaning that there is a lot of clutter in the raw text of an XML document, which makes it hard to discern both the structure and the content (and also very laborious to write).

To remove that clutter, many authors use XML editors that provide a graphical view of the content similar to that of a word processor. But while XML editors removes visual clutter, they also hide the structure. Even if the author is supposed to be working in the document domain or the subject domain, the editor is now displaying content in the media domain. This greatly reduces the functional lucidity of the document domain or subject domain language and encourages backsliding into the media domain.

And then there are the problems that arise when you try to edit the WYSIWYG view of an XML document. Underneath is a hierarchical XML structure, but all you can see is the flat media-domain view of the graphical editor. Editing or cutting and pasting structures

you can't see can be an exercise in futility. You can learn to do it, but it is frustrating and it takes time, and even when you learn, the process is still more complicated than it should be.

Concrete markup languages like Markdown, on the other hand, show you the structure you are creating and are simple to edit.

## Hybrid languages

There are significant advantages and significant disadvantages, then, in both concrete and abstract languages. Hybrid languages try to find a middle way.

By hybrid, I mean a language that combines both abstract and concrete markup in one language. A hybrid language has a base set of concrete syntax describing basic text structures but it also has abstract structures such as XML's elements and attributes that can be the basis of extensibility and constraint.

An example of a hybrid markup language is reStructuredText. Like Markdown, it has a basic concrete syntax for things like lists and paragraphs. But it also supports what it calls "directives", which are essentially named block structures. For example, a codeblock in reStructuredText looks like this:

```
.. code-block:: html
 :linenos:

 for x in range(10):
 print(x+1, "Hello, World")
```

reStructuredText provides an extension mechanism that allows you to add new directives. But while reST directives are similar to XML elements, reStructuredText predefines a core set of directives for common document structures. The `code-block` directive above is not an extension of reStructuredText, it is part of the core language.

Because it defines a large set of document-domain directives, reStructuredText is inherently a document domain language. You could, of course, add subject-domain directives to it. Most document-domain languages in use today include some subject-domain structures, reflecting the purpose they were originally designed to serve. Nonetheless, reStructuredText is inherently document domain.

Another important note about reST is that it has no constraint mechanism. You can add new directives, but you can't constrain their use, or the use of the predefined directives.

I have developed a hybrid markup language which is designed to be both extensible and constrainable. I call it SAM (which stands either for Semantic Authoring Markdown or

Semantic Authoring Markup, as you please). SAM is the language I have been using for most of the examples in this book.

Here is the *Moby Dick* passage written in SAM:

```
section: Moby Dick
```

```
Herman Melville's {Moby Dick} (novel) is a long book about a big w
```

In SAM, as in Markdown and most other concrete markup languages, a paragraph is just a block of text set off by whitespace. Thus there is no explicit structure named `p` or `para`.

At the beginning of a line, a single word without spaces and followed by a colon creates an abstract structure called a block. The word before the colon is the name of the block. Thus `section:` above creates a block structure named “section” just as in XML an element named `<section>` would create a structure named “section”. Blocks can contain blocks or text structures such as paragraphs and lists.

The hierarchy of a SAM document is indicated by indentation. Thus the paragraph in the sample is indented under the section block. This removes the need for end tags, which reduces verbosity and helps make the structure of the document visually clear.

Within a paragraph, curly braces markup a phrase, to which you can attach an annotation in parentheses. Here the phrase “Moby Dick” is annotated to indicate that it is a novel. SAM also supports decorations, like the underscores in the Markdown example, so in the media domain “Moby Dick” could have been written `_Moby Dick_`.

SAM is not intended to be nearly as general in scope as a purely abstract markup language like XML. It is meant for semantic authoring (which is to say, structured writing). As such it incorporates a number of shortcuts to make writing typical structured documents easier.

In a typical document, a block of text (larger than a paragraph) typically has a title. So in SAM, a string after a block tag is considered to be a title. That means that the markup above is equivalent to:

```
section:
 title: Moby Dick
```

```
Herman Melville's {Moby Dick} (novel) is a long book about a big w
```

Unlike RestructuredText, however, SAM does not have an extensive set of predefined blocks. It has just a few, which correspond to the basic text structures for which it provides shortcuts or concrete syntax. And SAM is designed to have a constraint

mechanism, allowing you to write a schema to define what blocks and annotation are allowed in a SAM document. This includes constraining the use of the concrete syntax as well. SAM thus represents a different type of hybrid from RestructuredText which is not extensible but not constrainable.

Also unlike RestructuredText, SAM is not intended to have its own publishing tool chain. SAM is really intended for creating subject domain languages, with just enough basic concrete document domain structures to make writing easier. The SAM parser outputs an XML document which can then be further processed by any existing publishing tool chains by transforming it into an appropriate document domain language. This book was written in SAM, using a simple document-domain language I created for the purpose, with a number of subject domain annotations. That language was transformed into a semantically equivalent XML document by the SAM parser. That XML document was then transformed into DocBook according to the publisher's specifications (the publisher has a number of constraints on the DocBook they use that are not expressed in DocBook itself). From that point on, the publisher's existing DocBook tool chain took over.

Most concrete markup languages, at least those designed for documents, try to make their marked-up documents look and read as much as possible like a formatted document. SAM is designed to be easy and natural to read, like a concrete markup language, but it is also designed to make the structure of the content as clear and explicit as possible while requiring the minimum of markup. This is why it uses indentation to express structure. Indentation shows structure clearly with a minimum of markup noise to distract the reader's eye.

Because it is meant specifically for authoring, a SAM parser outputs XML, which can then be processed by the standard XML tool chain. Below is how the SAM markup above would be output by a SAM parser:

```
<section>
 <title>Moby Dick</title>

 <p>Herman Melville's <phrase>
 <annotation type="novel"/>Moby Dick</annotation>
 </phrase> is a long book about a big whale.</p>
</section>
```

## Instances of hybrid markup languages

I said above that this book is written in SAM, but that is not quite accurate. As noted above, you can't write anything in an abstract or hybrid language directly. You write in instances of those languages. Thus DocBook is an instance of the abstract language XML. You can write documents in DocBook. We do say, of course, that we write documents in XML, but that statement is, if not wholly inaccurate, certainly non-specific. Saying the a

document is written in DocBook tells you what constraints it meets. Saying it is written in XML merely tells you which syntax it uses, which is a whole lot less informative.

So, to be more specific, this book is written in a markup language written in SAM, one that I created for the specific purpose of writing this book. That markup language was then transformed by a processing application into DocBook, which is the markup language that the publisher uses for producing books. From there it was processed through the publisher's regular DocBook-based tool chain to produce print and e-book output.

---

# Chapter 41. Lightweight Languages

The term “lightweight markup language” has arisen in recent years and seems to describe that set of markup languages which are designed to use a lightweight syntax, that is, one that imposes a minimal burden on the readability of the raw text of the document. The primary appeal of lightweight markup languages rests on two related phenomena.

- They have a high degree of functional lucidity at the syntactic level and often at the semantic level as well. It is usually possible to read the raw markup of a lightweight language more or less as if it were a conventional text document.
- They can be written effectively using a plain text editor (as opposed to an elaborate structured editor with a graphical editing view). This means that the editing requirements are lightweight as well.

Most examples also come with a simple processing applications that creates output directly in one or more output formats. This means that they have a lightweight tool chain that is easy and inexpensive to implement.

There are a number of lightweight markup languages. Some of the more prominent include:

## Markdown

The most prominent of the lightweight languages, and arguably the lightest-weight, is Markdown. Invented in 2004 by John Gruber as a way to quickly write simple web pages using syntax similar to that of an text-format email, it has spread to all kind of systems and now exists is multiple variants that have been adapted for different purposes.

Adapted for different purposes mostly means that people have created version with specific additional syntax and semantics in addition to those of Gruber’s first version. For instance, the code sharing site GitHub has adopted “Git Hub flavored markdown” as the standard format for user-supplied information on the site, such as project descriptions and issues, and has added syntax specific to tracking issue numbers and code commits for projects, allowing the automatic generation of links between commits and the issues that relate to them.

Markdown is a simple document-domain language. While its semantics are essentially a subset of HTML, it is more squarely in the document domain than HTML since it lacks any ability to specify formatting or even to create tables (though various MarkDown flavors have added support for tables).

One of the recurring patterns of technology development, and certainly markup language development, is that when some simple format becomes popular because of its simplicity, people start to add “just one more thing” to it, with the result that it either becomes

more complex (and thus less attractive) or more fragmented (and thus harder to build a tool chain for). Markdown is definitely going the route of fragmentation at the moment (though a standardization effort, in the form of CommonMark is also under way). There is even a project to add semantic annotation to MarkDown as part of the Lightweight DITA project.

None of this is a reason not to use MarkDown where its structures and syntax make it an appropriate source. MarkDown provides useful constraints on the basic formatting of a web page both by factoring out direct formatting features and by providing a very limited set of document domain features. These constraints help keep all contributors to a site from indulging in extravagant non-standard formatting or overly elaborate text structures.

It does not provide any kind of subject domain constraints at all. This may be a welcome feature when comparing it with more complex document domain languages, many of which do include some subject domain structures which can be confusing to some writers, or which writers may abuse to achieve formatting effects.

The inspiration for its syntax, text-format emails, has faded to obscurity, so it is not clear that everyone automatically knows how to write markdown, as was the original design intent, but a lot of it remains obvious and intuitive, meaning that, within its limits, MarkDown has good functional lucidity.

# Wiki markup

Another popular lightweight format is Wiki markup, introduced by Ward Cunningham in 1995 as the writing format for WikiWikiWeb, the first Wiki.<sup>1</sup> Wiki markup is similar to Markdown in many respects (most lightweight languages share the same basic syntax conventions, based on the imitation of formatted document features in plain text documents). What makes Wiki markup distinct is how it is tied into the operation of a Wiki. One of its most notable features is how linking is handled. In the original WikiWikiWeb markup, any word with internal capitals was considered a “WikiWord” and instantly became a link to a page with that WikiWord as the title. Such a page was created automatically if it did not already exist. This was an extremely simple implementation of a linking algorithm based on annotation rather than the naming of resources.

A wiki is a type of simple content management system which allows people to create and edit pages directly from a web browser. Wikipedia is by far the largest and most well known Wiki. Wikis are a preeminent example of a bottom-up information architecture. Anyone can add a page and that page is integrated into the overall collection by Wiki word style linking and by including itself in categories (conventionally by naming them on the page).

---

<sup>1</sup><https://en.wikipedia.org/wiki/WikiWikiWeb>

Cunningham described WikiWikiWeb as “The simplest online database that could possibly work.”<sup>2</sup> Like Markdown, its success has led to additional features, fragmentation, and growing complexity. Some commercial wikis are now complex content management systems. Indeed, it is somewhat difficult today to define the boundaries between Wikis, Blog platforms, and conventional CMSs.

If Wikis have a defining characteristic today it is probably the bottom-up architecture rather than the original novelty of in-browser editing which is now found across many different kinds of CMS. Cunningham designed Wikis to be collaborative platforms -- places where people could collaborate with people they did not even know to create something new without the requirement for central direction or control. The idea was not only architecturally bottom-up but editorially bottom-up. Most Wiki products today, however, include considerable features for exercising a degree of central control. Question and Answer sites like Stack Exchange with their distributed and democratic control systems may be closer today to Cunningham’s idea of a democratic creation space.

What Wikis illustrate for structured writing is that very simple markup innovations like the WikiWord can have revolutionary effects on how content is created and organized. Most Wikis today use ordinary words between double square brackets for WikiWords, rather than internal capitals, but the principle is the same. You can link to a thing merely by naming it.

Wiki words are also a case of subject domain annotation. Marking a phrase as a WikiWord says, “this is a significant subject”. It does not provide type information like most of the subject domain annotation examples shown in this book, but merely denoting a phrase as significant says that it names some subject of importance that deserves a page of its own. This illustrates the point about bottom-up information architectures, that structured writing, even in very simple form, can create texts that are capable of self-organization, that can be assembled into meaningful collections without the imposition of any external structures.

## reStructuredText

reSturcutedText is a lightweight hybrid markup language most often associated with the Sphinx documentation framework which was developed for documenting the Python programming language. We looked at reStructuredText briefly as an example of a hybrid markup language in Chapter 40, *Markup*.

Similar to MarkDown, ReStructuredText uses a plain text formatting approach to basic text structures. This part of the markup looks very natural because it uses characters and patterns that you might easily use to format a document if the only editor you had available were a plain text editor:

---

<sup>2</sup><http://www.wiki.org/wiki.cgi?WhatIsWiki>

Hard Boiled Eggs

=====

A hard boiled egg is simple and nutritious.

Prep time, 15 minutes. Serves 6.

Ingredients

-----

=====	=====
Item	Quantity
=====	=====
eggs	12
water	2qt
=====	=====

Preparation

-----

1. Place eggs in pan and cover with water.
2. Bring water to a boil.
3. Remove from heat and cover for 12 minutes.
4. Place eggs in cold water to stop cooking.
5. Peel and serve.

But reStructuredText also has a feature called directives which is used to create markup with more complex semantics. For this, reStructuredText takes the same approach as XML, using characters in a way that they are almost never used in a normal document. This approach simplifies parsing, because there is seldom any question about whether a particular pattern is intended to be markup of text, but it also makes it less natural to read and to write.

```
.. image:: images/biohazard.png
:height: 100
:width: 200
:scale: 50
:alt: alternate text
```

reStructuredText is therefore something of a syntactic hybrid as well as being a hybrid in the sense that it has both fixed concrete parts and extensible parts.

# ASCIIDoc

ASCIIDoc is a lightweight markup language based on the structure of DocBook. It is intended for the same sort of document types for which you might choose DocBook, but

allows you to use a lightweight syntax. In appearance it is very similar to MarkDown, as shown in this example from Wikipedia:

```
= My Article
J. Smith

http://wikipedia.org[Wikipedia] is an
on-line encyclopaedia, available in
English and many other languages.
```

```
== Software

You can install 'package-name' using
the +gem+ command:
```

```
gem install package-name
```

```
== Hardware
```

```
Metals commonly used include:
```

- \* copper
- \* tin
- \* lead

However, while MarkDown was designed for simple Web pages, ASCIIDoc was designed for complex publishing projects with support for a much wider array of document domain structures such as tables, definition lists, and tables of contents.

## LaTeX

LaTeX is a document domain markup language used extensively in academia and scientific publishing. It is not based on XML syntax but on the syntax of TeX, a typesetting system developed by Donald Knuth in 1978.<sup>3</sup> Here is an example of LaTex, from Wikipedia:

```
\documentclass[12pt]{article}
\usepackage{amsmath}
\title{\LaTeX}
\date{}
\begin{document}
```

---

<sup>3</sup><https://en.wikipedia.org/wiki/LaTeX>

```
\maketitle
\LaTeX{} is a document preparation system for
the \TeX{} typesetting program. It offers
programmable desktop publishing features and
extensive facilities for automating most
aspects of typesetting and desktop publishing,
including numbering and cross-referencing,
tables and figures, page layout,
bibliographies, and much more. \LaTeX{} was
originally written in 1984 by Leslie Lamport
and has become the dominant method for using
\TeX; few people write in plain \TeX{} anymore.
The current version is \LaTeXe.
```

```
% This is a comment, not shown in final output.
% The following shows typesetting power of LaTeX:
\begin{align}
E_0 &= mc^2 \\
E &= \frac{mc^2}{\sqrt{1 - \frac{v^2}{c^2}}} \\
\end{align}
\end{document}
```

Here is how that markup is rendered:<sup>4</sup>



  $\text{\LaTeX}$  is a document preparation system for the  $\text{\TeX}$  typesetting program. It offers programmable desktop publishing features and extensive facilities for automating most aspects of typesetting and desktop publishing, including numbering and cross-referencing, tables and figures, page layout, bibliographies, and much more.  $\text{\LaTeX}$  was originally written in 1984 by Leslie Lamport and has become the dominant method for using  $\text{\TeX}$ ; few people write in plain  $\text{\TeX}$  anymore. The current version is  $\text{\LaTeX} 2_{\varepsilon}$ .

$$E_0 = mc^2 \tag{1}$$

$$E = \frac{mc^2}{\sqrt{1 - \frac{v^2}{c^2}}} \tag{2}$$

It is the markup for the equation that shows why  $\text{\TeX}$  is popular for academic and scientific publishing. While not exactly transparent, the markup is compact and functionally lucid for anyone with a little experience with it.

---

<sup>4</sup>By The original uploader was Bakkedal at English Wikipedia - Own work, CC BY-SA 2.5, <https://commons.wikimedia.org/w/index.php?curid=30044147>

Wikipedia offers a comparison of various math markup formats which shows how big a difference syntax can make to the lucidity of markup language in some cases.

For the equation:

$$x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$

The LaTeX markup is:

```
x=\frac{-b \pm \sqrt{b^2 - 4ac}}{2a}
```

Whereas the XML-based MathML version looks like this:

```
<math mode="display" xmlns="http://www.w3.org/1998/Math/MathML">
<semantics>
 <mrow>
 <mi>x</mi>
 <mo>=</mo>
 <mfrac>
 <mrow>
 <mo form="prefix">−<!-- - --></mo>
 <mi>b</mi>
 <mo>±<!-- ± --></mo>
 <msqrt>
 <msup>
 <mi>b</mi>
 <mn>2</mn>
 </msup>
 <mo>−<!-- - --></mo>
 <mn>4</mn>
 <mo>⁢<!-- ⁢ --></mo>
 <mi>a</mi>
 <mo>⁢<!-- ⁢ --></mo>
 <mi>c</mi>
 </msqrt>
 </mrow>
 <mrow>
 <mn>2</mn>
 <mo>⁢<!-- ⁢ --></mo>
 </mrow>
 </mfrac>
 </mrow>
</semantics>
```

```
<mi>a</mi>
</mrow>
</mfrac>
</mrow>
</semantics>
</math>
```

Clearly MathML was not designed with the idea that anyone would ever try to write it raw. It is intended to be the output of a graphical equation editor.<sup>5</sup> You might well choose to use a graphical equation editor to create LaTeX math markup as well, but it is certainly possible to write it and read it in raw LaTeX.

LaTeX is not as lightweight a language as Markdown. Its markup is almost entirely explicit (except for paragraphs, which are delineated by blank lines just in Markdown). But it is certainly lighter weight in its syntax compared to XML-based languages and has much greater functional lucidity. There are a number of LaTeX editors, but they tend to be the same style of side-by-side style of editors also popular for MarkDown. That is, the author writes in LaTeX syntax and a preview is generated continuously in a separate panel. Sufficient functional lucidity to be able to write in raw markup rather than needing a graphical editor is the hallmark of a lightweight markup language.

## Subject Domain Languages

So far we have looked at languages that are primarily document domain oriented. The document domain is an obvious choice for a public language since the use of common document types like books and articles is widespread. But there are a number of public subject domain languages as well. One example that we have looked at before (Chapter 23, *Merge*) is JavaDoc. Here is the example we looked at there:

```
/**
 * Validates a chess move.
 *
 * Use {@link #doMove(int theFromFile, int theFromRank,
 * int theToFile, int theToRank)} to move a piece.
 *
 * @param theFromFile file from which a piece is being moved
 * @param theFromRank rank from which a piece is being moved
 * @param theToFile file to which a piece is being moved
 * @param theToRank rank to which a piece is being moved
 * @return true if the move is valid, otherwise false
 */
```

---

<sup>5</sup>Interestingly, MathML comes in two different flavors. Presentation MathML is a media domain language describing how an equation is presented. Content MathML is a subject domain language describing what it means.

```
boolean isValidMove(int theFromFile, int theFromRank,
int theToFile, int theToRank) {
 // ...body
}
```

Not only does JavaDoc have subject domain tags for parameters and return values, it effectively incorporates the Java code itself (all computer programs are structured text). Thus the JavaDoc processor will pull information from the function header itself to incorporate into the output.

There are a number of similar languages for documenting different programming languages, such as Doxygen which is used for multiple languages. Wikipedia maintains an extensive list at: [https://en.wikipedia.org/wiki/Comparison\\_of\\_documentation\\_generators](https://en.wikipedia.org/wiki/Comparison_of_documentation_generators).

It is difficult to find public subject-domain lightweight markup languages outside the realm of programming language and API documentation. This is probably because only programmers are likely to write their own parser in order to create a markup language. Most other people are going to choose an extensible language as a base, which today usually means XML. Part of my motivation for creating SAM is to provide a way to create subject domain languages with lightweight syntax.

# SAM

I've talked about SAM a couple of times now, and said that many of the examples in this book are written in SAM. That is true, in the same sense that some of the examples are written in XML. But this is not the same thing as saying that an example is written in reStructuredText. reStructuredText is a hybrid document domain markup language that predefines a large range of document domain structures. You can write a complete document in reStructuredText without defining any new structures. In other words, you can treat reStructuredText as a purely concrete language, and that is exactly what most people

SAM, on the other hand, cannot be used as a concrete language out of the box. While it has concrete lightweight syntax features, they need to be wrapped in a containing block, and SAM does not predefine any blocks. The most complex SAM document you can write without defining any blocks is a single paragraph or list. You need a block structure such as `topic` or `chapter` or `recipe` to contain everything and SAM requires you define those structures for yourself.

In this respect, SAM is much more like XML than reStructuredText. It is an extensible language that is intended to be used only when extended. It does not belong in this discussion of concrete lightweight markup languages, therefore. I'll get back to it in Chapter 43, *Extensible and Constrainable Languages*.

---

# Chapter 42. Heavyweight markup languages

I am using the term “heavyweight” here as an obvious contrast to the commonly used “lightweight”, even though the term “heavyweight” is not used commonly of markup languages. Nonetheless, it fits. Both the abstract language XML and the concrete languages like DocBook and DITA are heavyweights in the sense that they have a lot of capability that comes at the expense of a large footprint.

Having said that, I should make the distinction between the heavyweight syntax of XML and the heavyweight semantics of a DITA or DocBook. It would certainly be possible to create the semantics of DITA or DocBook in a more lightweight syntax. And it is certainly possible to create very simple markup languages (with semantics much more lightweight than something like reStructuredtext, for instance) using XML. Despite this, there is a definite connection between heavyweight syntax and heavyweight semantics, perhaps because the more heavyweight languages have more need of the capabilities of a fully abstract syntax of XML and the processing tools that go with it.

I’m going to briefly survey some of the heavyweight languages. One thing to note about heavyweight languages is that they often contain structures from more than one domain, though their core is usually in the document domain. But they typically contain some media domain structures for things like tables that are hard to abstract from the media domain in a generic way. They typically contain some subject domain structures, typically related to technology, since many heavyweight languages originated for documenting technical products. Finally, most contain some management domain structures, particularly for things like conditional text.

If the point of structured writing is improve the quality of content through the application of constraints, why is the structured writing landscape dominated by a few very large and quite loosely constrained markup languages?

- Partly because, as I pointed out at the beginning, all writing is structured and when someone says they are moving to structured writing, what they mean is that they are adding a more structure to their writing than they had before. DocBook may not be a highly constrained language, but it is lot more constrained than Microsoft Word or InDesign and somewhat more constrained than FrameMaker.
- Partly because a lot of the adoption of structured writing it not motivated primarily by content quality but by a desire to improve content management, particularly content reuse. While it is possible to do these things without resorting to structured writing for a format, structured writing formats ease the integration of the various parts. They also ease fears about having your content locked into the system of a single vendor.
- Partly it is because constraints are onerous if you don’t get them right, and the benefits of getting them right are often under-appreciated, especially in content management

applications where the consequences of a lack of constraints tends to show up years down the road (and it is all too easy to blame the problems that emerge on human failure rather than poor system design).

For all these reasons, it is worthwhile to look at where the big public languages fit in the structured writing picture.

For large systems like DocBook, DITA, and S1000D, there is not nearly enough space in this book to do them full justice or to fully characterize them in terms of the structured writing domains and algorithms described in this book. This chapter is therefore not to be taken as a buyers guide. Rather, this book as a whole is an attempt to provide an framework for thinking and talking about structured writing that will allow you to understand your requirements independently of any system, and then to evaluate, compare, and contrast systems in more or less neutral terms.

## DITA

There are two ways of looking at DITA. You can look at it as a complete structured writing system which can be used more or less out of the box. (Even packaged applications like Word or FrameMaker are not used completely out of the box for serious content creation: some customization of styles and output format is needed at least, and the same is true of DITA.)

Alternatively, you can look at as what it's name proclaims it to be: an information typing architecture. The acronym DITA stand for Darwin Information Typing Architecture, with the word "Darwin" representing DITA's approach to the extensibility of markup: specialization.

With out-of-the-box DITA, you get a fixed set of topic types provided by the DITA specification and implemented in the DITA Open Toolkit and other tools. With DITA as an information typing architecture, you get the capability to create a unbounded number of information types. I will discuss DITA as an information typing architecture in Chapter 43, *Extensible and Constrainable Languages*. Here I will look at out-of-the-box DITA.

Out-of-the-box DITA comes in three main forms.

1. The DITA Open Toolkit. You can download the DITA Open Toolkit for free and use it to produce content. The formatting stylesheets that come with the toolkit are very basic, so you will likely want to do some customization of the output as a minimum.
2. Packaged DITA tools. There are a variety of tools that package DITA. Most of these are essentially content management systems of one degree of sophistication or another. These may add additional capabilities over what is supplied by the DITA Open Toolkit and may hide the underlying DITA structures to one extent or another. I don't intend to say anything about any of these tools here.

3. Customized DITA systems. You may be handed a customized system created by a consultant or vendor, generally built using one of the available DITA CMS platforms. This may have been extended to provide new topic types, new output formats, or new management facilities. These could work in almost any way and it may not even be obvious that DITA has been used in their construction. Obviously there is not much that we can say about them here.

The key features of out-of-the-box DITA that will determine how well its fits with your needs are its topic model and its focus on the reuse algorithm. The description of the document-domain/management-domain approach in to reuse Chapter 13, *Reuse* is based on the DITA model, which provides comprehensive support in those domains.

As described in Chapter 36, *Blocks, fragments, paragraphs, and phrases*, the DITA topic model is based on the concept of information typing, which is the idea that information can be usefully broken down into different abstract types, and that there is value in clearly separating the different types. One of the problems with this theory, and consequently with the application of DITA's topic model, is that it is not clear how big an information type is, and whether an information type constitutes a narrative block or just a semantic block. Specializing DITA may allow you to be more specific on this point, but if you are using out-of-the-box DITA you are probably using the basic concept, task, and reference topic types (though out-of-the-box DITA now includes a number of other topic types such as Machine industry task and Troubleshooting). There is no clear answer to whether these topic types constitute semantic blocks or narrative blocks. In practice they are used both ways.

The principal thing that sets out-of-the-box DITA apart from other approaches to structured writing is its map and topic architecture. In most other systems, the unit that the writer writes and the unit that the reader reads are the same. For very long works, there may be a mechanism for breaking up and assembling pieces. For instance, in DocBook, you can write a book using a book document type in which you can include various chapter document types to create a complete book out of multiple files. In this model, a chapter is a narrative unit, though it may also be an element of a larger narrative unit which is the whole book. But DITA factors out this document assembly function into a separate file type, the map. In a DocBook book document, there is a lot of book content in addition to the included chapters. Indeed, you could write the entire book in one file if you wanted to. The content model of a DocBook document is described by a single schema and the content model of the chapters is simply part of the content model of the book. In other words, a DocBook book is a single document structure that just happens to be made up of individual files. A DITA map file, on the other hand, is an independent structure. It does not create a single logical document structure. It does not usually contain any actual content, and you can't write an entire book in a single map file. Instead, a map file is an instruction to a publishing tool chain about how to assemble a larger work out of component pieces.<sup>1</sup>

---

<sup>1</sup>DITA does also support the DocBook model of including one topic directly in another, but this practice is discouraged.

This distinction is very important. In the DocBook model, there is a continuity of constraint between the book and its chapters. In DITA, the constraints on the map and the constraints on the topics in the map, are completely separate. This means that in DITA, the topic is the largest unit of content to which content constraints can be applied (at least in the conventional way).

Maps are structured like tree so they can construct hierarchies an arbitrary number of layers deep. This means you have a choice about what parts of your structure you create using a map and what part you create inside a topic. If you have a list of four items, each of which needs two or three paragraphs of description, do you create one topic with the list of four items in it or do you create one topic for each item and then tie them together using a map? This is particularly important when we remember that the topic is the largest using of content constraint in DITA. If we break the content down to this fine a level, we lose the ability to apply constraints to it. On the other hand, we are now able to reuse each of these items separately.

This presents something of a dilemma. We have already talked about structure writing as dividing content into blocks and made a distinction between semantic blocks and narrative blocks. (Chapter 36, *Blocks, fragments, paragraphs, and phrases*) In the design of a markup language, narrative blocks are made up of semantic blocks which may be made up of smaller semantic blocks. This works fine for developing the structure of a narrative block, which is the work that will be presented whole to the reader. In that scenario, the unit that the writer writes is the narrative block. The semantic blocks are just elements of the model.

But things become more difficult when you attempt to do fine-grained reuse of content. Then you may want to write individual semantic blocks and combine them to produce narrative blocks. DITA will let you do this one of two ways. The first (which is frequently discouraged) is to nest one topic inside another. The second is to combine topics using a map, with the map representing the narrative block. However, DITA does not provide a high-level way to constrain the structure of a narrative block that is built this way.

If you want a constrained narrative block, you have to model it as a single DITA topic type. You can certainly do this by specializing from the base `topic` topic type, but in doing so you will probably move away from the “information typing” idea of keeping different types of information separate, as a full narrative topic often requires different types of information (as in the recipe example we have used so frequently).

This leads to the confusion about whether a DITA topic is a semantic block or a narrative block. For people who use out-of-the-box DITA this can be a problem because the default Web presentation of DITA places each topic on a separate page, which is not an appropriate presentation if your DITA topics are not narrative blocks. If your narrative block is made up of multiple DITA topics strung together by a map, and you want that narrative block to appear on a single page, you need to use a procedure called chunking,

which is not as straightforward as it should be. (Chunking is one of the things on the agenda to be fixed in DITA 2.0.<sup>2</sup>)

The idea that blocks are reusable is a very attractive one. But it is important to think through exactly what the reusable unit of content is. It is one thing to reuse narrative blocks whole (perhaps with some variations in the text). It is quite a different thing to reuse semantic blocks below the level of the narrative block, particularly if it is important to constrain the narrative block or to apply any of the other structured writing algorithms at the level of the narrative block.

Quality should be a serious concern with this model. If narrative blocks are being assembled out of smaller reusable units without proper attention to the narrative integrity or completeness and consistency of the result, quality can be seriously affected. If the author no longer sees, thinks, or works in the context of the narrative block, and if the structure of the narrative block is not constrained, content quality is very difficult to maintain.

DITA, as a technology, does not prevent you from working in whole narrative blocks, or from constraining your blocks in any way you want (using its information typing capabilities). But the block and map model (whether implemented by DITA or any other system) presents this inherent tension between creating smaller semantic blocks to optimize for reuse vs creating constrainable narrative blocks to optimize for content quality.

A related note here is that in a reuse scenario, your motives for constraining the semantic block may be different from your motives for constraining the narrative block. The reasons for constraining the semantic block might be to adhere to the theory of information typing, or to enhance the composability of the reusable blocks. The reasons for constraining the narrative block might be to ensure the quality or consistency of the information presented to the user.

But merely doing reuse of content blocks does not require either kind of constraint. The constraints may improve quality and reliability of the system if used correctly and consistently, but the actual act of composing larger blocks out of smaller blocks does not require them. This has led many organization to use DITA for its reuse capabilities without paying any particular attention to its constraint capabilities or its information typing roots. People taking this approach will sometime write their content in the base `topic` topic type rather than a more constrained specialization.

The growing popularity of this approach to reuse has led to the development of alternatives to DITA that provide the same reuse-management capabilities but remove the constraint mechanisms. One example of this trend is Paligo, a reuse-focused component content management system that uses DocBook as its underlying content

---

<sup>2</sup><http://docs.oasis-open.org/dita/dita-1.3-why-three-editions/v1.0/cn01/dita-1.3-why-three-editions-v1.0-cn01.html#future-of-dita>

format, specifically for the purpose of minimizing constraints on the content.<sup>3</sup> Such systems can reduce the up-front complexity of component-based content-reuse, though possibly at the expense of costs down the road due the failure to apply constraints up front.

DITA's sweet spot, therefore, would appear to be content reuse scenarios in which you want to place constraints on the reusable content units that follow the "information typing" model, but can live without placing constraints on the narrative blocks that are built from those units.

Note, though, that this analysis is only focused on DITA as a tool for structured writing. DITA is also a tool for content management and its role in making content management systems work effectively has to be evaluated separately. And, as pointed out above, DITA provides a high-level tools for information typing which can be used for things outside this sweet spot, and potentially having nothing to do with its reuse features. How broad this range of applicability is depends largely on where the sweet spots of other tools lie. Your needs may not coincide perfectly with the sweet spot of any one tool. At that point, the right choice is the one that can be adapted for your needs at the least cost. That calculation is outside the scope of this book. I will try to point out where costs lie in each alternative, but quantifying them for individual situations is an exercise left to the reader.

## DocBook

DocBook is an extensive, largely document-domain language with a long history and an extensive body of processing tools and support. As we have noted, DocBook is not a tightly constrained language. Instead it is focused on providing very broad capability for describing document structures.

Unlike DITA, DocBook does not ascribe to any information typing theory. It does not have an opinion about how content should be written or organized. It is very much about the structure of books, and leaves it to the author to decide what the rhetorical structure of the text should be. In other words, DocBook makes no attempt to constrain the rhetorical structure of a work, and in fact makes every attempt to avoid constraining it.

The result, however, is an extremely complex system that can be quite challenging to learn and use. Because of this, writers often use simplified subsets of DocBook. (Where DITA is sometimes customized by the addition of elements, DocBook is often customized by their subtraction.) However, DocBook remains popular with many for its lack of constraint combined with its rich feature set.

Because of its lack of constraint, DocBook is not a particularly great fit with the idea of structured writing as a means improve content quality through the application of constraints. However, it can play a very useful role in a structured writing tools chain as

---

<sup>3</sup><http://idratherbewriting.com/2016/08/01/paligo-the-story-xml-ccms-in-the-cloud/>

a language for the presentation algorithm. This is exactly how it is used in the production of this book. The book is written in SAM is a small, constrained language developed just for the purpose, which is then transformed by the presentation algorithm into DocBook, which then feeds the publisher's standard publishing tools. The DocBook created by this method matches the publisher's exact specifications as required to make the tools work correctly.

This is a more reliable process than if I had written the book in DocBook directly. I wrote my previous book in DocBook (an experience that contributed to my decision to develop SAM) but it took a lot of revision to get the DocBook I wrote into the form that the publication process required. In other words, the publishing process has a set of constraints that are not enforced by DocBook itself, and have to be imposed by human oversight and editing when an author writes in DocBook. But in my highly constrained SAM-based markup language, all those constraints were factored out, which enabled me to translate it reliably into the DocBook that the publisher needed.

## S1000D

S1000D is a specification developed in the aviation and defense industries specifically for the complex documentation tasks of those industries, and intended to support the development of the Interactive Electronic Technical Manuals (IETMs) that are typically required in that space. While it obviously has a fair amount of subject domain structures for the target domains, it also has media domain structures targeted at the production of IETMs and extensive management domain structures designed to support the common source database (CSDB), the content management architecture which is part of the S1000D specification. S1000D, in other words, is much more than a structured writing format. It is a specification for a complete document production system for a specific industrial sector.

## HTML

HTML is widely used as an authoring format for content. For the most part this is a pure media domain usage: people writing for the web in its native format, often using a WYSIWYG HTML editor.

But HTML is still a document domain language, and efforts have been made over the years to factor out the media domain aspects of the languages and leave the formatting to CSS stylesheets. This makes HTML a legitimate document domain markup language. In particular, people interested in using HTML this way often use XHTML, the version of HTML that is a valid instance of XML. Being an instance of XML is important because it means you can write XHTML in an XML editor and process it with XML processing tools. This means that you can potentially publish content written in XHTML by processing it into other formats or by modifying its structure for use in different HTML-based media such as the Web and ebooks.

# Subject domain languages

There are hundreds of subject domain languages written in XML. Most of these are more data oriented than content oriented, but you might be able to derive content from some of them using the extract and merge algorithm. Wikipedia maintains an extensive list at [https://en.wikipedia.org/wiki/List\\_of\\_markup\\_languages](https://en.wikipedia.org/wiki/List_of_markup_languages).

One interesting example is BeerXML:

```
<?xml version="1.0" encoding="ISO-8859-1"?>
<RECIPES>
 <RECIPE>
 <NAME>Dry Stout</NAME>
 <VERSION>1</VERSION>
 <TYPE>All Grain</TYPE>
 <BREWER>Brad Smith</BREWER>
 <BATCH_SIZE>18.93</BATCH_SIZE>
 <BOIL_SIZE>20.82</BOIL_SIZE>
 <BOIL_TIME>60.0</BOIL_TIME>
 <EFFICIENCY>72.0</EFFICIENCY>
 <TASTE_NOTES>Nice dry Irish stout with a warm body but low st...
 <RATING>41</RATING>
 <DATE>3 Jan 04</DATE>
 <OG>1.036</OG>
 <FG>1.012</FG>
 <CARBONATION>2.1</CARBONATION>
 <CARBONATION_USED>Kegged</CARBONATION_USED>
 <AGE>24.0</AGE>
 <AGE_TEMP>17.0</AGE_TEMP>
 <FERMENTATION_STAGES>2</FERMENTATION_STAGES>
 <STYLE>
 <NAME>Dry Stout</NAME>
 <CATEGORY>Stout</CATEGORY>
 <CATEGORY_NUMBER>16</CATEGORY_NUMBER>
 <STYLE_LETTER>A</STYLE_LETTER>
 <STYLE_GUIDE>BJCP</STYLE_GUIDE>
 <VERSION>1</VERSION>
 <TYPE>Ale</TYPE>
 <OG_MIN>1.035</OG_MIN>
 <OG_MAX>1.050</OG_MAX>
 <FG_MIN>1.007</FG_MIN>
 <FG_MAX>1.011</FG_MAX>
 <IBU_MIN>30.0</IBU_MIN>
```

```

<IBU_MAX>50.0</IBU_MAX>
<COLOR_MIN>35.0</COLOR_MIN>
<COLOR_MAX>200.0</COLOR_MAX>
<ABV_MIN>3.2</ABV_MIN>
<ABV_MAX>5.5</ABV_MAX>
<CARB_MIN>1.6</CARB_MIN>
<CARB_MAX>2.1</CARB_MAX>
<NOTES>Famous Irish Stout. Dry, roasted, almost coffee ...
</STYLE>
<HOPS>
 <HOP>
 <NAME>Goldings, East Kent</NAME>
 <VERSION>1</VERSION>
 <ALPHA>5.0</ALPHA>
 <AMOUNT>0.0638</AMOUNT>
 <USE>Boil</USE>
 <TIME>60.0</TIME>
 <NOTES>Great all purpose UK hop for ales, stouts, por...
 </HOP>
</HOPS>
<FERMENTABLES>
 <FERMENTABLE>
 <NAME>Pale Malt (2 row) UK</NAME>
 <VERSION>1</VERSION>
 <AMOUNT>2.27</AMOUNT>
 <TYPE>Grain</TYPE>
 <YIELD>78.0</YIELD>
 <COLOR>3.0</COLOR>
 <ORIGIN>United Kingdom</ORIGIN>
 <SUPPLIER>Fussybrewer Malting</SUPPLIER>
 <NOTES>All purpose base malt for English styles</NOT...
 <COARSE_FINE_DIFF>1.5</COARSE_FINE_DIFF>
 <MOISTURE>4.0</MOISTURE>
 <DIASTATIC_POWER>45.0</DIASTATIC_POWER>
 <PROTEIN>10.2</PROTEIN>
 <MAX_IN_BATCH>100.0</MAX_IN_BATCH>
 </FERMENTABLE>
 ...
</FERMENTABLES>
<MISCS>
 <MISC>
 <NAME>Irish Moss</NAME>
 <VERSION>1</VERSION>
 <TYPE>Finishing</TYPE>

```

```

<USE>Boil</USE>
<TIME>15.0</TIME>
<AMOUNT>0.010</AMOUNT>
<NOTES>Used as a clarifying agent during the last few
 </MISC>
</MISCS>
<WATERS>
 <WATER>
 <NAME>Burton on Trent, UK</NAME>
 <VERSION>1</VERSION>
 <AMOUNT>20.0</AMOUNT>
 <CALCIUM>295.0</CALCIUM>
 <MAGNESIUM>45.0</MAGNESIUM>
 <SODIUM>55.0</SODIUM>
 <SULFATE>725.0</SULFATE>
 <CHLORIDE>25.0</CHLORIDE>
 <BICARBONATE>300.0</BICARBONATE>
 <PH>8.0</PH>
 <NOTES> Use for distinctive pale ales strongly hopped
 </NOTES>
 </WATER>
</WATERS>
<YEASTS>
 <YEAST>
 <NAME>Irish Ale</NAME>
 <TYPE>Ale</TYPE>
 <VERSION>1</VERSION>
 <FORM>Liquid</FORM>
 <AMOUNT>0.250</AMOUNT>
 <LABORATORY>Wyeast Labs</LABORATORY>
 <PRODUCT_ID>1084</PRODUCT_ID>
 <MIN_TEMPERATURE>16.7</MIN_TEMPERATURE>
 <MAX_TEMPERATURE>22.2</MAX_TEMPERATURE>
 <ATTENUATION>73.0</ATTENUATION>
 <NOTES>Dry, fruity flavor characteristic of stouts.
 <BEST_FOR>Irish Dry Stouts</BEST_FOR>
 <FLOCCULATION>Medium</FLOCCULATION>
 </YEAST>
</YEASTS>
<MASH>
 <NAME>Single Step Infusion, 68 C</NAME>
 <VERSION>1</VERSION>
 <GRAIN_TEMP>22.0</GRAIN_TEMP>
 <MASH_STEPS>
 <MASH_STEP>

```

```
<NAME>Conversion Step, 68C </NAME>
<VERSION>1</VERSION>
<TYPE>Infusion</TYPE>
<STEP_TEMP>68.0</STEP_TEMP>
<STEP_TIME>60.0</STEP_TIME>
<INFUSE_AMOUNT>10.0</INFUSE_AMOUNT>
</MASH_STEP>
</MASH_STEPS>
</MASH>
</RECIPE>
</RECIPES>
```

As you can see, this is a recipe language, but a much more detailed and precise recipe language than any that we have looked at before. It is a recipe language for beer and for beer only. It takes a lot of information that might normally be written in paragraphs and breaks it up into precisely labeled fields. This means, of course, that it is far more constrained than generic recipe markup. You can do far more conformance testing on this recipe than you could on a normal one, and manipulate and query the information in far more ways that you could with a conventional recipe format.

---

# Chapter 43. Extensible and Constraintable Languages

The languages we have looked at to this point are publicly specified and have existing tool chains. Some of them are more constrained than others, and some of them support different structured writing algorithms. Choosing one of them makes sense if the constraints they express and the algorithms they support are the ones that matter most to your organization, and they support them sufficiently well to meet your needs. If not, you will need to create your own structures.

There are essentially three options for this:

- Create your own language entirely from scratch, creating both the syntax and the semantics.
- Use an existing definition of markup syntax, such as XML or SAM, and create your own semantics by defining named structures using that syntax.
- Take an existing markup language with extensible and/or constraintable semantics, such as DITA or DocBook and extend and/or constrain it to meet your needs.

Each of these approaches has its merits and its drawbacks. For instance, creating a new language entirely from scratch may enable to achieve exceptional functional lucidity for a particular type of information, extending/constraining an existing language can save you a lot of tool development costs, while defining your own semantics based on an existing syntax may let you find the right balance between functional lucidity and development costs.

Also note that this is not an all or nothing choice. Your system might use a combination of these techniques. You might even combine them in one markup language by defining the content on a particular element or block as a pattern with its own particular syntax for a specific purpose, as discussed in Chapter 39, *Patterns*.

## XML

The X in XML stands for eXtensible, but, as we noted in Chapter 40, *Markup* XML is an abstract language that does not define any document structures itself. Extension in XML, therefore, is extension from zero. Starting from scratch. You can define the structure of a new XML tagging language using one of the several available schema languages. We will look at schema languages in Chapter 44, *Constraint Languages*.

# DITA

DITA is somewhat unique among markup languages in that it was designed for extension from the beginning. In fact, it is something of a misnomer to call DITA a markup language. DITA actually calls itself an information typing architecture. What is an information typing architecture? DITA is really the only thing that calls itself by this name, so to a certain extent we have to derive the definition from the properties of this one example.

XML schema languages are information typing languages. Their sole function is to define information types. So what does an information typing architecture provide over and above what an information typing language provides?

There are plenty of precedents for this distinction. The programming world makes use of architectures and frameworks to abstract certain types of operation to a higher level. You could program these functions from scratch, but the architecture or framework is designed to save time and possibly avoid errors.

In DITA's case, the architecture consists of a set of predefined tagging languages that are intended as a basis for extension through a mechanism known as specialization. It also includes an extensive set of management domain markup and the specification of the behavior it should produce, as well as a facility (maps) for assembling information products, and a facility (subject schema) for managing metadata.

In other words, it predefines a range of structures, semantics, and operations that you might need in establishing an information architecture and then provides a way for you to build from there.

As with any other architecture, its usefulness depend on how well the predefined structures, semantics, and operations suit your needs, how easy the extension mechanism is to use, and how reliable the available implementations are.

It is common in the software world for there to be many competing architectures with different sweet spots. Because an architecture is essentially a series of guesses about what a variety of systems may have in common, different architectures may be constructed very differently to cover different sets of commonalities among diverse projects and you may not see equivalent architectural features from one architecture to another.

There are not a lot of information typing architectures. The only other one I am aware of is the one I am developing myself, which is called SPFE. SPFE, however, is a very different kind of architecture from DITA. DITA has grown into a very large and complex architecture over the year, and it is out of scope for this book to attempt to describe its architecture in full. All I shall attempt to do here is to do a basic mapping of some key features of the DITA architecture to the structured writing concepts explored in this book.

Inherent in the process of constructing an architecture is that you constrain the field in certain ways. Architectures move functionality to a higher level by choosing some options and rejecting others. I said that an XML schema is an information typing language. But a schema can define a markup language for any purpose at all, such as recording transfers between banks or storing the configuration options of an editor. Describing banking transactions or storing configuration options is not within the scope of the information typing that the DITA architecture was designed for. DITA therefore has a more restricted definition of “information typing”. The DITA specification defines “information typing” this way:

Information typing is the practice of identifying types of topics, such as concept, reference, and task, to clearly distinguish between different types of information.

—<http://docs.oasis-open.org/dita/dita/v1.3/csd01/part3-all-inclusive/archSpec/base/information-typing.html>

Unfortunately this definition is largely circular. But it does help establish a scale. Information typing is about defining topic types. The spec goes on to define the purpose of information typing:

Information typing is a practice designed to keep documentation focused and modular, thus making it clearer to readers, easier to search and navigate, and more suitable for reuse.

DITA information typing then, is not as general as structured writing. It is focused on information at a particular scale and on a subset of the structured writing algorithms. (That does not mean that it makes it impossible to work at other scales or implement other algorithms, it just means that these are the areas that the architecture supports at a higher level.)

Out-of-the-box DITA is commonly associated with the idea that there are just three information types, task, concept, and reference. The DITA spec makes it clear that this is not the intention of DITA as an information typing architecture.

DITA currently defines a small set of well-established information types that reflects common practices in certain business domains, for example, technical communication and instruction and assessment. However, the set of possible information types is unbounded. Through the mechanism of specialization, new information types can be defined as specializations of the base topic type (`<topic>`) or as refinements of existing topics types, for example, `<concept>`, `<task>`, `<reference>`, or `<learningContent>`.

As we have noted many times, many of the structured writing algorithms work best with more specific markup, particularly markup in the subject domain. The ability to create an unbounded set of information types is therefore very relevant to getting the most out of structured writing.

Clearly, though, one does not need an information typing architecture to define an information type. You can, as John Gruber did with MarkDown, sit down and sketch out a set of structures and a syntax to represent them, and then write a program to process them. With an abstract language like XML or SAM, you can create a new information type by defining a set of named elements and attributes using a schema language. How does using a higher level “information typing architecture” like DITA change this process?

First and foremost, it means that you don’t start from scratch. All topic types in DITA are derived from a base topic type called `topic` by a process called specialization.

What is specialization? We noted that XML is an abstract language, meaning that its syntax defines abstract structures that do not occur in documents: elements, attributes, etc. To create a markup language in XML, you define named elements and attributes for the structures you are creating. Thus in DocBook `para` is a type of element. `para` has what is called an “is-a” relationship to elements. This is a type of specialization. `para` “is-an” element, but it is a special type of element. An XML parser will process it generically as an element, reporting its name to the application layer. The application layer will have a rule that processes just this specialized `para` element (and not the also specialized but different `title` element).

DITA specialization follows the same principle, but moves it up a level. The base `topic` topic type is the abstract structure. More specific types like `knitting-pattern` or `ingredients-list` are specializations of `topic` (or of other topic types that are specializations of `topic`). A generic DITA processor can process them as a instance of `topic`, but it would require additional code to process them specifically as `knitting-pattern` or `ingredient-list`. Each of these specialized types has an “is-a” relationship with the type it was specialized from. So `knitting-pattern` “is-a” `topic`.

But DITA specialization is different from simply creating new named elements in XML in a number of ways.

First, the base DITA topic type is not an abstraction like an XML element. You cannot create an XML element without inventing a name for it. The base DITA topic type, on the other hand, is a fully implemented topic type that you can use directly. You can, and people do, write directly in the base topic type without inventing anything new. We noted in the discussion of rhetorical structure Chapter 33, *Rhetorical Structure* that it is sometimes easy to treat what is intended as a meta model as a generic model. This is the case here. All topic types in DITA are derived by specialization from the generic `topic` type. They all have an “is-a” relationship to this generic type.

One consequence of this is that while the set of topic types you can create with DITA may be unbounded, it is not necessarily universal. The generic topic type has specific characteristics and if a specialized topic has an is-a relationship to the generic topic, the structures in the specialized topics have a corresponding is-a relationship to structures in the generic type. Thus there can be information types that cannot reasonably be said to

have an is-a relationship to a DITA generic topic. That is, there can be information types such that processing them using the code of the type they are specialized from would produce no meaningful result. For example, an information type that factored out most of the text that would appear in the published version would not process meaningfully as a generic topic because the factored-out text would not be restored.

To specialize a topic type, you specialize the root element and any child elements or attributes that you need to define your new topic type. Each specialized element or attribute should have an is-a relationship to the element it specializes. Thus a procedure element might be a specialization of an ordered list element and its step elements might be specializations of a list item element. (You can see that in this case, processing a procedure as an ordered list would produce meaningful output, but that you might also want to specialize the output of steps in a procedure, perhaps by prefixing each step with "Step 1:" rather than as "1.".)

The second way in which specialization differs from giving names to abstract elements is that specialization is recursive. That is, suppose you have a topic type `animal-description`, which is a specialization of `topic`. You want to impose additional constraints on the description of different types of animal, so you create more specialized types `fish-description` and `mammal-description` which are specializations of `animal-description` (and would be processed like an `animal-description` if no other processing were specified for them). Then you might decide to impose still more constraints on the description of different kinds of mammals, so you create a type, `horse-description` that is a specialization of `mammal-description`. This type will be processed as a `mammal-description` if no specific processing is provided for `horse-description`, as `animal-description` if no specific `mammal-description` processing is provided, and as `topic` if no specific `animal-description` is provided.

The second way in which information typing in DITA differs from doing it from scratch is that DITA information types share a common approach to processing and to information architecture. In particular, they inherit a common set of management domain structures and their associated management semantics. It is possible to ignore all of these things, but if you do so, the value of using DITA for information typing is reduced because you then have to invent your own ways of doing these things.

As a generality, the fewer pieces of an architecture you use, the less value there is to basing your work on that architecture, both because you have more work to do, and because you take less advantage of the infrastructure or tools and expertise surrounding that architecture, and create a system that is less understandable to people versed in the architecture. All architectures come with overheads and if you don't use their features, you still have to live with their overheads, which adds cost and complexity to your system. Thus while you can use DITA and depart from the default DITA way of doing things, the value of using DITA diminishes the further you depart from the DITA way. The same would be true of any other information typing architecture.

# Specializing between domains

We have talked a lot about the benefits of moving content from the document domain to the subject domain as a means to factor out constraints in the document domain and as a way to enable multiple structured writing algorithms and improve functional lucidity.

DITA's base types clearly belong to the document domain. Clearly you can use specialization to create more specialized document domain structures. But can you use specialization to create subject-domain structures that factor out aspects of the document domain?

In formal terms the answer is no. Subject domain information does not have an is-a relationship to document domain information precisely because it is the document domain structures that you factor out when you move to the subject domain.

Take the list of ingredients in a recipe. In the document domain, they could be presented as a list or as a table. In subject domain terms they are actually more of a table (database sense) than a list. That is, a set of records with a defined semantic structure:

```
ingredients:: ingredient, quantity, unit
eggs, 3, each
salt, 1, tsp
butter, .5, cup
```

How do you create this record structure by specializing document domain elements? What is the best starting point? A table is the most obvious candidate because it is structured like a set of records. However, the more common presentation of an ingredient list is as a list. But there is no structure in a list that supports dividing a list item into three named fields. The ingredient structure above is neither a list nor a table; it is a record set, a data structure whose contents could be presented in many different ways in the document domain, but is not a specialization of any of them.

I have made a distinction between two main types of subject domain languages, the rhetorical subject domain, were the focus is on how to order information on a particular subject, and the data-oriented subject domain, where the focus is on capturing the various pieces of information that are required on a subject while factoring out the presentation. Both present problems for specialization.

One of the elementary things that the rhetorical subject domain does is to factor out the titles of various sections of a topic, replacing the generic structure of sections and titles with specific names blocks whose titles can be supplied consistently by the presentation algorithm. These named sections do have a kind of is-a relationship to generic sections, in that they are sections. But this is not a true is-a relationship because a generic section requires a title, and a rhetorical block cannot have one. It therefore does not therefore have all the characteristics of the thing it is specialized from, and therefore it is not a section. It

is a named block of text that, in the document domain, would be expressed as a section, but that is not the same things as falling back to being a section. Processing such a block as a section would result in the text following on from that of the previous section with no title, thus lacking the defining document and media domain characteristic of a section.

For the data-oriented subject domain, which, let us remember, can include material recorded in databases and source code, there is no necessary relationship between the order of content in the source and the order of content in the document domain. Information that is in separate fields in the subject domain may need to be combined to form sentences by the presentation algorithm before it is even readable as a piece of content. Such structures are clearly have no is-a relationship to a DITA generic topic.

Does this mean that it is impossible to subject-domain structured writing using DITA specialization? No. Even if you are breaking the is-a relationship when you do, you can still use specialization as a mechanism for writing a schema for your subject domain language. But consider this: the value of an architecture depends on how much of it you can use in your project without modification. Modify too much of it and you will end up doing more work, and dealing with more complexity than if you had built from scratch. Why might the move to the subject domain force you to modify so much of the DITA architecture? As we have seen, the nature of almost all structured writing algorithms changes profoundly when you move your content from one domain. For instance, consider the differences between the linking and content reuse algorithms in the document domain and the subject domain. As long as you are creating document domain structures in DITA, much of its substantial linking and content reuse architecture will fully apply to what you are creating. But if you move to the subject domain, and intend to apply the subject domain approaches to linking and content reuse, that architecture will be of no use to you.

## DocBook

DocBook is not really extensible in the same sense as the other languages mentioned here, but it still deserves a mention. DocBook does not provide an extension mechanism like DITA's specialization. What it does provide is a deliberately modular construction that makes it easy to create new schemas that include elements from DocBook. In short, DocBook takes full advantage of the extensibility features built into XML schema languages.

Does the fact DocBook that does not invent its own extension mechanism means that it is not as extensible as DITA? No. By relying on XML's own extensibility features, which are both more comprehensive and lower level than DITA's specialization and constraint mechanisms, DocBook is as extensible as it is possible for any XML vocabulary to be.

Where it differs from DITA is that DocBook extensions are not DocBook and cannot be processed by standard DocBook tool chains. DITA's specialization mechanism means that a specialized topic will always pass through the DITA publication process, though

whether it will be presented in a useful or comprehensible way very much depends on how well the is-a relationship between specialization and base was maintained. If you would rather ensure that topics always pass through the publication process, even if the results are gibberish, DITA will support that. If you want to ensure that errors are raised if any structure is not recognized by the publishing tool chain (thus avoiding accidental gibberish, but potentially holding up the entire production process) then DocBook's extension mechanism will give you that.

Another aspect of DocBook customization deserves to be mentioned here even though it is not strictly speaking extension. DocBook has a huge tag set and it is quite conceivable that if you want a small constrained document domain markup languages that you can create one by sub-setting DocBook. DocBook provided for just about every document structure out there, so if you are building a document domain language, chances are the pieces you need are in there.

The great advantage of creating a new language as a subset of DocBook is that the result is also a valid DocBook document and can therefore be published by the DocBook tool chain. You will not have to write any algorithms at all if you take this approach. Creating a subset of DocBook can therefore allow you to impose more constraints and improve functional lucidity significantly compared to standard DocBook without having to write any processing code at all.

Technically speaking, any XML-based markup language is extensible in the same way that DocBook is. However, DocBook's structure, and the implementation of its schemas, was designed deliberately to support both extension and sub-setting of DocBook, something which is not true for many markup languages.

## RestructuredText

Restructured text defines blocks using directives.

```
.. image:: images/biohazard.png
:height: 100
:width: 200
:scale: 50
:alt: alternate text
```

It is extensible by adding new directives to the language. However, there is no schema language for RestructuredText. To create a new directive, you have to create the code that processes it.

There is an important distinction to be made between languages that are extensible by schema and those that are extensible by writing code to process the extension. If a language is extended by writing processing code for the extension, the only way to know if the input is valid is by processing it. If it raises a processing error, it is invalid.

If you have only one processor for a language, you can treat that processor as normative. That is, the definition of a correct file is any file that can be successfully be processed by the normative processor. The language, in other words, is defined by the processor. But if you have multiple processors, how do you determine who is at fault when a processor fails to process a given input file? Is the processor incorrect or the source file?

A schema creates a language definition that is independent of any processor. It is the schema that is normative, not any of the processors. If the source file is valid per the schema, the processor is at fault if it does not process that file correctly. If the source file is not valid per the schema, the blame lies with the source file.

In the case of RestructuredText, the capacity of the processor to be extended in this way is built into the processor architecture. It is not like you have to hack around in the code to add your extensions. There is a specific and well documented way to do it. But while RestructuredText allows you to extend it by adding new directives, it does not have a constraint mechanism. There is no mechanism (other than by hacking into the code) to restrict the use either of new directives or the existing directives and structures.

## TeX

TeX is a typesetting system invented by Donald Knuth in 1978. As a typesetting language it is a concrete media domain language. But Knuth also included a macro language in TeX which allows users to define new commands in terms of existing commands. (I say commands because that is the term used in TeX. Markup in the media domain tends to be much more imperative than markup in the subject domain, which is entirely descriptive, so “commands” is an appropriate name for TeX’s tags.) This macro language has been used to extend TeX, most notably in the form of LaTex, a document-domain language that we looked at in Chapter 42, *Heavyweight markup languages*.

As we noted with RestructuredText, extension of a language is not the same thing as constraint. Introducing new commands does not create a constraint mechanism.

## Sam

While lightweight languages provide great functional lucidity, they suffer from limited extensibility, which generally requires writing code, and a general lack of constraint mechanisms. I believe that a fully extensible, fully constrainable lightweight markup language would be a valuable addition to the structured writing toolkit. This is why I have been developing SAM, the markup language used for most of the examples in this book. (The book was also written in SAM.)

As described in Chapter 40, *Markup*, SAM is a hybrid markup language which combines implicit syntax similar to MarkDown with an explicit syntax for defining abstract structures called blocks, recordsets, and annotations, and with specific concrete markup for common features such as insertions, citations, and variable definitions.

SAM, like XML, for defining specific tagging languages. However, all languages defined in SAM share a small common base set of text structures for which SAM provide concrete syntax. This allows SAM to combine lightweight syntax of the most common text structures with the ability to define specific constrained markup languages for particular purposes, particularly subject domain languages.

So when I say that this book is written in SAM, what I mean was that it is written in a specific tagging language based on SAM that I created for the purpose of writing this book. It is a mostly document-domain language, because discursive texts like this one don't offer a lot of scope for useful subject domain constraints, but it make significant use of subject domain annotation, which was used for auditing, linking, cross referencing, formatting, and indexing purposes.

SAM is designed to be extensible and constrainable through a schema language (this is not complete at time of writing, but hopefully will be available by the time you read this). The intent is that the schema language should be able not only to define and constrain new block structures, but to constrain the use of the concrete structures as well, and to constrain the values of fields using patterns.

SAM is not designed to be nearly as general as XML in its applications. As a result, its syntax is simple and more functionally lucid and it schema language should also be simpler and make it much easier for writers to develop their own SAM-based markup languages.

I use SAM for the majority of the examples in this book because SAM is designed to make structure clear and that is all I have needed to do in most examples. Naturally, to write in SAM you would need to know more about the rules of the language, but you should be able to read a typical SAM document and understand its structure with little or no instruction.

This is similar, but not identical, to the aim of mainstream concrete and hybrid languages such as Markdown and Restructured Text, which is to have the source file be readable as a document. In other words, they strive to make the document structure clear from the markup. They are document domain languages, and they strive to make sure the the markup expresses the document structure they create in a way that is readable.

SAM has the same goal, except that SAM was designed for creating subject domain languages. As such, it is designed to make the subject domain structure of the document clear to the reader. A SAM document may not look as much like a finished document as a Markdown or reST document. For example, it does not use underlines to visually denote different levels of header. Instead, it focuses on creating a hierarchy of named blocks and fields. In doing so, it uses the kind of markup people commonly use to create named blocks of text and to express a hierarchical relationship between them. Blocks are introduced with a name followed by a colon, and hierarchy is expressed through indentation.

examples: Basic SAM structures

example: Paragraphs

The is a sample paragraph. It is inside the {block}(structure) called `example`. It contains two {annotations}(structure), including this one. It ends with a blank line.

This is another paragraph.

example: Lists

Then there is a list:

1. First item.
2. Second item.
3. Third item.

example: Block quote

Next is a block quote with a {citation}(structure).

"" "[Mother Goose]  
Humpty Dumpty sat on a wall.

SAM is an open source project. A description of the language and a set of associated tools are available from <https://github.com/mbakeranalecta/sam>.

## SPFE

SPFE is another project of mine. It is designed to be a framework for implementing structured writing algorithms and its structure followed the model I laid out in Chapter 11, *Processing Structured Text*. It is tempting to compare it to DITA as an information typing architecture, but as I commented before, architectures are not necessarily parallel to each other and often differ in their emphasis. SPFE is principally designed for subject domain markup. As such, it does not start with a generic document domain topic type like DITA. SPFE does not require any particular schema, though it does require that schemas meet certain constraints.

But SPFE does not leave it entirely to you to develop schemas from scratch. Instead, it supports building schemas from pre-built components. The pre-built components include a collection of semantic blocks and the default processing code for each stage of the publishing algorithm. SPFE also allows you to define your own reusable structure components with processing code. This is, essentially, extensibility through composition,

rather than extensibility through specialization (as in DITA) or extensibility through processor extension (and in reStructuredText). Constraints are supported through normal schema mechanisms and by selecting the minimal required structural components for the individual case.

By strictly segregating the presentation and formatting layers, SPFE reduces the effort required to process custom markup formats. Custom format are processed to a common document-domain markup language which it then processed to all required media-domain output formats. The SPFE Open Tool Kit includes a basic document domain language for this purpose, but you can also use DocBook or DITA in this role, allowing you to take advantage of their existing publishing capabilities. This also allows you to install SPFE as an authoring layer on top of an existing DITA or DocBook tool chain.

To create a subject-domain markup language in SPFE, therefore, all you have to define for yourself are the key subject-domain fields and blocks that are essential to your business. All the other elements you need, such as paragraphs, lists, tables, and common annotations, you can include from the pre-built components, along with their default processing code.

Among its default processing steps, the SPFE process includes the subject-based linking algorithms described in Chapter 17, *Linking* and the subject-based composition and architecture algorithms described in Chapter 14, *Composition* and Chapter 24, *Information Architecture*, including bottom-up information architecture. The conformance and audit algorithms are well-supported as well.

While it has support for reuse, SPFE is not as focused on content reuse or content management as DITA. While it can produce books and top-down information architectures, its main focus is hypertext and bottom-up information architectures. SPFE does not define or require maps as an assembly mechanisms, though you could implement maps in SPFE if you wanted them, nor does it have a standardized metadata architecture like DITA's subject schema. SPFE's processing model is modeled on a software build architecture and it is designed to work well with a version control system system as a repository rather than a content management system. One of its key design objectives is that author should have to know little or nothing about how SPFE works.

Both SAM and XML are supported as markup syntax for SPFE, and you can freely mix and match SAM and XML content.

SPFE is an open source project available from <http://spfeopentoolkit.org>.

---

# Chapter 44. Constraint Languages

Structured writing is about applying constraints to content and recording the constraints that the content follows, both to constrain what writers write and to constrain how algorithms interpret the content. This requires some way to express constraints in a formal and machine readable way. The way you do this is with a schema language. Schema languages are, quite simply, languages for expressing constraints.

For a concrete markup languages like MarkDown, the content constraints are established in the code of the MarkDown processor. They are validated when the processor parses the MarkDown file. (In practice, though, MarkDown does no meaningful validation. Anything it does not recognize as markup, it simply outputs as text.)

For abstract markup languages like XML you define structures yourself (or use an existing definition). Basic XML syntax is validated by the parser, but the definition of constraints is the business of a schema language. The validation of those constraints is the business of a piece of software called a “validator”.

A schema language is a structured language for defining structured languages. The schema for a markup languages says what structures are allowed and in what order and relationship. A given document either conforms to those constraints or it does not.

Here is an example of a schema in a schema language called RelaxNG, which is one of several schema languages available for defining XML-based markup languages:<sup>1</sup>

```
<element name="book" xmlns="http://relaxng.org/ns/structure/1.0">
 <oneOrMore>
 <element name="page">
 <text/>
 </element>
 </oneOrMore>
</element>
```

This example defines two elements and three constraints. The first element is called book and the second is called page. The constraints are:

- The page element must occur inside the book element. (Because the page element is defined inside the book element structure.)
- There must be at least one page element inside the book element, and there can be more. (Because the page element is defined inside an RNG oneOrMore element.)

---

<sup>1</sup>[https://en.wikipedia.org/wiki/RELAX\\_NG#XML\\_syntax](https://en.wikipedia.org/wiki/RELAX_NG#XML_syntax)

- Text can occur inside the page element, but not directly inside the book element. (Because the RNG text element occurs inside the definition of the page element, but not as a direct child of the book element definition.)

Thus if an author wrote:

```
<book>Moby Dick
 <page>Call me Ishmael. Some years ago- never mind how long
 precisely- having little or no money in my purse, and nothing
 particular to interest me on shore, I thought I would sail
 about a little and see the watery part of the world.</page>
</book>
```

the validator would report an error because the words “Moby Dick” are directly inside the book element and text is not allowed in that position.

There are several different schema languages for XML each of which is capable of expressing and enforcing different sets of constraints. It is not unusual to combine different schema languages to more completely constrain a markup language. In particular, it is not uncommon to use a schema language called Schematron in concert with other schema languages like RNG or XSD. While most schema languages work by modeling the structure of a document, as in the RNG example above, Schematron works by making assertions about the structure in a language called XPath. A Schematron schema would make a lousy guide for authoring, but it can test assertions, and therefore enforce constraints, that no other schema language can.

Here is a simple Schematron example. It defines one of the constraints listed above, namely that a book element must contain at least one page element.

```
<schema xmlns="http://purl.oclc.org/dsdl/schematron" queryBinding="xsl">
 <pattern>
 <title>Book constraint</title>
 <rule context="book">
 <assert test="page">A book must contain at least one page
 </rule>
 </pattern>
</schema>
```

The rule says that in the context of the element book the assertion that there is an element page must be true. If it is not, the message within the assert element will be displayed.

In the RNG schema example, the schema is essentially a template that models the hierarchy of the document. The constraints are consequences of the structure of that hierarchy. (The book element cannot contain text because there is no place for text in that part of the hierarchy.) Only documents that conform to the template are valid.

In the Schematron example, the schema is a series of constraint statements. The hierarchy of the document is a consequence of meeting all the constraints. Any document is valid as long as it conforms to all the stated constraints.

There are essentially two ways to describe constraints. One is to start from the basis that nothing is allowed unless there is a specific rule allows it. The other is to start from the basis that everything is allowed unless there is some specific rule that forbids it. Schematron is based on the latter doctrine. It says anything is valid as long as it passes a given set of tests. Other schema languages take the former approach. If you validate an XML document against an empty XSD schema, for instance, the validator will say it is invalid because it can't find the definition of the root element. In essence, therefore, their constraints are described as a set of permissions. The writer is then constrained to stay within the boundaries of what is permitted.

Starting on the basis that everything is allowed is not generally a good model for structured writing. You have to write algorithms that can handle everything that is included in a document, so you pretty much have to start with a definitive list of what is allowed and in what combinations. This is why Schematron is seldom used as a standalone schema language. But it can be very useful for qualifying the use of certain structures permitted by other in other schema languages, in particular the kinds of constraints that cannot be expressed in the main schema language.

The wider point here are that there are many occasions in the course of writing and publishing content in which to express and enforce constraints. Some constraints can only be expressed during the publishing process, for instance. An example would be the resolution of a key reference. We can only tell if there is a value or resource corresponding to the key when the synthesis algorithm is compete and all the elements of the published content set are assembled and resolved.

Constraint, therefore, is not something that exists at a single point in the writing and publishing process. It is something you need to consider across the entire publishing process design. In particular, when errors occur at any point in the publishing process, you should trace them back to the point where the constraint that could have prevented them was violated, and ask how that constraint could have been better implemented or enforced.

One of the most basic design principles of any process is that errors should be caught as early in the process as possible, and that any error-prone process should be designed to make errors less likely. Factoring out constraints by moving content to a different domain, as we have noted several times, is an excellent way to reduce errors in structured writing, and to ensure that errors are caught earlier in the process.

It is tempting, of course, to avoid having to design constraints for yourself by adopting an off-the-shelf structured writing system. But you should reflect on exactly how effective such a system can be in enforcing the constraints that matter to your business, and the how much it costs you to handle the errors that the lack of such constraints cause throughout your publishing process.

---

## **Part V. Choices**

---

## Table of Contents

45. Choosing .....	362
Complexity .....	362
Scale .....	364
Tool driven vs. structure driven decision making .....	365

---

# Chapter 45. Choosing Complexity

Content creation, particularly content creation at scale and in a dynamic environment is not simple. Structured writing can deliver many benefits in quality and efficiency, but it cannot magically make the whole process simple.

In fact, structured writing systems don't make any of the complexity of content creation go away. What they do is move that complexity from one place to another. For instance, separating content from formatting moves the complexity of formatting text away from the author and places it on a formatting algorithm, and therefore on the person who writes and maintains the formatting algorithm. Using the subject-domain approach to linking moves the complexity of finding content to link to away from the author of the source content and places it partly on an algorithm, but also on the author of the content that is being linked to, since they must accurately index their content to indicate what subjects it covers.

The point of moving complexity from one place to another is to move each type of complexity to the person or process who is best placed to handle it. Of course, the machinery we use to move this complexity -- structures and algorithms -- also introduces its own complexity into the process. This means that structured writing is actually adding to the total amount of complexity in the system, which makes it all the more important to ensure that that complexity is moved to the right place. It is all too easy for the introduction of a structured writing system to make the content process less efficient and the resulting content of lower quality, simply because it does not move the complexity to the right place or introduces more complexity into critical functions than it removes from them.

For example, separating content from formatting requires that we factor out the formatting and replace it with structures from which the formatting can be derived. But if we turn to a complex document domain markup language like DocBook for this, we will often find that we have made the author's task more complex than it was before.

The very fact the structured writing has the ability to move the complexity around the system, imposing it on one role at the expense of another, makes the task of choosing a structured writing approach all the more complex. For instance, there has been a long history of IT departments choosing content management system on the basis that they were easy for the IT department to install and administer, only to have them be deeply unpopular with users because they pushed complexity out to authors and other users. On the other hand, some groups of authors want to write in uber-simple formats like Markdown, despite the difficulties that its limited structure and capabilities create for an overall publishing process.

Inconsistency, duplication, delay, error, and failure are escape valves for the complexity the no part of the organization is willing to accept. Traditionally, a style guide has become the dumping ground for complexity. Every new rule and requirement gets written into an ever growing style guide.

Change management is another scapegoat for complexity that no one will own. It is very common to blame the failure of structured writing systems on unwillingness of writers to change. The cure, the backers of the system assure us, is more training and more generalization of the change. But what is usually going on in these cases is that the system design has dumped new and unmanageable complexity on whatever group is rebelling -- almost always the writers, either the full time writers or the occasional contributors.

Here is the inescapable fact: complexity is irreducible. It can be moved, not destroyed. When you simplify one processes, therefore, you have added complexity somewhere else in the system. It is all too easy to get caught up in the gains you expect from removing complexity from A and forget to pay attention to what the effect will be of the complexity you have added to B. The right way to design a structured writing system, therefore, is not to focus solely on where you can create simplicity, but where you can move complexity to, and how well, and how reliably, that complexity can be handled in that place.

Complexity that is not handled in the organization is dumped onto the customer. This can take the form of sub-optimal information architectures, such as vast hierarchical organizational schemes that make no sense to the reader, or a lack of linking to allow the reader to easily find their way from one subject to another. It can also manifest itself in poor information design, in topics that do not adequately define their content, leading to users not being able to identify them as the place to find the information they need. It can also manifest itself in simpler things like lack of information, information expressed in the wrong way, or a general lack of consistency that makes navigation and comprehension difficult.

It is natural for designers and implementors to distribute complexity away from themselves.

There are knock-on effects to complexity as well. If authoring is complex, collaboration becomes complex. So does conformance.

The right mix and distribution of complexity and simplicity is likely to be different from one organization to another, even if all parties have a good understanding of where complexity lies and how it can be distributed or minimized. It is beyond the scope of this book to walk you through this decision making process. I hope that what this book has contributed to that process is an understanding of what structured writing is capable of bringing to an organization, and a tool-neutral language for talking about it.

# Scale

The question of whether it is worthwhile to transfer complexity from one part of a process to another often depends on scale. Transferring complexity from a human to a algorithm, for instance, creates a complex task of writing, testing, and maintaining the algorithm. But algorithms don't cost any more to write, test, and maintain based on the amount of content you process with them. You have to pay a human every time they do the task by hand. You only have to pay them once to write an algorithm to do it automatically.

All of the algorithms we have discussed so far can be carried out either by humans or by algorithms. Algorithms are much faster than humans, but algorithms have to be designed, written, and maintained, and they don't have the same capacity as humans to adapt on the fly when conditions warrant it. They also require much more precisely structured inputs than humans, meaning that they require more work from human writers, at least for original content creation. (They may save writers work in all sorts of other ways, as we have seen.)

How do you determine when it is worth the investment to design structures and algorithms and have your writers write structured documents, rather than having them just execute all the algorithms themselves. Issues of consistency and quality are important considerations here. There are also some information designs that are very costly to develop without the aid of algorithms. However, one of the key factors in making this decision is scale.

The issue of scale applies to many systems and many activities. A home kitchen does not work like a restaurant kitchen. The restaurant kitchen has multiple work stations as divides the work up among multiple cooks. They prepare ingredients in advance so that when an order comes in, they can prepare it very quickly. All the equipment, planning, and preparation costs money, but it pays off when you have to process a lot of food orders during a busy dinner shift. By contrast, a cook in a home kitchen starts with basic ingredients and does all the steps themselves to prepare a single meal for themselves or a small group.

The restaurant could not possibly keep up with orders if it worked the way a home kitchen worked. That method of preparation does not scale up to the throughput of a restaurant. Equally, though, the approach of a restaurant kitchen would not work for a home kitchen. The overhead of all the stations, the planning, and the advanced preparation would not be economical for preparing a single small meal. The approaches of a large-scale operation do not scale down to a small-scale operation anymore than the approaches of a small-scale operation scale up for a large organization.

The same is true of writing.

# Tool driven vs. structure driven decision making

I do want to say one thing about the decision making process, though. It is important to decide up front if you are going to have a tool-driven decision process or a process driven by algorithms you want to use.

Of course, your decision making process should ultimately be driven by business requirements. But even a process driven by business requirements tend to be heavily influenced by what we think the available tools are capable of. We don't tend to write requirements for systems unless we have an idea that there are systems out there that are capable of meeting those requirements.

Over the years I have seen many requirements documents for proposed structured writing systems that essentially said that the proposed system must work exactly like Microsoft Word. This is not surprising when the people writing the requirements have used nothing but Word to create content for years. The tools you know shape how you work and what you think of as possible. As Henry Ford is supposed to have said of the Model T, "If I asked customers what they wanted, they would have said faster horses." Even when we are dissatisfied with our current tools, we tend to want the same basic tool only more so. Faster horses. This is why so many structured writing tool vendors literally advertise that their editor looks and feels "just like Microsoft Word". (Not to mention those vendors who create tools that modify Word itself.)

But Microsoft Word is a tool that sits on the boundary between the Media and Document domains. Using Word itself, or something that looks like Word, is usually an attempt to move its use slightly more into the document domain, but as we have seen, the WYSIWYG authoring interface invites a slide back into the media domain by hiding the structured that is supposed to be created and showing only the formatting that is supposed to have been factored out in adopting the document domain.

It is little to be wondered then that the structured writing tools that have been popular in the market to date have been predominantly document domain tools, and have tended, like DocBook, to be very loosely constrained. (It is much easier to write an XML document in a WYSIWYG editor if the underlying structures are minimally constrained, since it lets you insert whatever bit of formatting you want anywhere you want.)

Even with tools like DITA, which, while it is still fundamentally a document domain system, is more constrained, and capable of being constrained further, tend to be used in its generic out of the box form and with a Word-like WYSIWYG interface.

Thus even when a decision-making process is based on business requirements rather than specific tools, it is often tacitly driven by existing tool sets and ways of doing things, because those existing tools and processes shape our view of what the business

requirements actually are. We don't ask for a way to get from DesMoins to Albuquerque, we ask for a faster horse that eats few oats.

By laying out the structured writing algorithms as algorithms rather than as features of a tool, and by showing the different ways in which those algorithms can work based on content in the different structured writing domains, I hope I have provided you with a way to break free of the dead hand of current tools and processes in thinking about your structured authoring decision.

If so, this will allow you to have a decision making process based on algorithms and a knowledge of how the algorithms and the structure that support them can work together to address your business problems, while putting the complexity in the right place for your organization.

If you can come up with a set of requirements that say that you want to execute a given set of algorithms with a given level of reliability and to distribute the complexity of the system to an appropriate set of roles, then you will have a much more neutral basis on which to approach vendors or consultants.

Every tool in the market place, from end to end systems, to individual tools like XML editors, to frameworks and tools kits, represents an encapsulation of algorithms and the structures that support those algorithms. Not only do they represent a view about which algorithms are most important, but also a view about how they should be implemented -- which ultimately means how and to whom the complexity attendant of the algorithm should be distributed. Every tool, therefore, can be represented as a set of algorithms and structures, just like your set of requirements.

The decision making process then comes down to seeing if the best match between capabilities and requirements is to be found in a single end to end system, in a combination of separate tools, in modifying or building upon an existing framework, or in building some or all of the components of your system from scratch. When talking to any tool vendor, you can ask them to demonstrate how their tool supports each algorithms you care about, how reliable that implementation is, and where it distributes the complexity.

While individual circumstances can vary greatly, there are general patterns in how the use of different domains to implement algorithms distribute complexity within a system. Media domain systems are simple for small things, but provide no help for any kind of management or single sourcing algorithm, distributing all the complexity of those functions to the author and content manager. (Which results in complex content management systems.) They simplify ad hoc formatting of individual items but distribute all the complexity of conformance and consistency to authors. In short, they are simple in themselves but do nothing to alleviate the complexity of any other part of the publishing process.

The document domain distributes formatting complexity away from authors to the creators of publishing routines. However, a different kind of complexity can easily take

its place if the number of document domain structures that people have to remember and use grows in order to support complex output requirements. If single sourcing is required, the document domain again distributes some of the complexity away from authors to the creators of publishing algorithm, with the limits we have noted in regard to differential single sourcing. By itself, the document domain, like the media domain, does little to alleviate other parts of the publishing process.

The management domain, in concert with the document domain, enables a range of content management features that would otherwise be either tedious or impossible. Content reuse is probably the prime example of this. Content reuse in pure media-domain or pure document-domain content is generally so complex that it would not be attempted without the introduction of management domain structures. However, management domain structures distribute a huge amount of complexity to authors, effectively shutting out contributors without specialized training and tools.

The subject domain distributes a huge amount of complexity away from authors, and, by making so many of the structured writing algorithms more reliable, it also distributes a lot of complexity and effort away from content managers. However it distributes that complexity to the information architects and tools people who design the subject domain languages and write the algorithms that process them.

Complexity is by no mean the same thing as effort. Structured writing systems that are well designed to support appropriate algorithms can reduce overall effort considerably, which significantly improving quality. But where they place complexity matters. Even if a task requires less effort, adding complexity to it changes how the person assigned to that task works, and how they need to be qualified and trained. It is important to appreciate how the distribution of complexity and effort in the system you choose affects the dynamics and composition of your team.

As is no doubt apparent to you by this point, I am a strong believer in distributing the complexity of a structured writing system away from writers. The reason for this is simple. When structured writing systems distribute complexity toward writers, they don't merely add a new and complex task that must be learned, they impose that complexity directly on the activity of writing itself. It is not really possible for the writer to segment the writing process from the process of creating structure. They are too bound up with each other. The whole point of structured writing is that they should be bound up with each other. Yet writing is an activity requiring the whole of ones attention. Any unavoidable division of that attention directly detracts from the quality of writing.

Any avoidable complexity should be purged from the writing process itself. Not all markup represents an addition of complexity, however. Effective subject domain markup can work to direct the writers attention into the appropriate channels, distributing design complexity away from the writing task and leaving more attention available for the task of writing itself. (This is the basis on which we use forms for all the things we use forms for.)

Unfortunately, structured writing systems are often designed with other priorities in mind. In particular, they are often designed with content management priorities in mind. Structures are designed to support content management algorithms with little attention paid to how those structures distribute complexity towards authors.

In the end, the question of where complexity is distributed in your system is at least as important, if not more so, than the question of how much effort is avoided. The wrong distribution of complexity can not only undermine quality, it can also undermine the attempt to reduce effort. Complexity in the wrong place not only undermines the productivity of those saddled with it, it also undermines the reliability of every other algorithm, thus undermining the attempt to reduce effort and cost in those algorithms.

The economics of this decision are clearly complex. You may decide that the cost of creating and maintaining the most appropriate algorithms and structures is not worth the cost or quality improvements they promise. But hopefully that decision can be made with a full appreciation of the benefits that those algorithms are capable of delivering. But whatever you decide, make sure you understand how complexity is being distributed in the systems that you implement, and very conscious of the ability of those you are distributing it to handle it, and the effect it will have on their productivity and the reliability of their work.

---

# **Index**

---

# Colophon

## About the Author

Mark Baker is a twenty-five-year veteran of the technical communication industry, with particular experience developing task-oriented, topic-based content, and in designing and implementing structured authoring systems. He is also a frequent speaker on matters related to technical communications and structured authoring, and contributes to several publications in the field. Mark is currently President and Principal Consultant for Analecta Communications, Inc. [<http://analecta.com/>] in Ottawa, Canada.

Mark's blog, Every Page is Page One [<http://everypageispageone.com>] is focused on the idea that, in the context of the Web, Every Page is Page One, that the future of technical communication lies on the Web, and that to be successful on the Web, technical communicators cannot simply publish traditional books or help systems, they must create content that is native to the Web.

## About XML Press

XML Press (<http://xmlpress.net>) was founded in 2008 to publish content that helps technical communicators be more effective. Our publications support managers, social media practitioners, technical communicators, and content strategists and the engineers who support their efforts.

Our publications are available through most retailers, and discounted pricing is available for volume purchases for business, educational, or promotional use. For more information, send email to [orders@xmlpress.net](mailto:orders@xmlpress.net) [<mailto:orders@xmlpress.net>] or call us at (970) 231-3624.

---

# **Appendix A. Copyright and Legal Notices**

Structured Writing for Content Quality

Mark Baker

Printed in the United States of America.

Copyright © 2017 Mark Baker

All rights reserved. No part of this book may be reproduced or transmitted in any form or by any means without the prior written permission of the copyright holder, except for the inclusion of brief quotations in a review.

## **Credits**

Cover Image:

Cover Background:

Foreword:

## **Disclaimer**

The information in this book is provided on an “as is” basis, without warranty. While every effort has been taken by the author and XML Press in the preparation of this book, the author and XML Press shall have neither liability nor responsibility to any person or entity with respect to any loss or damages arising from the information contained in this book.

This book contains links to third-party web sites that are not under the control of the author or XML Press. The author and XML Press are not responsible for the content of any linked site. Inclusion of a link in this book does not imply that the author or XML Press endorses or accepts any responsibility for the content of that third-party site.

## **Trademarks**

XML Press and the XML Press logo are trademarks of XML Press.

All terms mentioned in this book that are known to be trademarks or service marks have been capitalized as appropriate. Use of a term in this book should not be regarded as affecting the validity of any trademark or service mark.